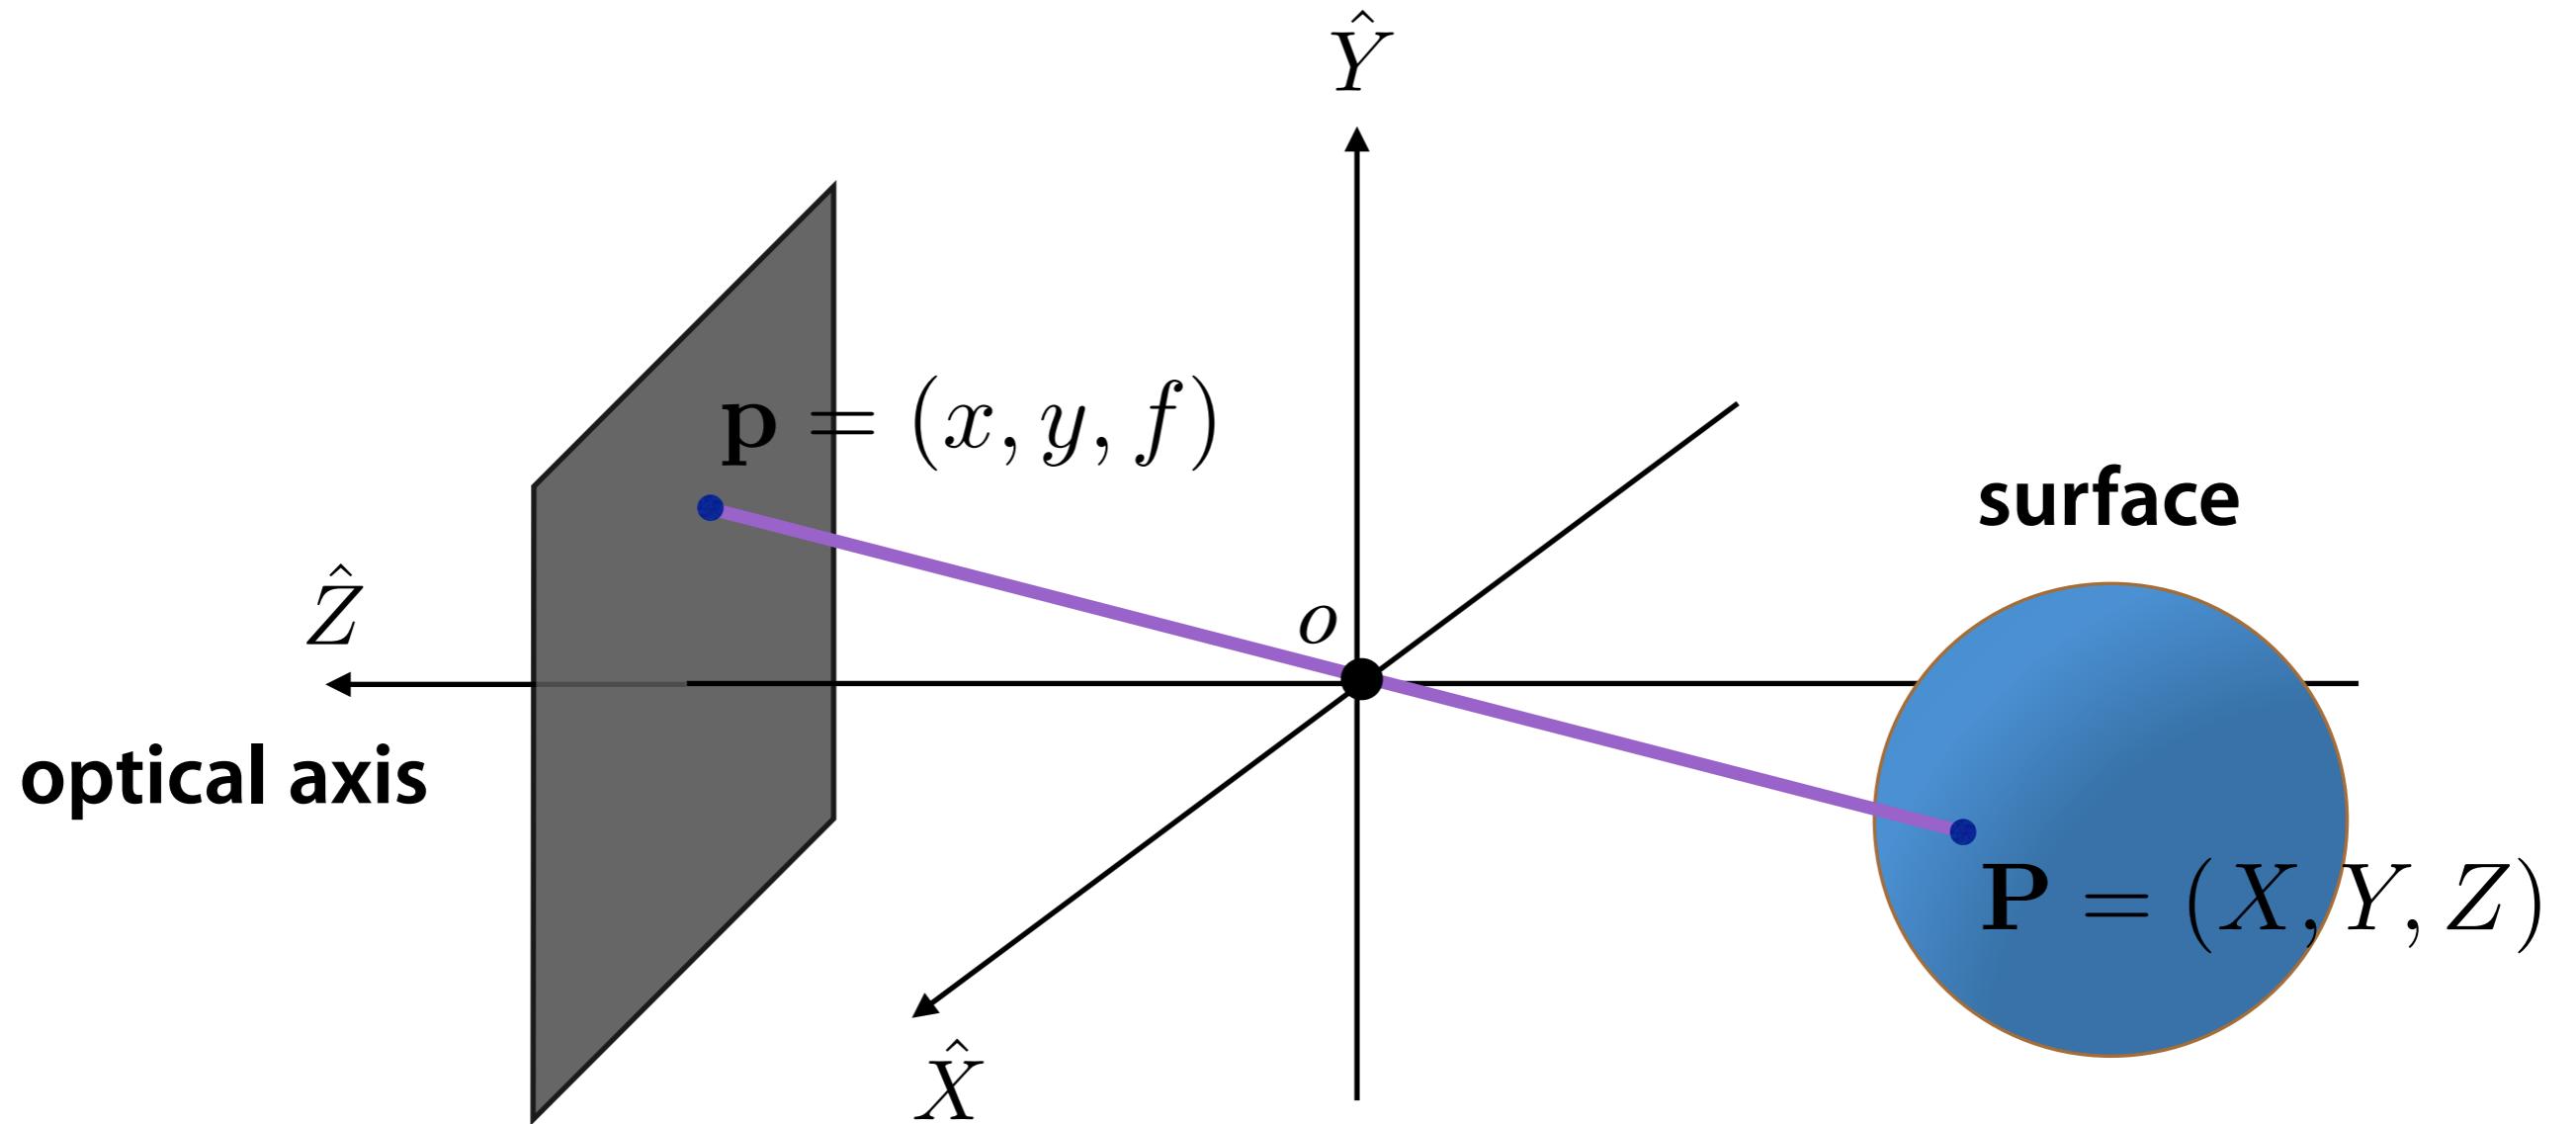


Intro to

# Computer Vision

with Prof. Kosta Derpanis

## Image filtering



# images as functions

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}$$

images as functions

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$



# images as functions

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$\bullet \begin{pmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{pmatrix}$$

**filter**

$I$

**image**

$\phi$

$\phi(I)$

**image**

Image

**NOISE**



original



**Gaussian noise**

image  
noise

$$I(x, y) = I_{\text{ideal}}(x, y) + \eta(x, y)$$

where

$$\eta(x, y) \sim \mathcal{N}(\mu = 0, \sigma)$$

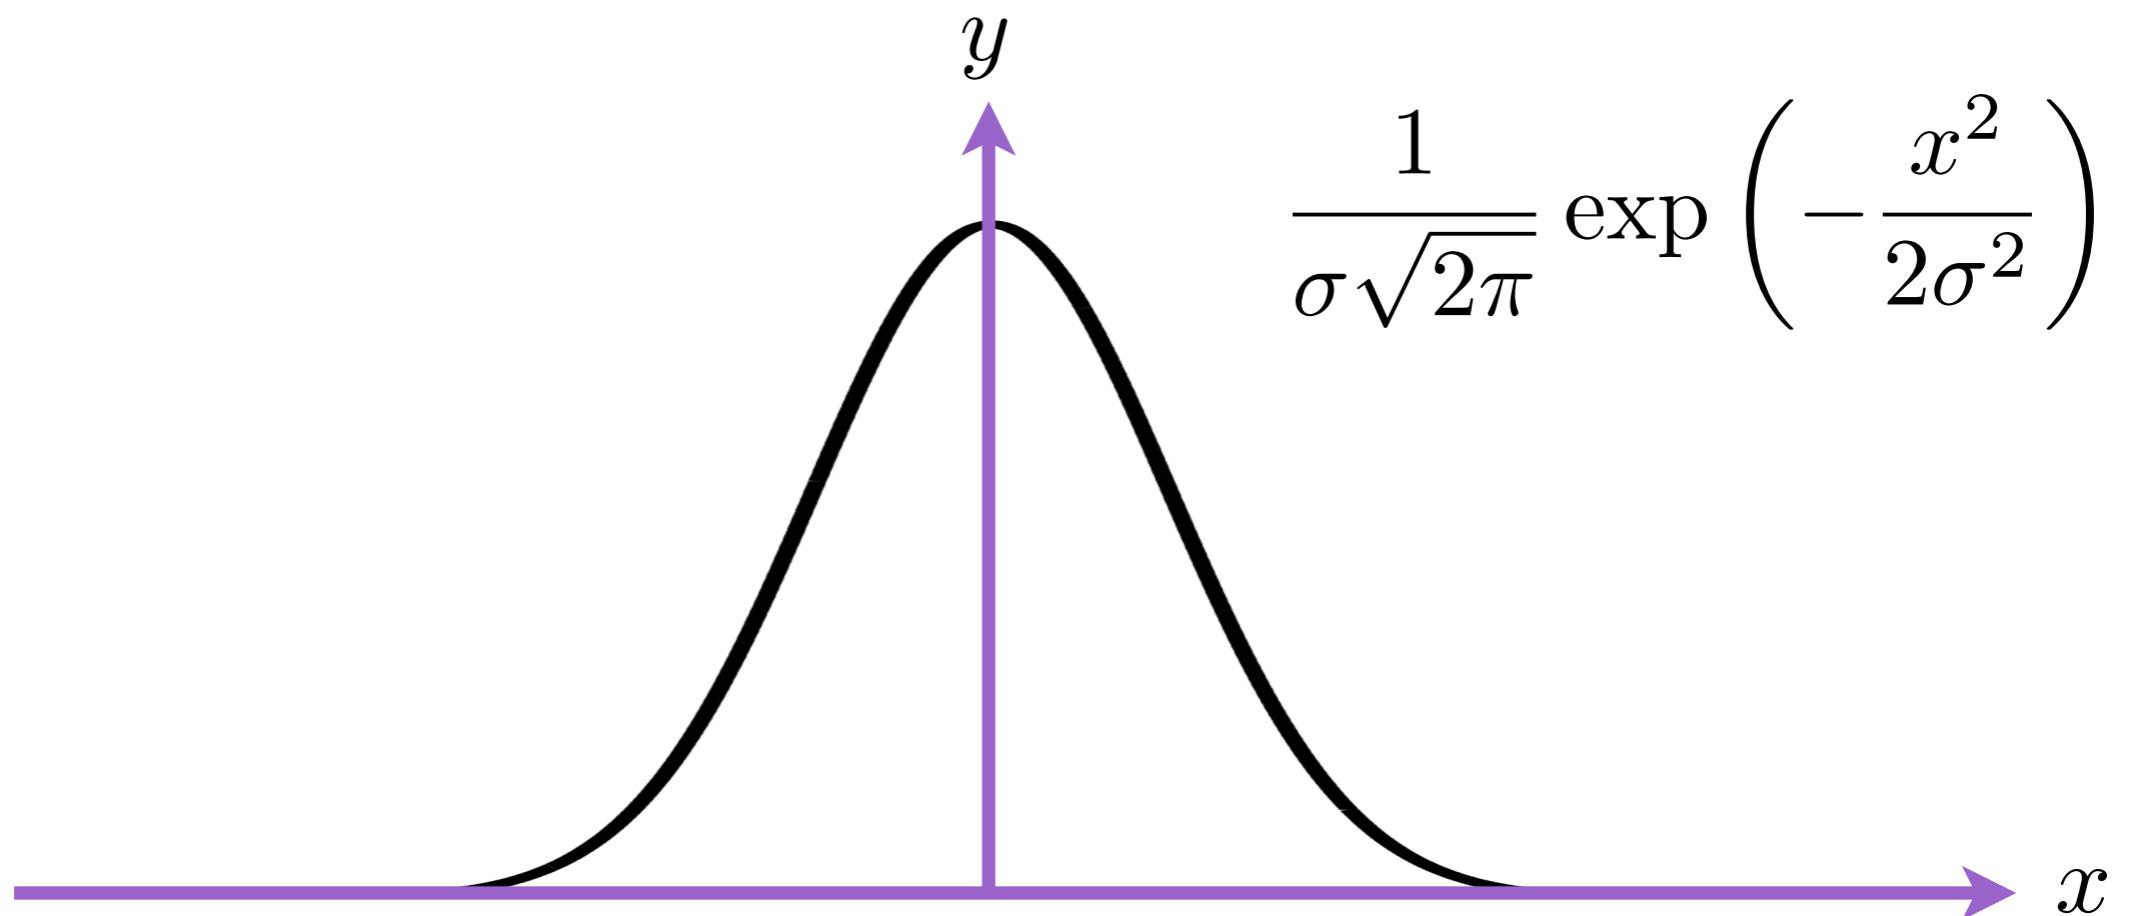
assumed **independent** and **identically distributed**

I.I.D.

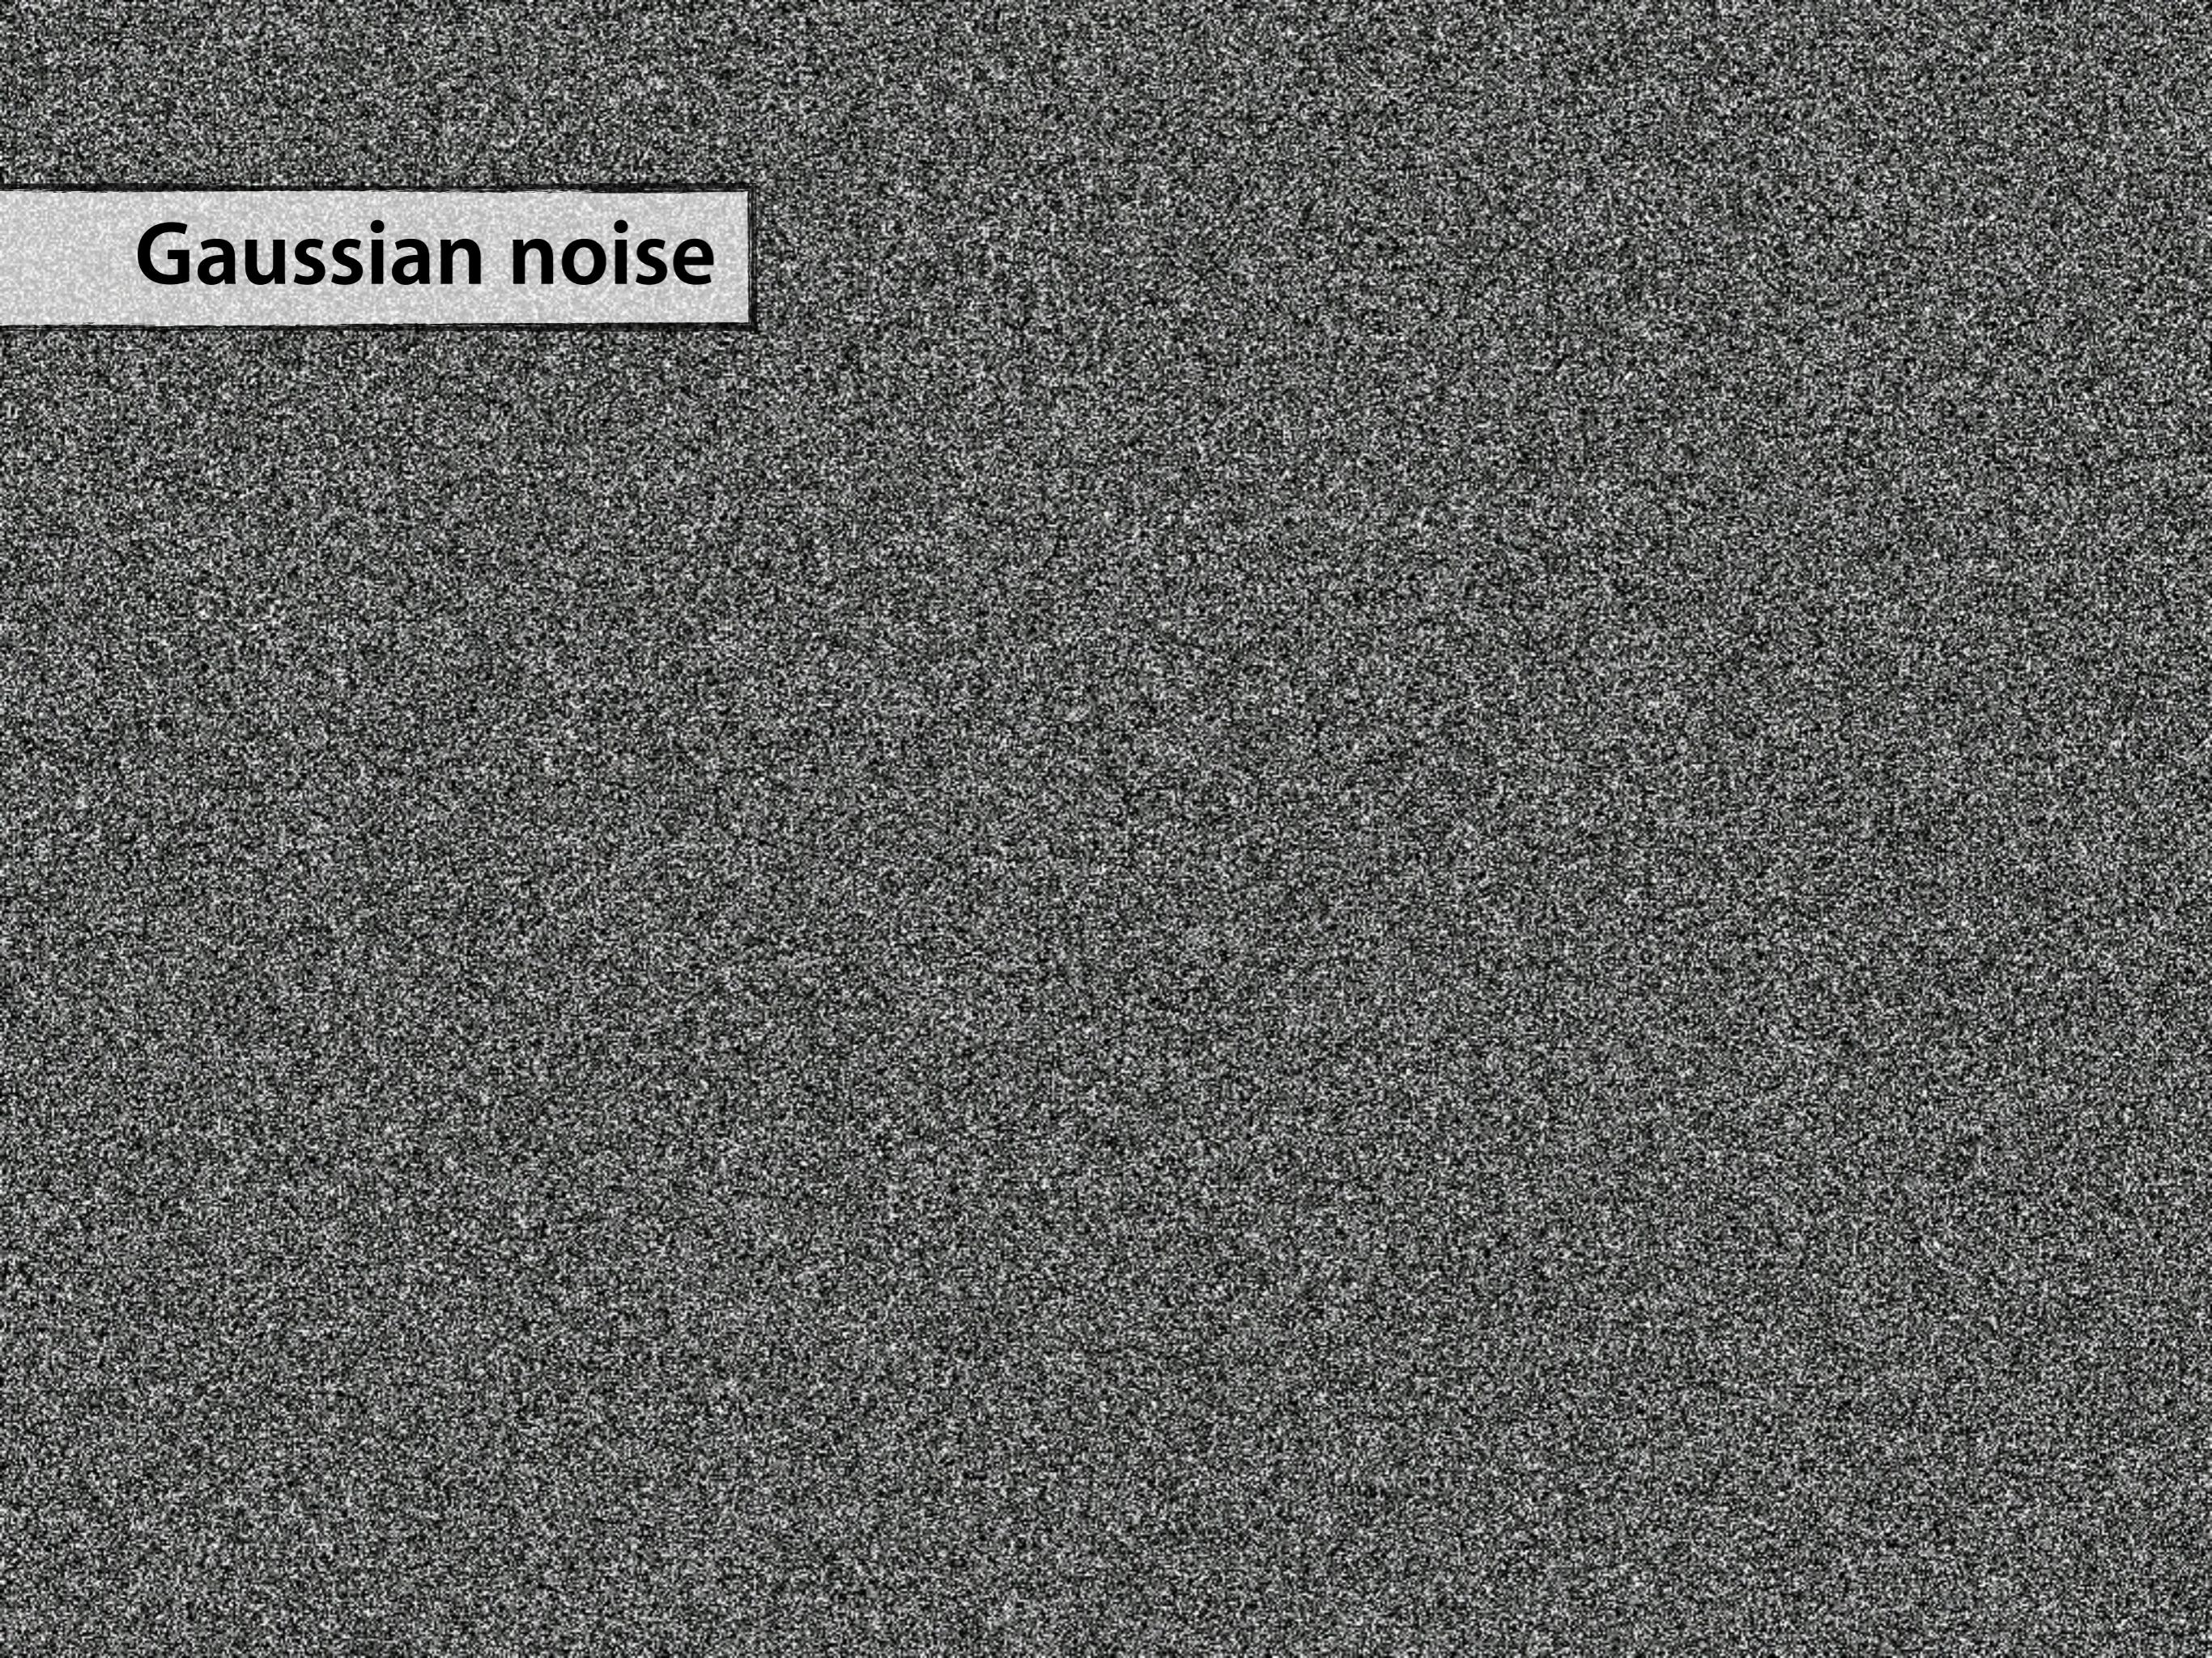
$$I(x, y) = I_{\text{ideal}}(x, y) + \eta(x, y)$$

where

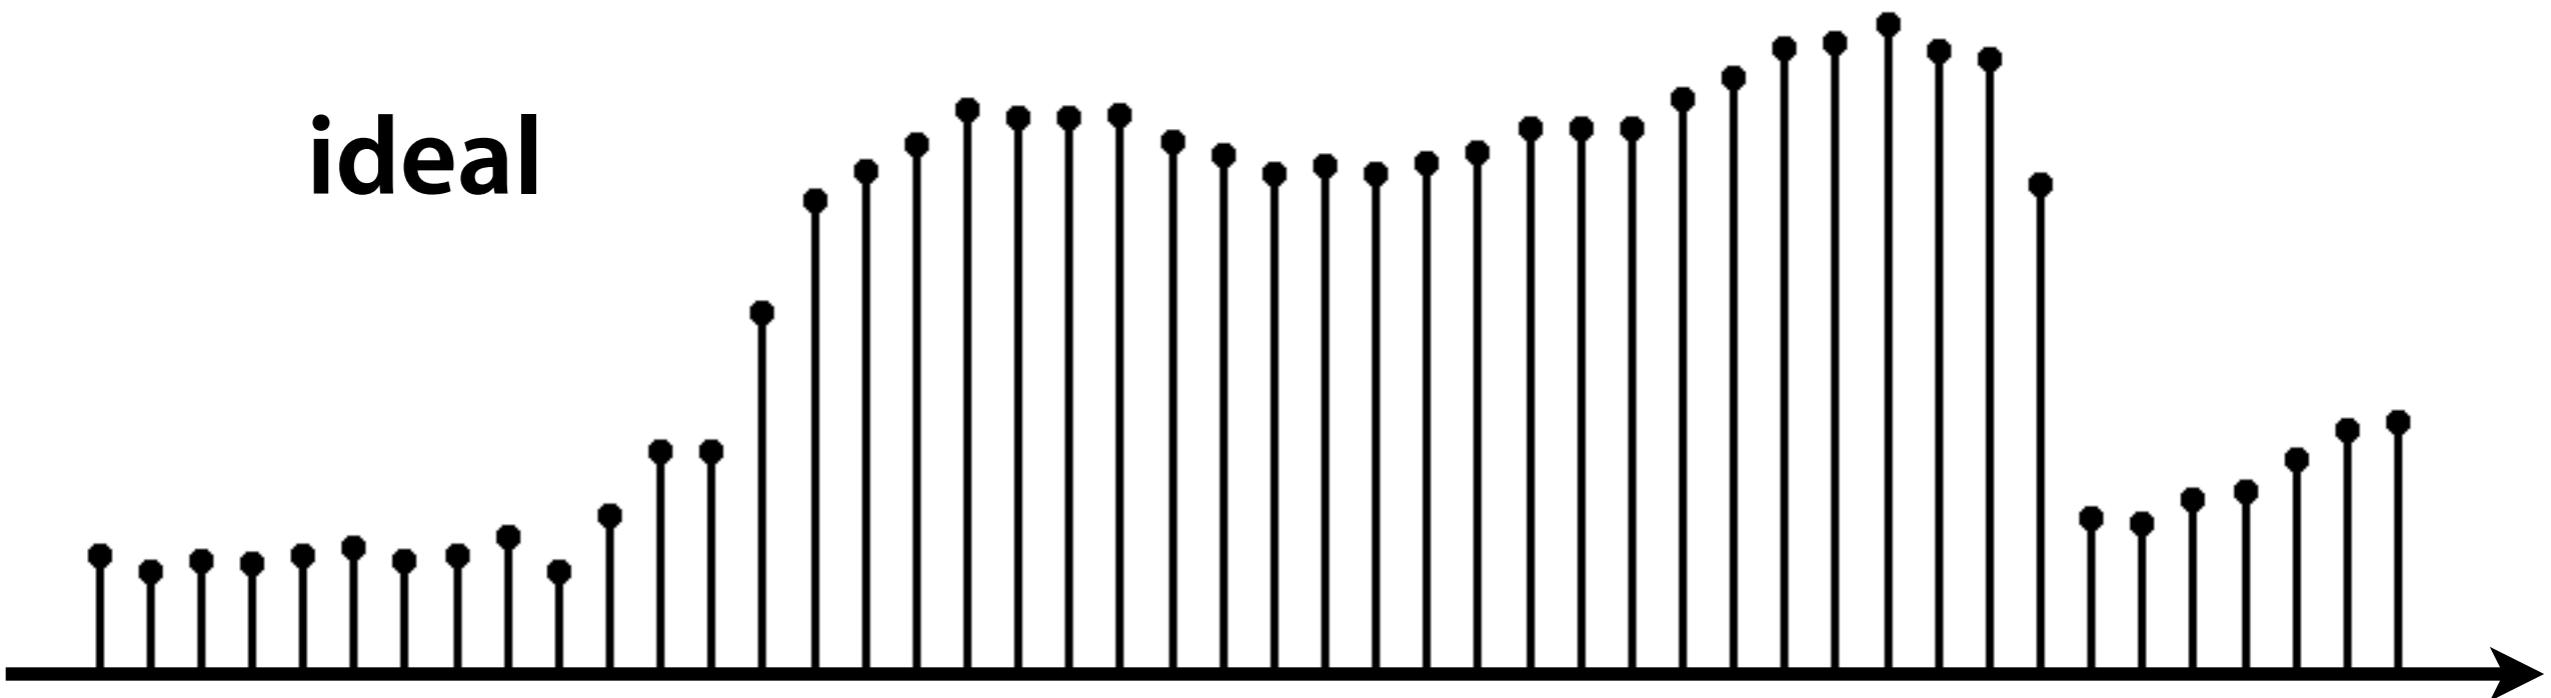
$$\eta(x, y) \sim \mathcal{N}(\mu = 0, \sigma)$$



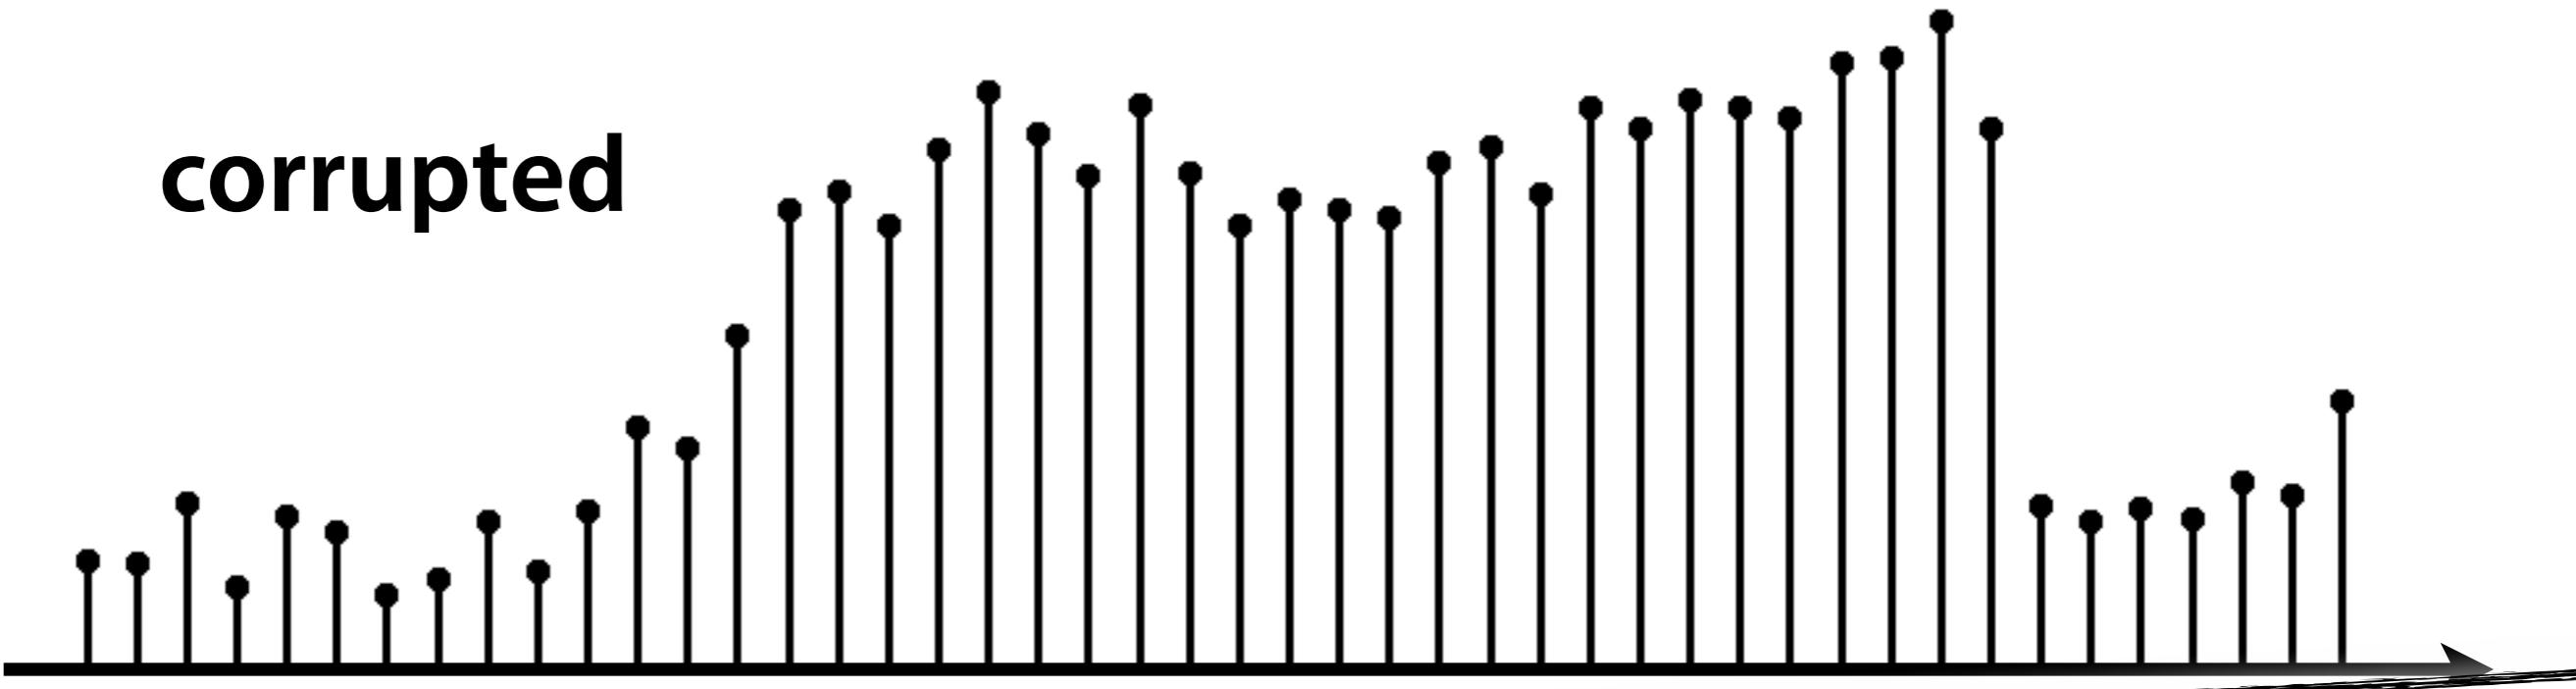
# Gaussian noise



**ideal**

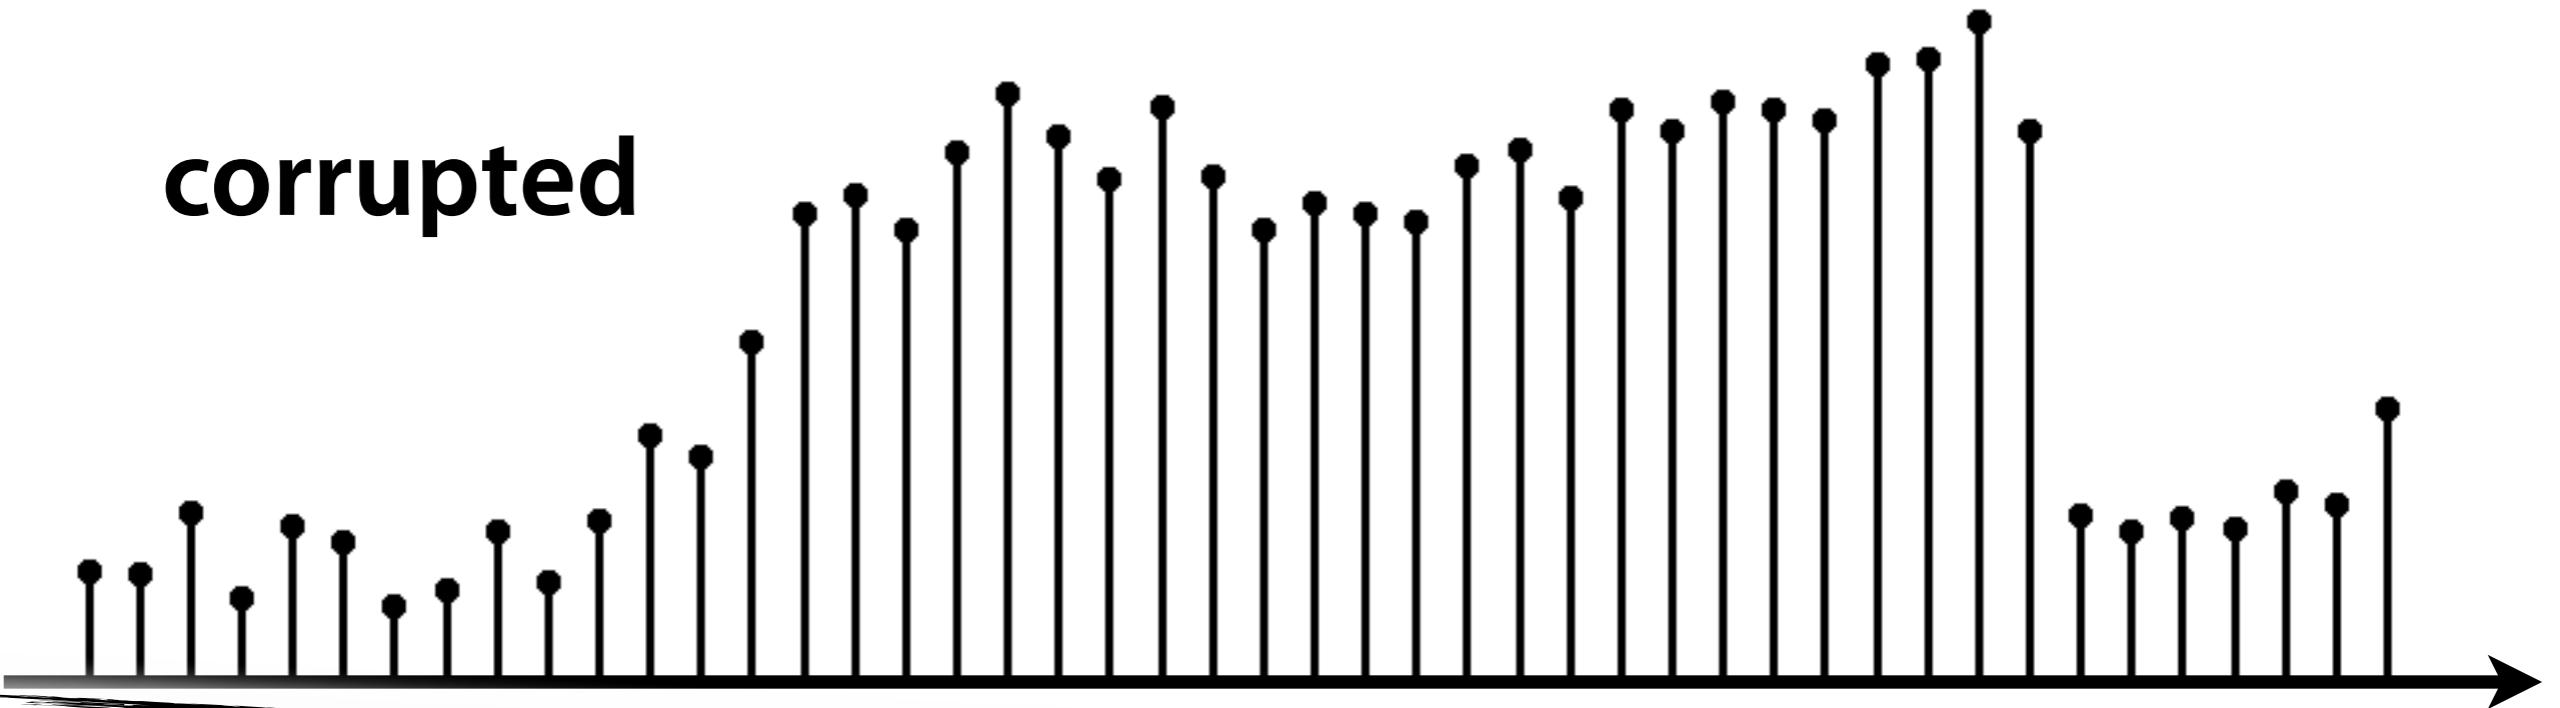


**corrupted**



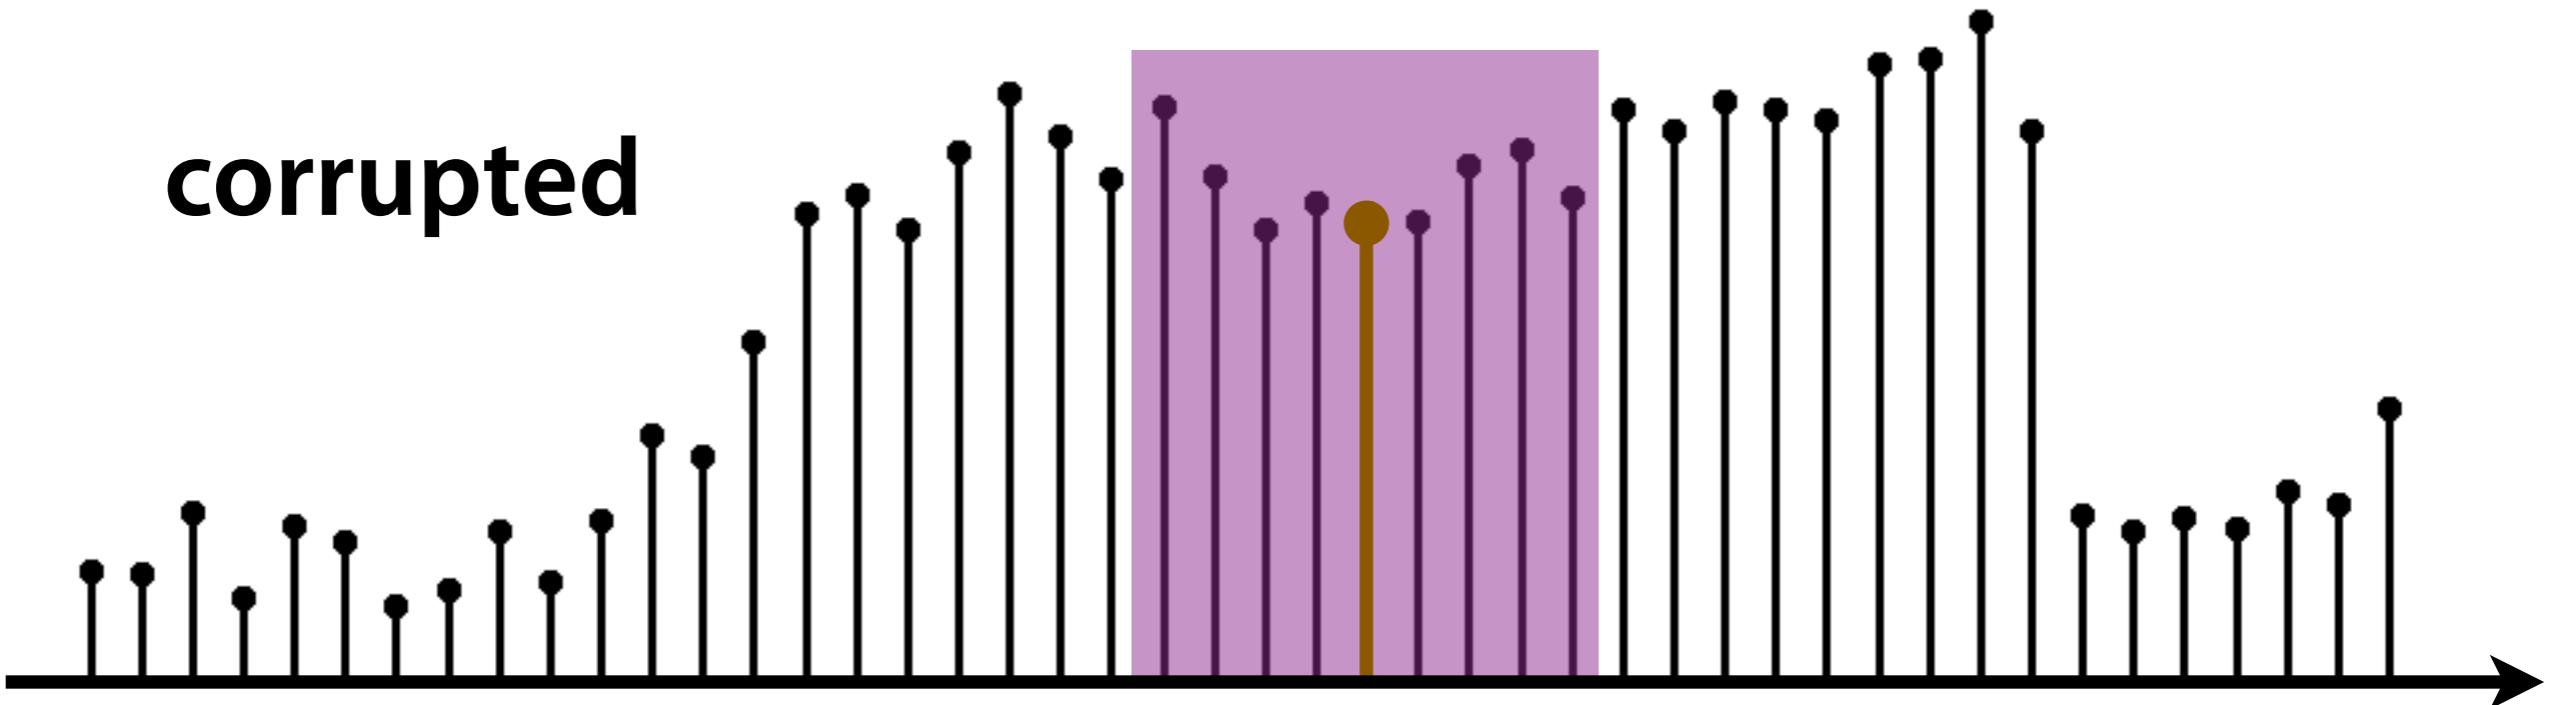
**How can we remove the noise?**

**corrupted**



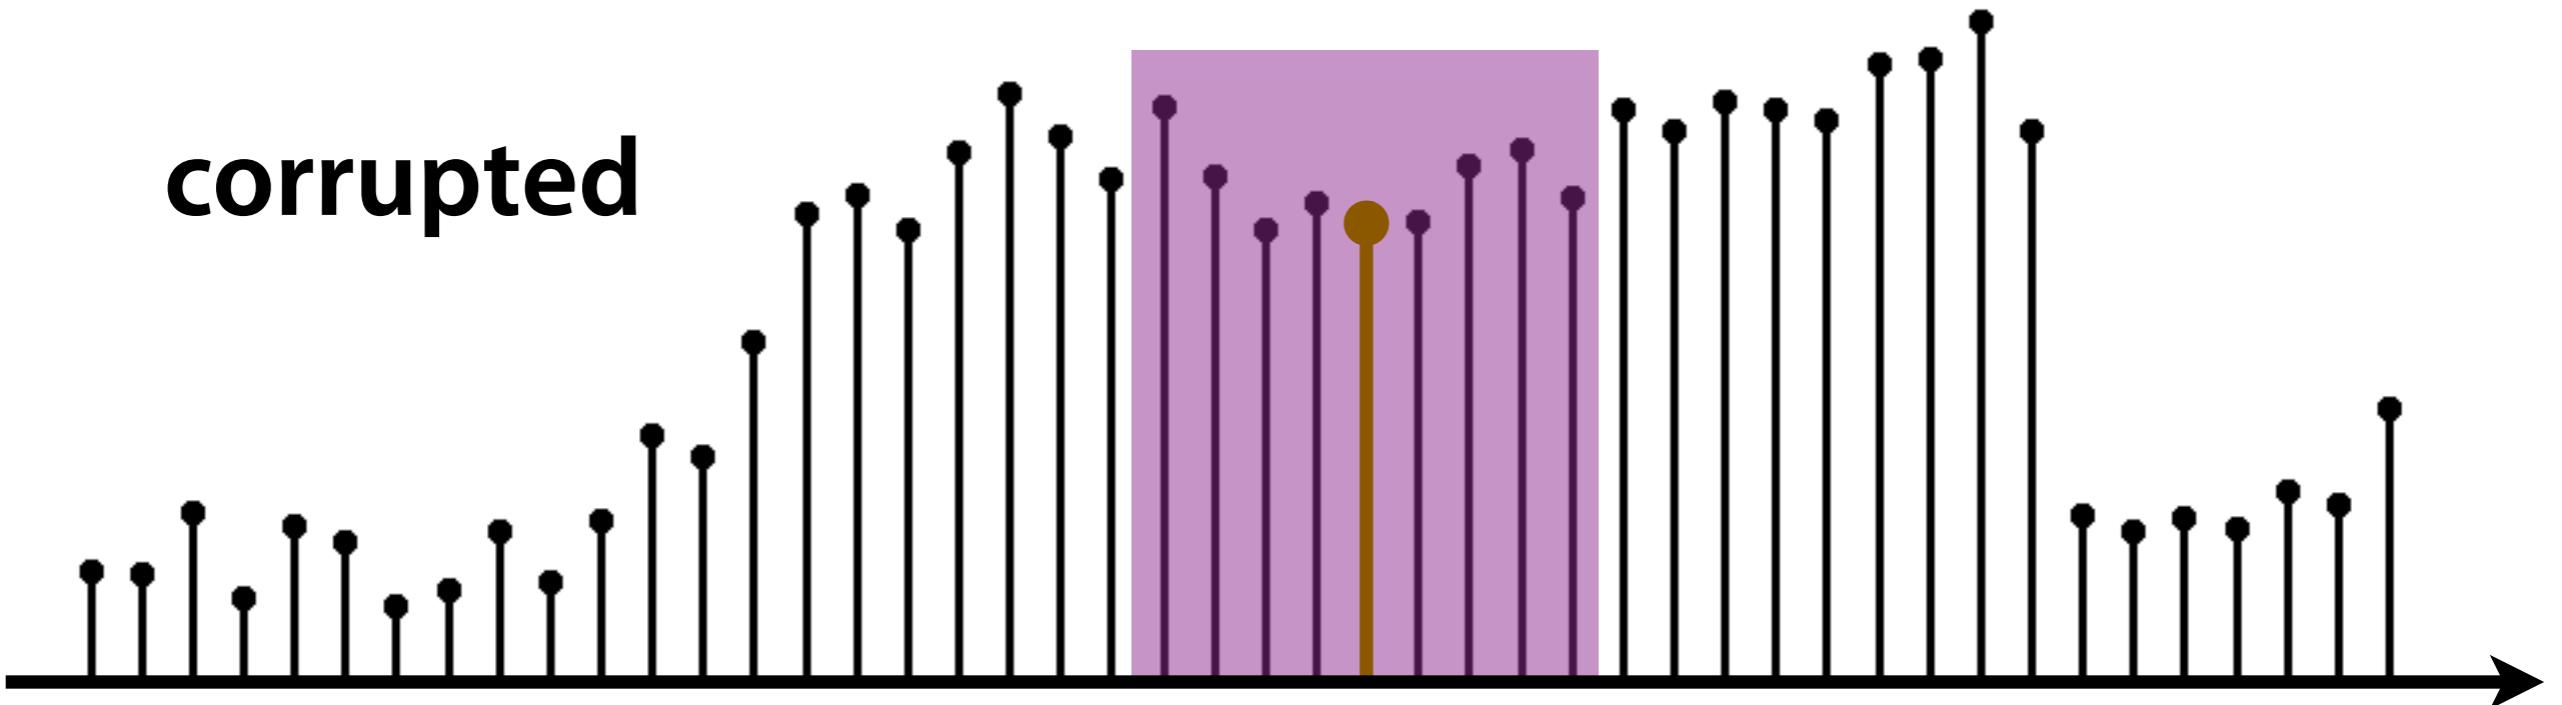
**assume noise is I.I.D. and pixel neighbours are similar**

**corrupted**



replace each pixel with an average of its neighbours

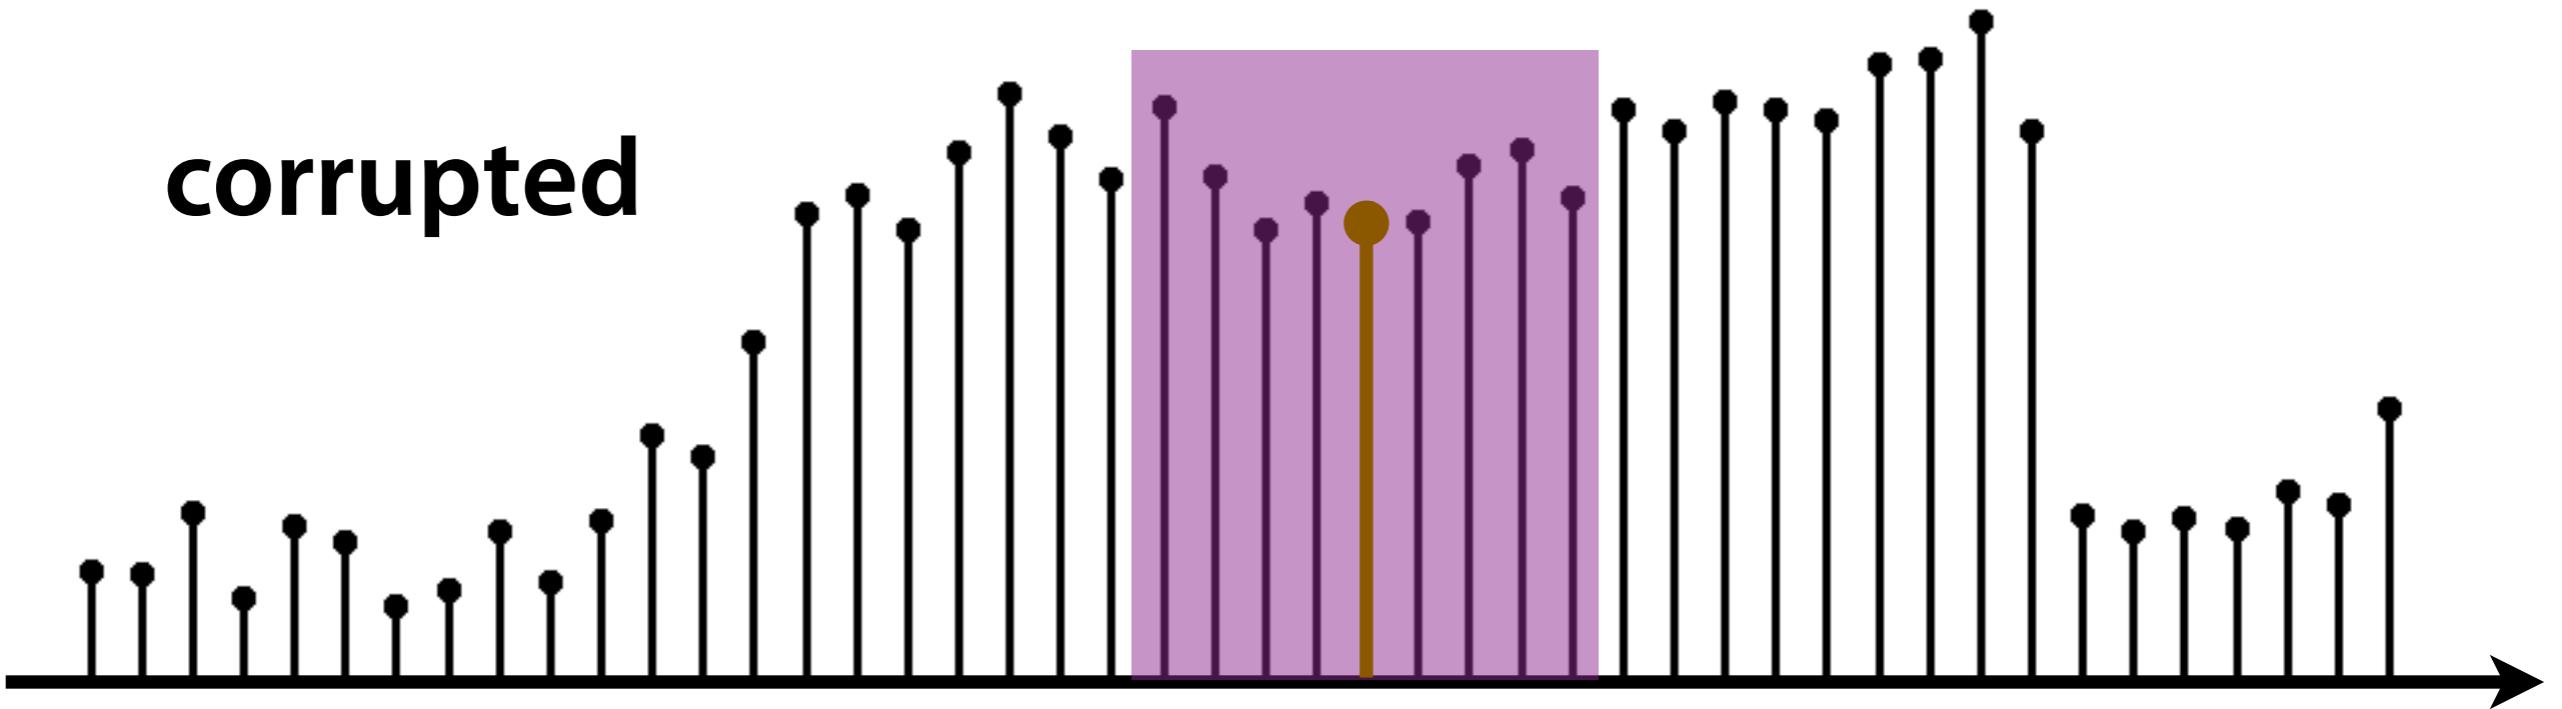
**corrupted**



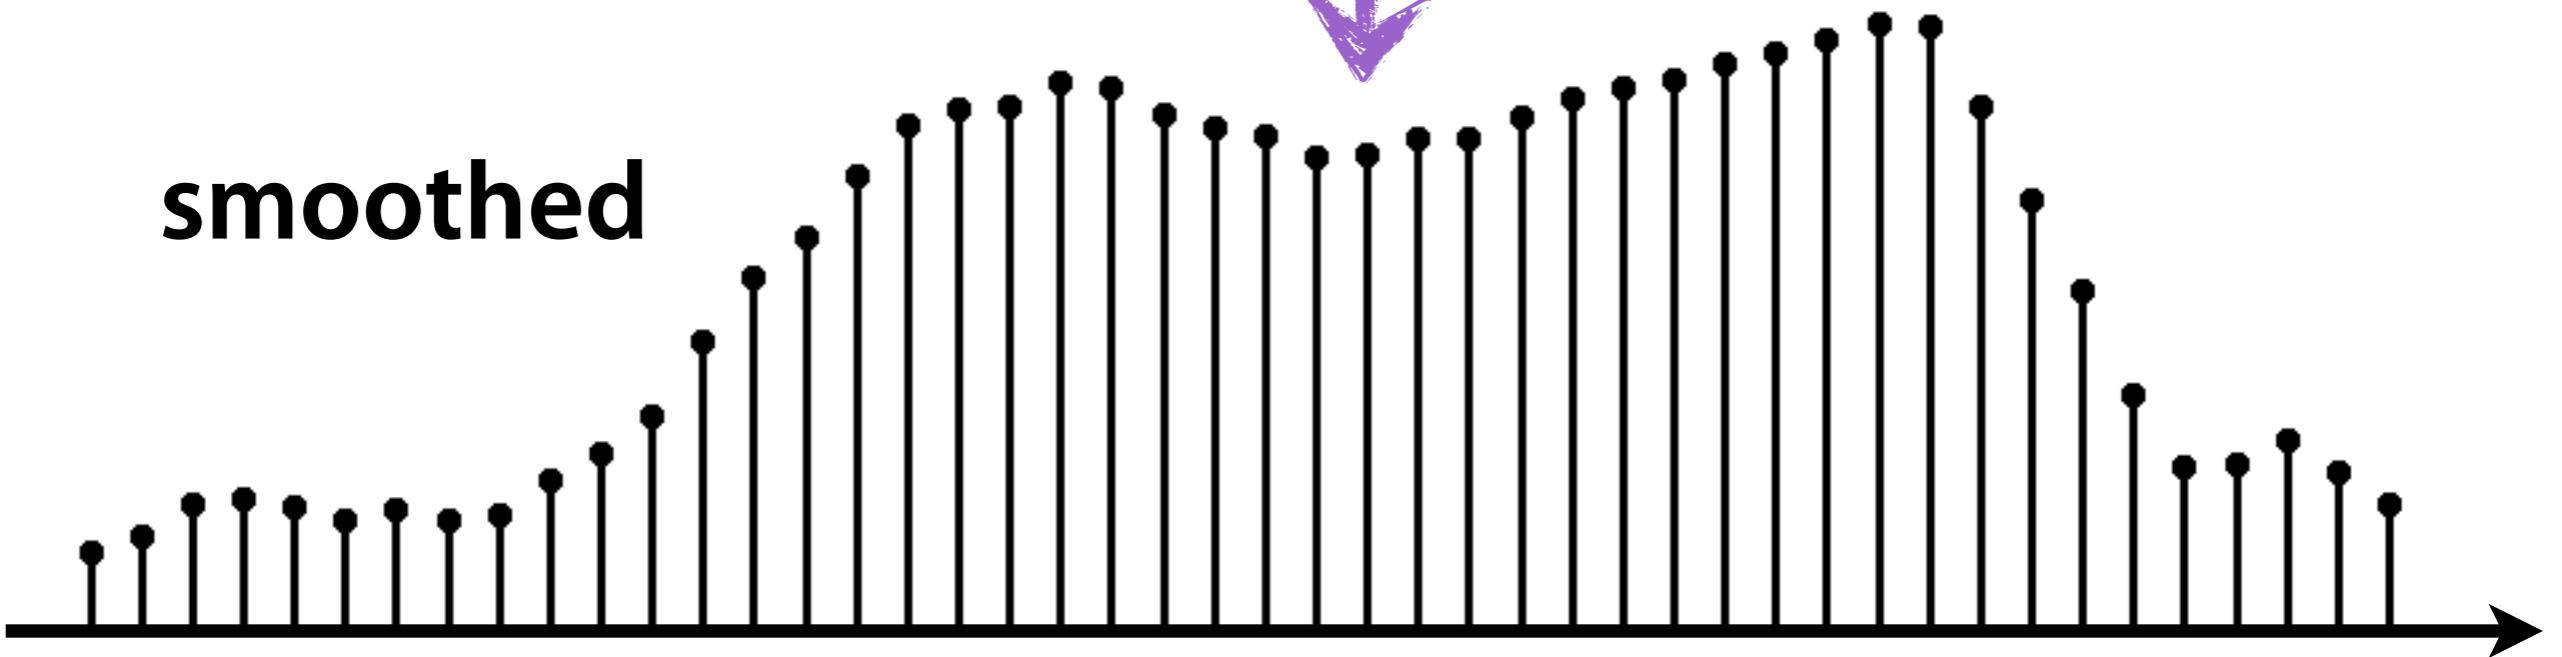
**smoothed**

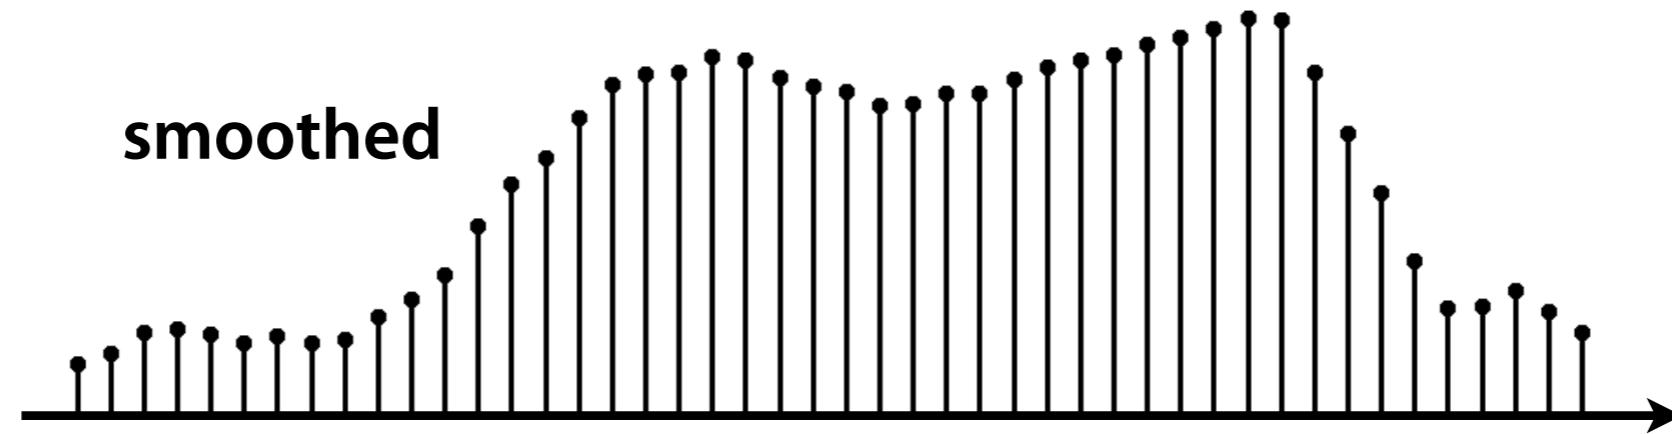
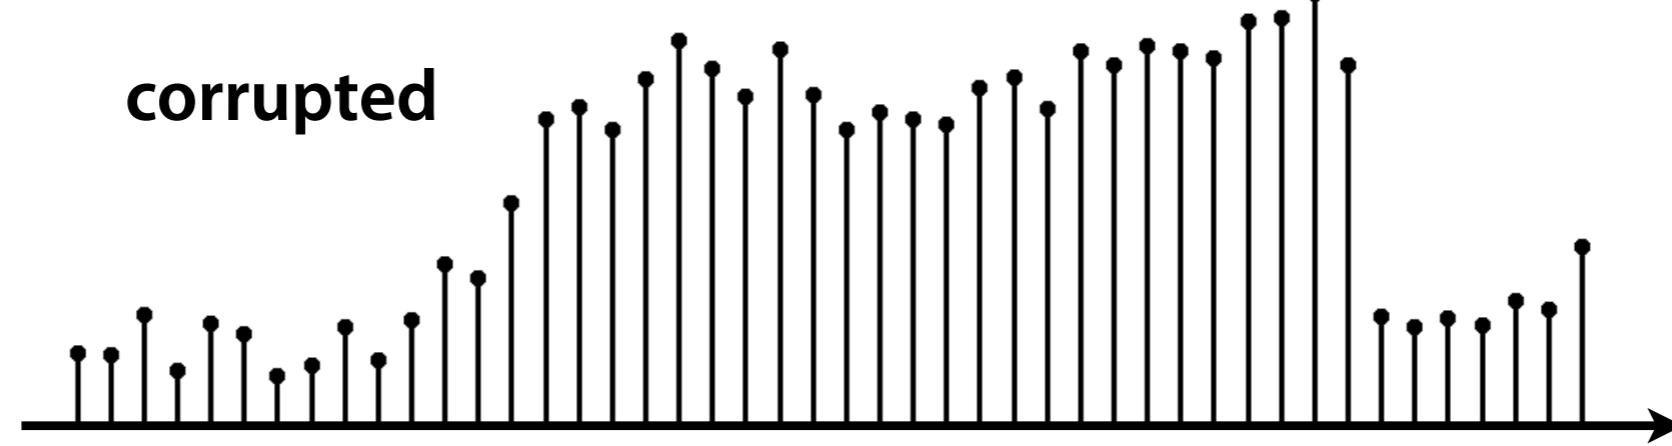
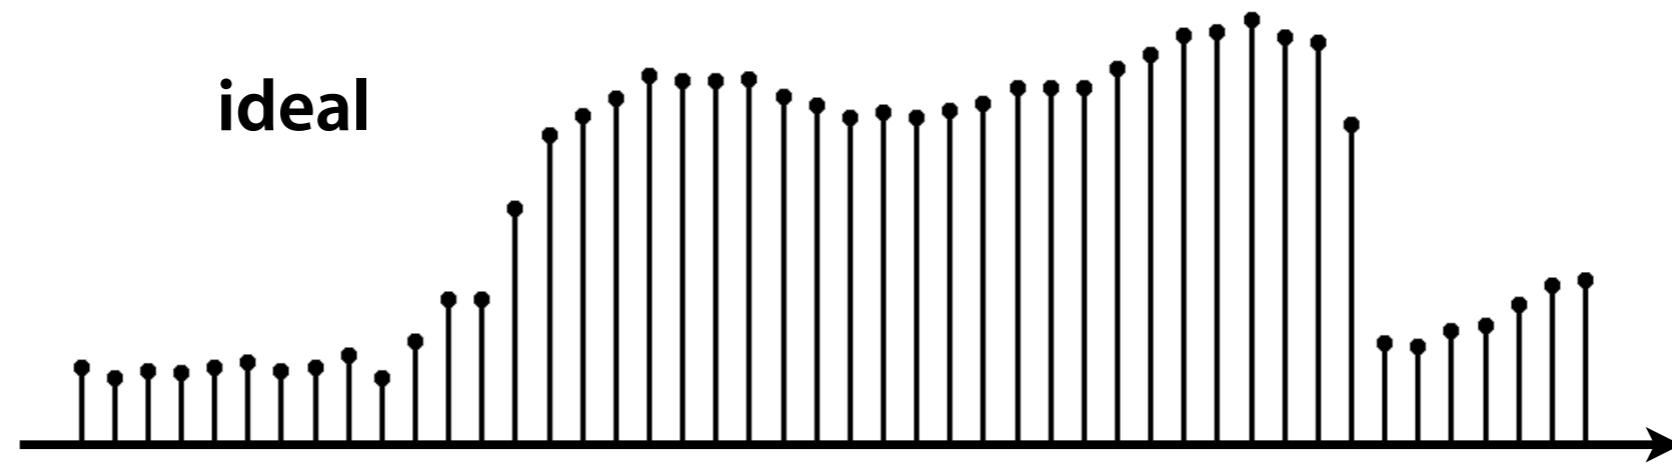


**corrupted**



**smoothed**





|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

**input**

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**output**

|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

**input**

|  |  |    |    |  |
|--|--|----|----|--|
|  |  |    |    |  |
|  |  | 10 | 15 |  |
|  |  |    |    |  |
|  |  |    |    |  |
|  |  |    |    |  |

**output**

|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

**input**

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**output**

|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

**input**

|  |  |  |    |    |
|--|--|--|----|----|
|  |  |  |    |    |
|  |  |  |    |    |
|  |  |  | 10 | 15 |
|  |  |  | 10 | 15 |
|  |  |  | 0  | 2  |

**output**

$$(2K + 1) \times (2K + 1)$$

**averaging window size**

$$H[x, y] = \frac{1}{(2K + 1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x + u, y + v]$$

**output**

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x + u, y + v]$$

**input**

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x + u, y + v]$$

loop over pixels in  
neighbourhood

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x+u, y+v]$$

weight

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x+u, y+v]$$

**How can the weights vary per neighbourhood position?**

**mask, kernel or filter**

$$H[x, y] = \frac{1}{(2K + 1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x + u, y + v] G[u, v]$$

correlation

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x+u, y+v] G[u, v]$$

**replace each pixel with a linear combination of neighbours**

$$H = G \otimes F$$

**cross correlation**

|   |   |   |   |   |
|---|---|---|---|---|
| 3 | 3 | 2 | 1 | 0 |
| 0 | 0 | 1 | 3 | 1 |
| 3 | 1 | 2 | 2 | 3 |
| 2 | 0 | 0 | 2 | 2 |
| 2 | 0 | 0 | 0 | 1 |

\*

|   |   |   |
|---|---|---|
| 0 | 1 | 2 |
| 2 | 2 | 0 |
| 0 | 1 | 2 |

correlation

|        |        |        |   |   |
|--------|--------|--------|---|---|
| 3<br>0 | 3<br>1 | 2<br>2 | 1 | 0 |
| 0<br>2 | 0<br>2 | 1<br>0 | 3 | 1 |
| 3<br>0 | 1<br>1 | 2<br>2 | 2 | 3 |

|    |  |  |
|----|--|--|
| 12 |  |  |
|    |  |  |
|    |  |  |

**pointwise multiplication and sum**

|   |   |   |   |   |
|---|---|---|---|---|
| 2 | 0 | 0 | 0 | 1 |
|---|---|---|---|---|

|   |     |     |     |   |
|---|-----|-----|-----|---|
| 3 | 3 0 | 2 1 | 1 2 | 0 |
| 0 | 0 2 | 1 2 | 3 0 | 1 |
| 3 | 1 0 | 2 1 | 2 2 | 3 |
| 2 | 0   | 0   | 2   | 2 |
| 2 | 0   | 0   | 0   | 1 |

==

|    |    |  |
|----|----|--|
| 12 | 12 |  |
|    |    |  |
|    |    |  |

|   |   |     |     |     |
|---|---|-----|-----|-----|
| 3 | 3 | 2 0 | 1 1 | 0 2 |
| 0 | 0 | 1 2 | 3 2 | 1 0 |
| 3 | 1 | 2 0 | 2 1 | 3 2 |
| 2 | 0 | 0   | 2   | 2   |
| 2 | 0 | 0   | 0   | 1   |

==

|    |    |    |
|----|----|----|
| 12 | 12 | 17 |
|    |    |    |
|    |    |    |

|        |        |        |   |   |
|--------|--------|--------|---|---|
| 3      | 3      | 2      | 1 | 0 |
| 0<br>0 | 0<br>1 | 1<br>2 | 3 | 1 |
| 3<br>2 | 1<br>2 | 2<br>0 | 2 | 3 |
| 2<br>0 | 0<br>1 | 0<br>2 | 2 | 2 |
| 2      | 0      | 0      | 0 | 1 |

=

|    |    |    |
|----|----|----|
| 12 | 12 | 17 |
| 10 |    |    |
|    |    |    |

|   |        |        |        |   |
|---|--------|--------|--------|---|
| 3 | 3      | 2      | 1      | 0 |
| 0 | 0      | 1      | 3      | 1 |
| 3 | 1<br>0 | 2<br>1 | 2<br>2 | 3 |
| 2 | 0<br>2 | 0<br>2 | 2<br>0 | 2 |
| 2 | 0<br>0 | 0<br>1 | 0<br>2 | 1 |

=

|    |    |    |
|----|----|----|
| 12 | 12 | 17 |
| 10 | 17 | 19 |
| 9  | 6  |    |

*correlation*

|   |   |     |     |     |
|---|---|-----|-----|-----|
| 3 | 3 | 2   | 1   | 0   |
| 0 | 0 | 1   | 3   | 1   |
| 3 | 1 | 2 0 | 2 1 | 3 2 |
| 2 | 0 | 0 2 | 2 2 | 2 0 |
| 2 | 0 | 0 0 | 0 1 | 1 2 |

==

|    |    |    |
|----|----|----|
| 12 | 12 | 17 |
| 10 | 17 | 19 |
| 9  | 6  | 14 |



original



**box filter**

**smoothing by averaging**

|   |   |   |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

$\frac{1}{9}$

$$G[x, y]$$



|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

$$F[x, y]$$

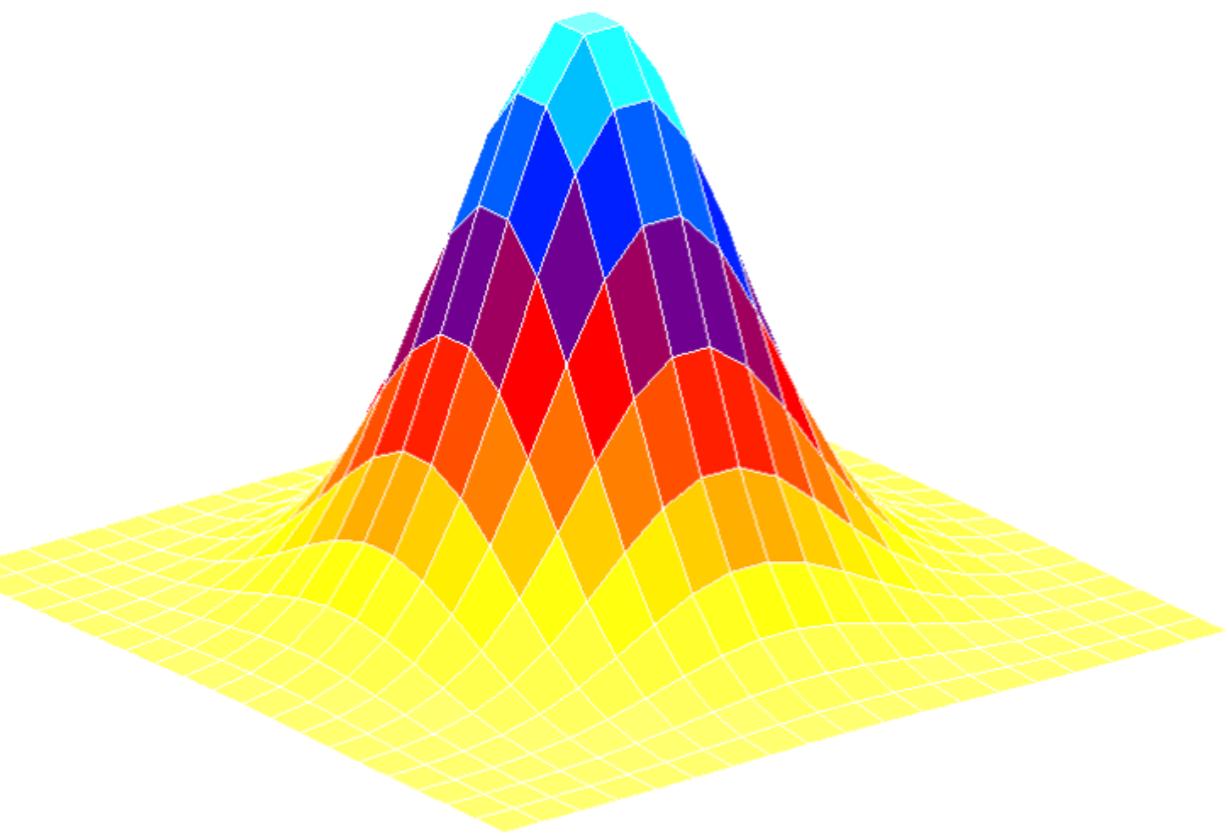
$\frac{1}{16}$ 

|   |   |   |
|---|---|---|
| 1 | 2 | 1 |
| 2 | 4 | 2 |
| 1 | 2 | 1 |

 $G[x, y]$ 

|   |    |   |    |    |
|---|----|---|----|----|
| 0 | 0  | 0 | 9  | 18 |
| 0 | 90 | 0 | 27 | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 9  | 0  |
| 0 | 0  | 0 | 0  | 18 |

 $F[x, y]$



$\approx$

$$\frac{1}{16}$$

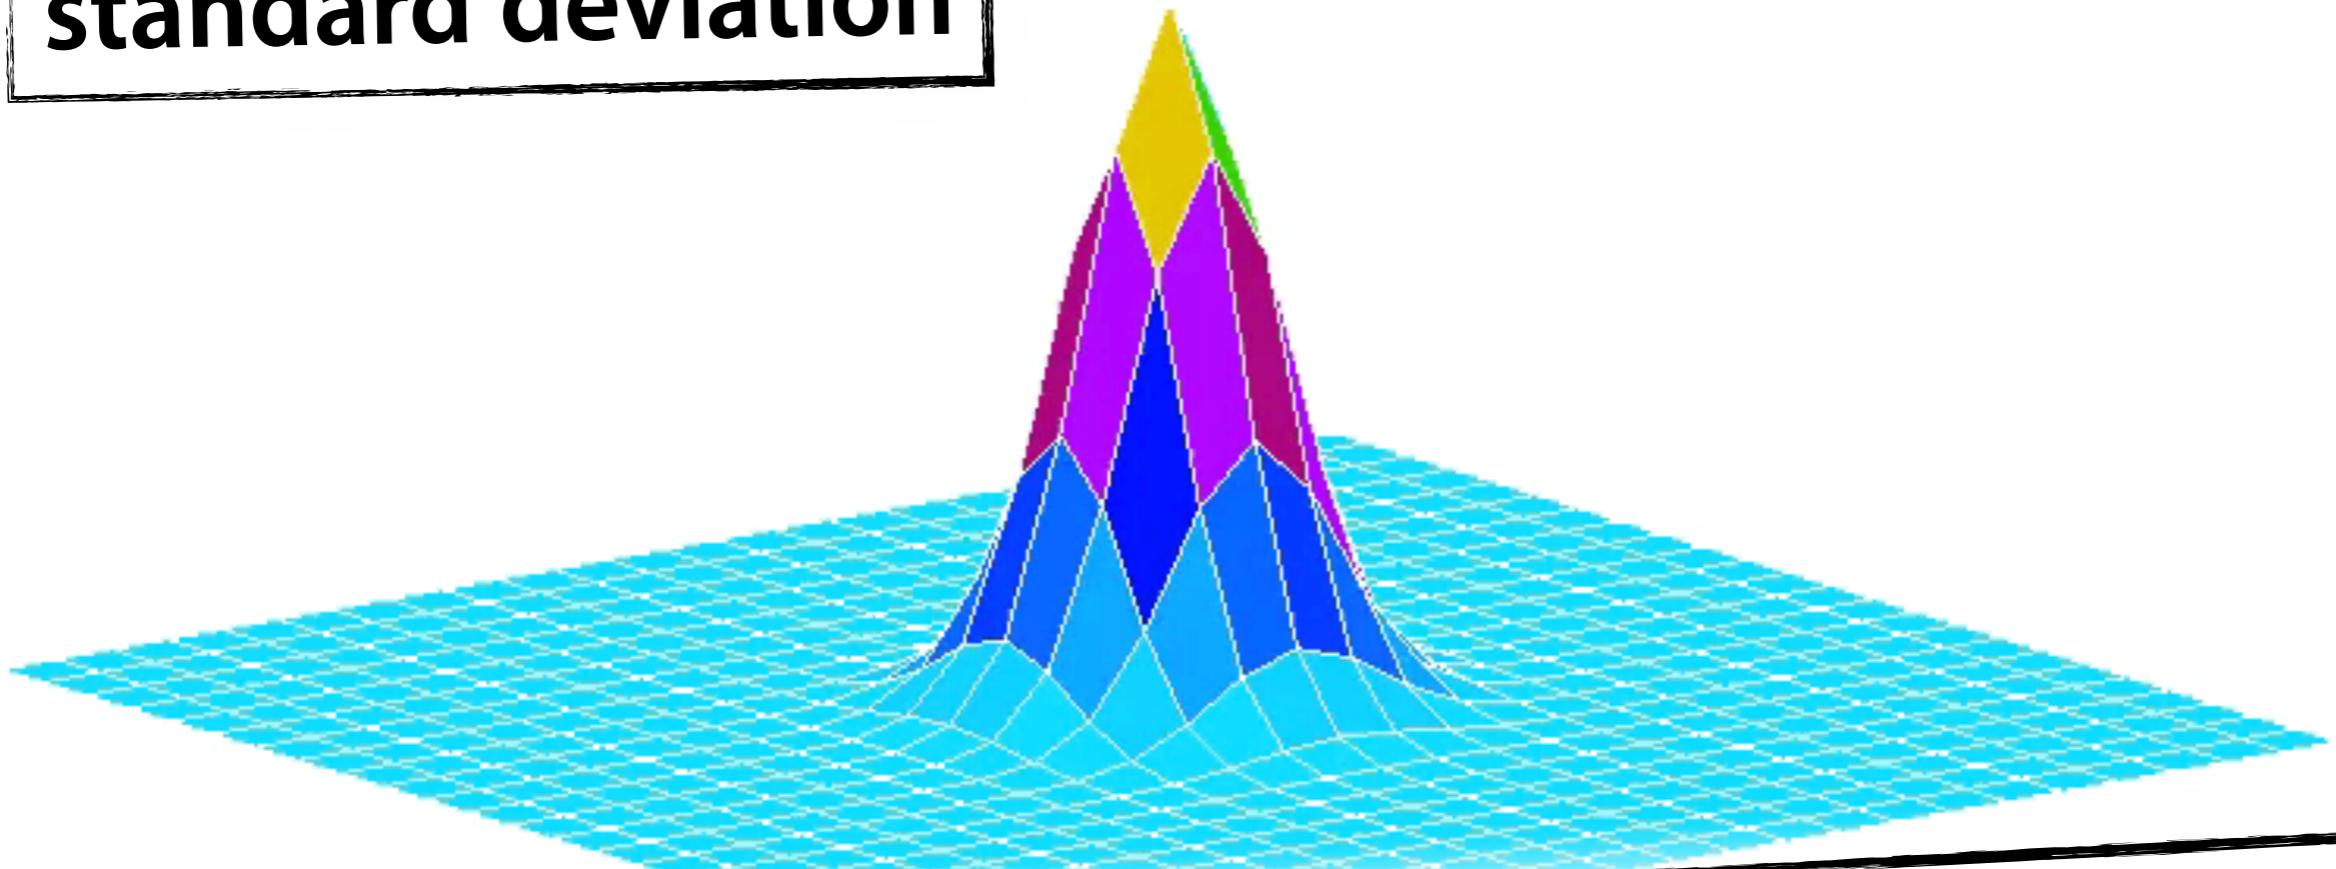
|   |   |   |
|---|---|---|
| 1 | 2 | 1 |
| 2 | 4 | 2 |
| 1 | 2 | 1 |

$$\frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

$$G[x, y]$$

$$\frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

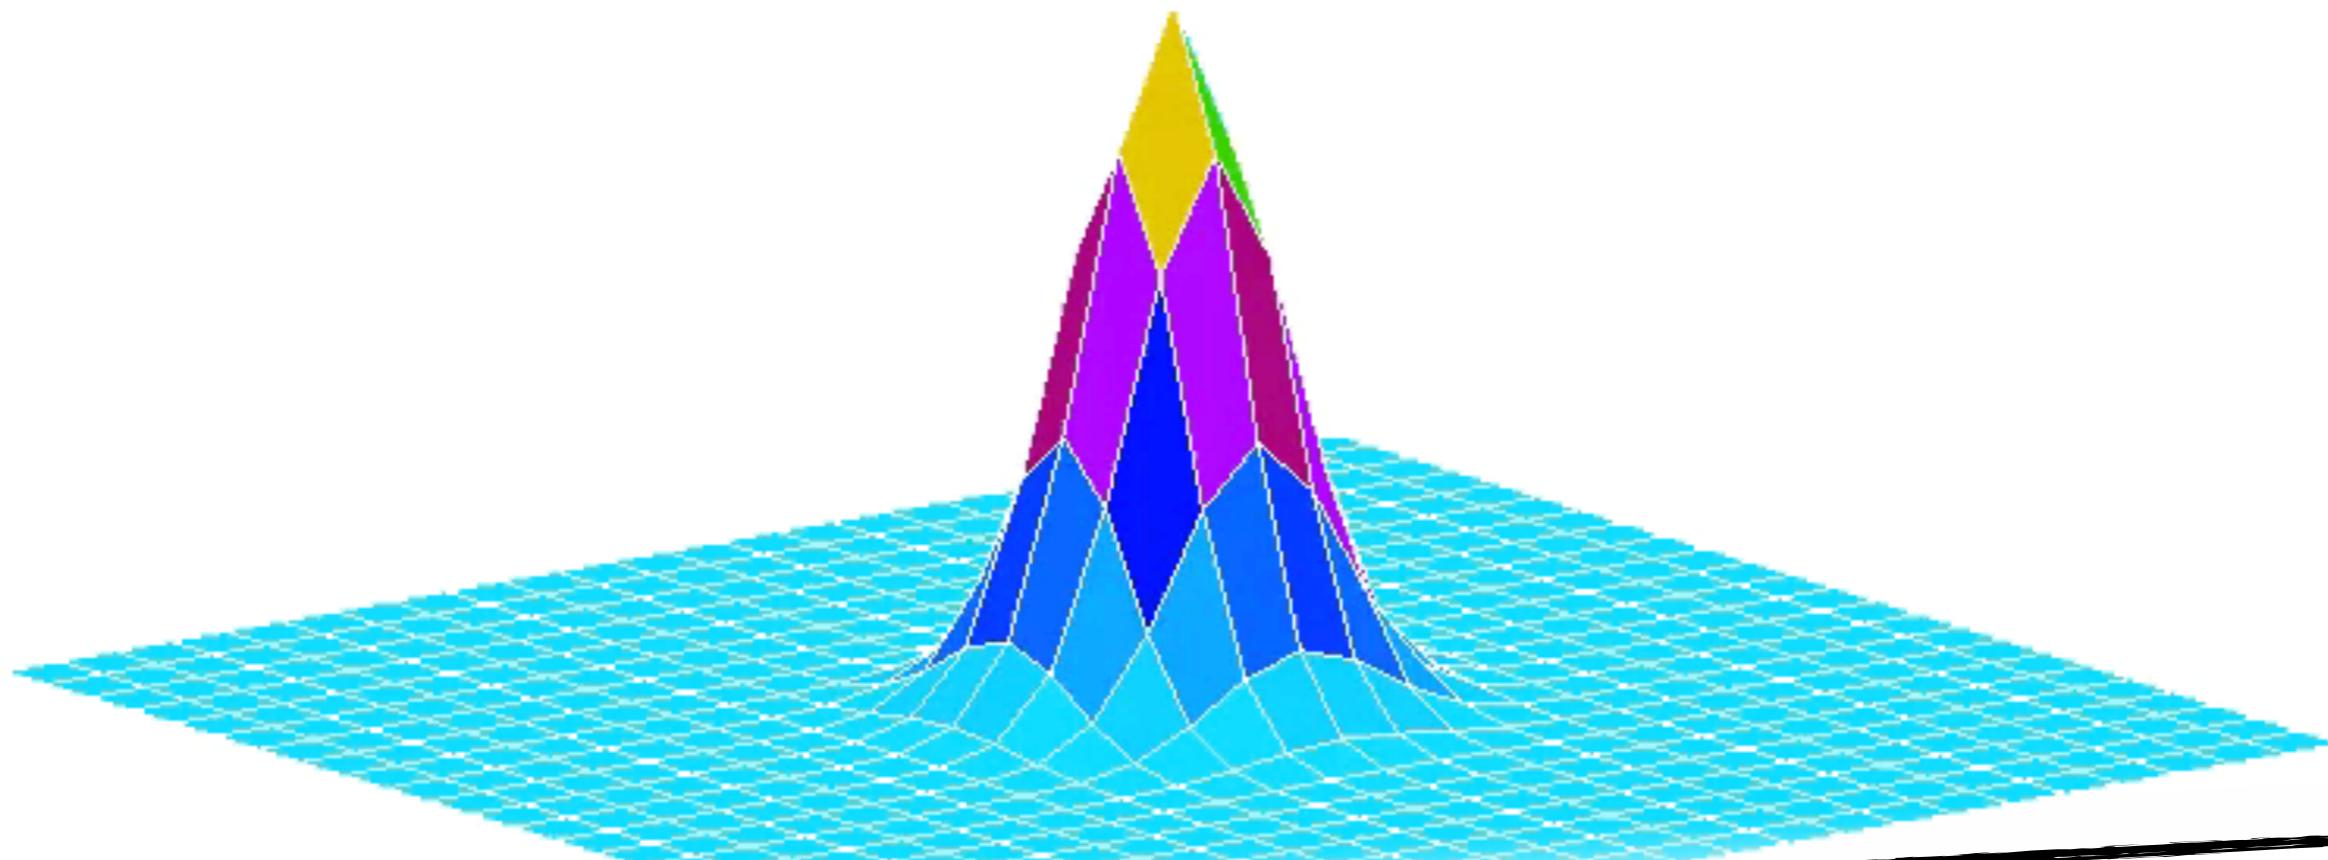
**standard deviation**



**How does the Gaussian vary with  $\sigma$  ?**

## normalization

$$\frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

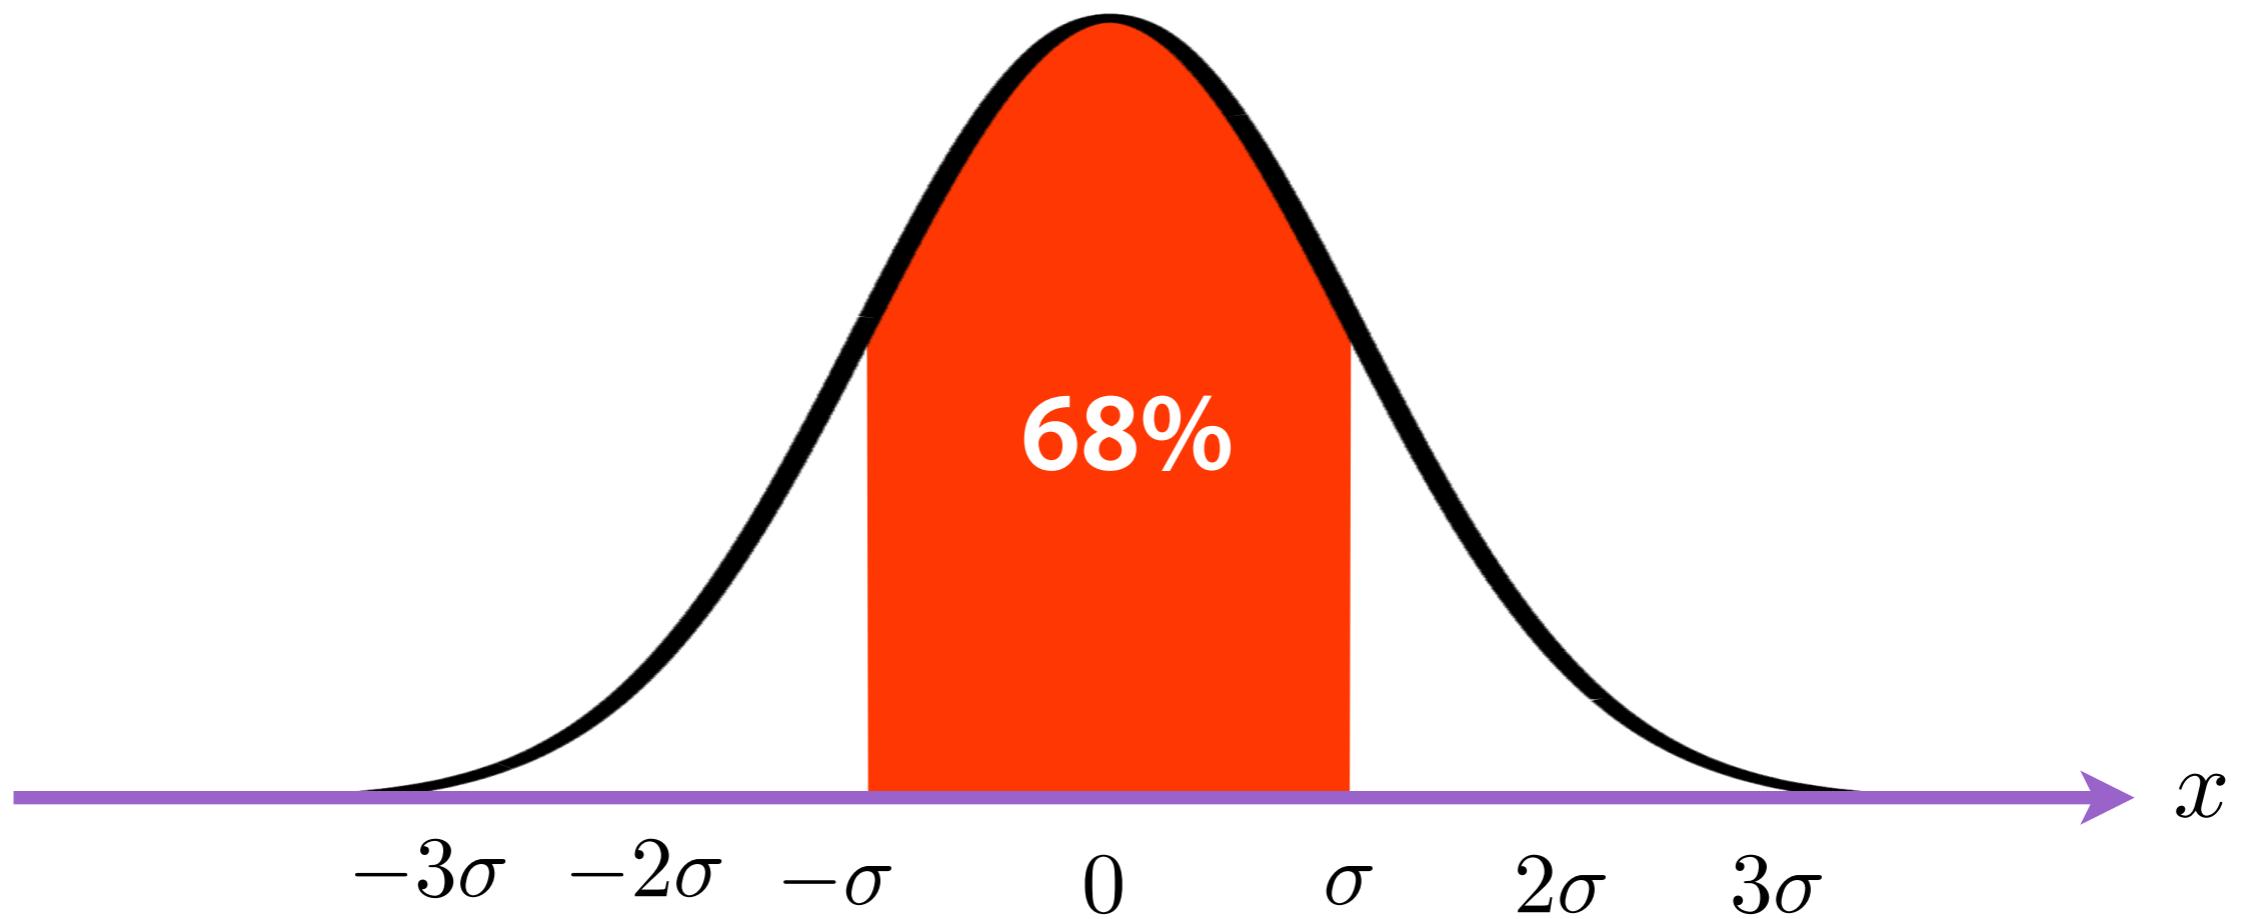


How does the Gaussian vary with  $\sigma$  ?

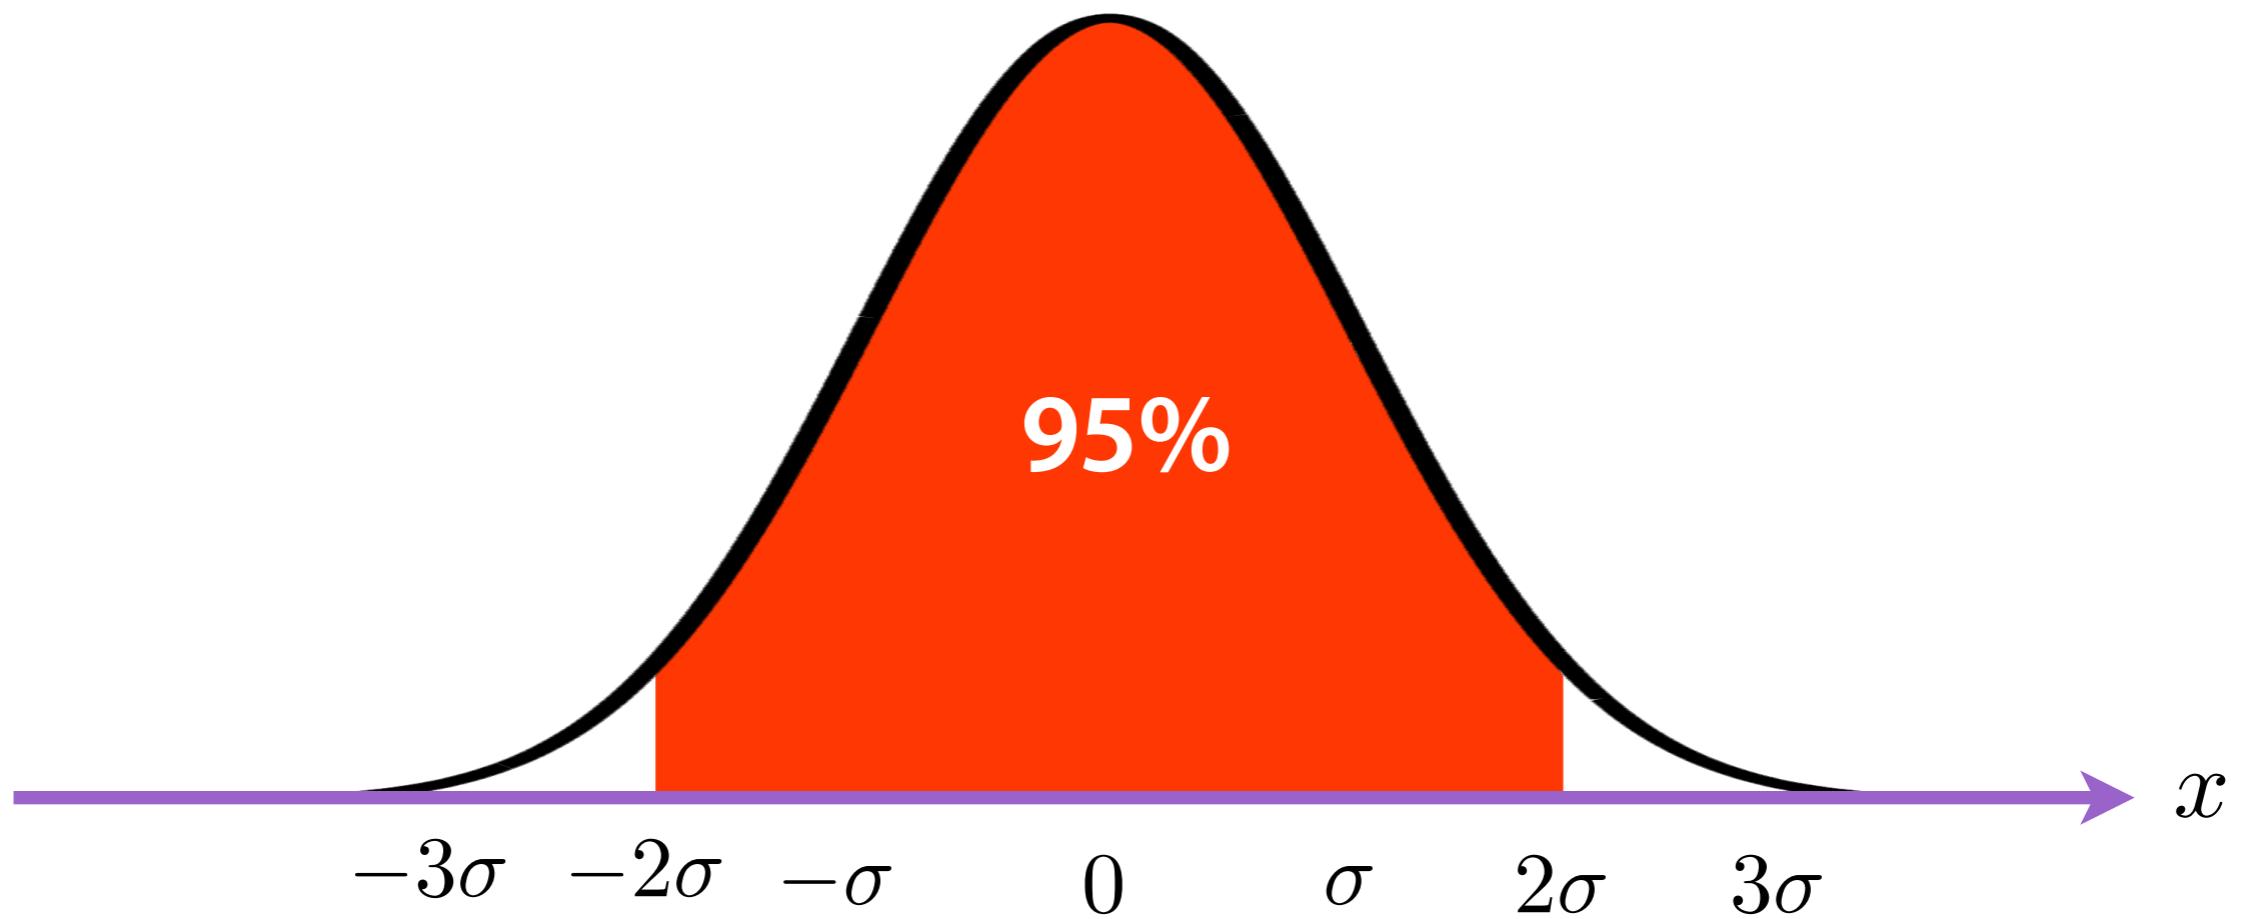
**Gaussian filters have infinite support**

**Discrete filters use finite kernels**

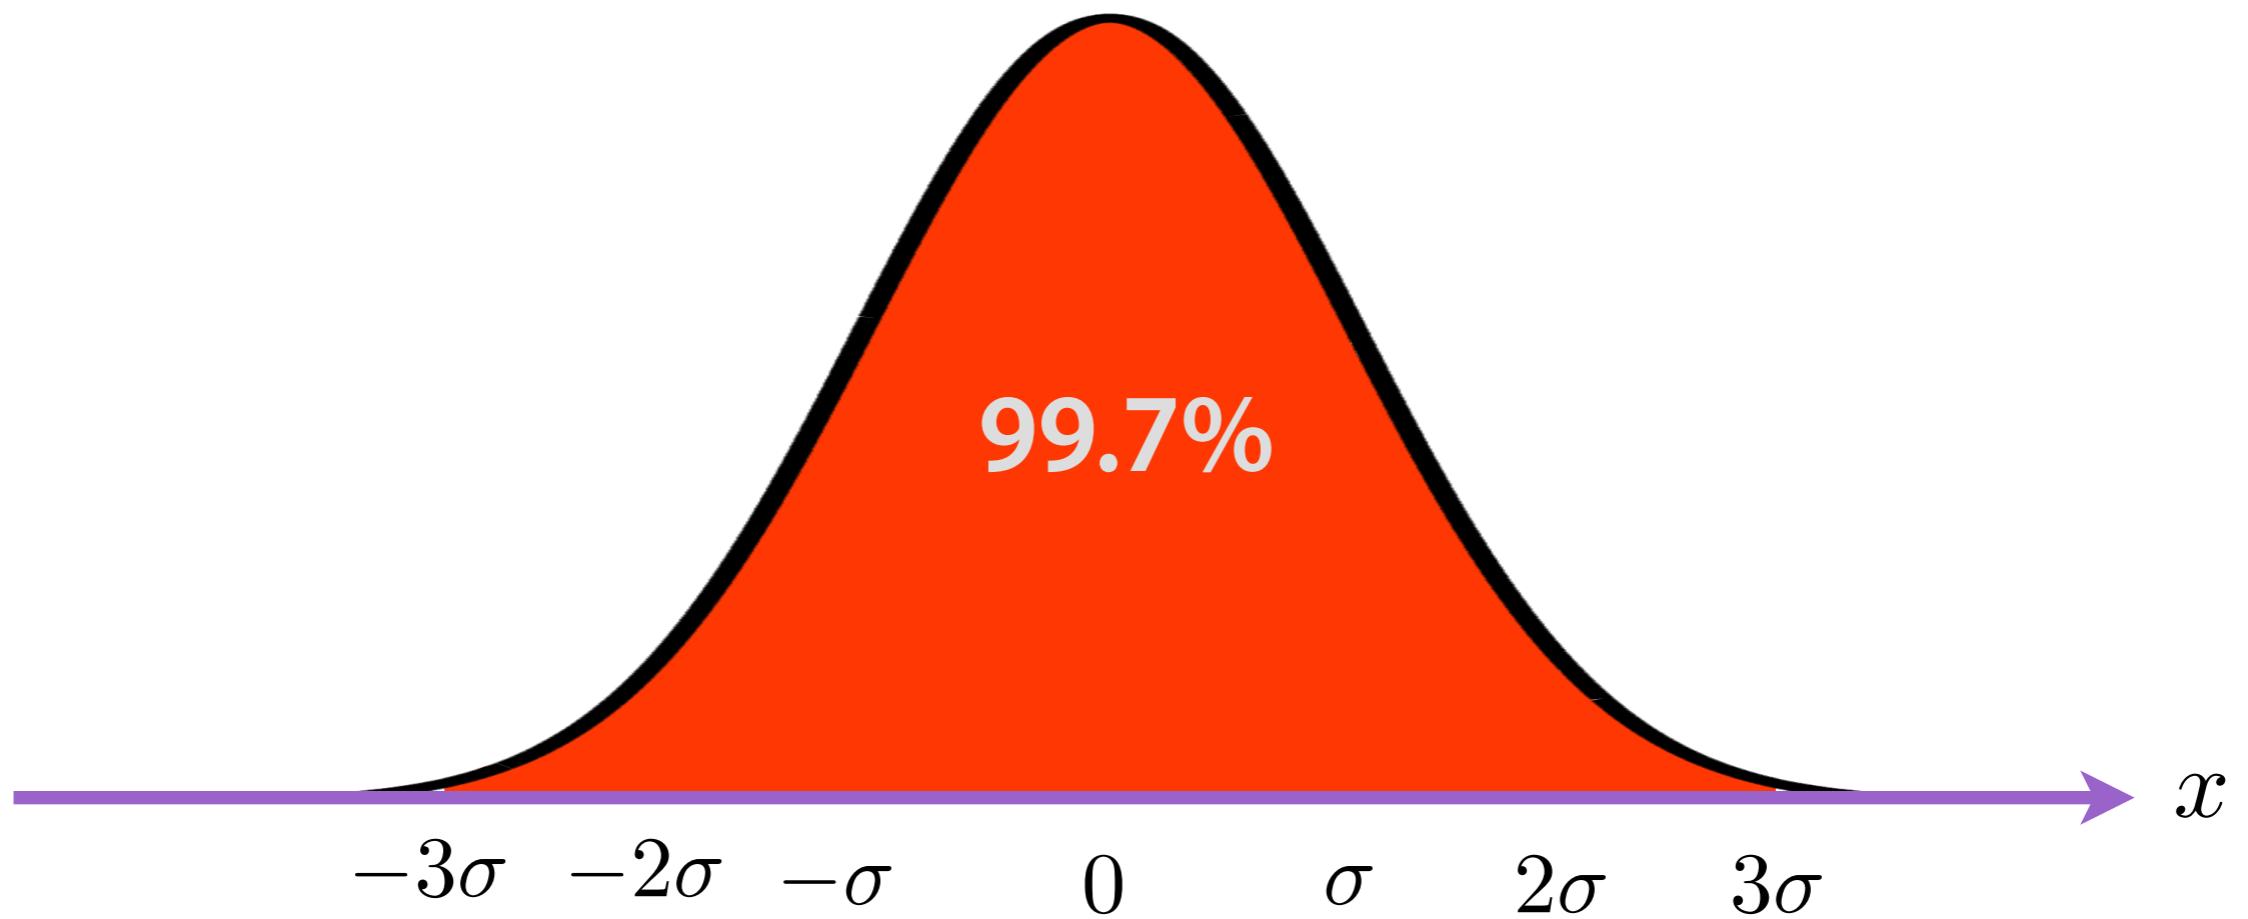
three-sigma  
rule



three-sigma  
rule



three-sigma  
rule



**original**

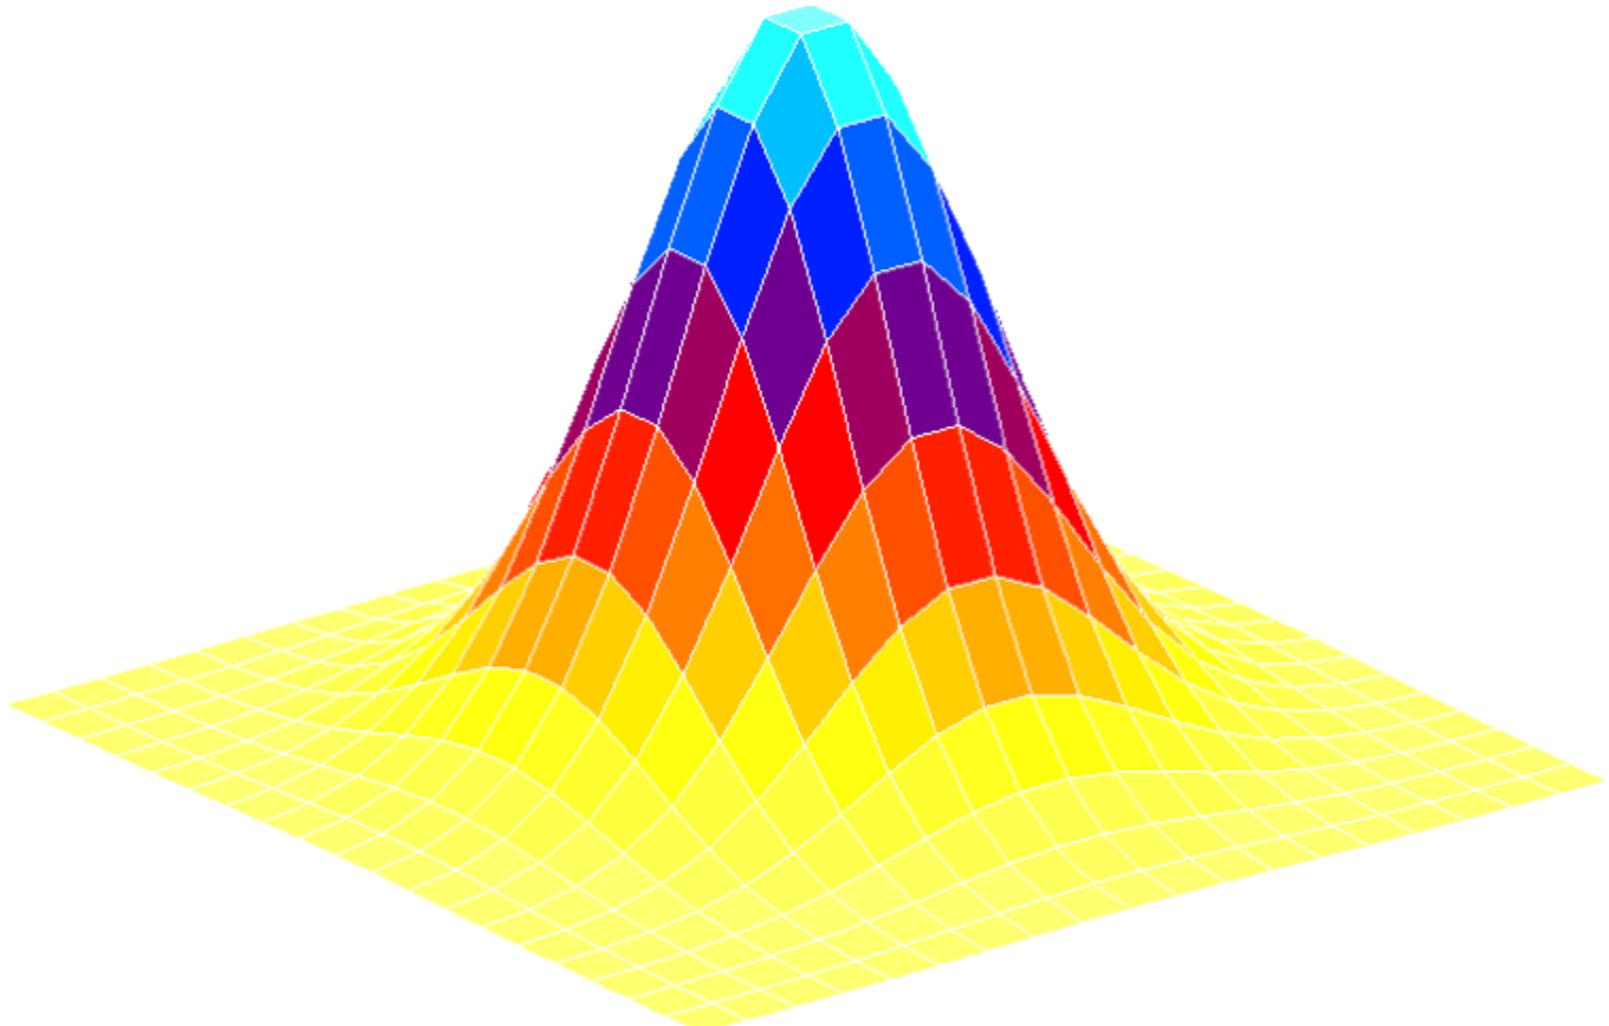
**Gaussian filter**



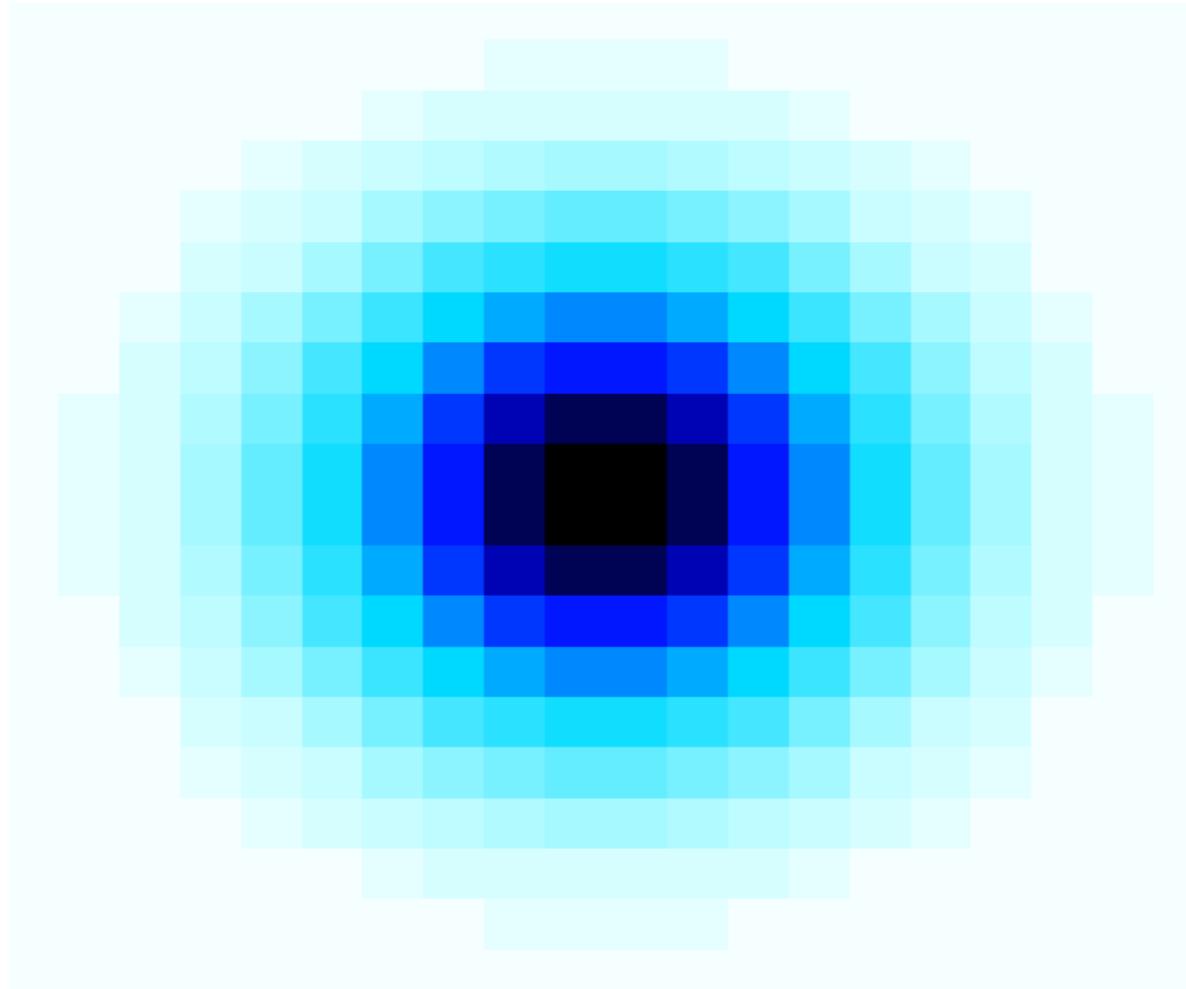
**Gaussian filter**



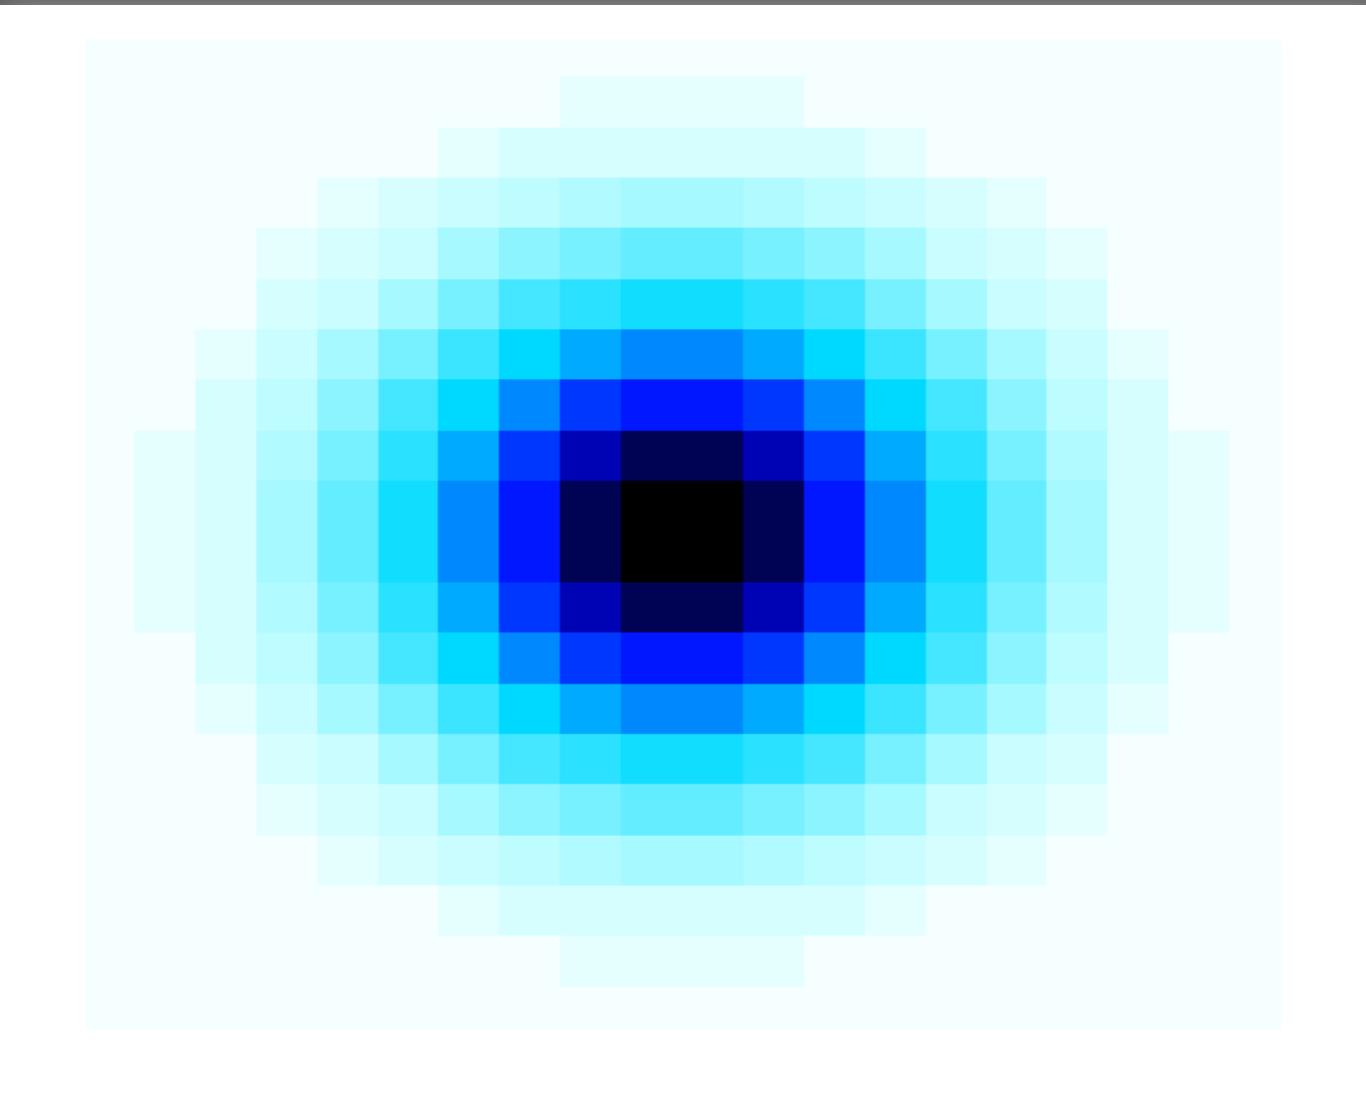
**Gaussian blur**



```
>> width = 19;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> surf(h)  
>> set(gca, 'Color', [0 0 0]); % optional
```

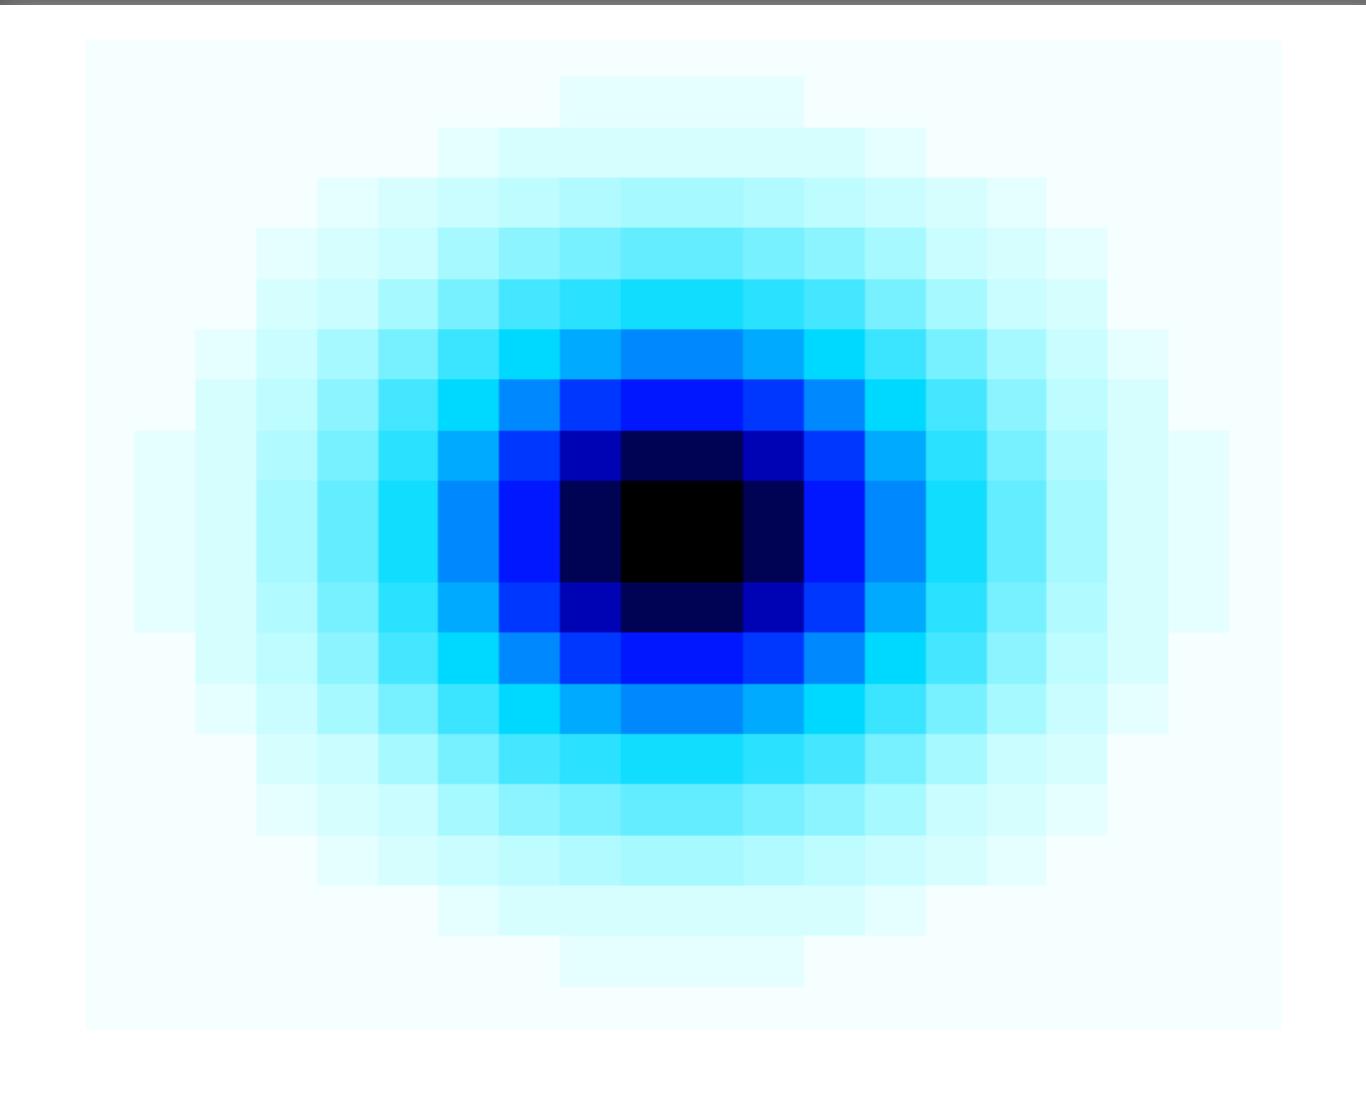


```
>> width = 19;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> imagesc(h);  
>> colormap hot % optional
```



```
>> width = 19;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> imagesc(h);
```

display image with scaled colours



```
>> width = 19;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> imagesc(h);  
>> colormap hot % optional
```



```
>> width = 21;
>> sigma = 3;
>> h = fspecial('gaussian', width, sigma);
>> I = imfilter(im, h, 'symmetric', 'same');
>> imshow(I, [ ]);
```



```
>> width = 21;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> I = imfilter(im, h, 'symmetric', 'same');  
>> imshow(I, [ ]);
```



```
>> width = 21;  
>> sigma = 3;  
>> h = fspecial('gaussian', width, sigma);  
>> I = imfilter(im, h, 'symmetric', 'same');  
>> imshow(I, [ ]);
```



```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



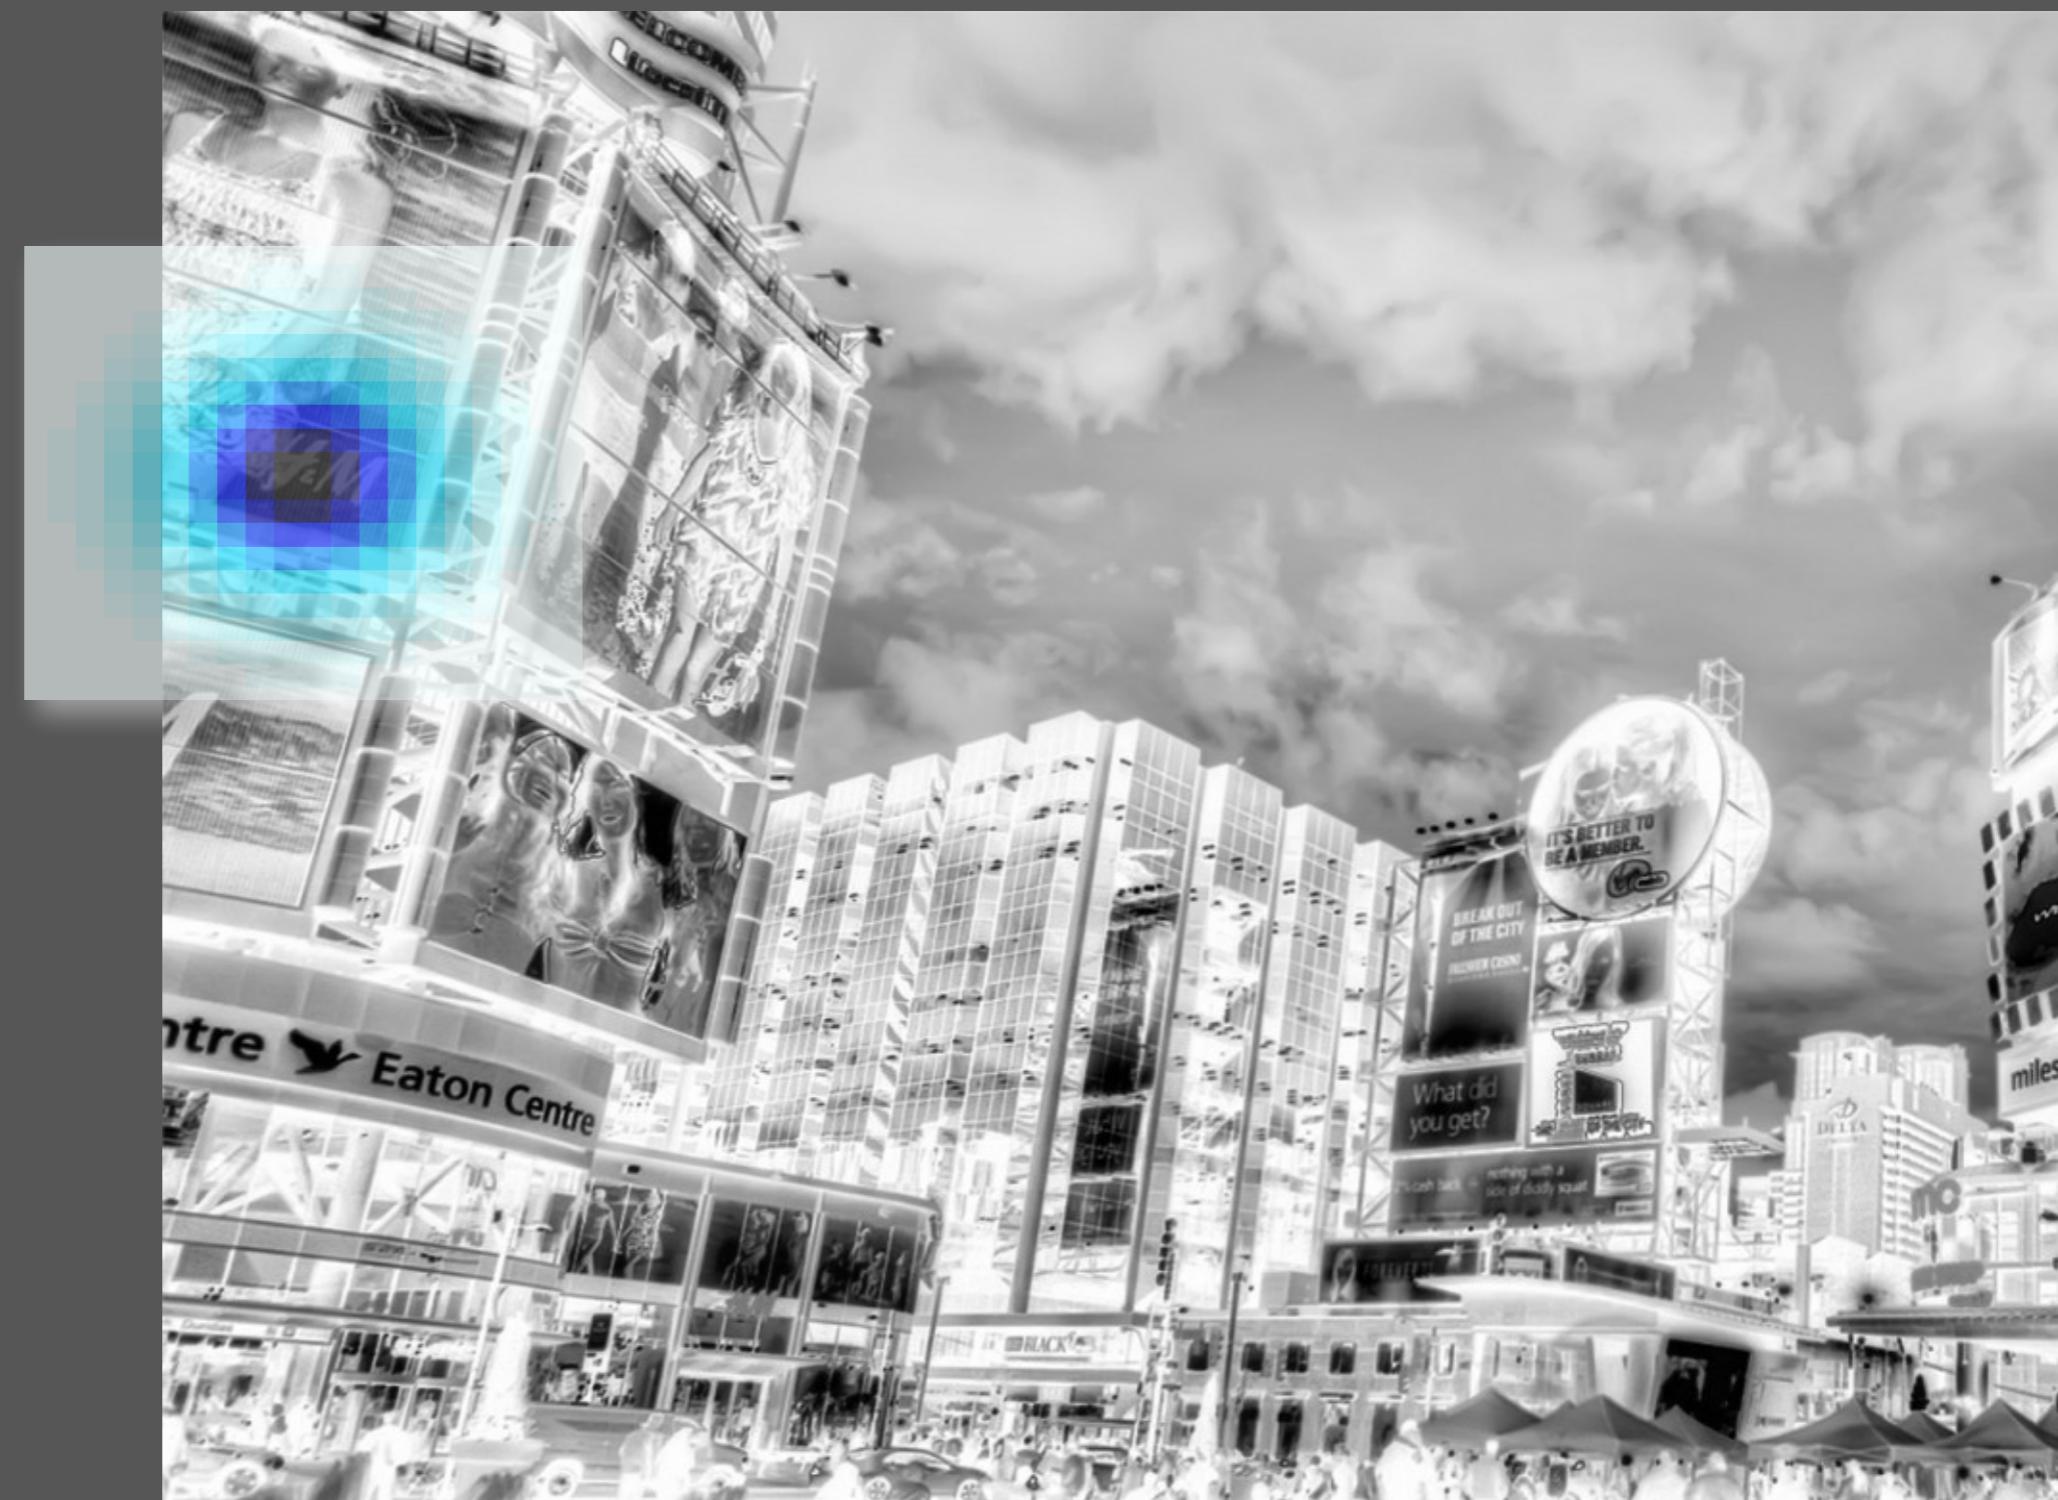
```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



```
>> sigma = 16;  
>> noise = randn(size(im))*sigma;  
>> I = im + noise;  
>> imshow(I, [ ]);
```



**What about near the image edges?**





clip padding

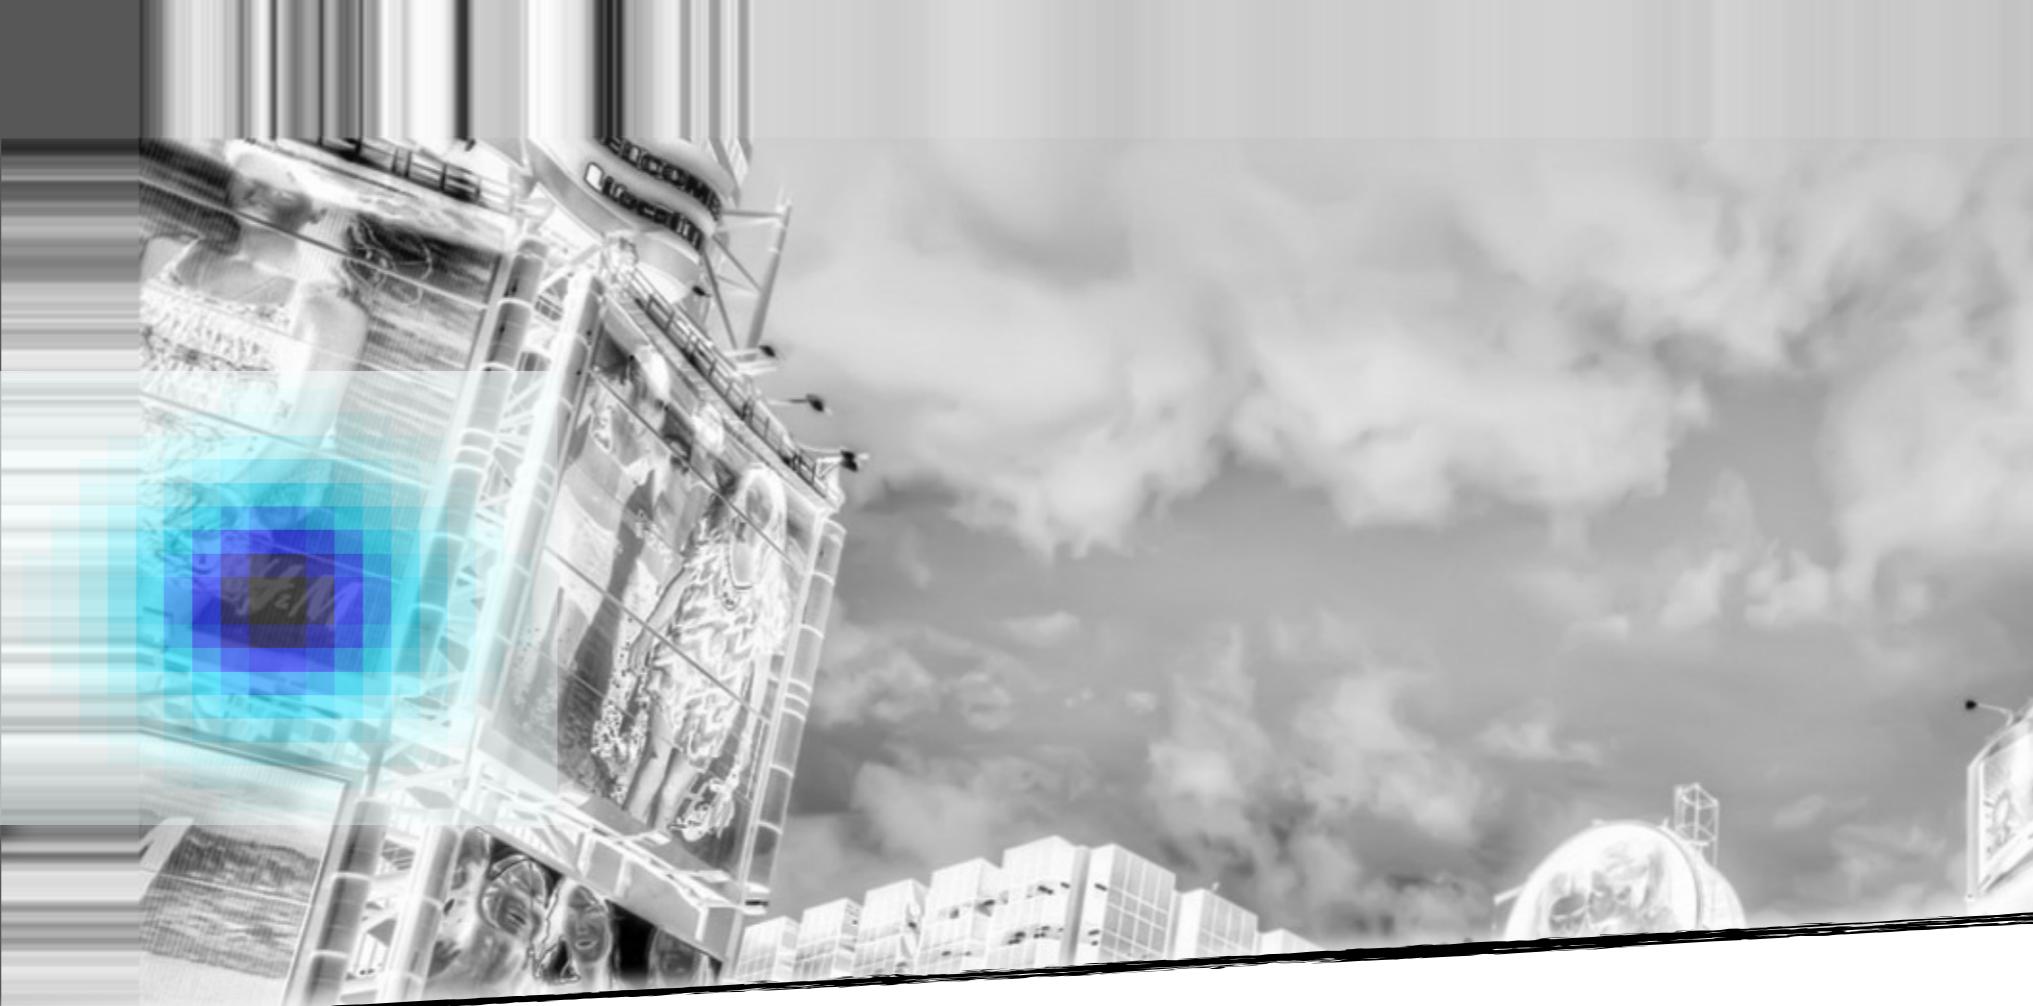


|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 3 | 3 | 2 | 1 | 0 | 0 |
| 0 | 0 | 0 | 1 | 3 | 1 | 0 |
| 0 | 3 | 1 | 2 | 2 | 3 | 0 |
| 0 | 2 | 0 | 0 | 2 | 2 | 0 |
| 0 | 2 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |



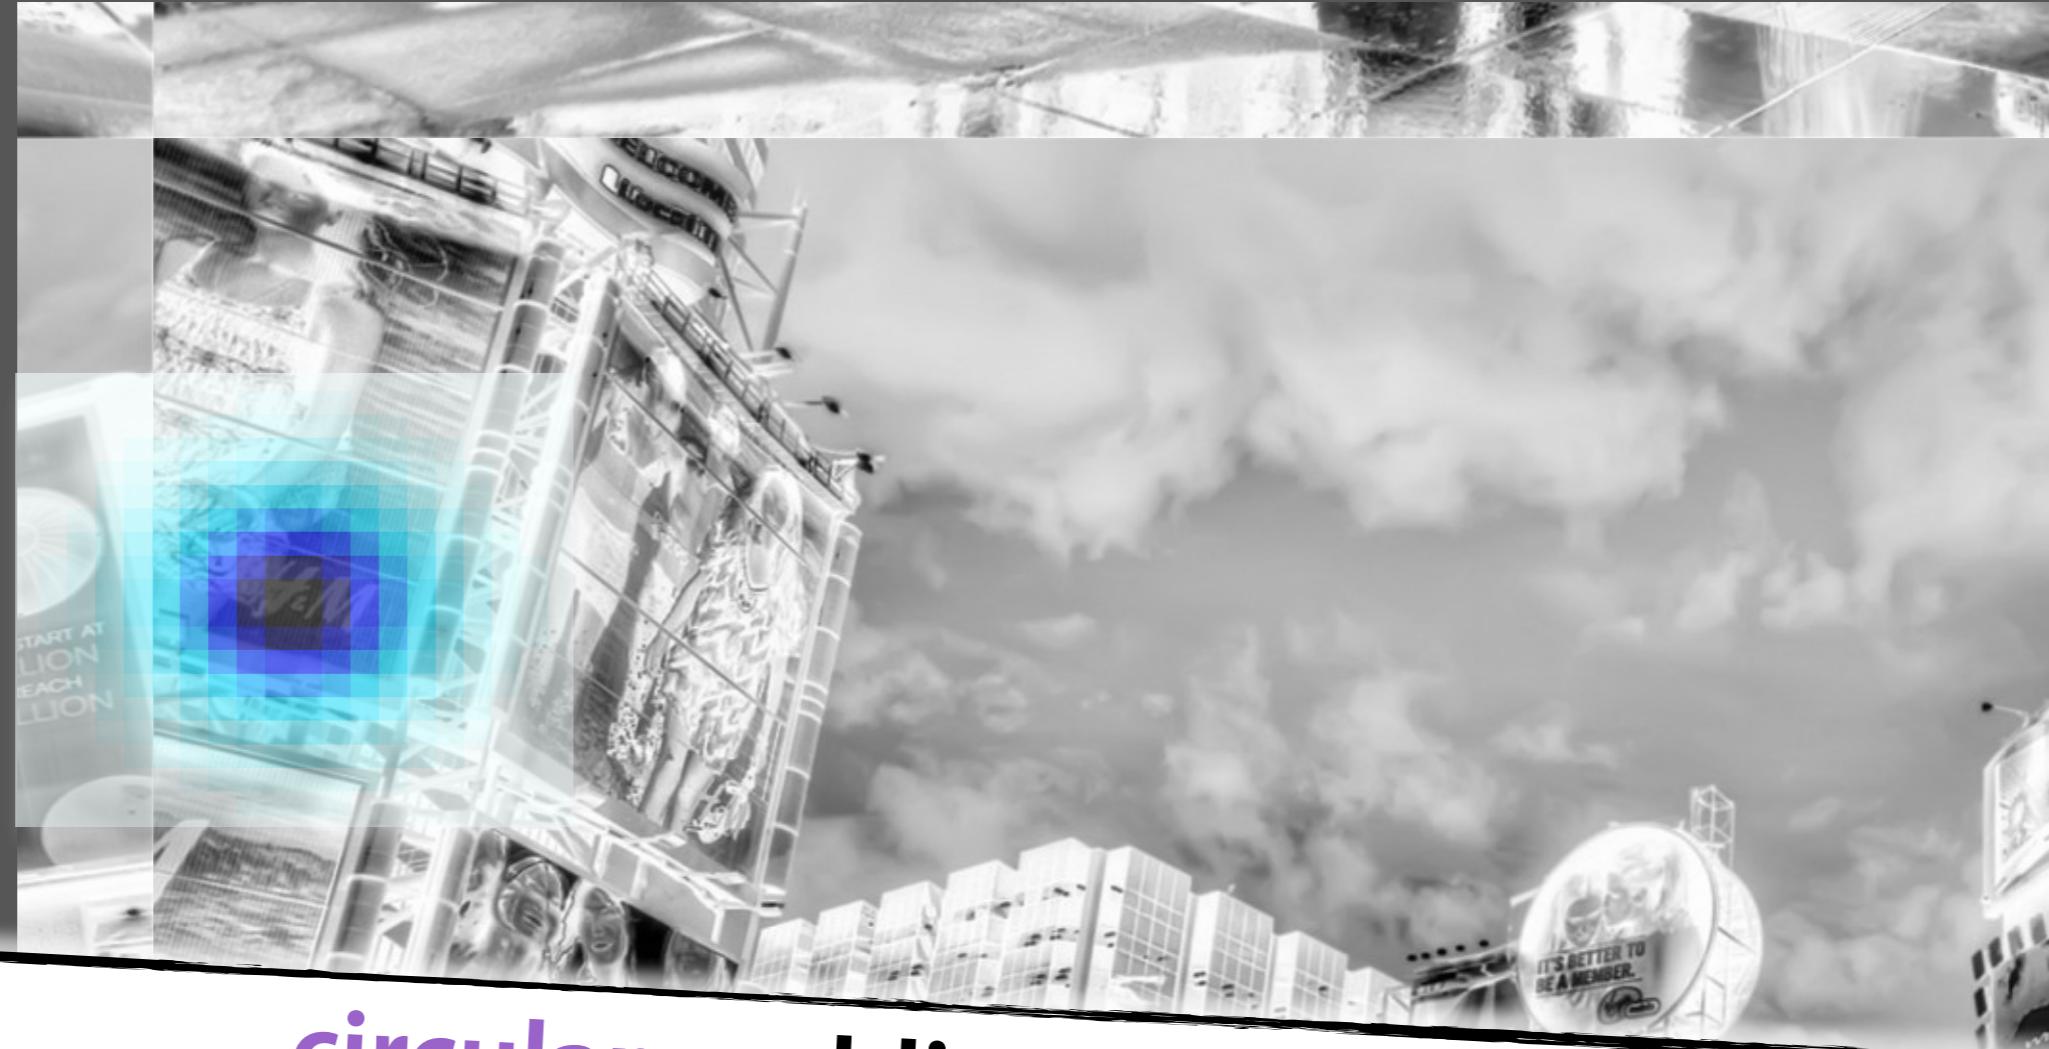
## symmetric padding

```
I = imfilter(im, h, 'symmetric');
```



## replicate padding

```
I = imfilter(im, h, 'replicate');
```

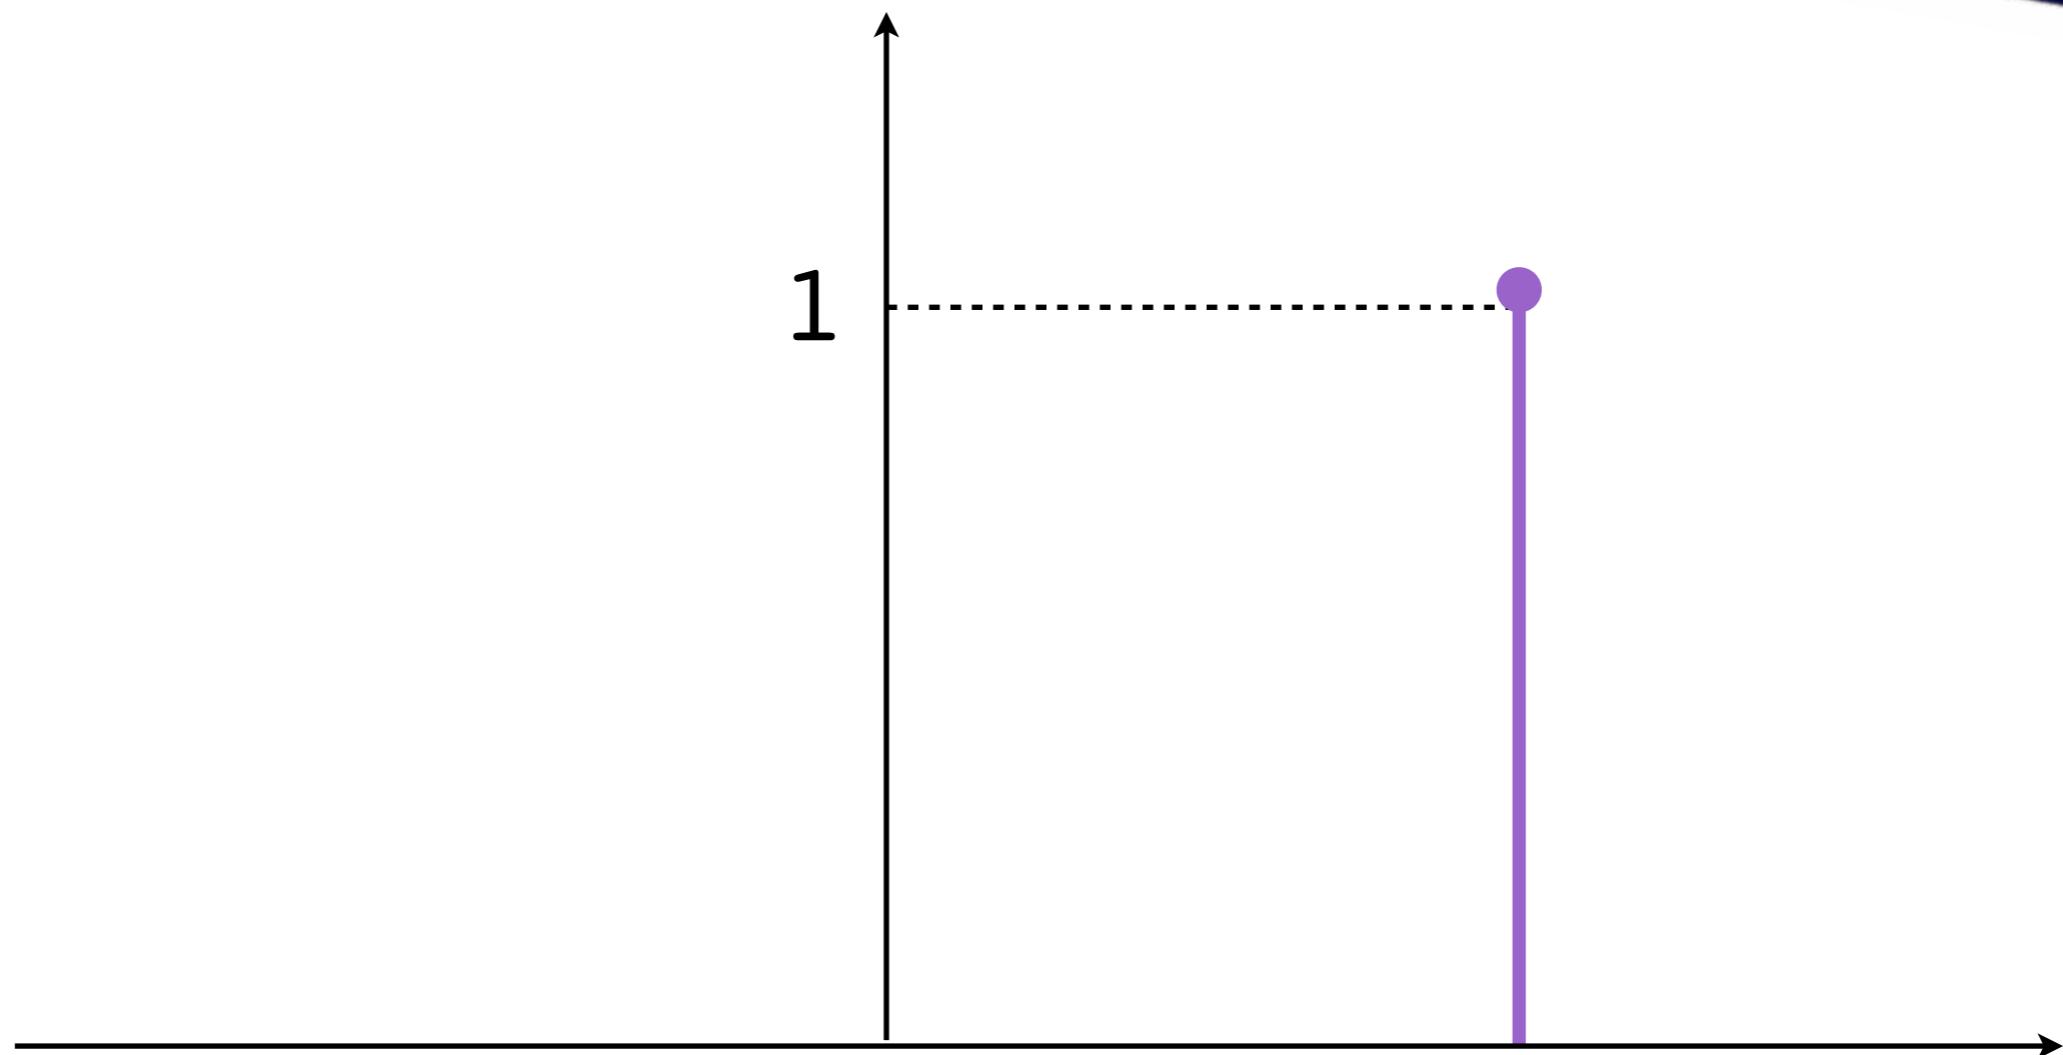


## circular padding

```
I = imfilter(im, h, 'circular');
```

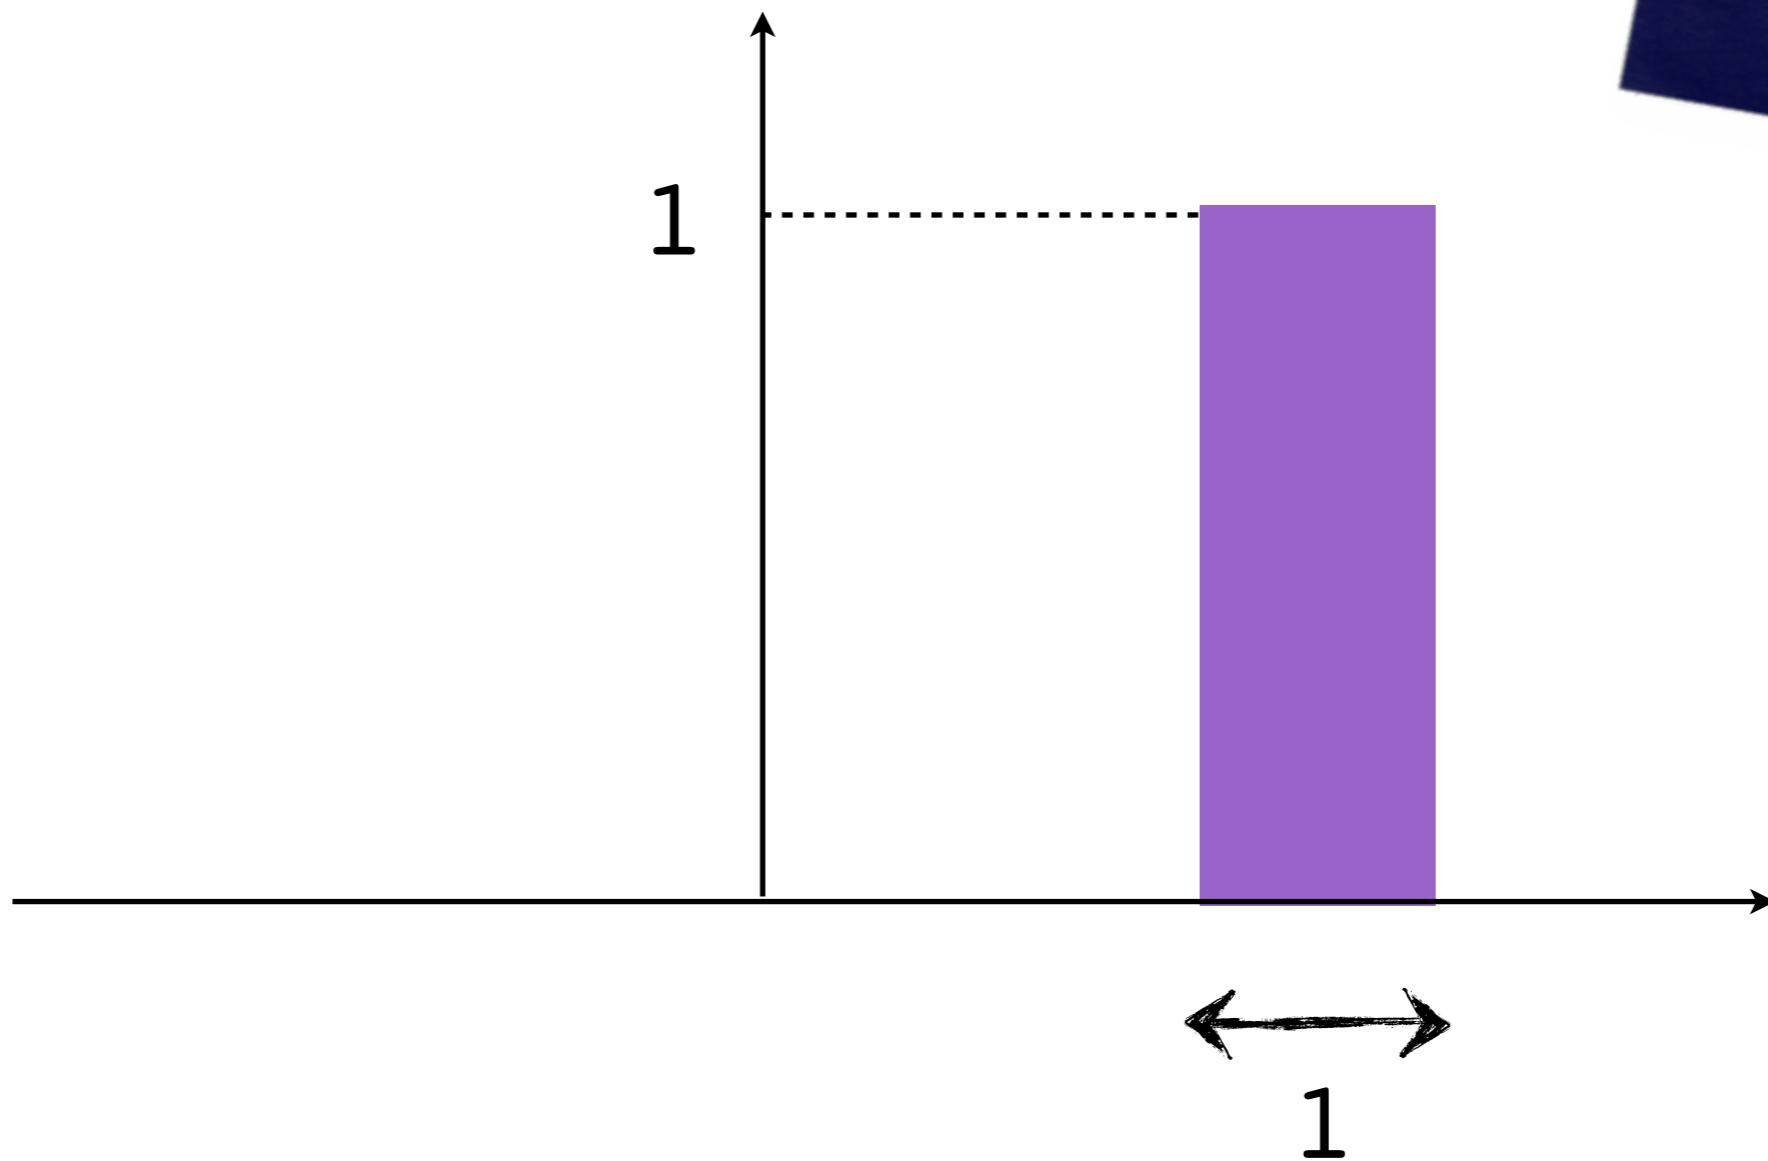
# Convolution

*discrete  
impulse*



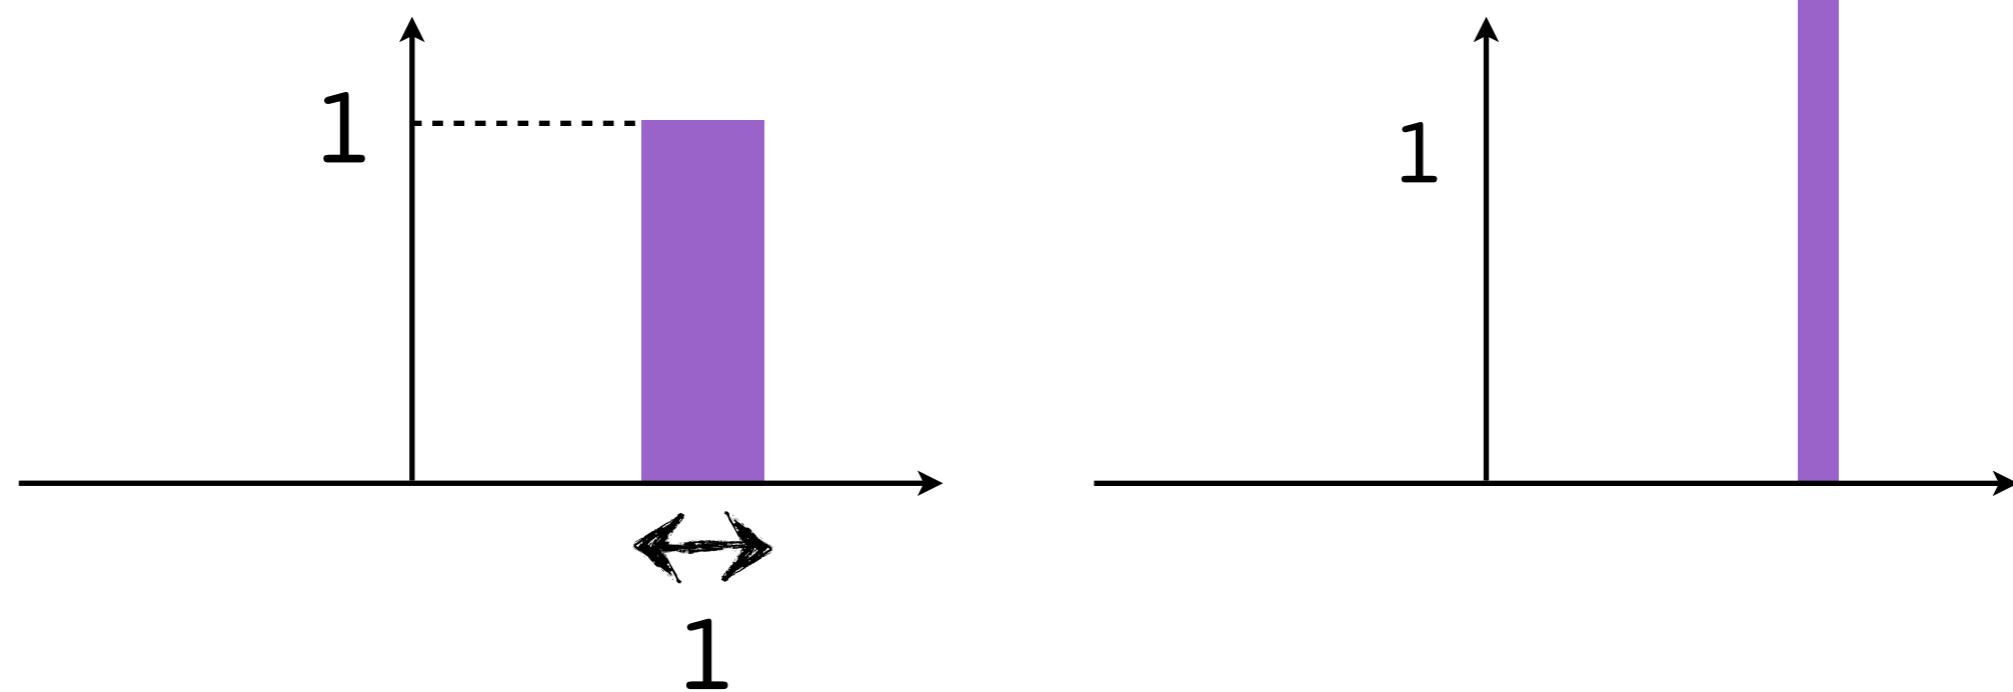
**A function with a value of 1 at a single location**

*continuous  
impulse*



**A function that is very narrow and very tall such that at the limit it has a unit area**

*continuous  
impulse*



**A function that is very narrow and very tall such that at the limit it has a unit area**

|   |        |        |        |   |
|---|--------|--------|--------|---|
| 0 | 0<br>a | 0<br>b | 0<br>c | 0 |
| 0 | 0<br>d | 0<br>e | 0<br>f | 0 |
| 0 | 0<br>g | 1<br>h | 0<br>i | 0 |
| 0 | 0      | 0      | 0      | 0 |
| 0 | 0      | 0      | 0      | 0 |

==

|   |   |  |
|---|---|--|
| i | h |  |
|   |   |  |
|   |   |  |

|   |   |     |     |     |
|---|---|-----|-----|-----|
| 0 | 0 | 0 a | 0 b | 0 c |
| 0 | 0 | 0 d | 0 e | 0 f |
| 0 | 0 | 1 g | 0 h | 0 i |
| 0 | 0 | 0   | 0   | 0   |
| 0 | 0 | 0   | 0   | 0   |

=

|   |   |   |
|---|---|---|
| i | h | g |
|   |   |   |
|   |   |   |

|        |        |        |   |   |
|--------|--------|--------|---|---|
| 0      | 0      | 0      | 0 | 0 |
| 0<br>a | 0<br>b | 0<br>c | 0 | 0 |
| 0<br>d | 0<br>e | 1<br>f | 0 | 0 |
| 0<br>g | 0<br>h | 0<br>i | 0 | 0 |
| 0      | 0      | 0      | 0 | 0 |

=

|   |   |   |
|---|---|---|
| i | h | g |
| f |   |   |
|   |   |   |

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | a | b | c |
| 0 | d | 1 | e | f |
| 0 | g | 0 | h | i |
| 0 | 0 | 0 | 0 | 0 |

=

|   |   |   |
|---|---|---|
| i | h | g |
| f | e |   |
|   |   |   |

|   |   |     |     |     |
|---|---|-----|-----|-----|
| 0 | 0 | 0   | 0   | 0   |
| 0 | 0 | 0   | 0   | 0   |
| 0 | 0 | 1 a | 0 b | 0 c |
| 0 | 0 | 0 d | 0 e | 0 f |
| 0 | 0 | 0 g | 0 h | 0 i |

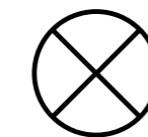
=

|   |   |   |
|---|---|---|
| i | h | g |
| f | e | d |
| c | b | a |

|   |   |   |
|---|---|---|
| i | h | g |
| f | e | d |
| c | b | a |

**Filter output is REVERSED**

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |



|   |   |   |
|---|---|---|
| a | b | c |
| d | e | f |
| g | h | i |

How would you modify the kernel to avoid the flip?

Flip the kernel

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |



|   |   |   |
|---|---|---|
| c | b | a |
| f | e | d |
| i | h | g |

How would you modify the kernel to avoid the flip?

Flip the kernel

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |



|   |   |   |
|---|---|---|
| i | h | g |
| f | e | d |
| c | b | a |

How would you modify the kernel to avoid the flip?

Flip the kernel

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 0 | i | 0 | h | 0 | g | 0 | 0 |
| 0 | f | 0 | e | 0 | d | 0 | 0 |
| 0 | c | 0 | b | 1 | a | 0 | 0 |
| 0 |   | 0 |   | 0 |   | 0 | 0 |
| 0 |   | 0 |   | 0 |   | 0 | 0 |

==

|   |  |  |
|---|--|--|
| a |  |  |
|   |  |  |
|   |  |  |

|   |     |     |     |   |
|---|-----|-----|-----|---|
| 0 | 0 i | 0 h | 0 g | 0 |
| 0 | 0 f | 0 e | 0 d | 0 |
| 0 | 0 c | 1 b | 0 a | 0 |
| 0 | 0   | 0   | 0   | 0 |
| 0 | 0   | 0   | 0   | 0 |

==

|   |   |  |
|---|---|--|
| a | b |  |
|   |   |  |
|   |   |  |

|   |   |            |            |            |
|---|---|------------|------------|------------|
| 0 | 0 | 0 <i>i</i> | 0 <i>h</i> | 0 <i>g</i> |
| 0 | 0 | 0 <i>f</i> | 0 <i>e</i> | 0 <i>d</i> |
| 0 | 0 | 1 <i>c</i> | 0 <i>b</i> | 0 <i>a</i> |
| 0 | 0 | 0          | 0          | 0          |
| 0 | 0 | 0          | 0          | 0          |

=

|   |   |   |
|---|---|---|
| a | b | c |
|   |   |   |
|   |   |   |

|   |   |            |            |            |
|---|---|------------|------------|------------|
| 0 | 0 | 0          | 0          | 0          |
| 0 | 0 | 0          | 0          | 0          |
| 0 | 0 | 1 <i>i</i> | 0 <i>h</i> | 0 <i>g</i> |
| 0 | 0 | 0 <i>f</i> | 0 <i>e</i> | 0 <i>d</i> |
| 0 | 0 | 0 <i>c</i> | 0 <i>b</i> | 0 <i>a</i> |

=

|          |          |          |
|----------|----------|----------|
| <i>a</i> | <i>b</i> | <i>c</i> |
| <i>d</i> | <i>e</i> | <i>f</i> |
| <i>g</i> | <i>h</i> | <i>i</i> |

$$(2K + 1) \times (2K + 1)$$

**filter dimensions**

convolution

$$H[x, y] = \sum_{u=-K}^{K} \sum_{v=-K}^{K} G[u, v] F[x - u, y - v]$$

correlation

$$H[x, y] = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K F[x + u, y + v] G[u, v]$$

*convolution*

$$H[x, y] = \sum_{u=-K}^K \sum_{v=-K}^K G[u, v] F[x - u, y - v]$$

$$H = G * F$$

**convolution**

*convolution*

$$H[x, y] = \sum_{u=-K}^{K} \sum_{v=-K}^{K} G[u, v] F[x - u, y - v]$$

$$\begin{matrix} G & * & F \\ \otimes & & \otimes \end{matrix} \equiv \begin{matrix} \mathcal{D} \end{matrix}$$

# Convolution Properties

*identity*

$$F[x,y] * \delta[x,y] = F[x,y]$$

*distributive*

$$E[x, y] * (F[x, y] + G[x, y])$$

$$= (E[x, y] * F[x, y]) + (E[x, y] * G[x, y])$$

linearity

$$\begin{aligned} G * (\alpha F_1[x, y] + \beta F_2[x, y]) \\ = \alpha H_1[x, y] + \beta H_2[x, y] \end{aligned}$$

shift  
invariant

$$G[x, y] * F[x - \alpha, y - \beta]$$

$$= H[x - \alpha, y - \beta]$$

associative

$$\begin{aligned}(E[x, y] * F[x, y]) * G[x, y] \\= E[x, y] * (F[x, y] * G[x, y])\end{aligned}$$

**sequential filtering is equivalent to applying one filter**

$$G(\mathbf{x}, \sqrt{\sigma_1^2 + \sigma_2^2}) = G(\mathbf{x}, \sigma_1) * G(\mathbf{x}, \sigma_2)$$

**convolution of Gaussians is a Gaussian**

associative

$$\begin{aligned}(E[x, y] * F[x, y]) * G[x, y] \\= E[x, y] * (F[x, y] * G[x, y])\end{aligned}$$

**sequential filtering is equivalent to applying one filter**

correlation

$$\left( \begin{array}{c|c} -1 & 1 \end{array} \otimes \begin{array}{c|c} 1 & -1 \end{array} \right) \otimes \begin{array}{c|c} -1 & 1 \end{array}$$

correlation

$$1 \Big) \otimes \begin{array}{|c|c|} \hline -1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline -1 & 3 & -3 & 1 \\ \hline \end{array}$$

correlation

$$\left( \begin{array}{c|c} -1 & 1 \end{array} \otimes \begin{array}{c|c} 1 & -1 \end{array} \right) \otimes \begin{array}{c|c} -1 & 1 \end{array}$$

correlation

$$\begin{pmatrix} -1 & 1 \end{pmatrix} \otimes \left( \begin{pmatrix} 1 & -1 \end{pmatrix} \otimes \begin{pmatrix} -1 & 1 \end{pmatrix} \right)$$

correlation

$$\begin{bmatrix} 1 \\ -3 \\ 3 \\ -1 \end{bmatrix} = \begin{bmatrix} -1 \\ 1 \end{bmatrix} \otimes \left( \begin{bmatrix} 1 \\ -1 \end{bmatrix} \right)$$

correlation

$$\begin{array}{|c|c|c|c|} \hline 1 & -3 & 3 & -1 \\ \hline \end{array} \neq \begin{array}{|c|c|c|c|} \hline -1 & 3 & -3 & 1 \\ \hline \end{array}$$

correlation

$$\begin{matrix} 1 & -3 & 3 & -1 \end{matrix} \neq \begin{matrix} -1 & 3 & -3 & 1 \end{matrix}$$

**correlation is not associative**

commutative

$$F[x,y]*G[x,y] = G[x,y]*F[x,y]$$

commutativity  
proof

By definition

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

Let  $p = x - u$  and  $q = y - v$

substitute

$$H * F = \sum_{q=-\infty}^{\infty} \sum_{p=-\infty}^{\infty} H[x - p, y - q]F[p, q]$$

rearrange

$$= \sum_{q=-\infty}^{\infty} \sum_{p=-\infty}^{\infty} F[p, q]H[x - p, y - q]$$

$$= F * H$$



motion  
blur

MAPLE

*motion  
blur*

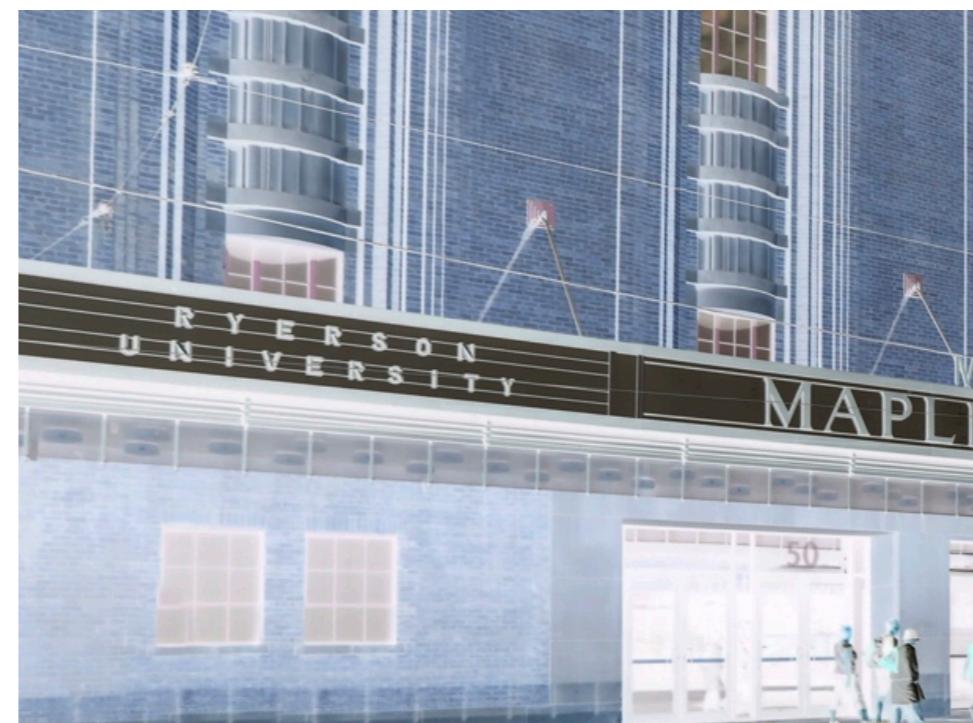


**blurred image**

*motion  
blur*



**blurred image**

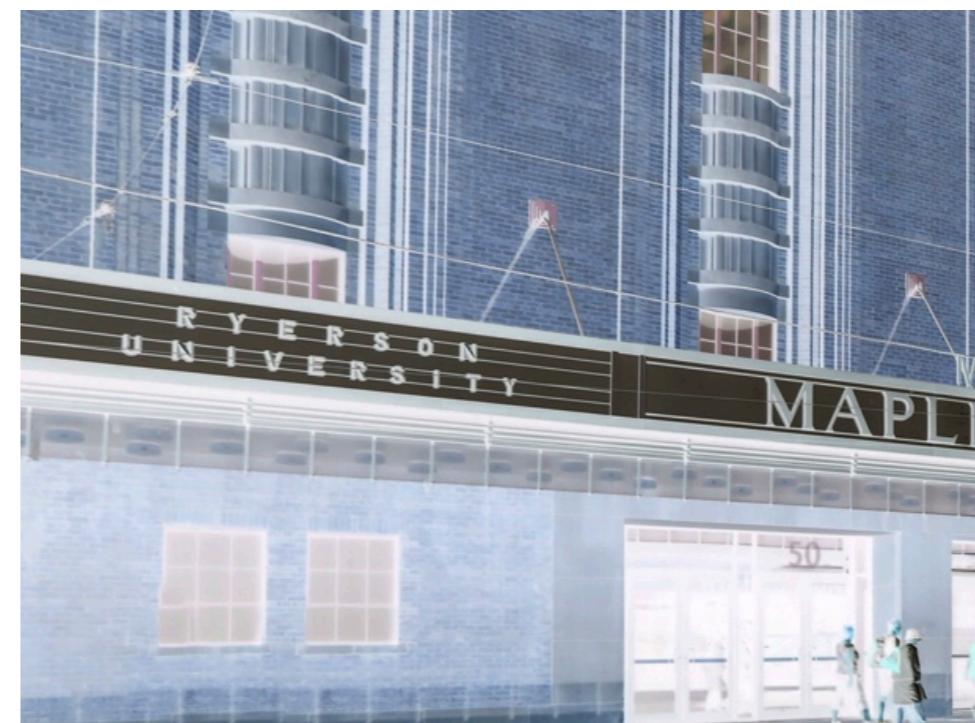


**ideal image**

*motion  
blur*

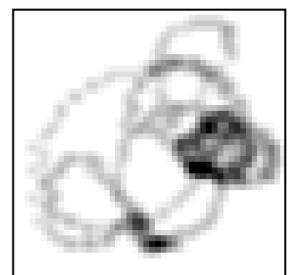


**blurred image**



**ideal image**

\*



**kernel**

$$H[x, y] = \sum_{u=-K}^{K} \sum_{v=-K}^{K} G[u, v] F[x - u, y - v]$$

How many operations does convolution require per pixel?

$$(2K + 1)^2$$

*separable  
filters*

$$F[x, y] = U[x]V[y]$$

**Assume the filter can be written as**

$$F[x, y] = U[x]V[y]$$

---

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

**Assume the filter can be written as**

$$F[x, y] = U[x]V[y]$$

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

*factor out*

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

Assume the filter can be written as

$$F[x, y] = U[x]V[y]$$

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

*factor out*

1D horizontal convolution

Assume the filter can be written as

$$F[x, y] = U[x]V[y]$$

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

*factor out*

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

1D vertical convolution

**Assume the filter can be written as**

$$F[x, y] = U[x]V[y]$$

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

*factor out*

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

**Filter horizontally then vertically**

Assume the filter can be written as

$$F[x, y] = U[x]V[y]$$

$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

*insert filter definition*

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

*factor out*

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

How many operations does convolution require per pixel?

$$2(2K + 1)$$

convolution

$$H[x, y] = \sum_{u=-K}^{K} \sum_{v=-K}^{K} G[u, v] F[x - u, y - v]$$

How many operations does convolution require per pixel?

$$(2K + 1)^2$$

**Assume the filter can be written as:**

$$F[x, y] = U[x]V[y]$$

separable  
filtering

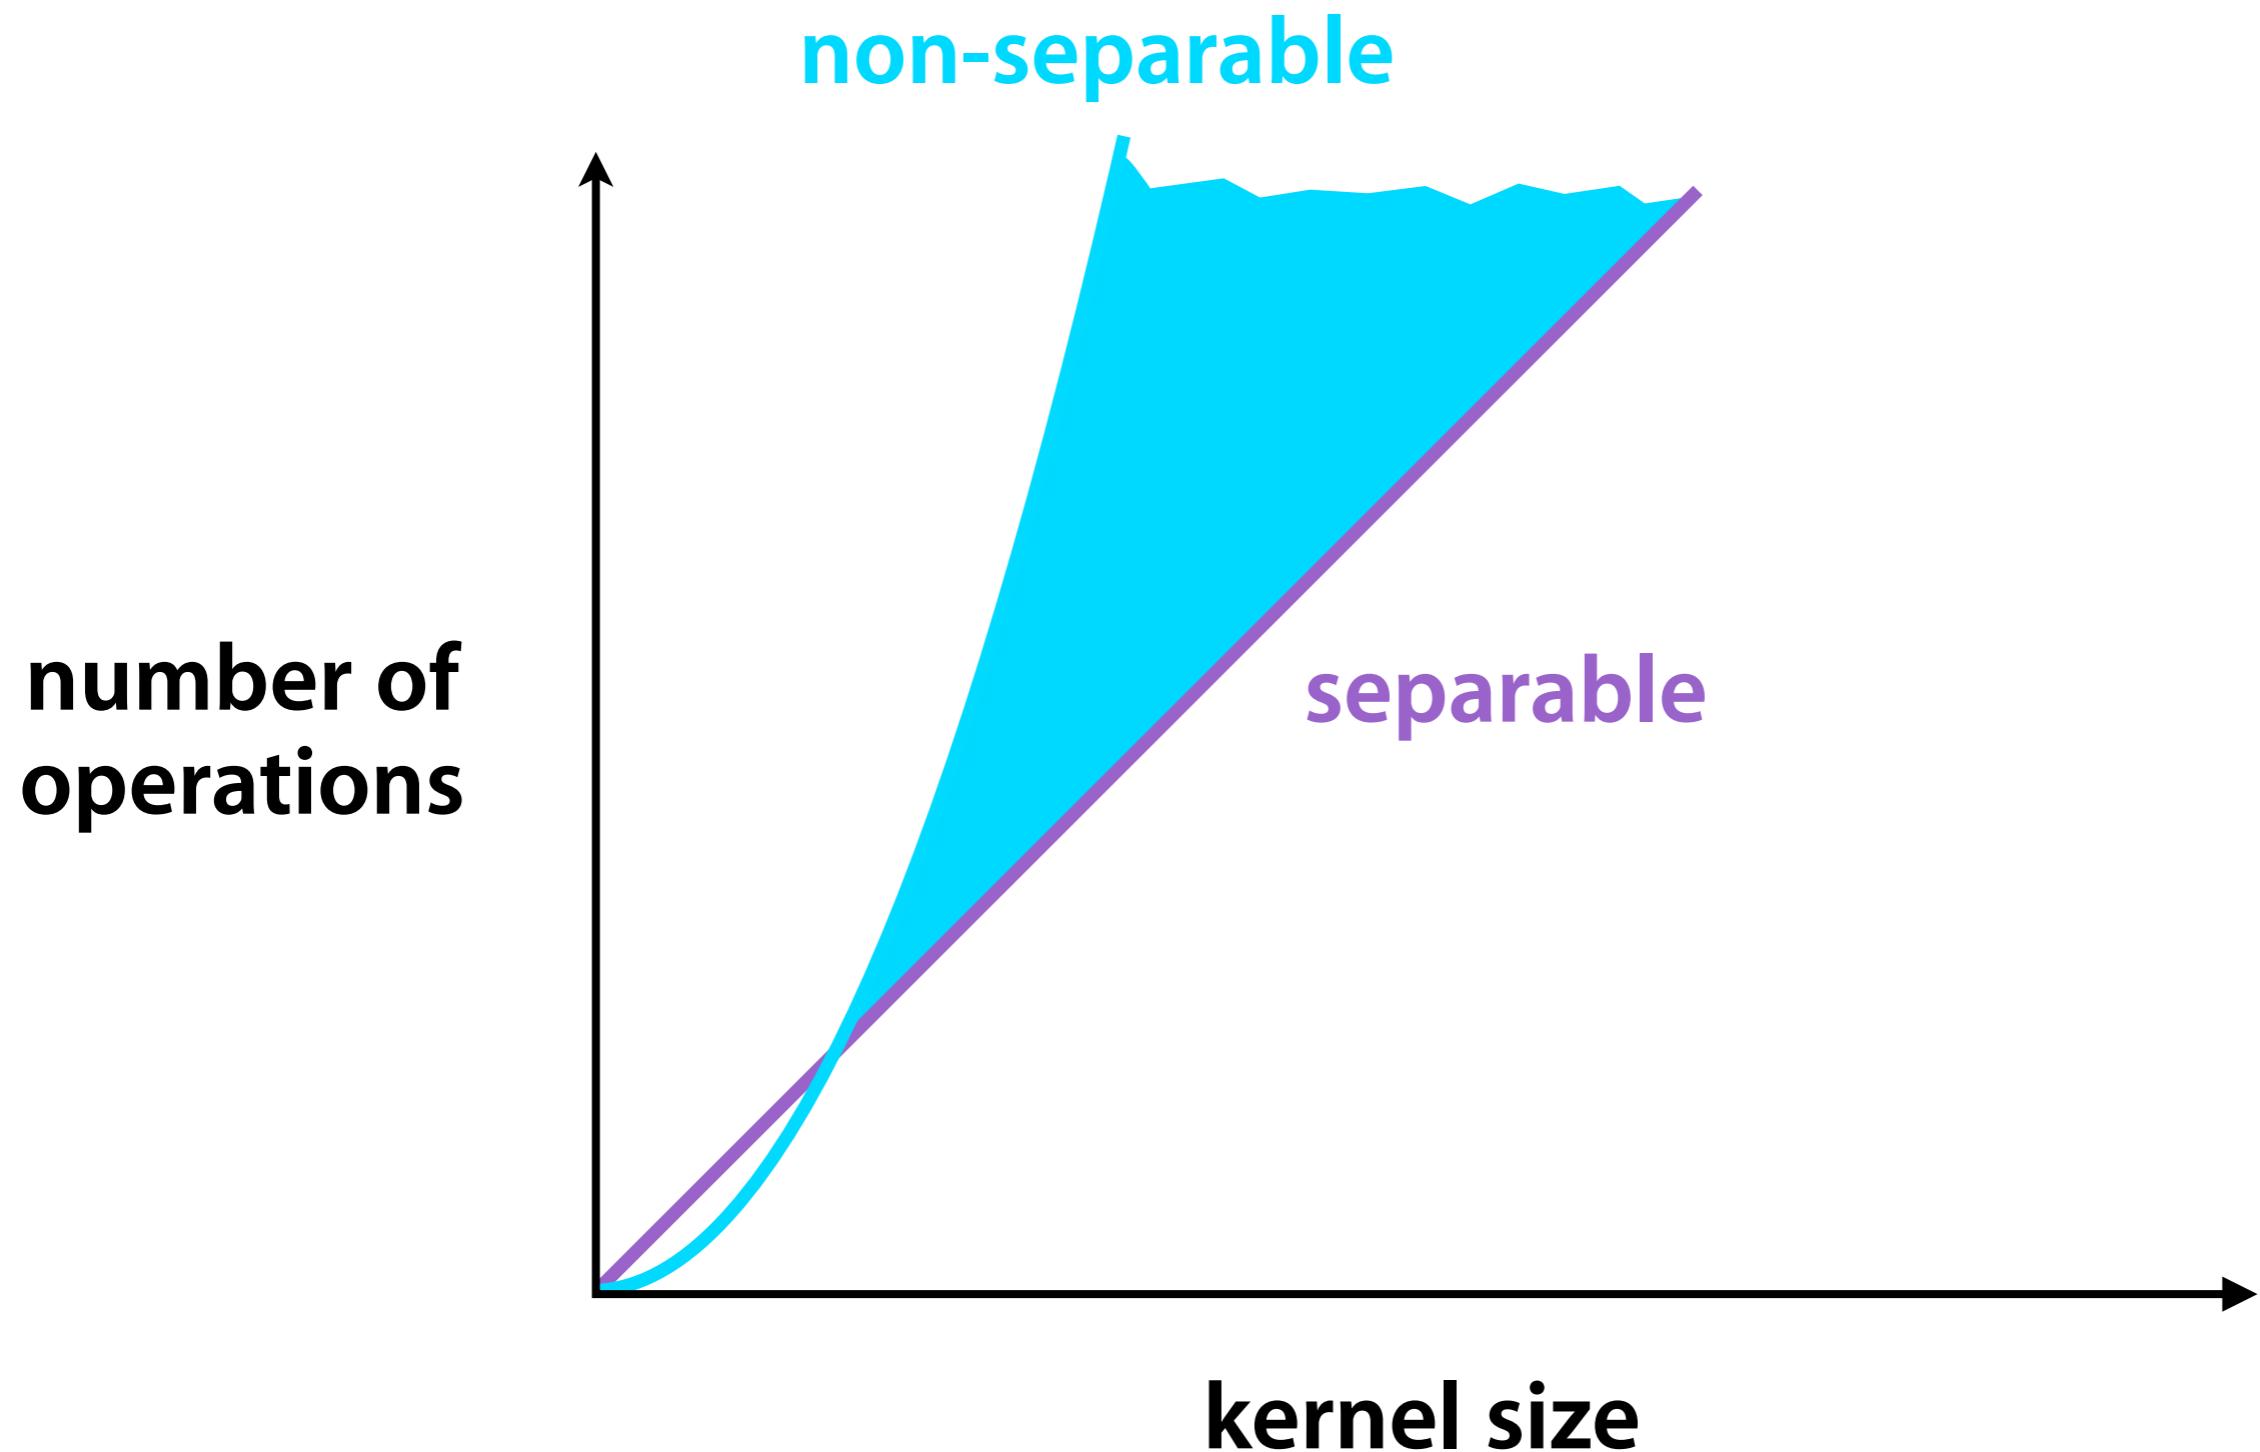
$$H * F = \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]F[x - u, y - v]$$

$$= \sum_{v=-\infty}^{\infty} \sum_{u=-\infty}^{\infty} H[u, v]U[x - u]V[y - v]$$

$$= \sum_{v=-\infty}^{\infty} V[y - v] \left( \sum_{u=-\infty}^{\infty} H[u, v]U[x - u] \right)$$

**How many operations does convolution require per pixel?**

$$2(2K + 1)$$



**Show that the 2D Gaussian filter is separable.**

$$\begin{aligned} G(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \\ &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \exp\left(-\frac{y^2}{2\sigma^2}\right) \end{aligned}$$

**Show that the 2D Gaussian filter is separable.**

$$\begin{aligned} G(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \\ &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \exp\left(-\frac{y^2}{2\sigma^2}\right) \\ &= \left[ \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{x^2}{2\sigma^2}\right) \right] \left[ \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{y^2}{2\sigma^2}\right) \right] \\ &= G_1(x)G_1(y) \end{aligned}$$

**For a symmetric filter how will the outputs differ?**

**no difference between outputs**

**For an impulse signal, how will the outputs differ?**

**output of convolution is flipped**

**For an impulse filter, how will the outputs differ?**

**no difference between outputs**

image  
sharpening



**input**



**blurred**



**"sharp-stuff"**

image  
sharpening



**blurred**

+



**“sharp-stuff”**

=



**sharpened**

|    |    |    |
|----|----|----|
| -1 | -1 | -1 |
| -1 | 9  | -1 |
| -1 | -1 | -1 |

# Sharpening Filter



**input**

$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} & = \end{matrix}$$

?



input

$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} & = \end{matrix}$$



no change to input image



input

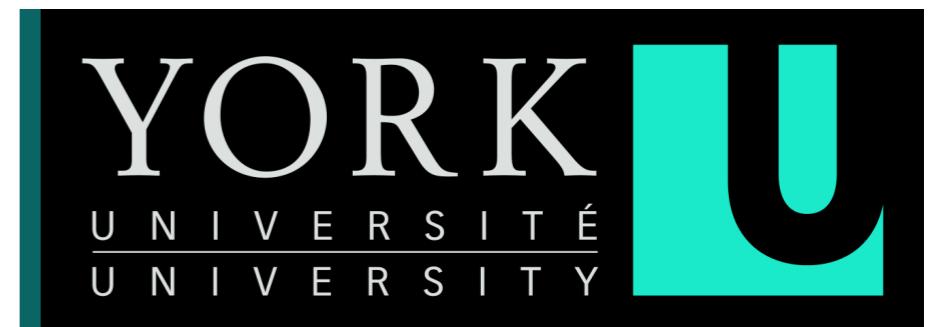
$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} & = \end{matrix}$$

?



input

$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} & = \end{matrix}$$





$$\begin{matrix} \mathbf{U} \\ * \end{matrix} \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{pmatrix} =$$



shift input image one pixel to the right



**input**

$$* \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$

?



input

$$* \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



input image blurred by box filtering

|     |     |     |     |     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 108 | 76  | 98  | 14  | 246 | 146 | 56  | 236 | 14  | 12  |
| 20  | 44  | 180 | 62  | 244 | 8   | 20  | 156 | 150 | 122 |
| 224 | 144 | 230 | 80  | 108 | 76  | 140 | 114 | 44  | 144 |
| 174 | 234 | 26  | 56  | 34  | 234 | 228 | 40  | 100 | 34  |
| 152 | 194 | 90  | 30  | 202 | 22  | 200 | 226 | 130 | 50  |
| 164 | 112 | 40  | 52  | 200 | 0   | 8   | 44  | 252 | 102 |
| 186 | 34  | 150 | 212 | 102 | 236 | 48  | 72  | 228 | 54  |

How can we compute the sum within the box?

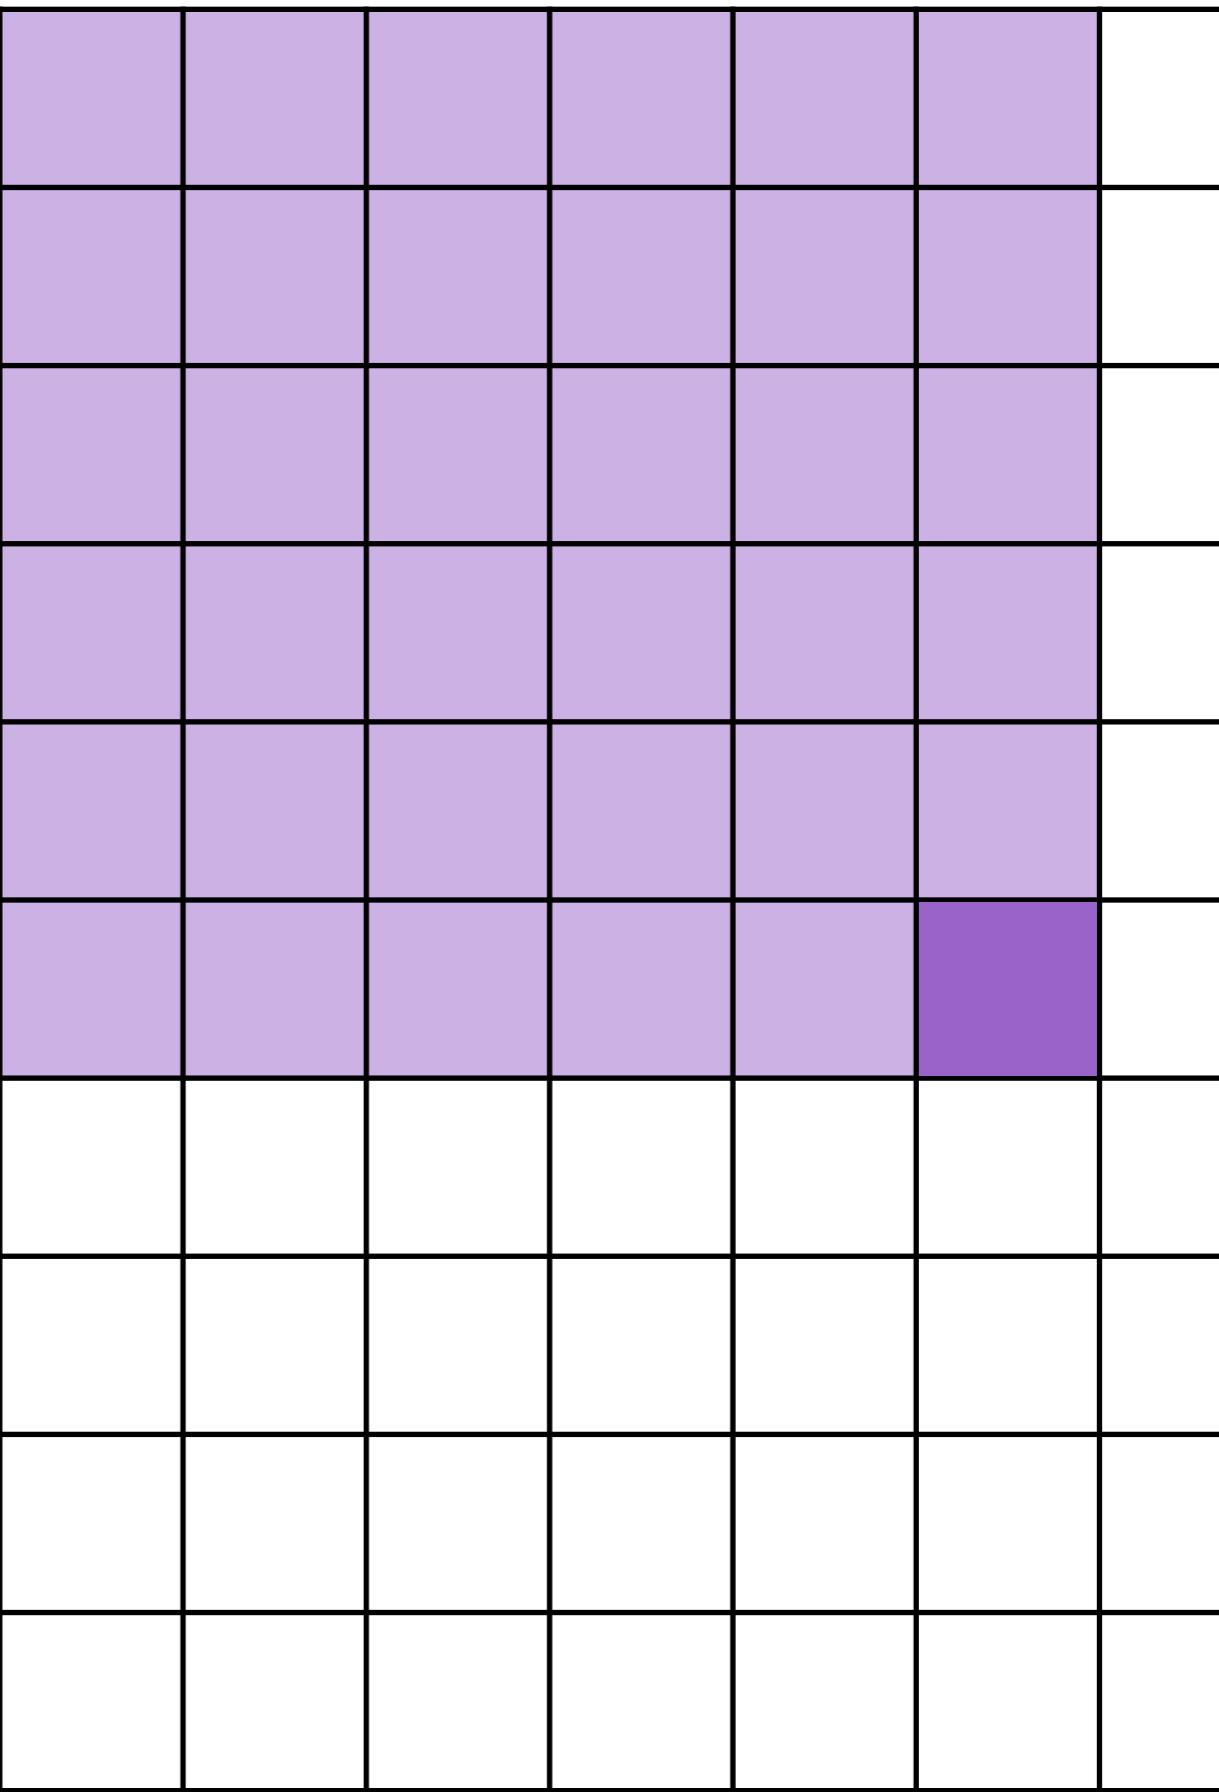
|     |     |     |    |     |    |     |     |     |     |
|-----|-----|-----|----|-----|----|-----|-----|-----|-----|
| 118 | 200 | 196 | 28 | 94  | 34 | 32  | 52  | 104 | 160 |
| 52  | 34  | 18  | 90 | 202 | 52 | 110 | 218 | 252 | 74  |

A large grid of squares, mostly light purple, with a dark blue corner containing white text.

# Summed Area Table

elements contain the pixel sums in the upper-left

$$\begin{aligned}I_{\text{S.A.T}}[x, y] &= \\& I[x, y] + \\& I_{\text{S.A.T.}}[x - 1, y] + \\& I_{\text{S.A.T.}}[x, y - 1] - \\& I_{\text{S.A.T.}}[x - 1, y - 1]\end{aligned}$$

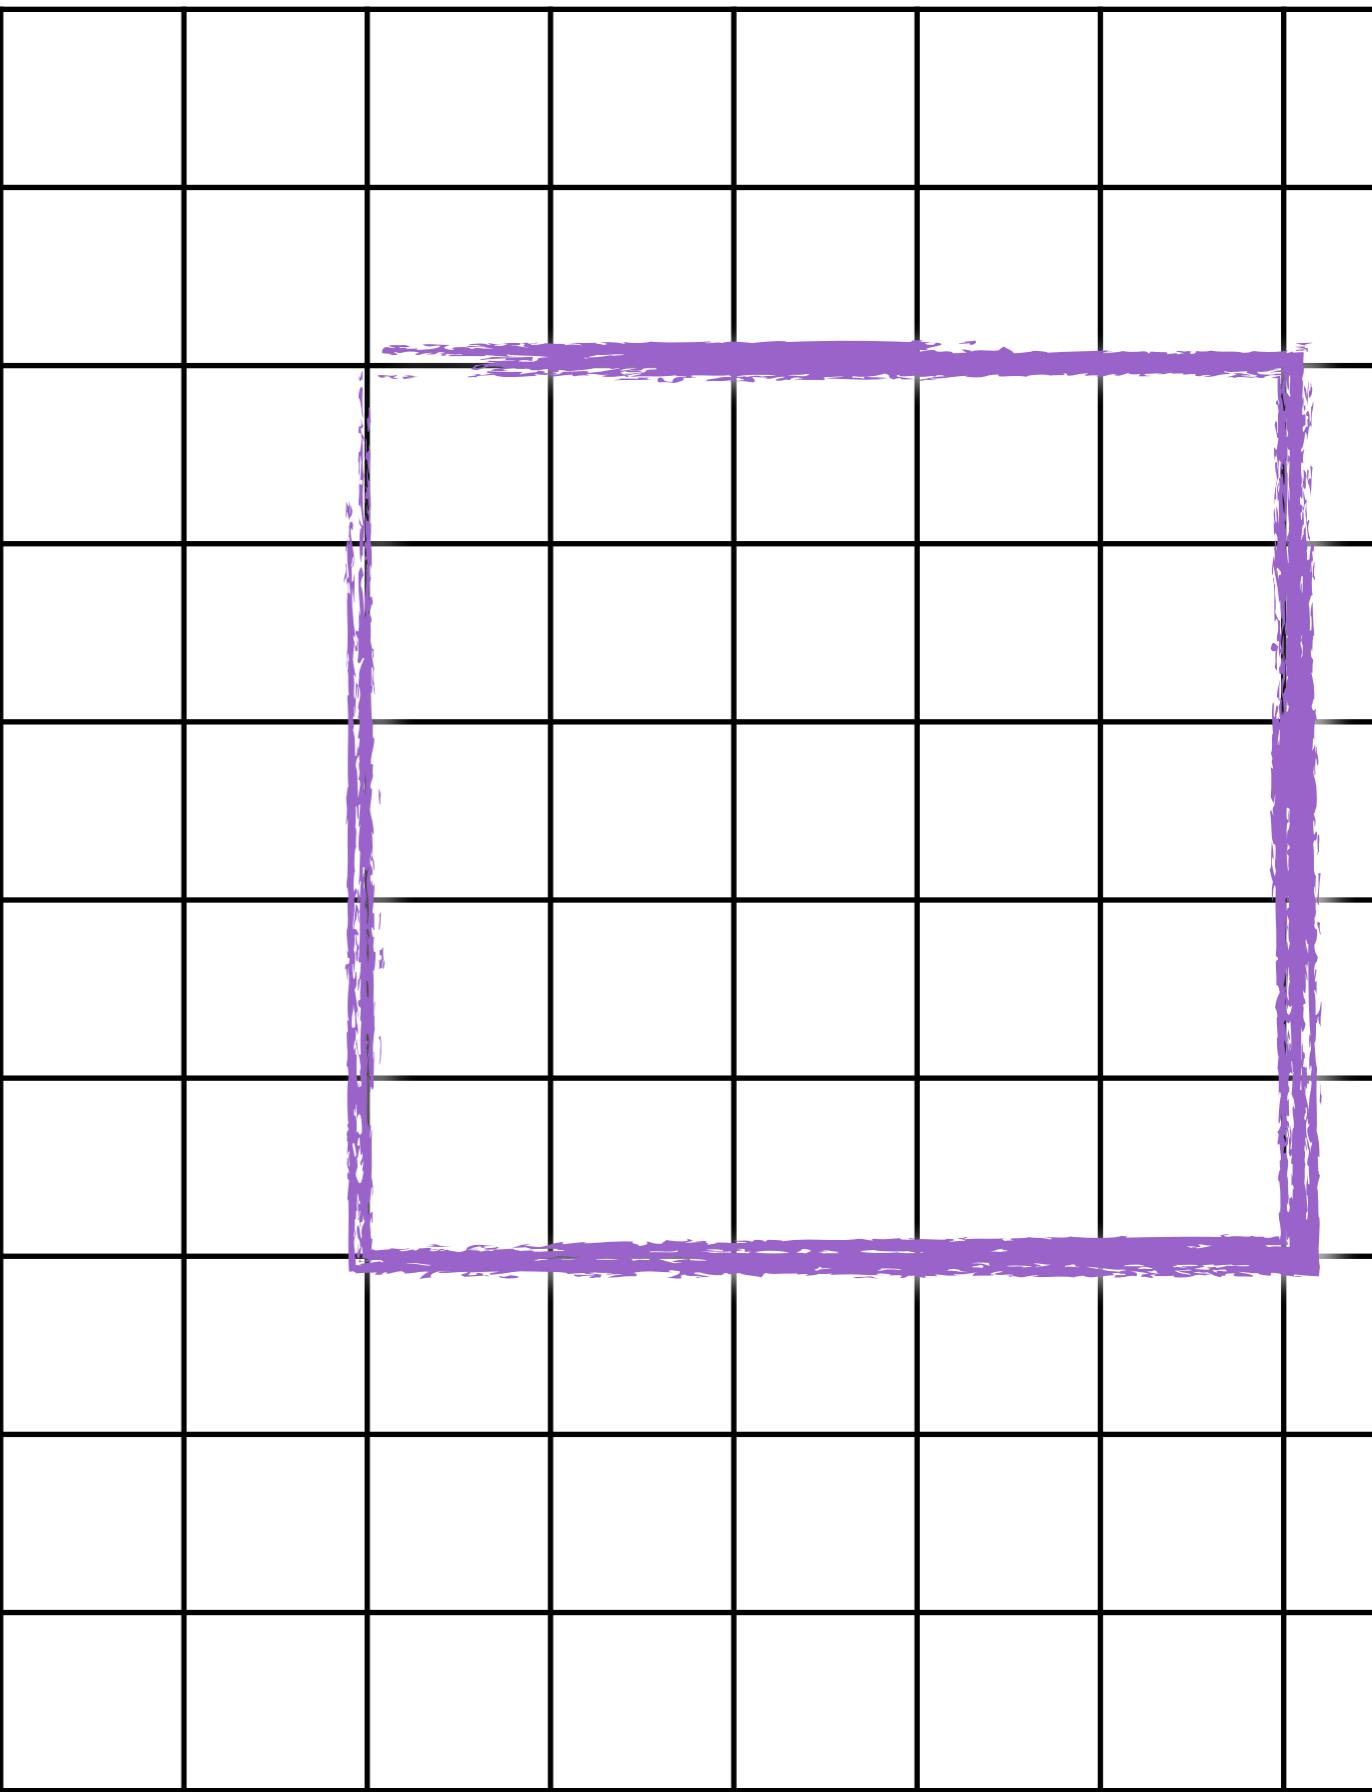


|  |  |  |  |  |  |  |  |  |
|--|--|--|--|--|--|--|--|--|
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |

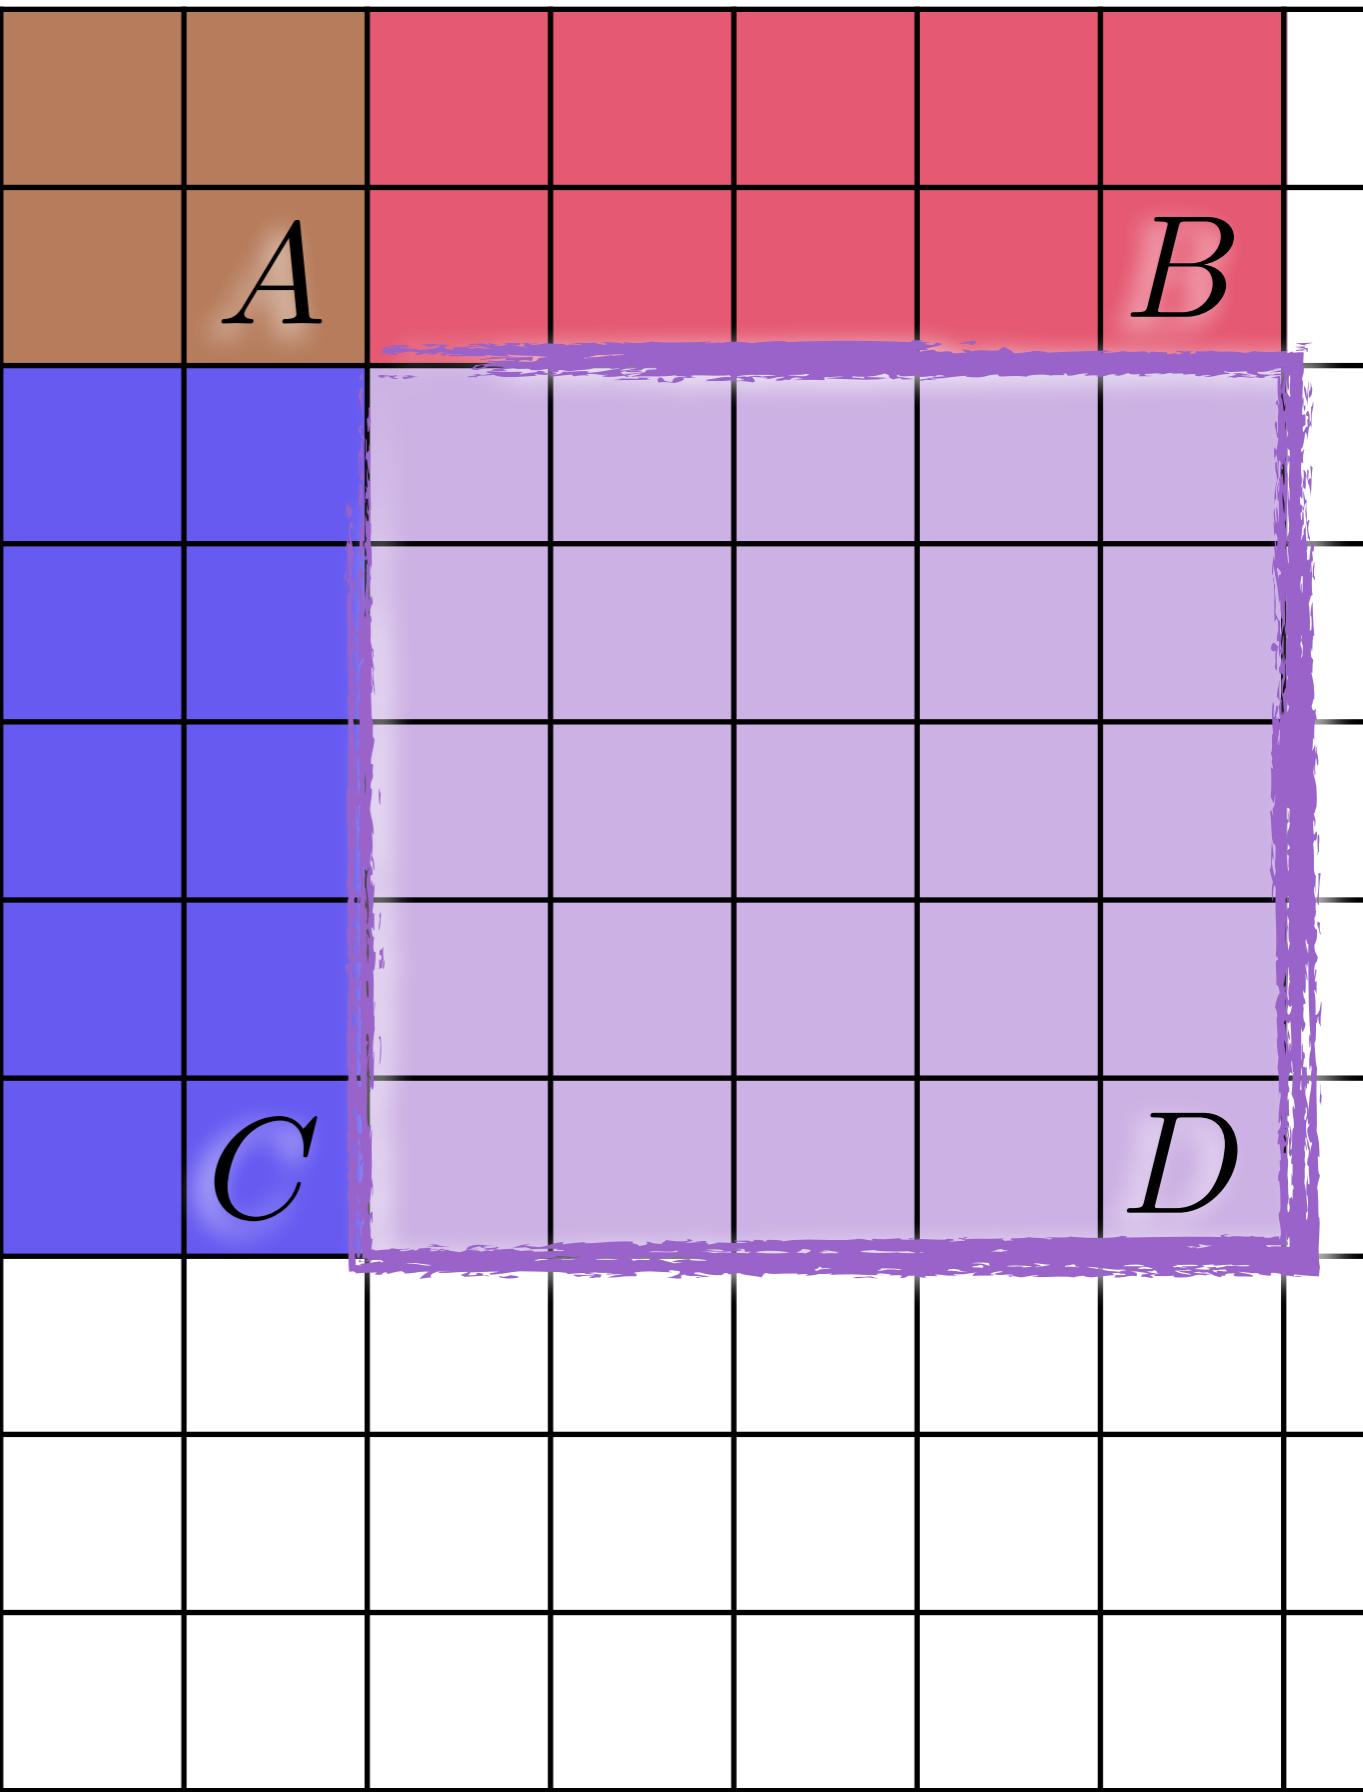
**only three additions are required to compute box sum**

|  |  |  |  |  |  |  |  |  |
|--|--|--|--|--|--|--|--|--|
|  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |

$$D - B - C + A$$



$$D - B - C + A$$

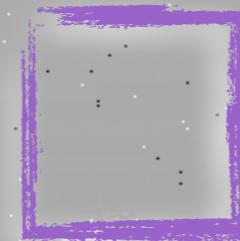


# **Non-linear Filtering**



**salt and pepper noise**





|    |    |    |
|----|----|----|
| 10 | 15 | 20 |
| 23 | 90 | 27 |
| 33 | 31 | 30 |

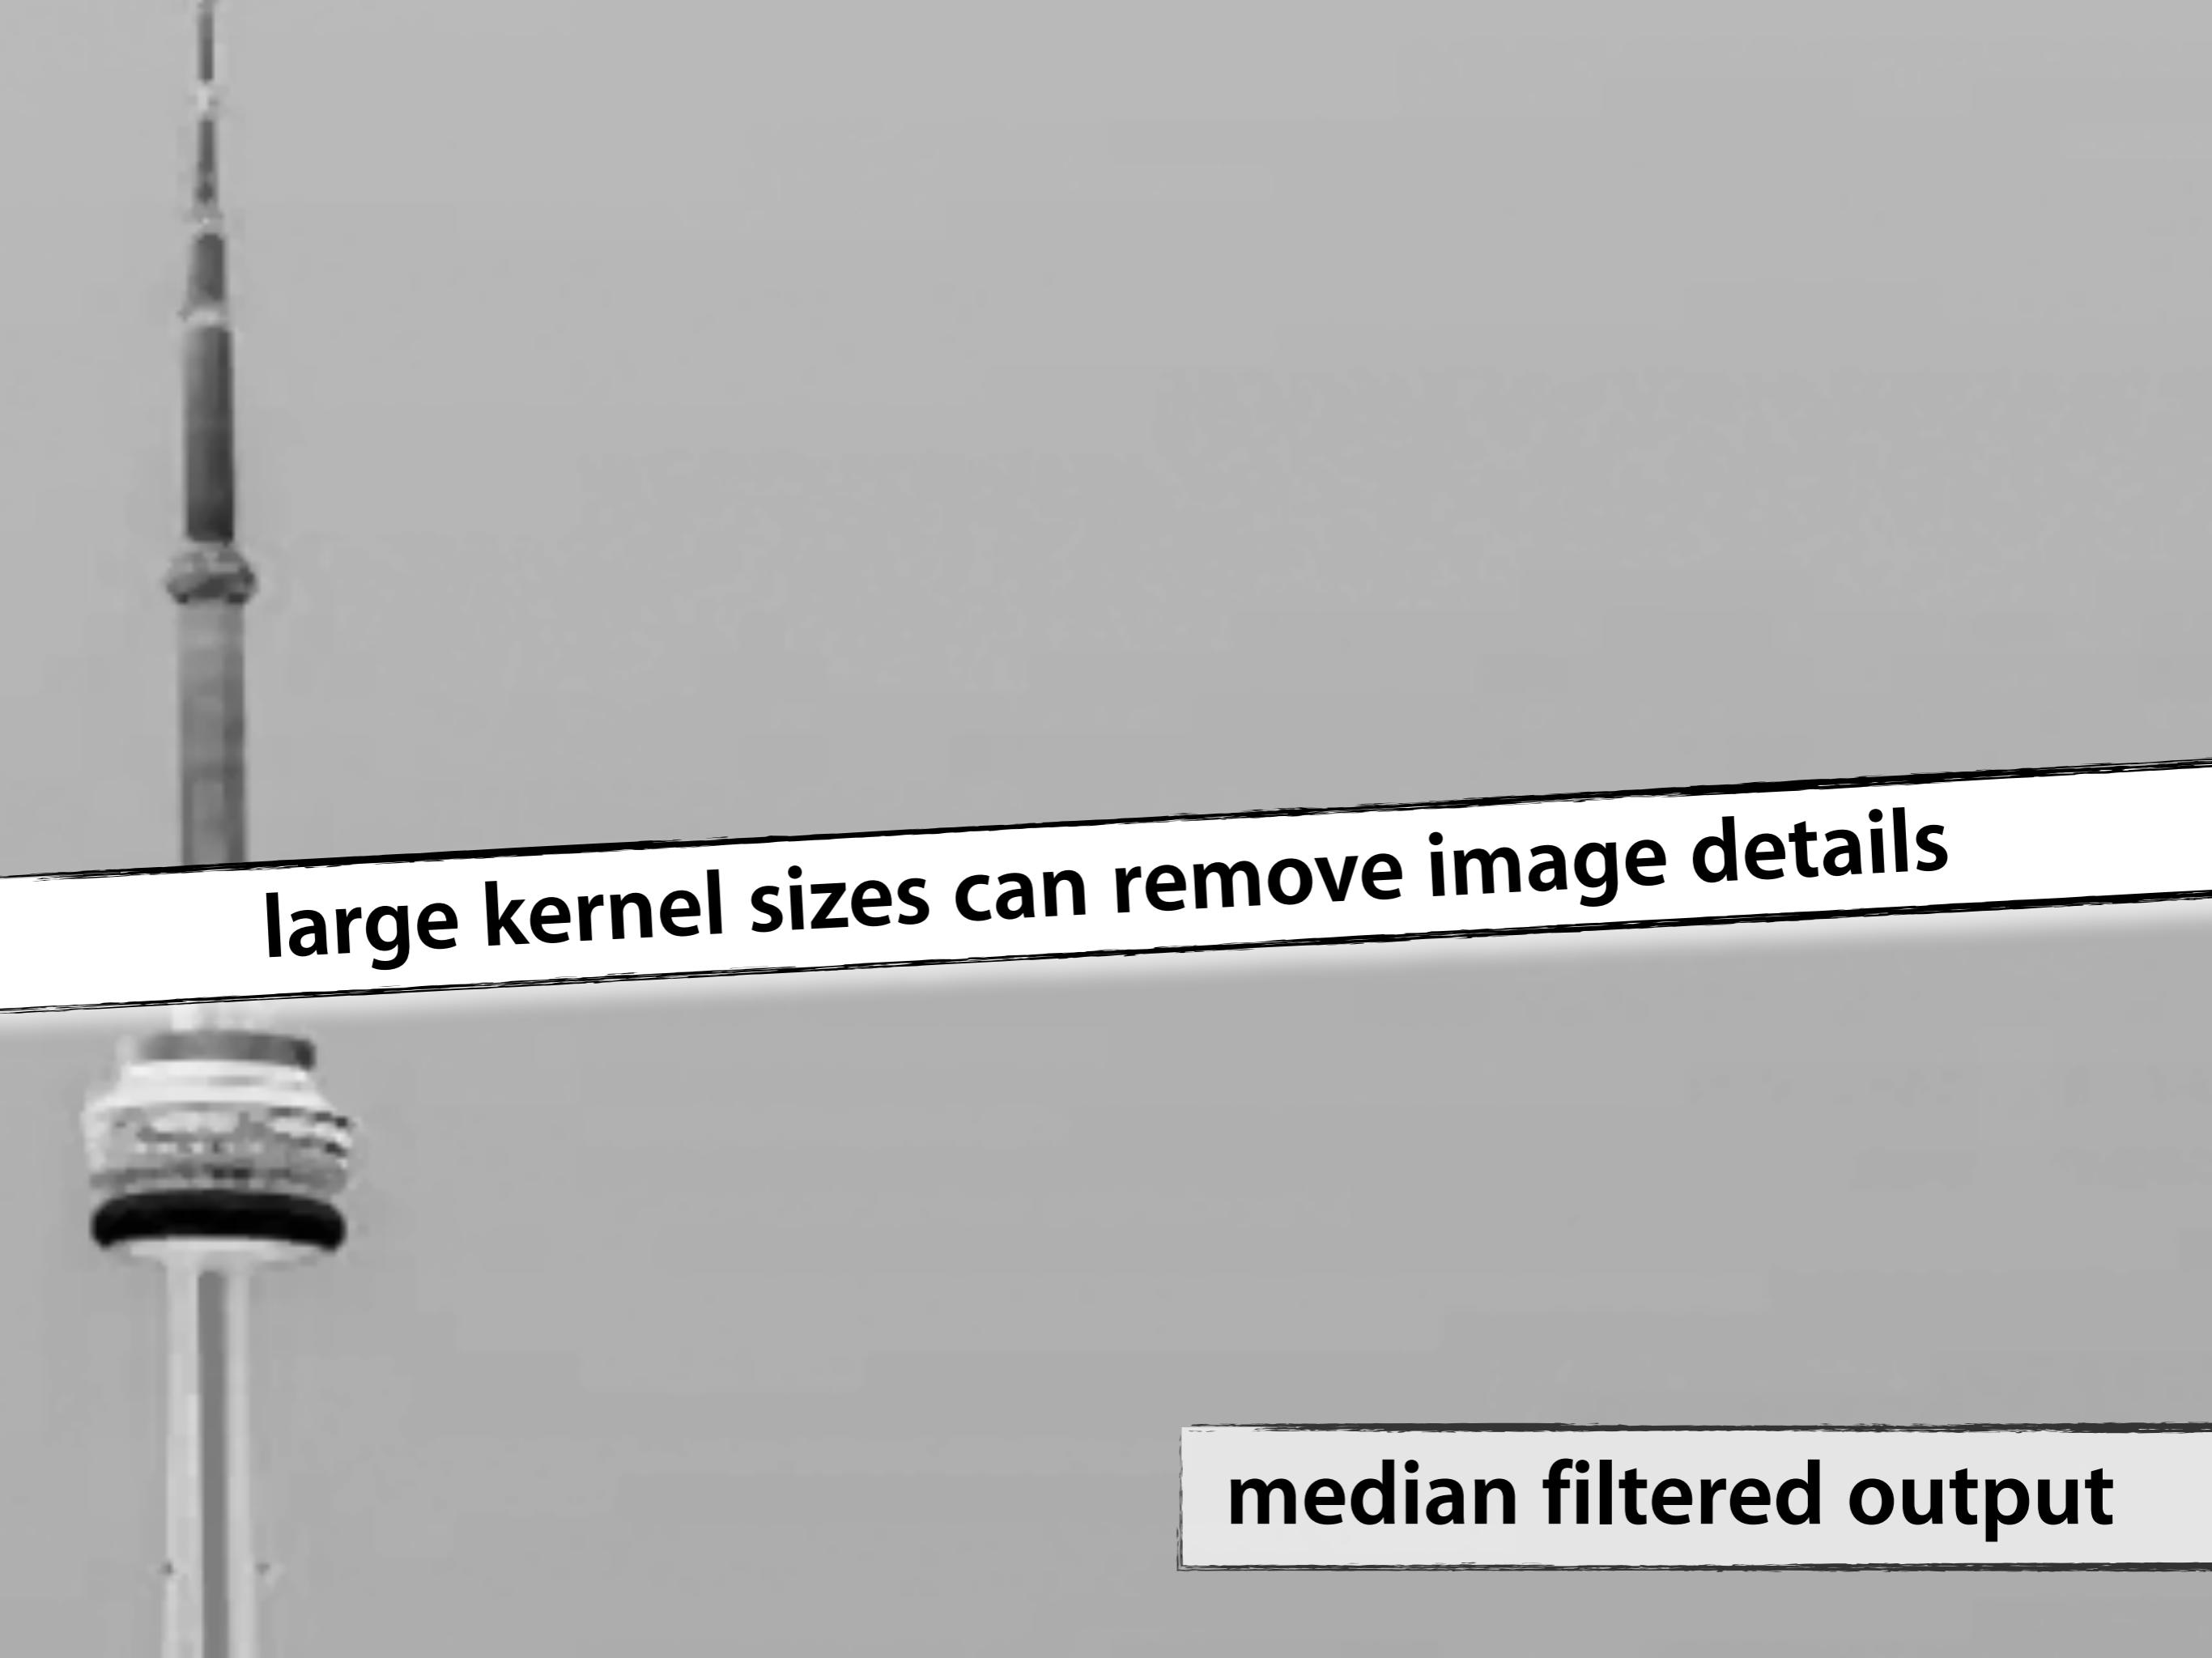
sort pixel values

10 15 20 23 27 30 31 33 90

replace pixel with median value



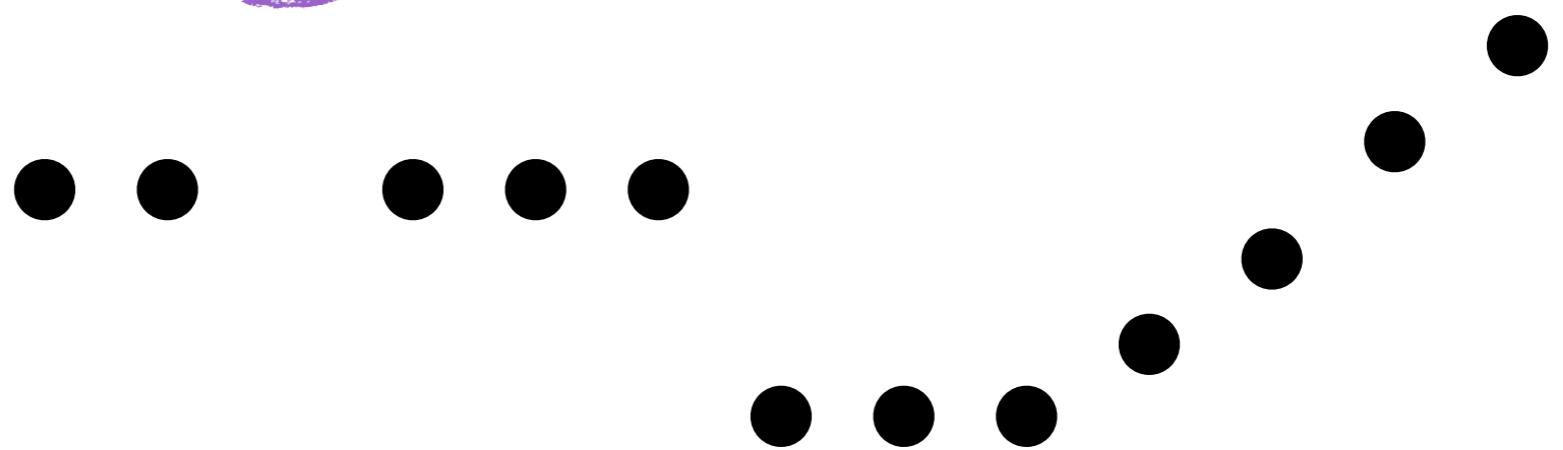
**median filtered output**



**large kernel sizes can remove image details**

**median filtered output**

input



median  
filtering

median  
filtering

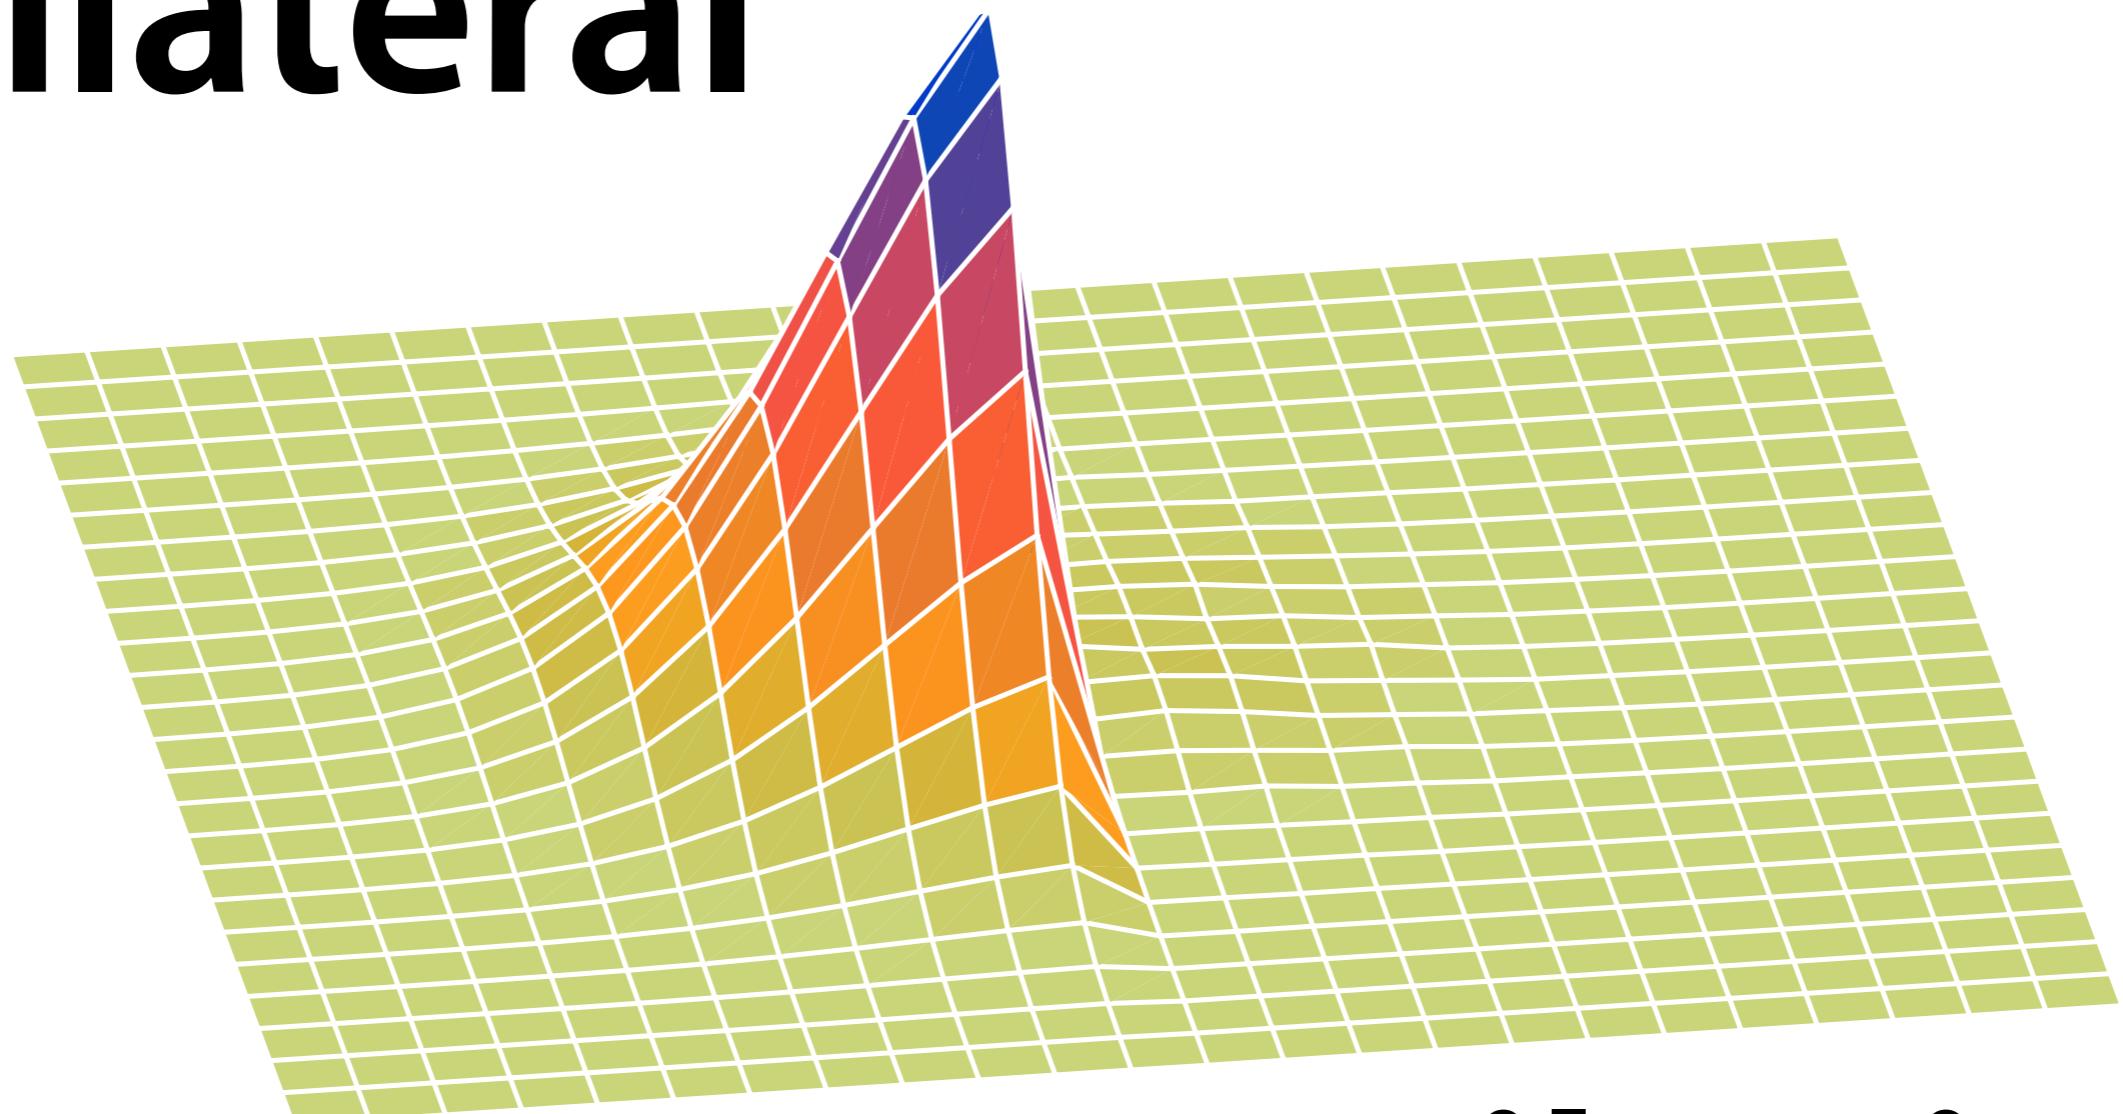
input



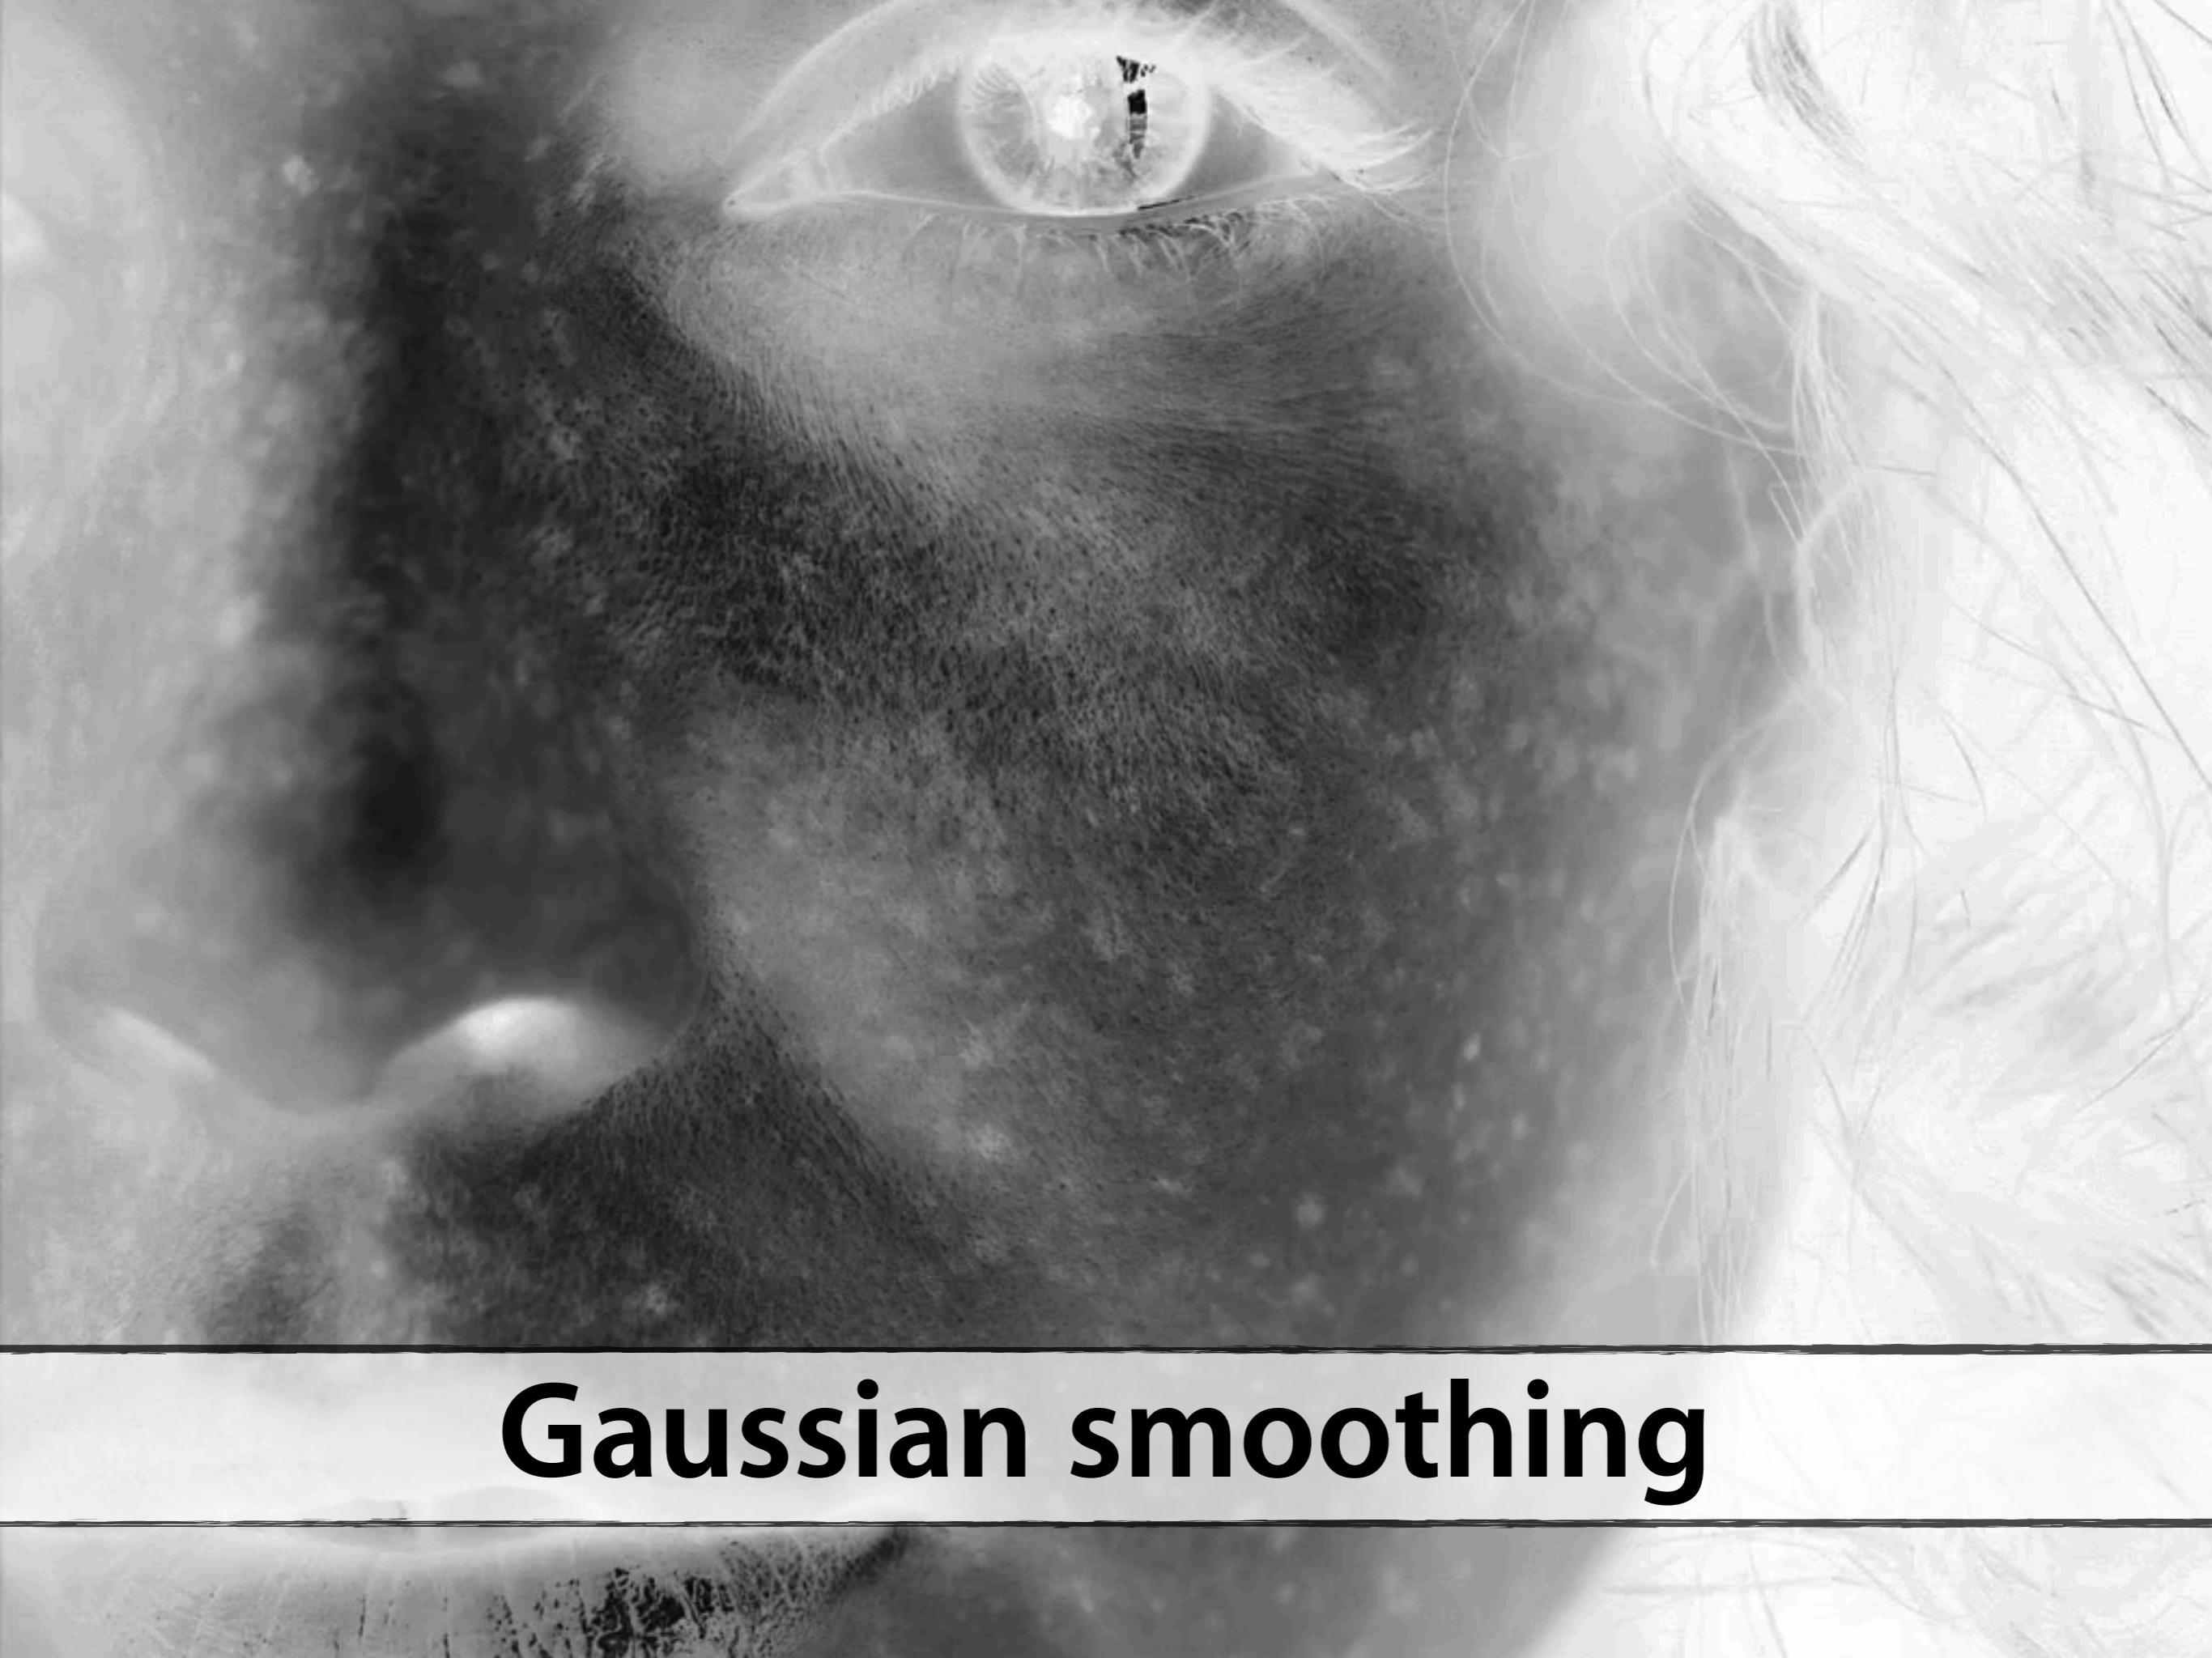
median  
filtered

mean  
filtered

# Bilateral

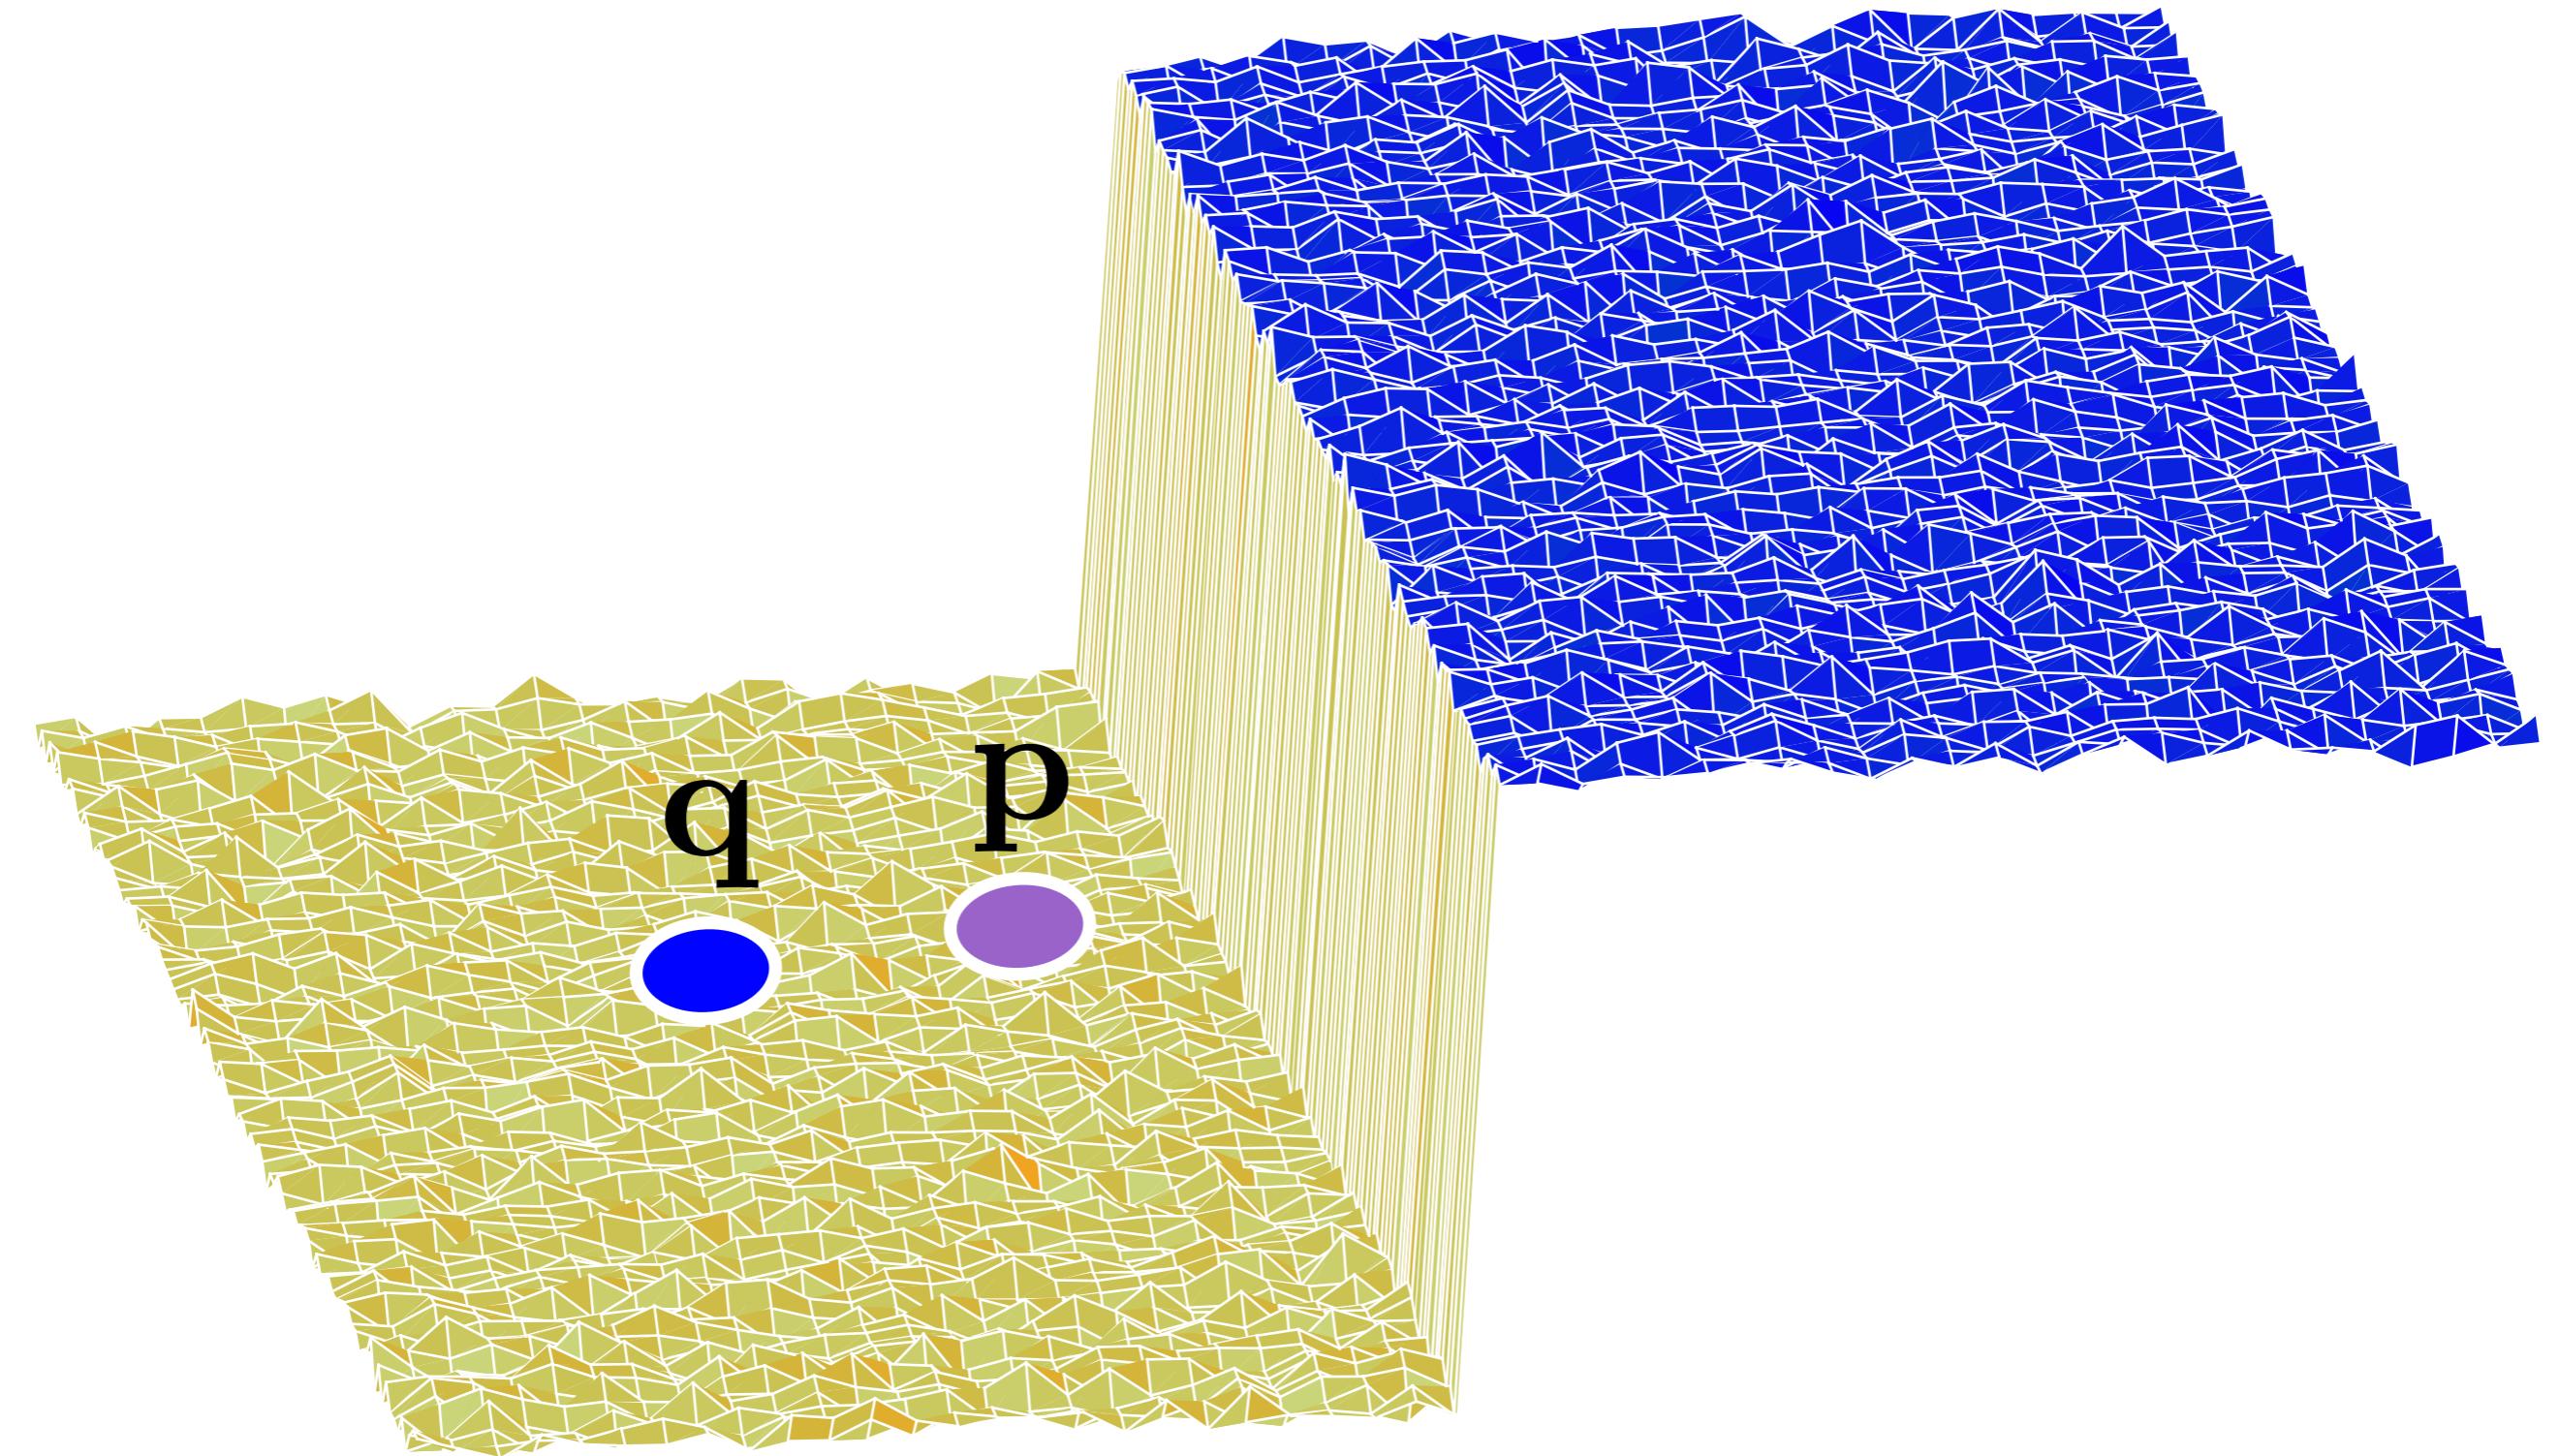


# Filtering



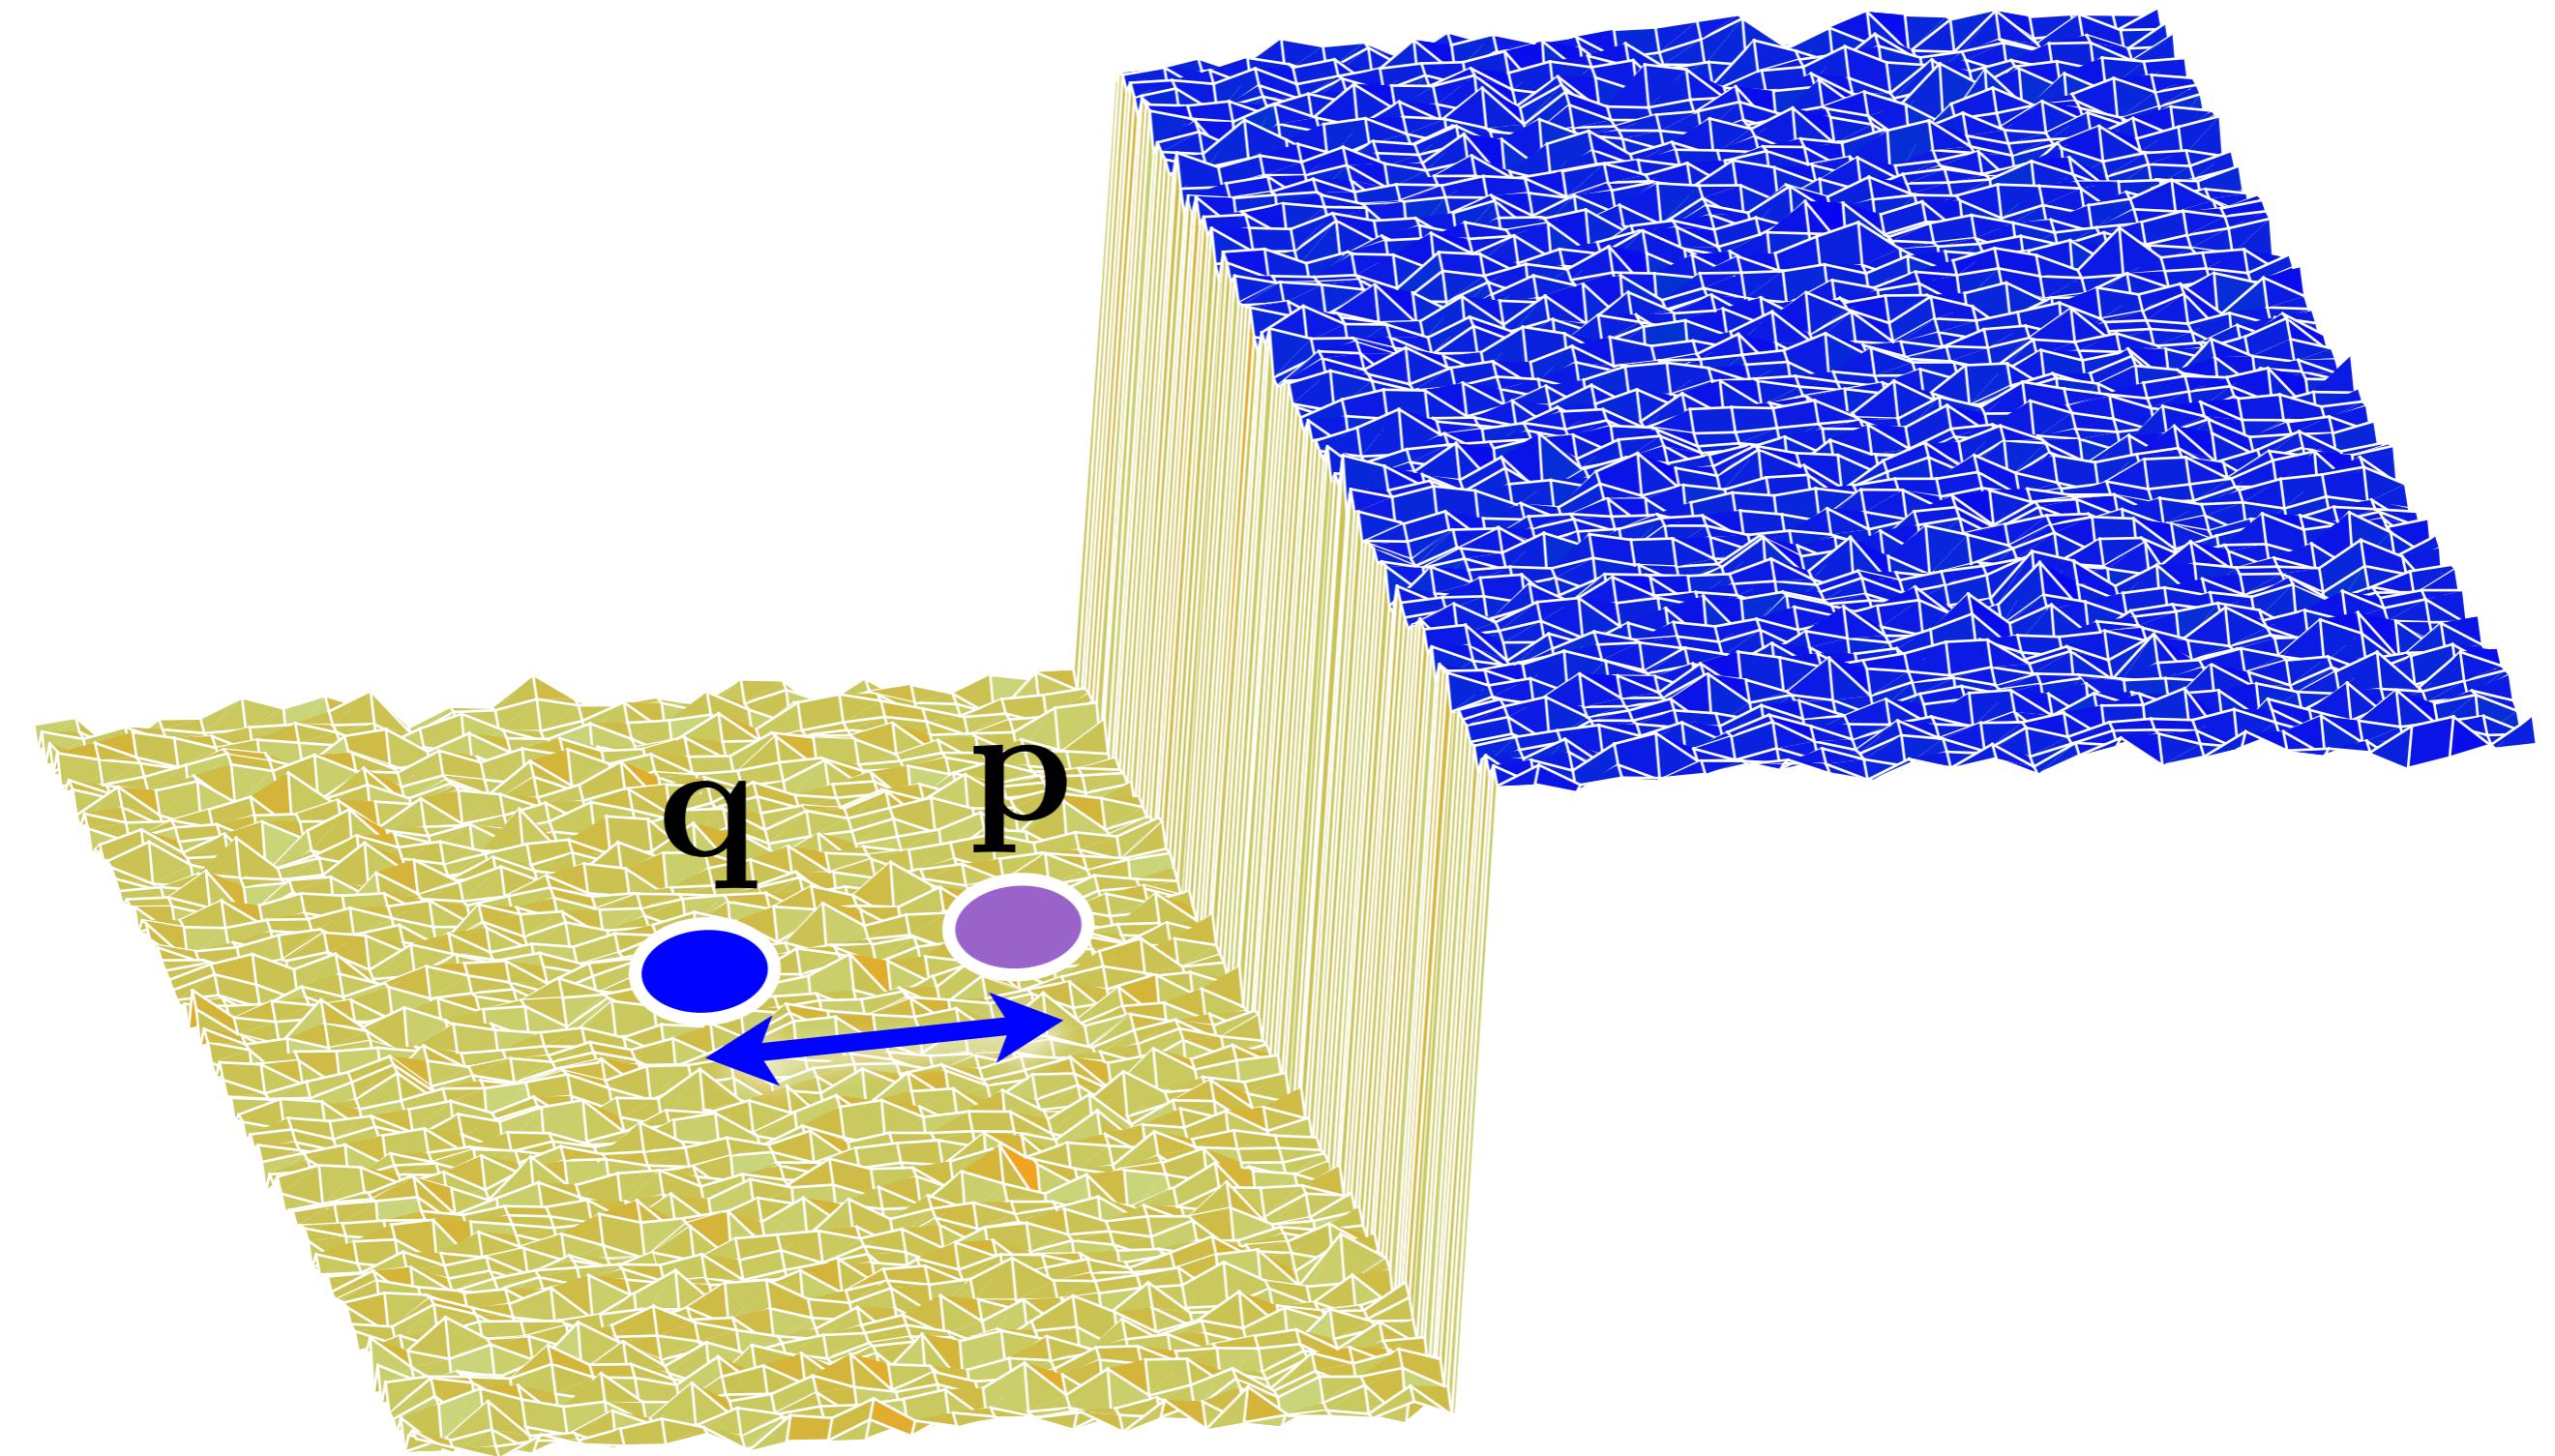
# Gaussian smoothing

$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$



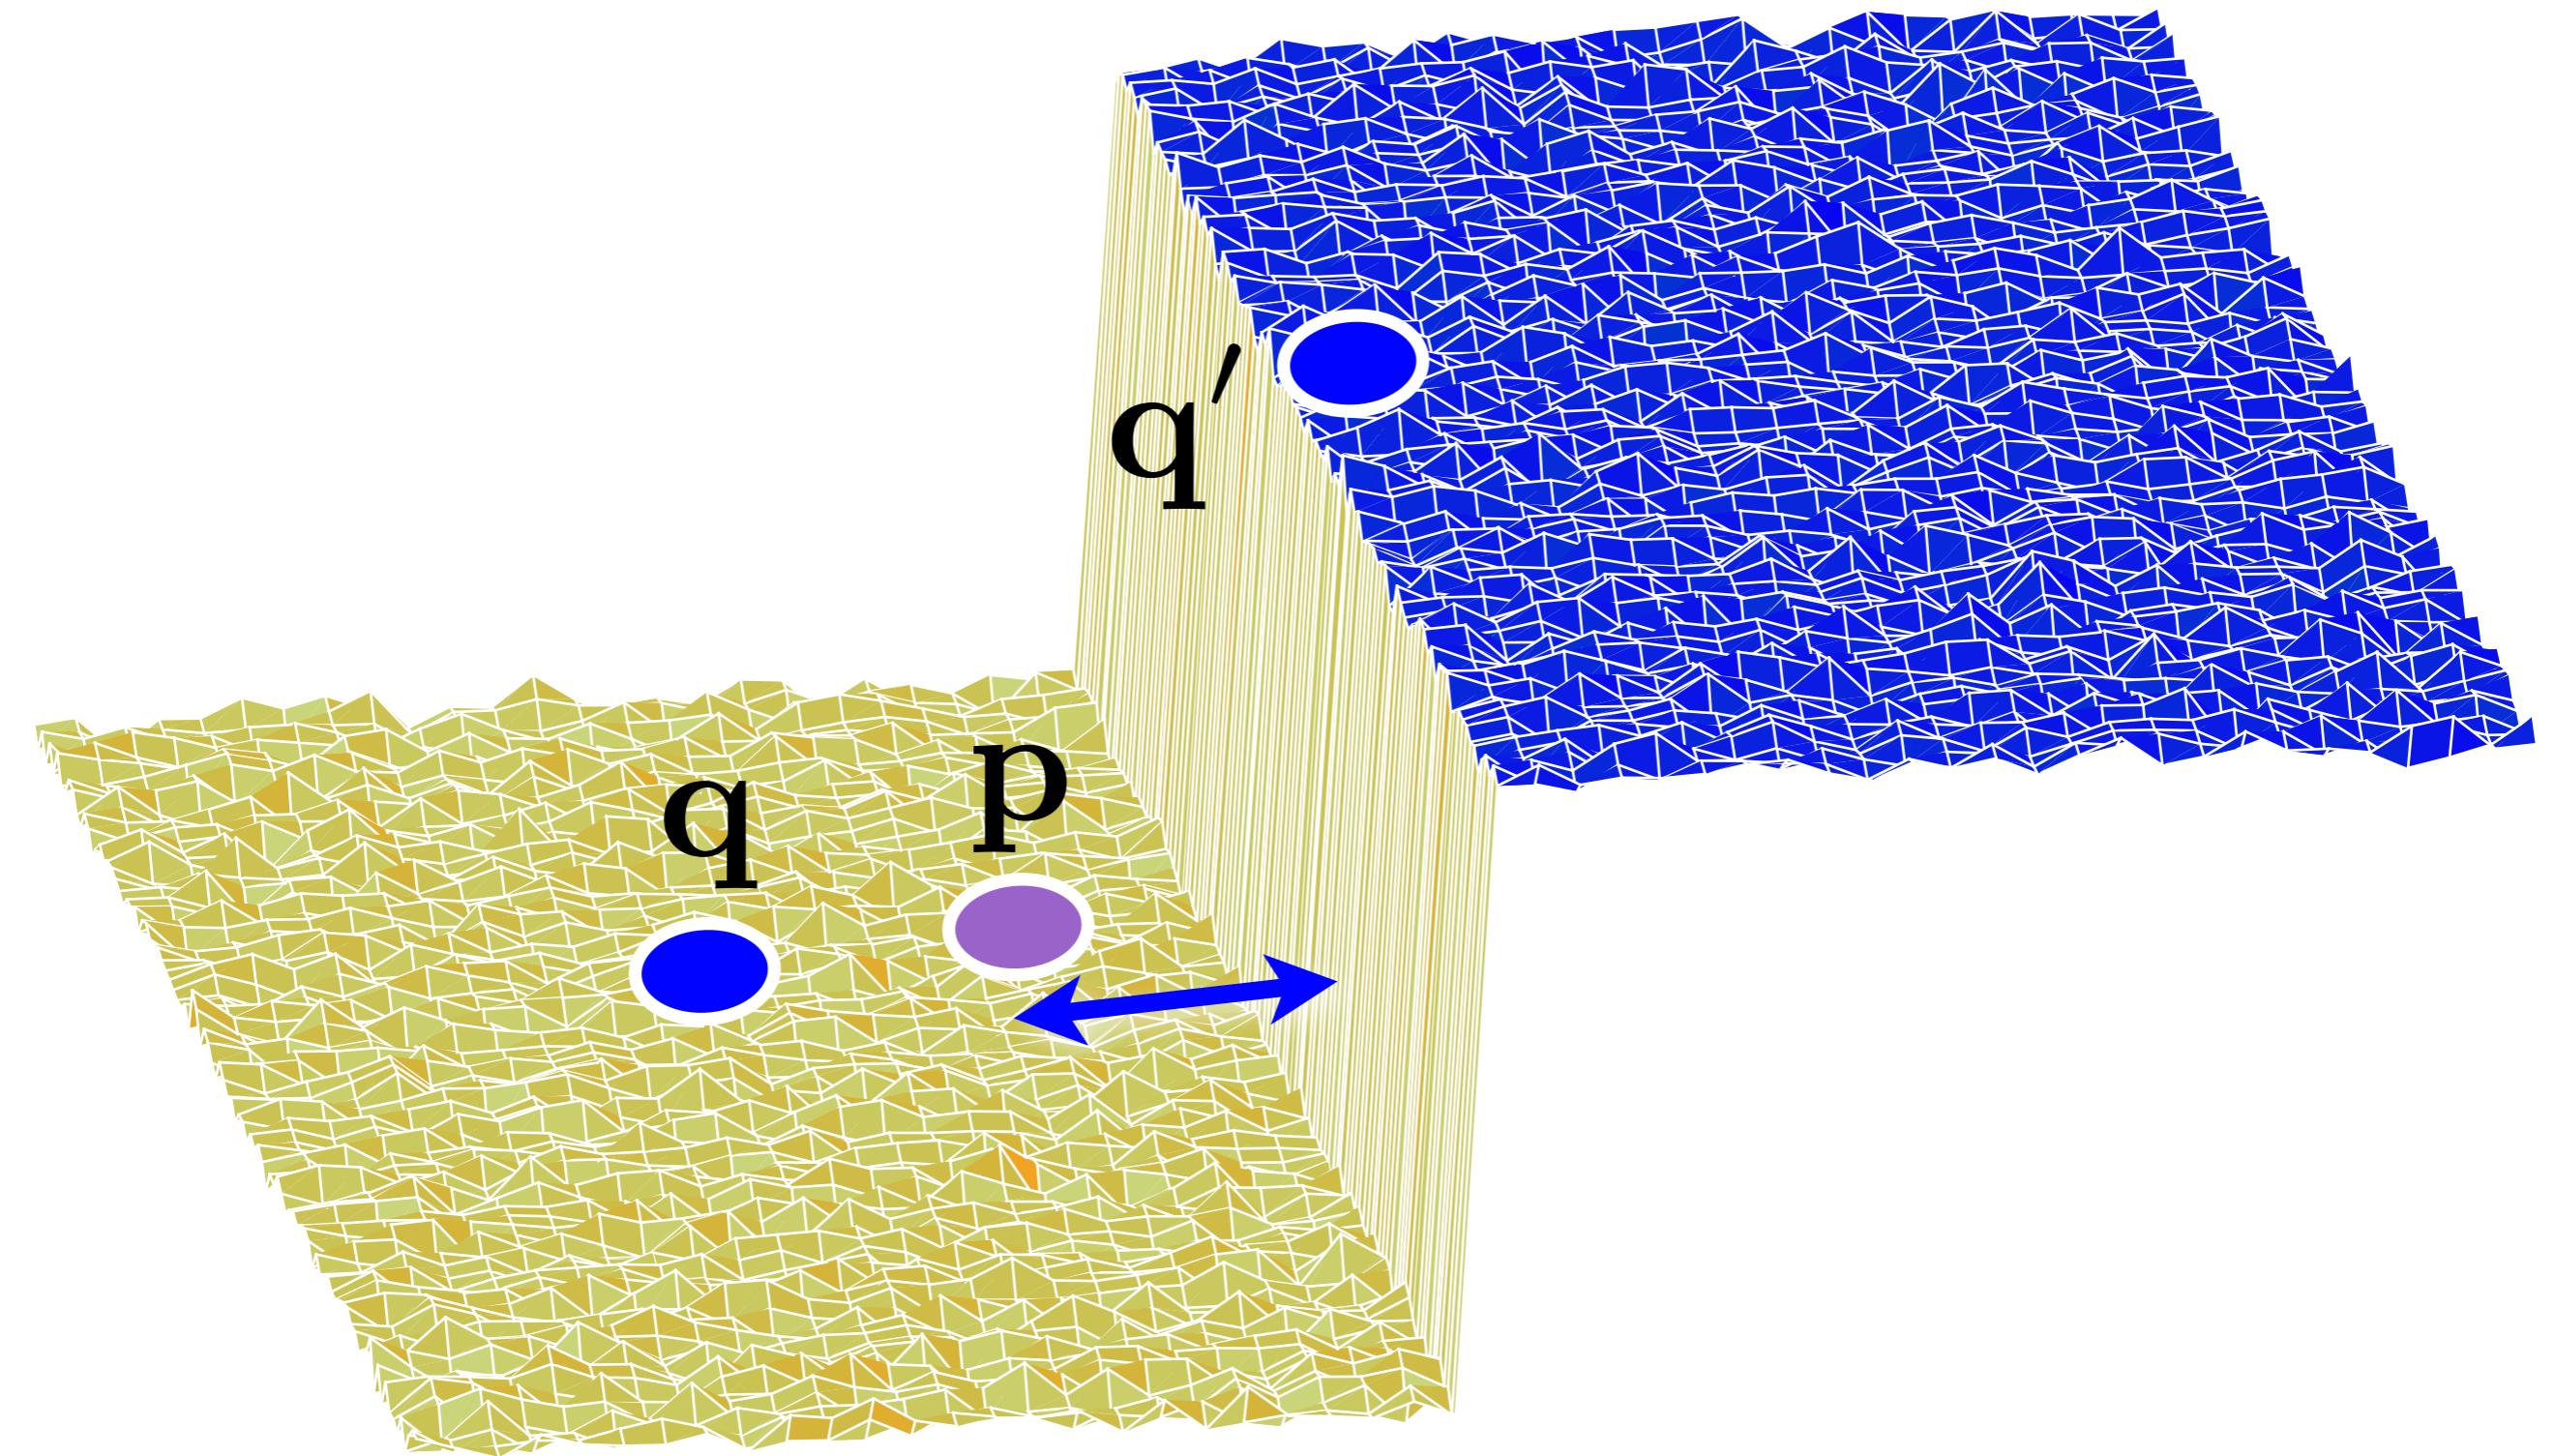
# spatial weight

$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$

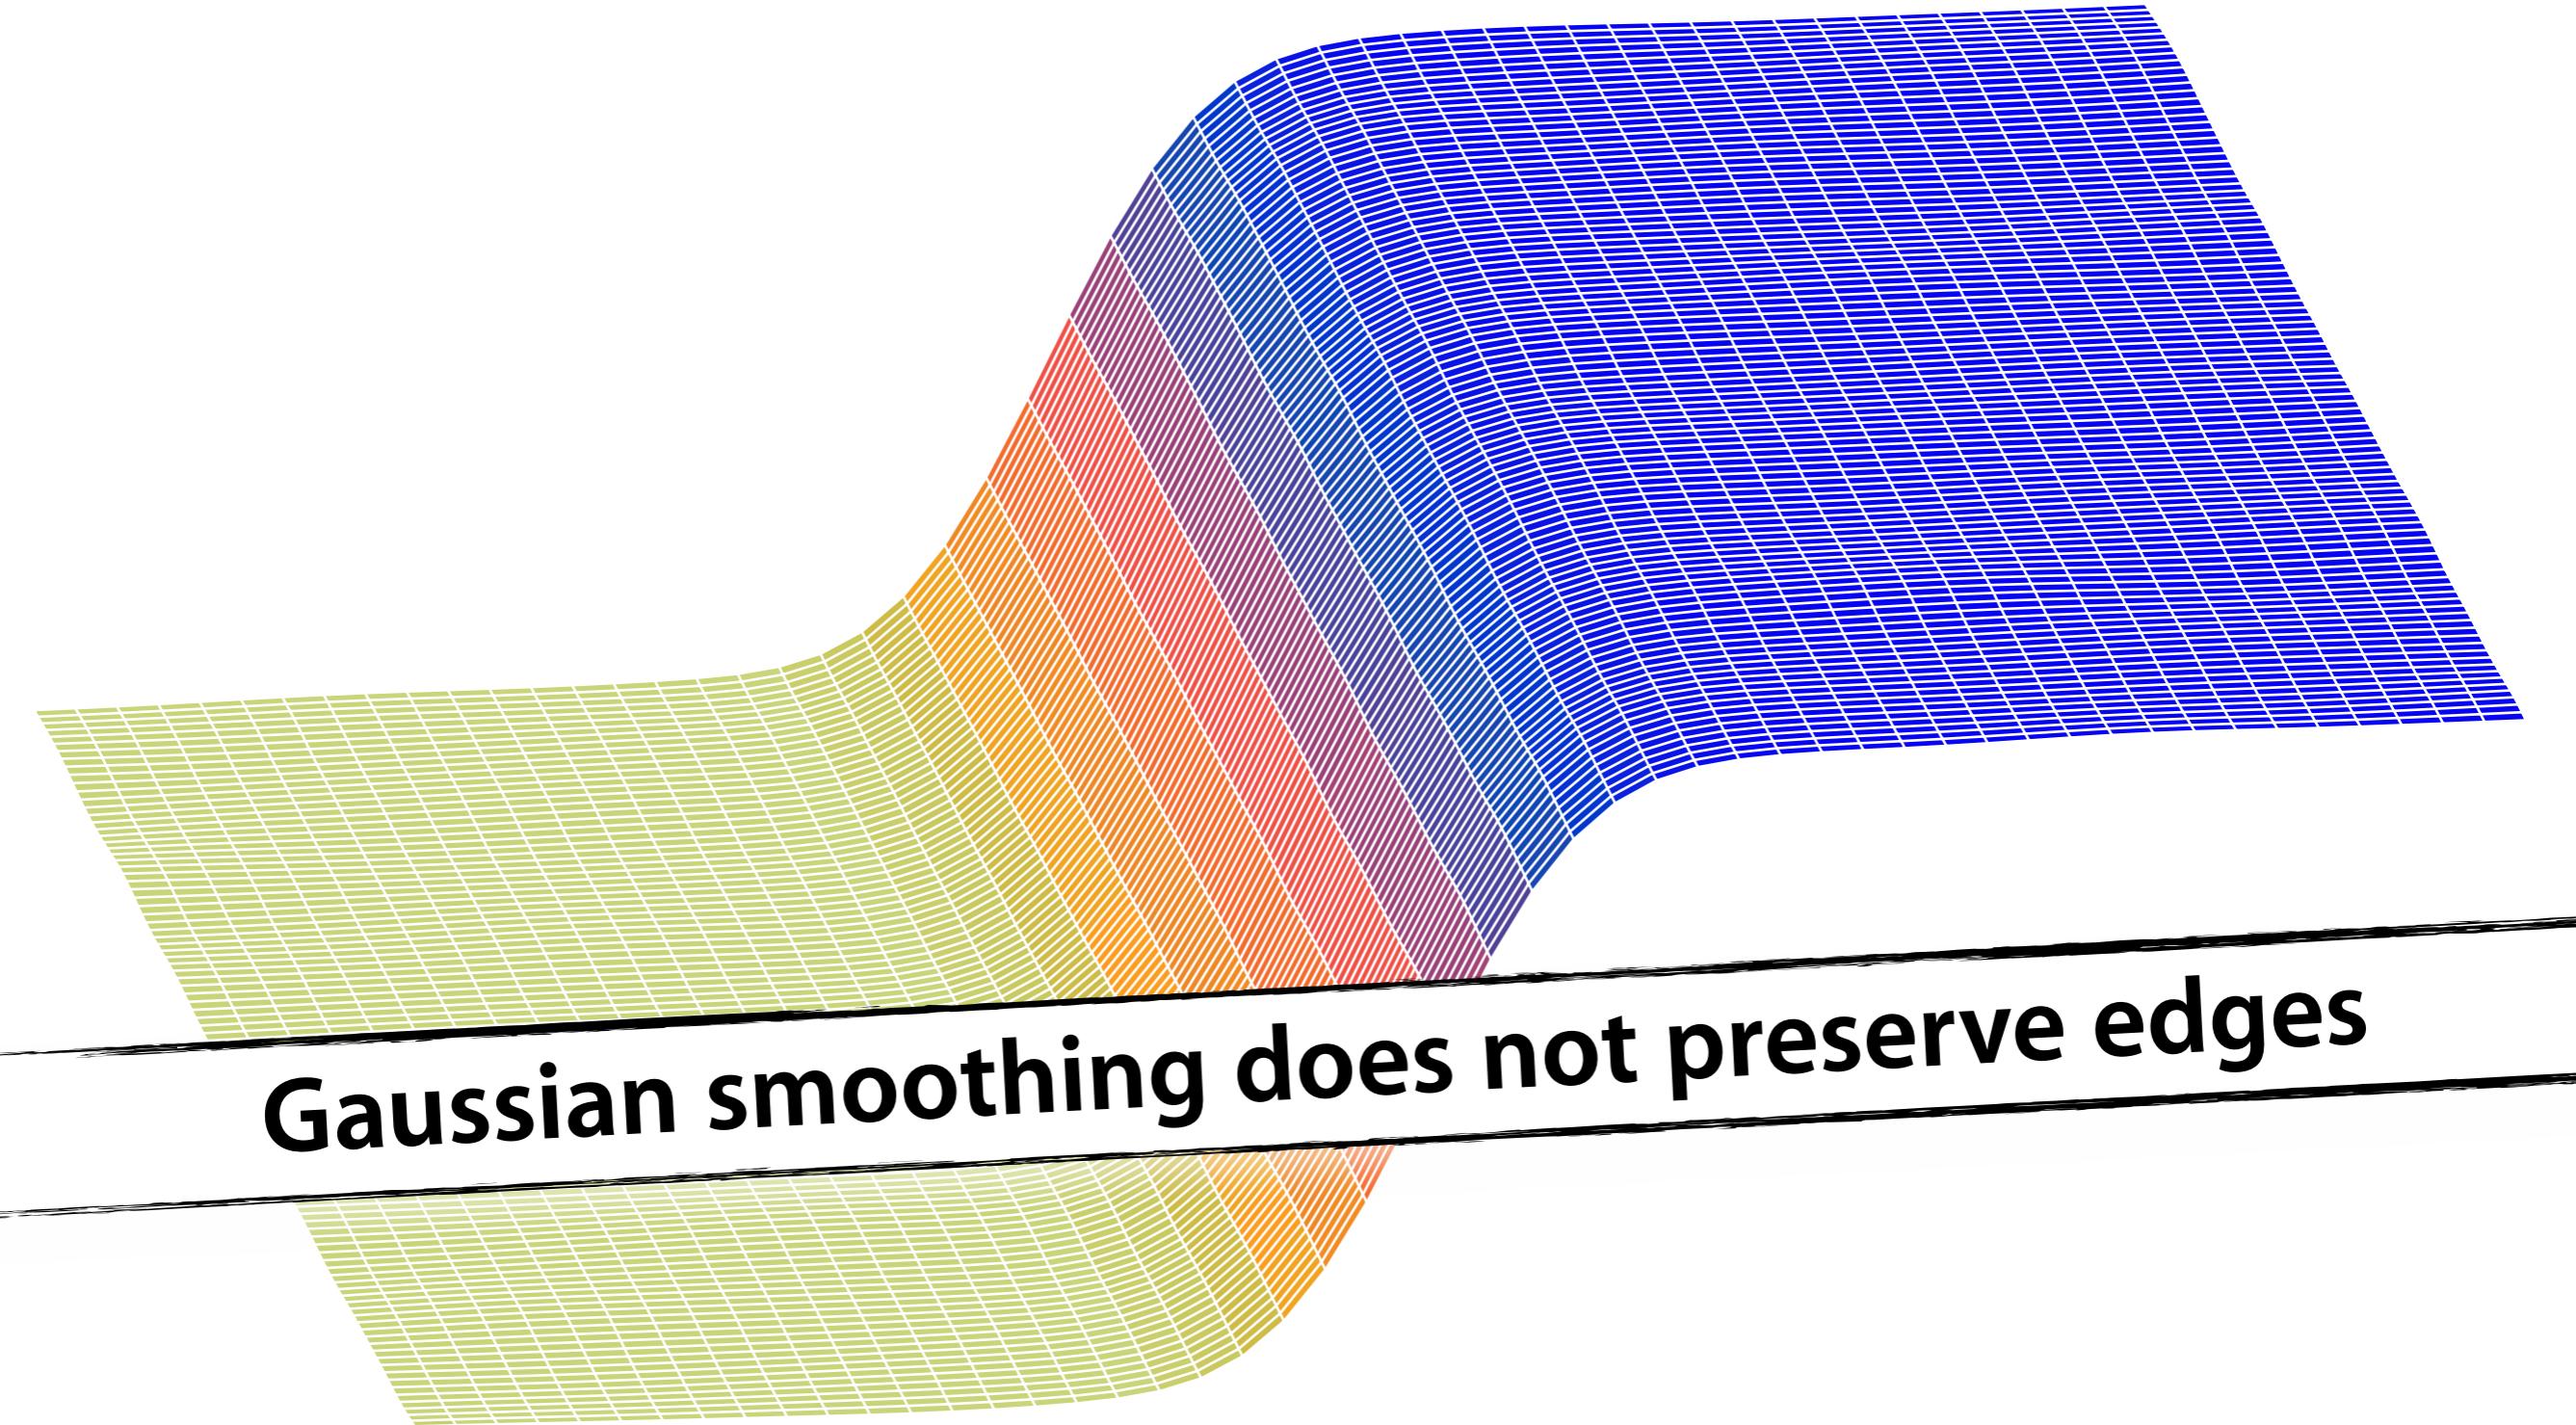


# spatial weight

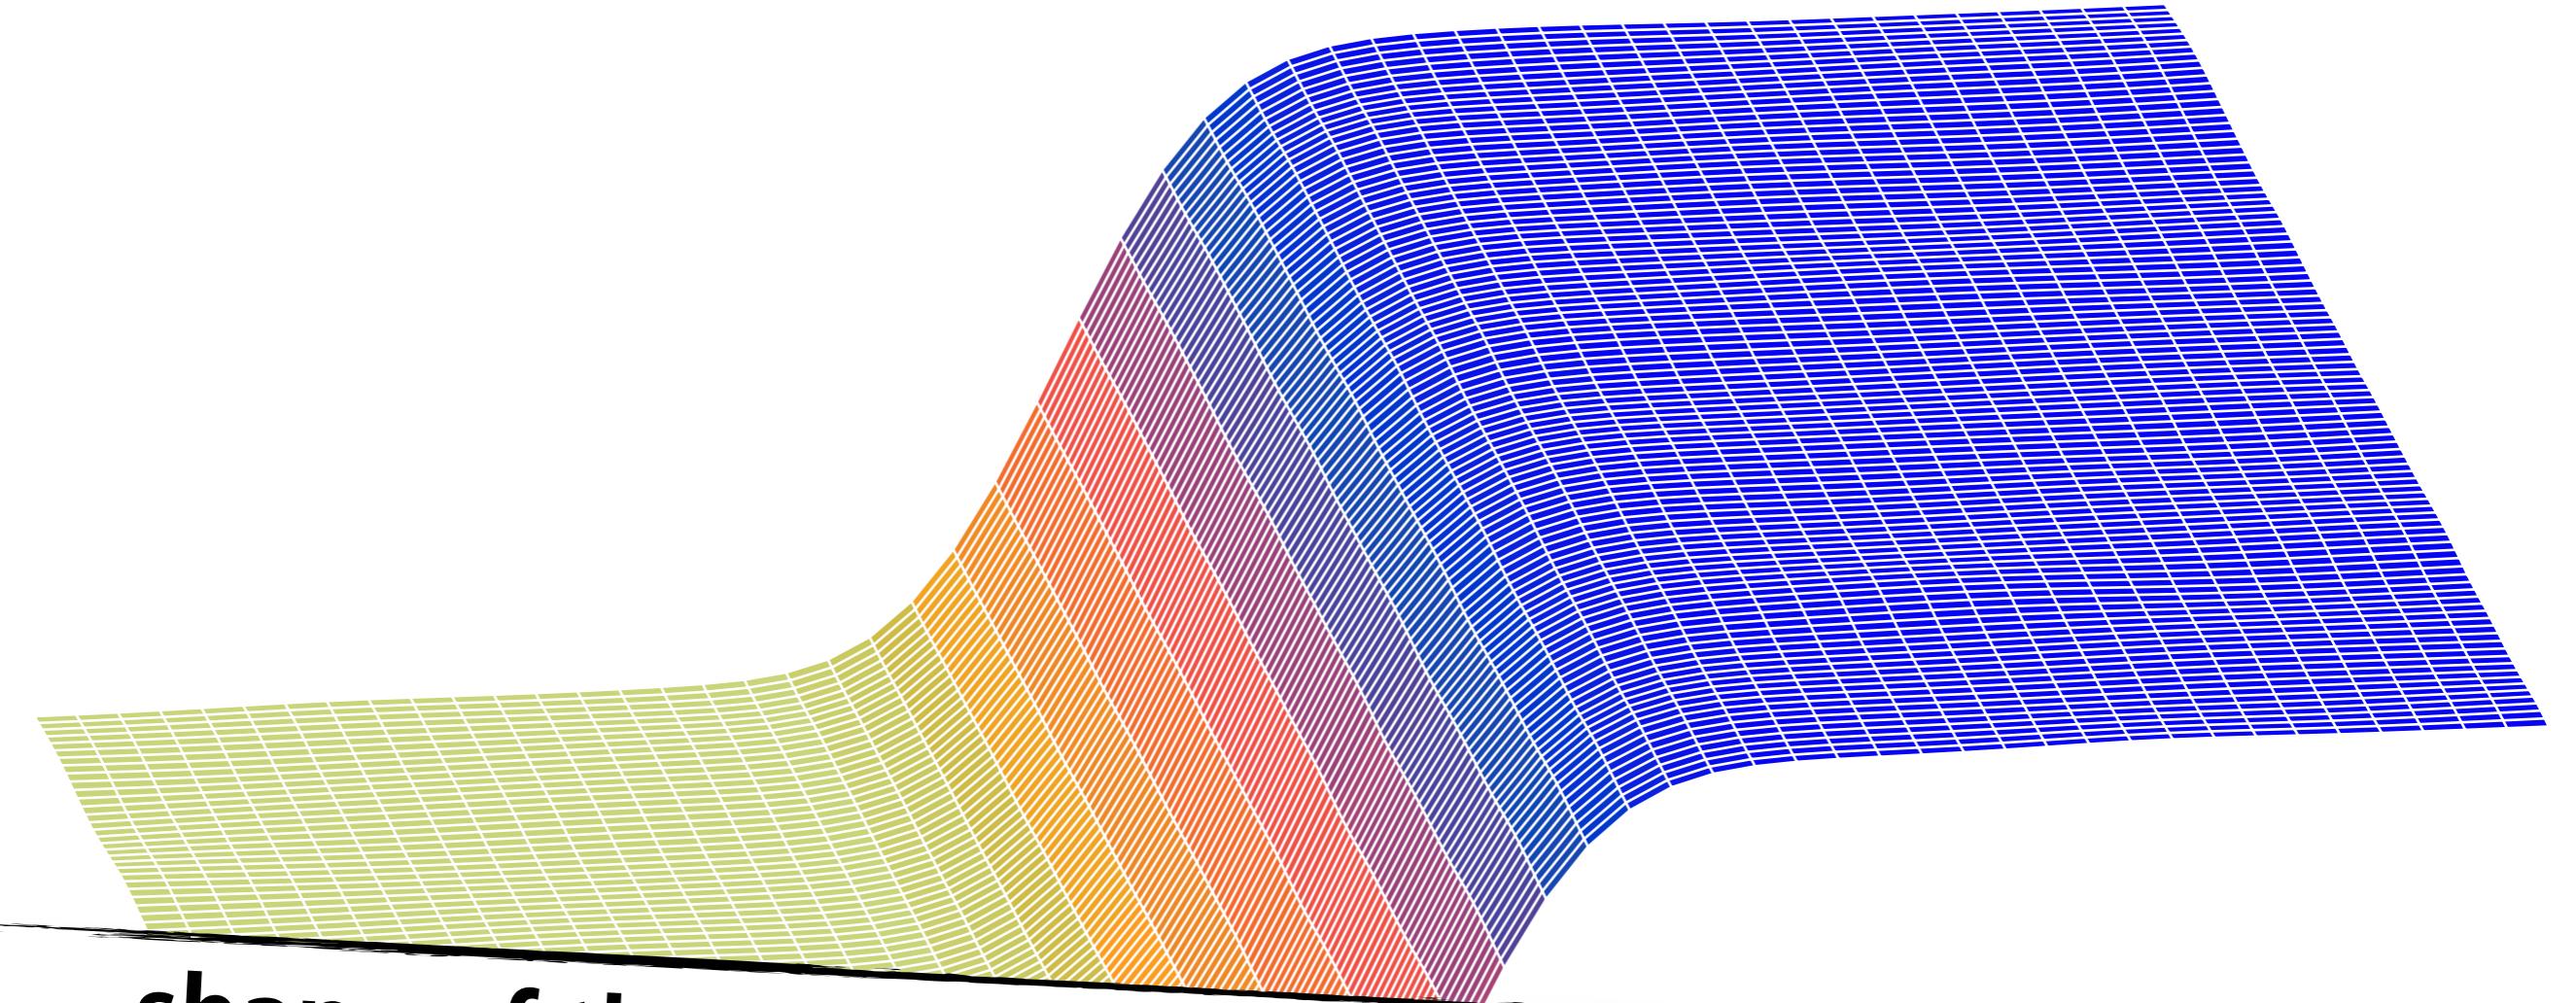
$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$



$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s} (\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$

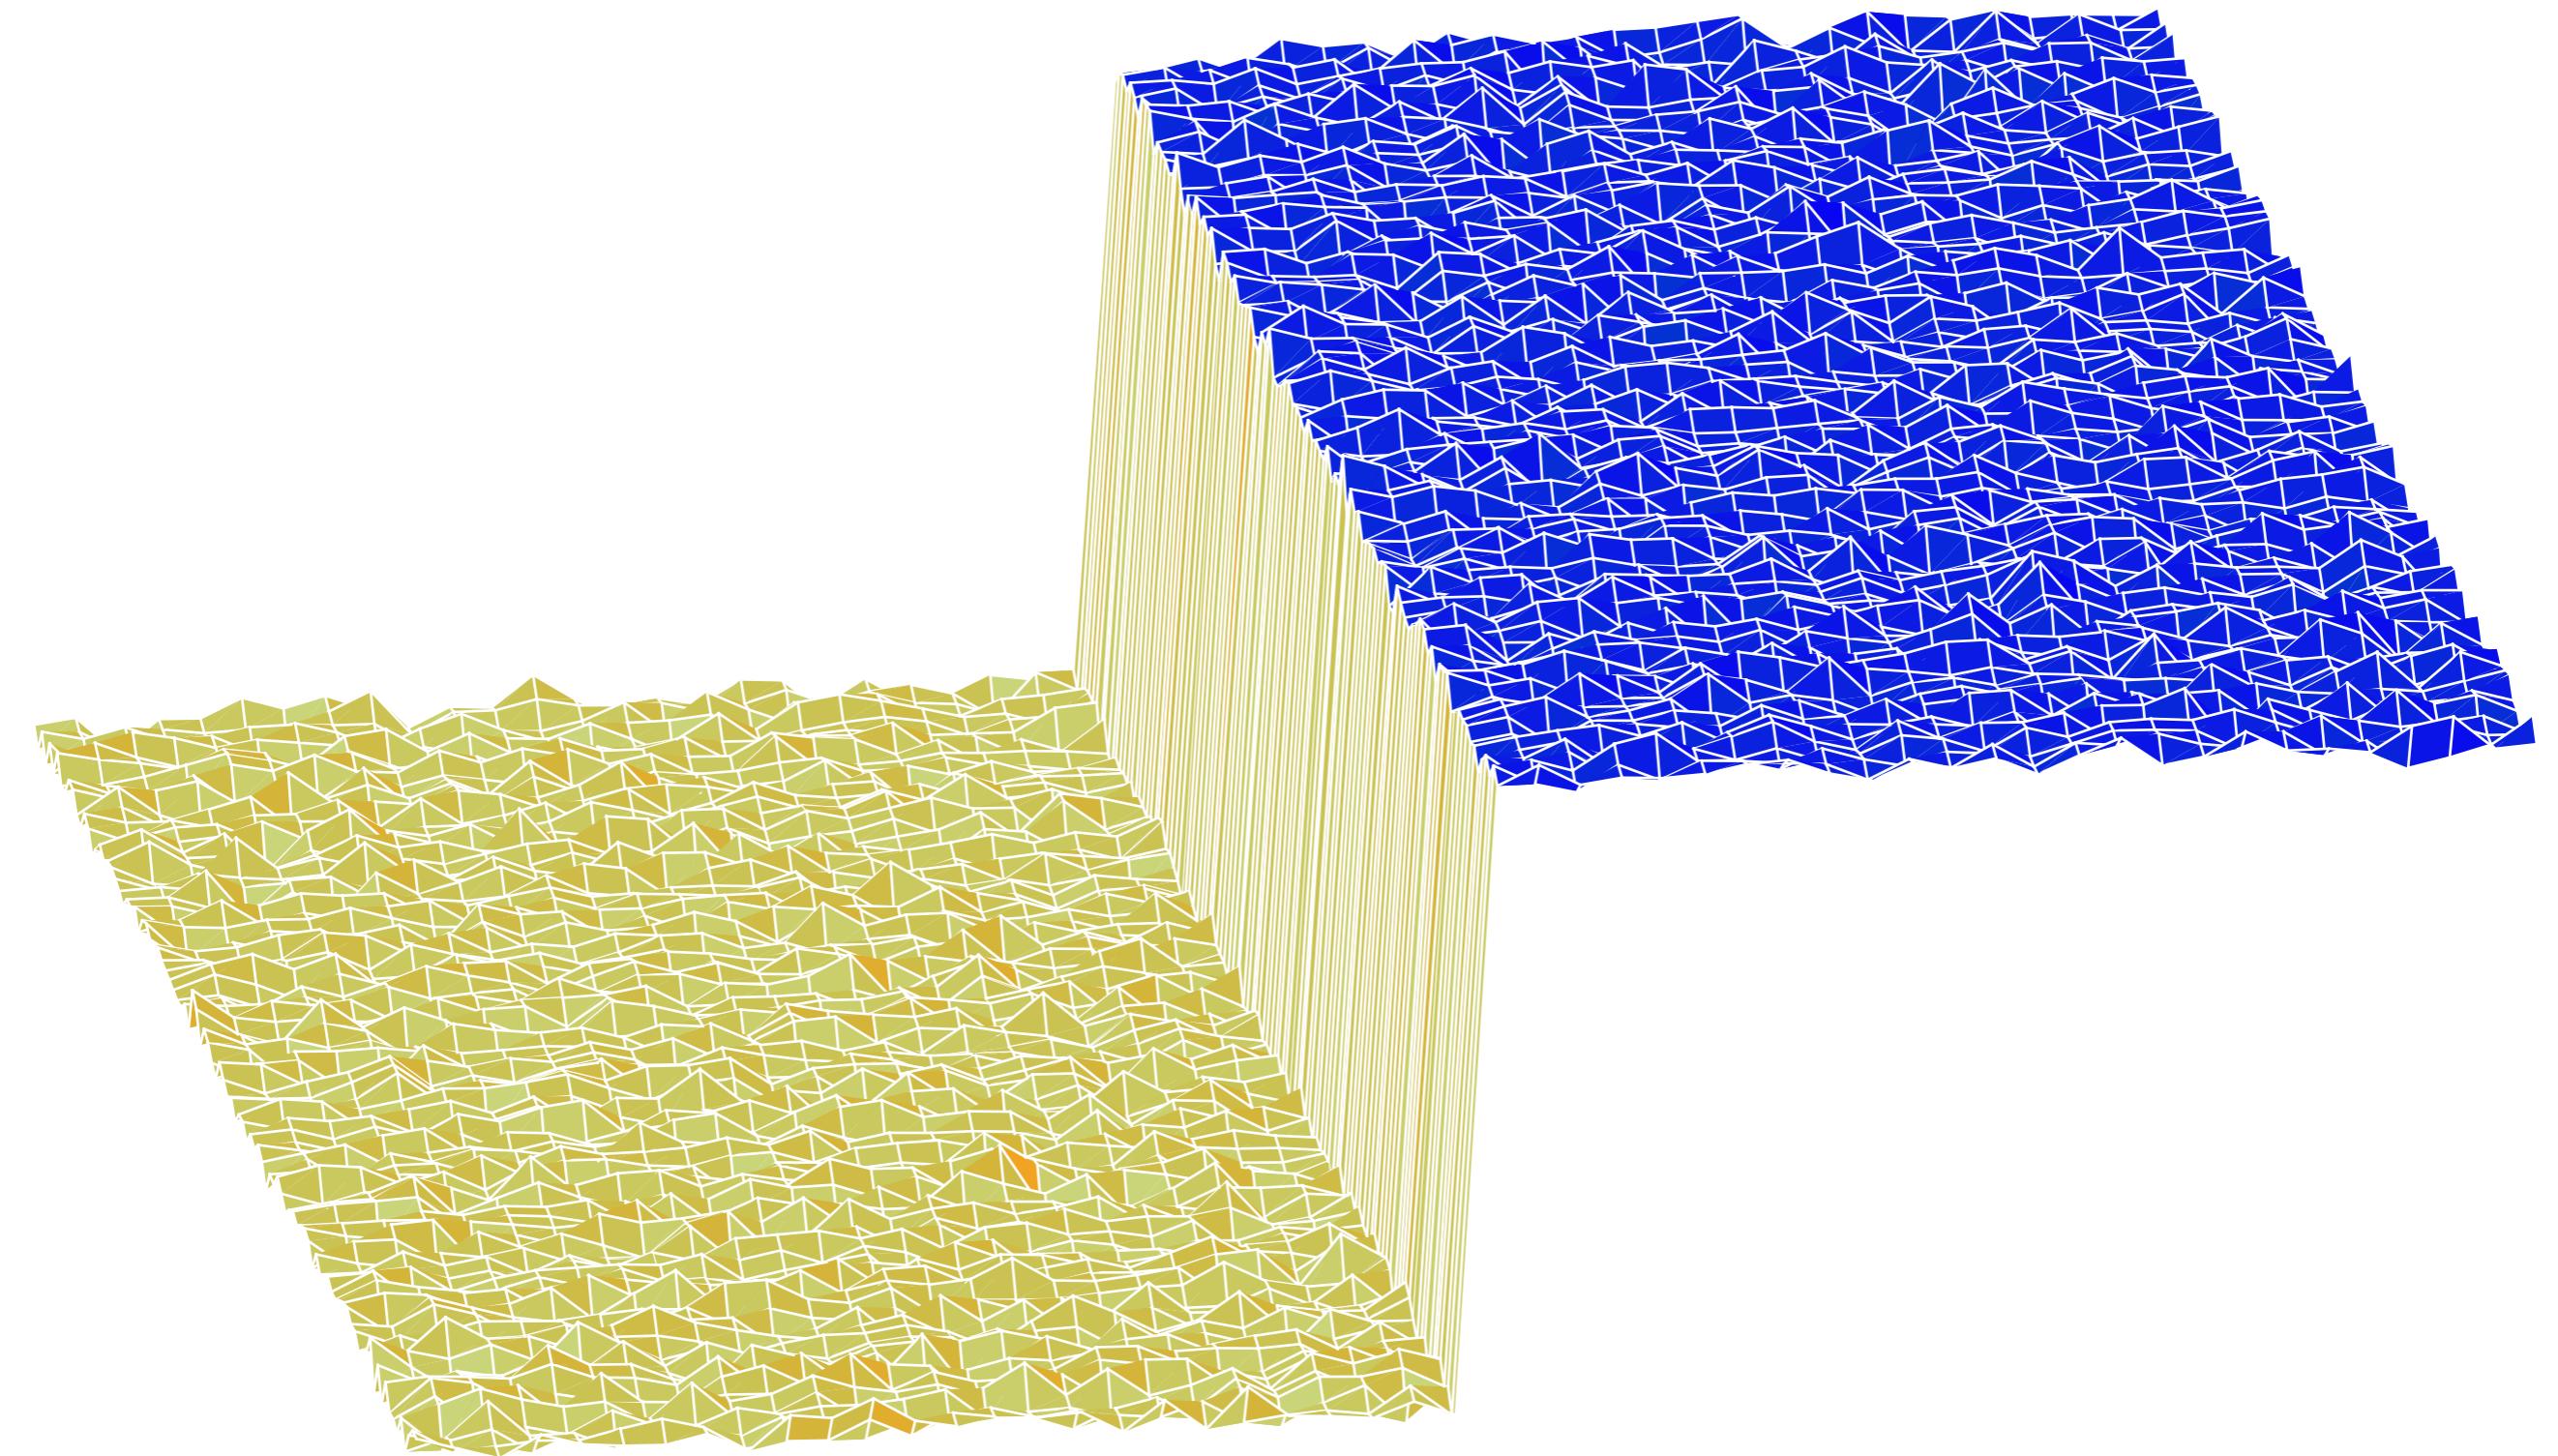


$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s} (\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$

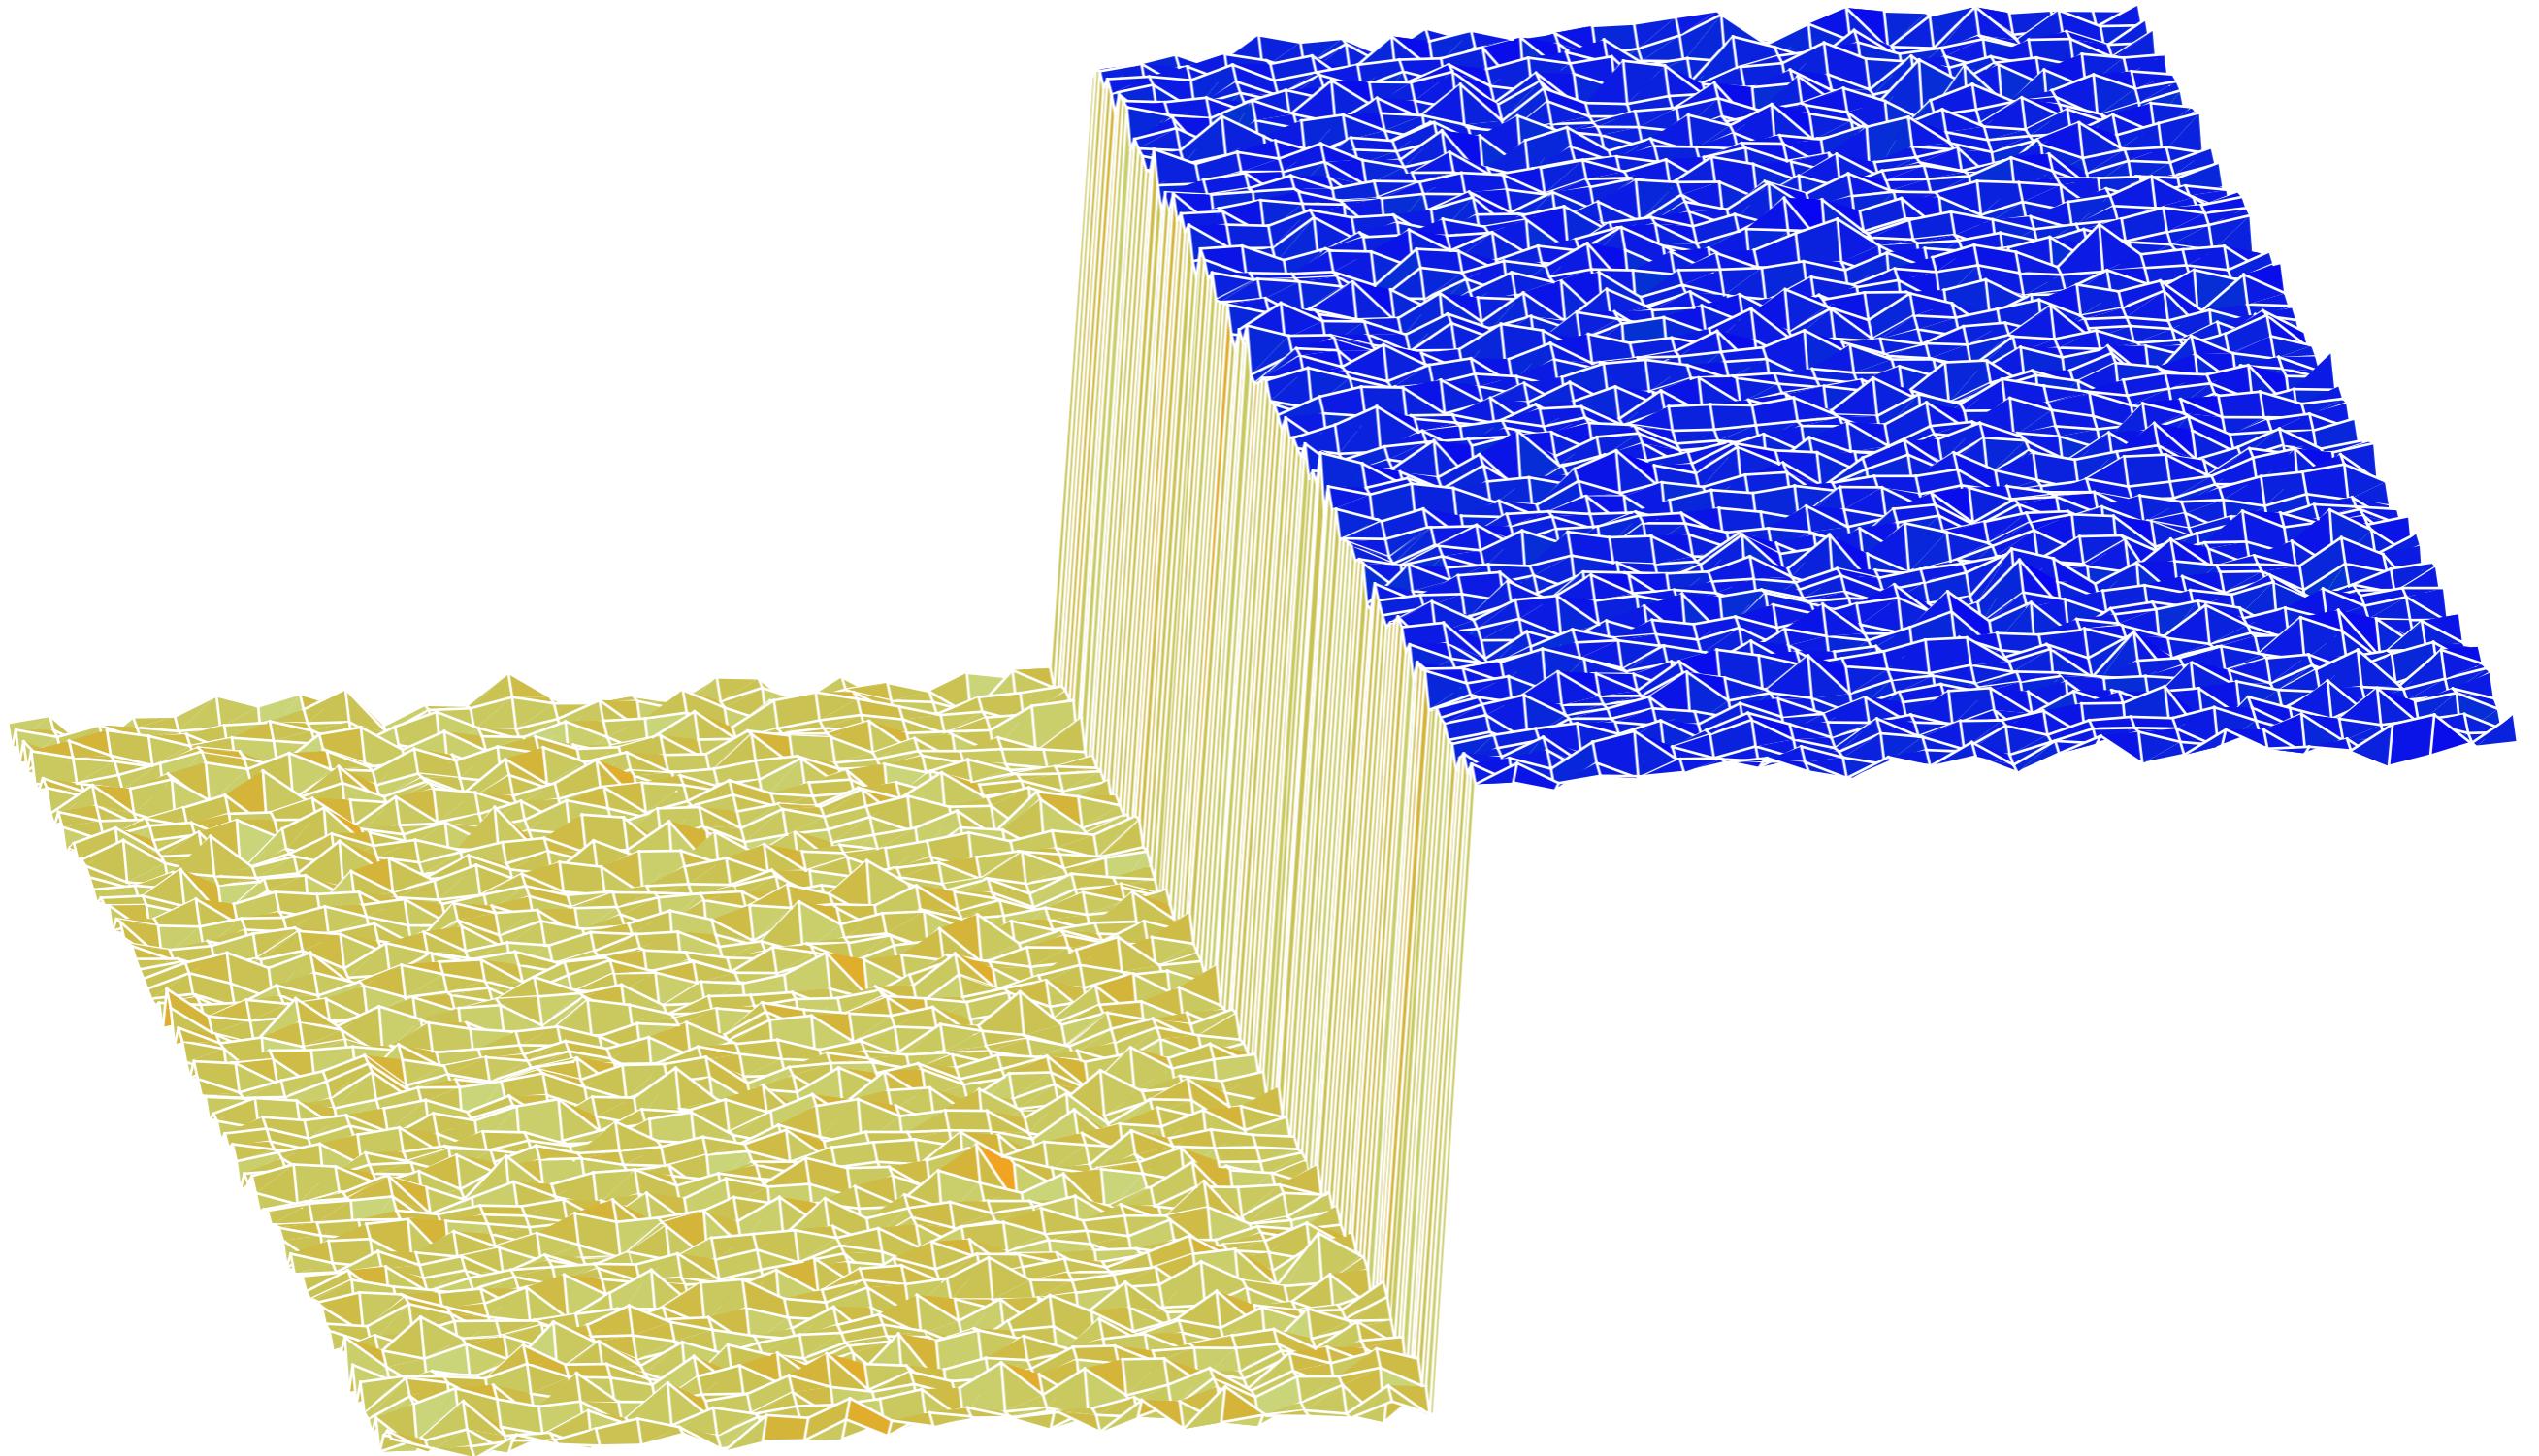


**shape of the spatial weight is the same everywhere**

$$I'(\mathbf{p}) = \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) I(\mathbf{q})$$

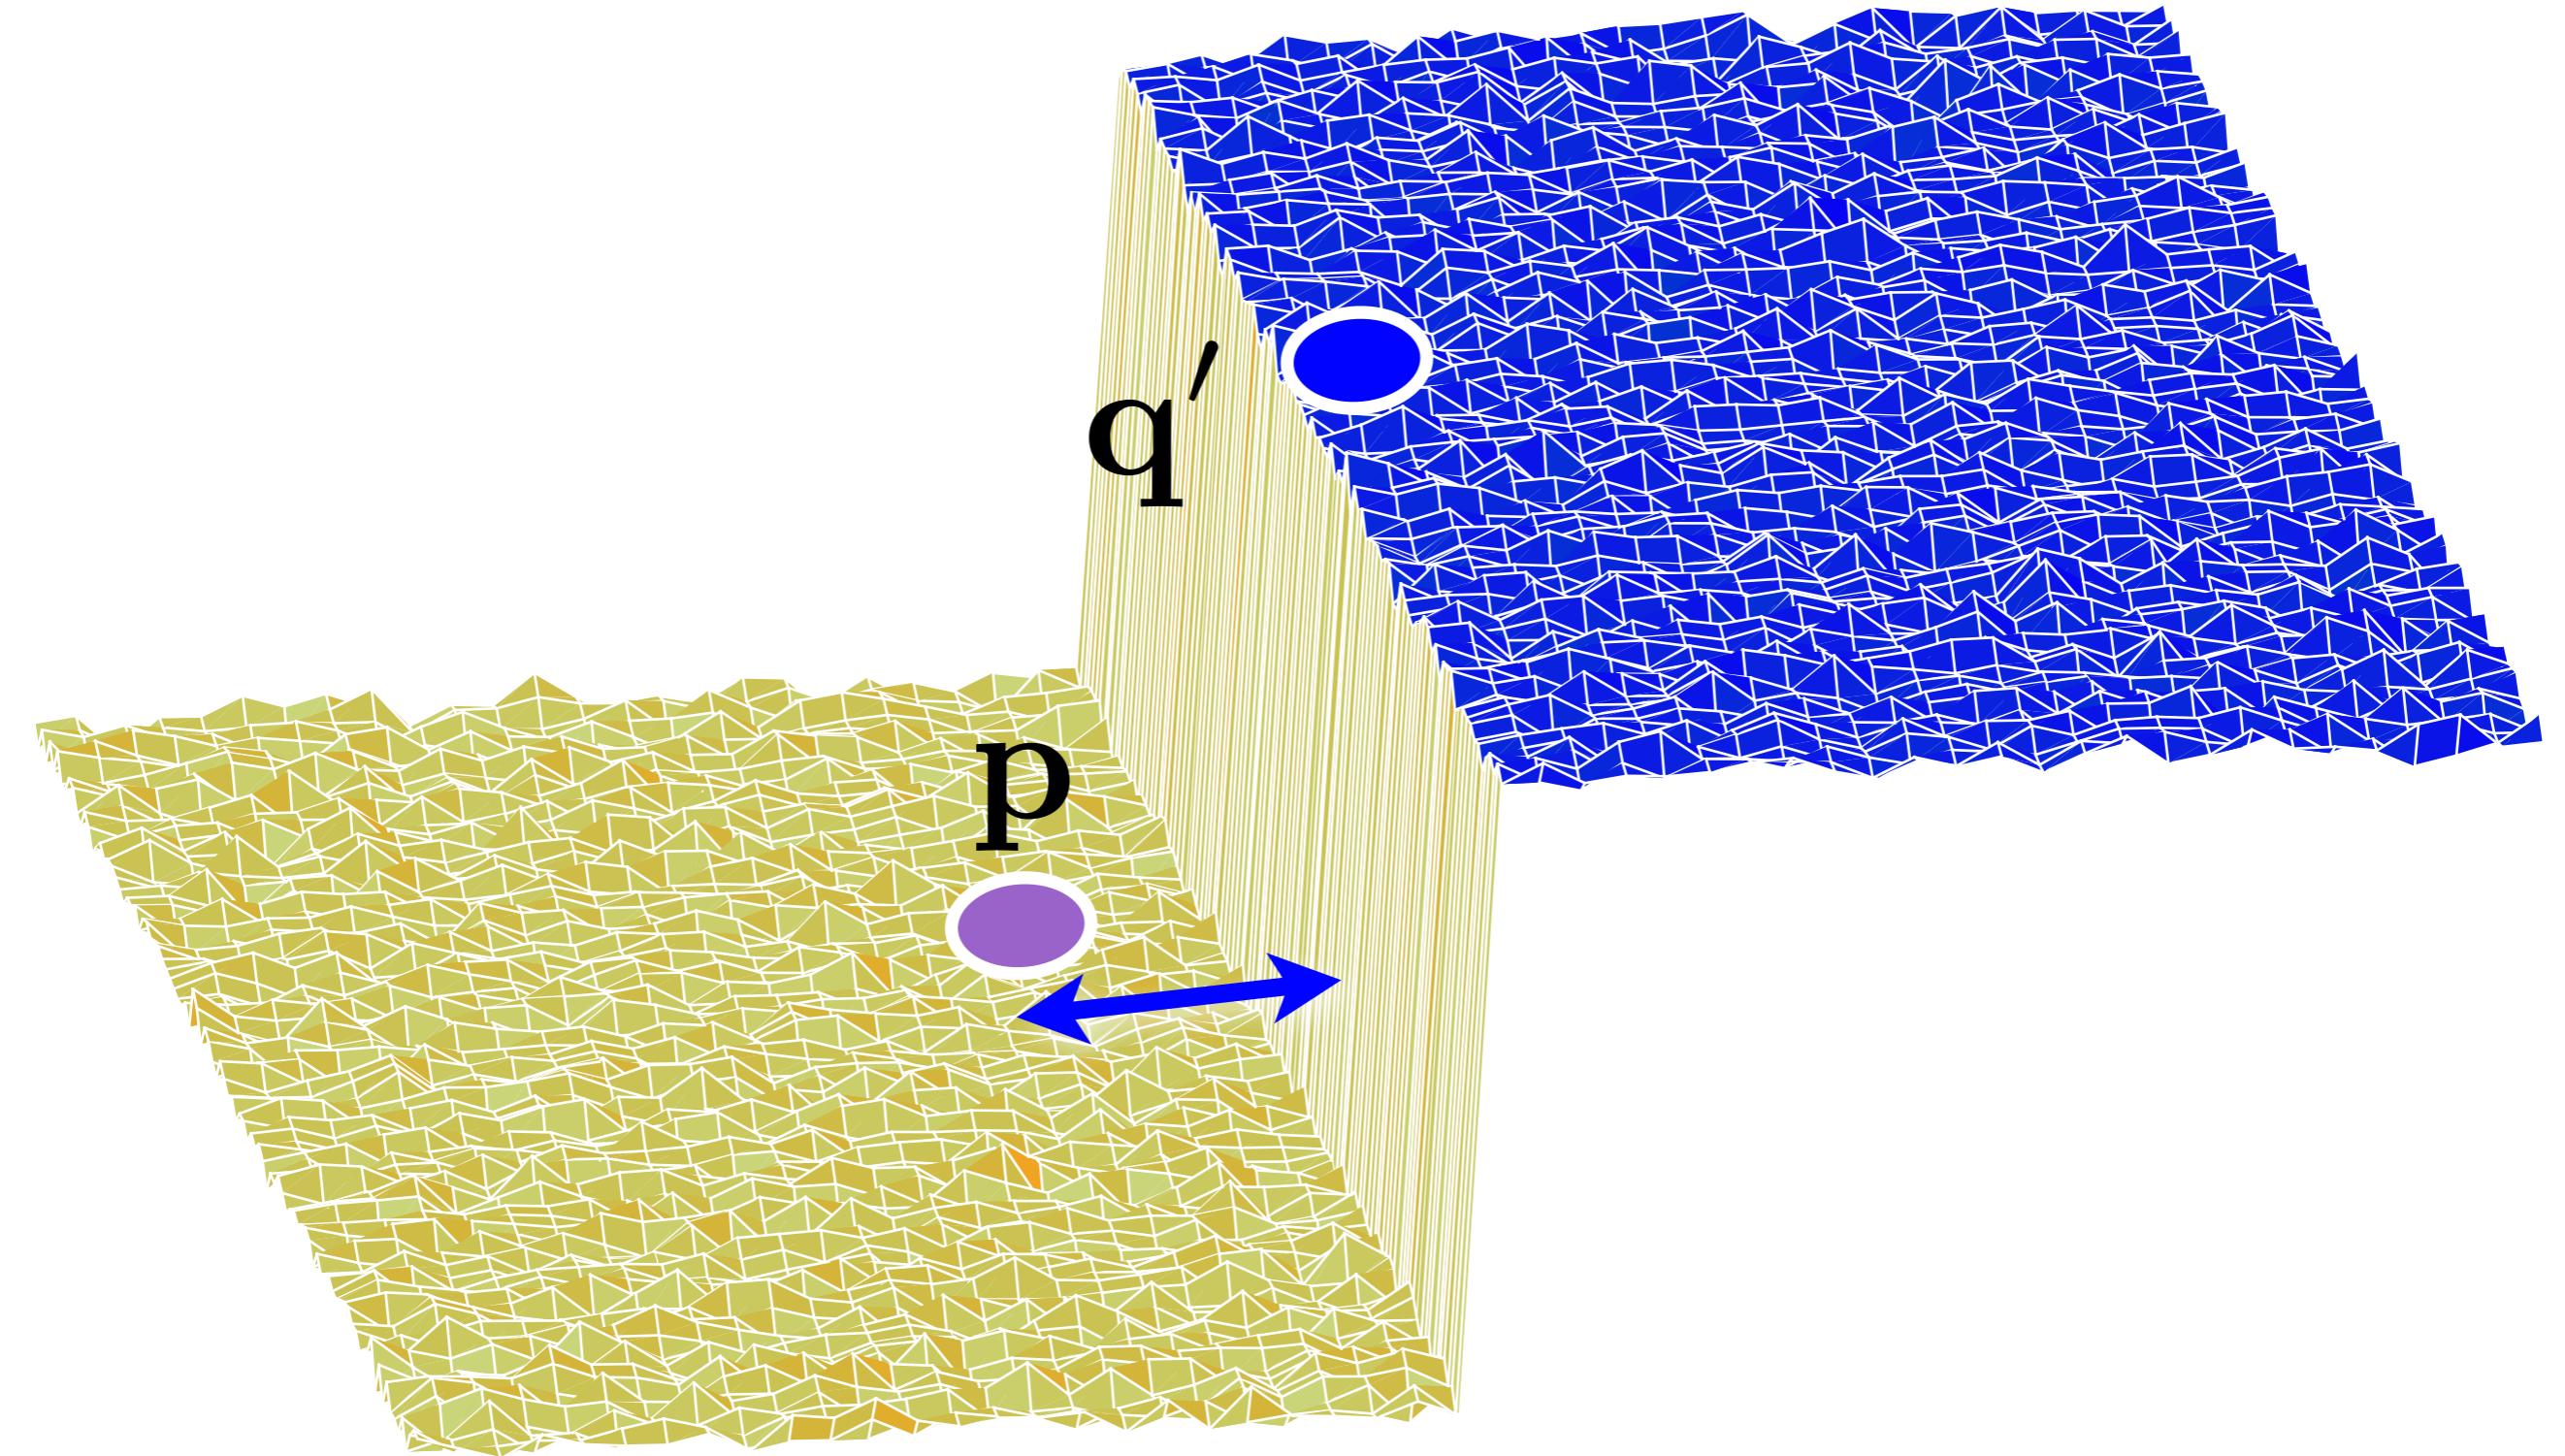


$$I'(\mathbf{p}) = w(\mathbf{p}) \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|) I(\mathbf{q})$$



# spatial weight

$$I'(\mathbf{p}) = w(\mathbf{p}) \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|) I(\mathbf{q})$$

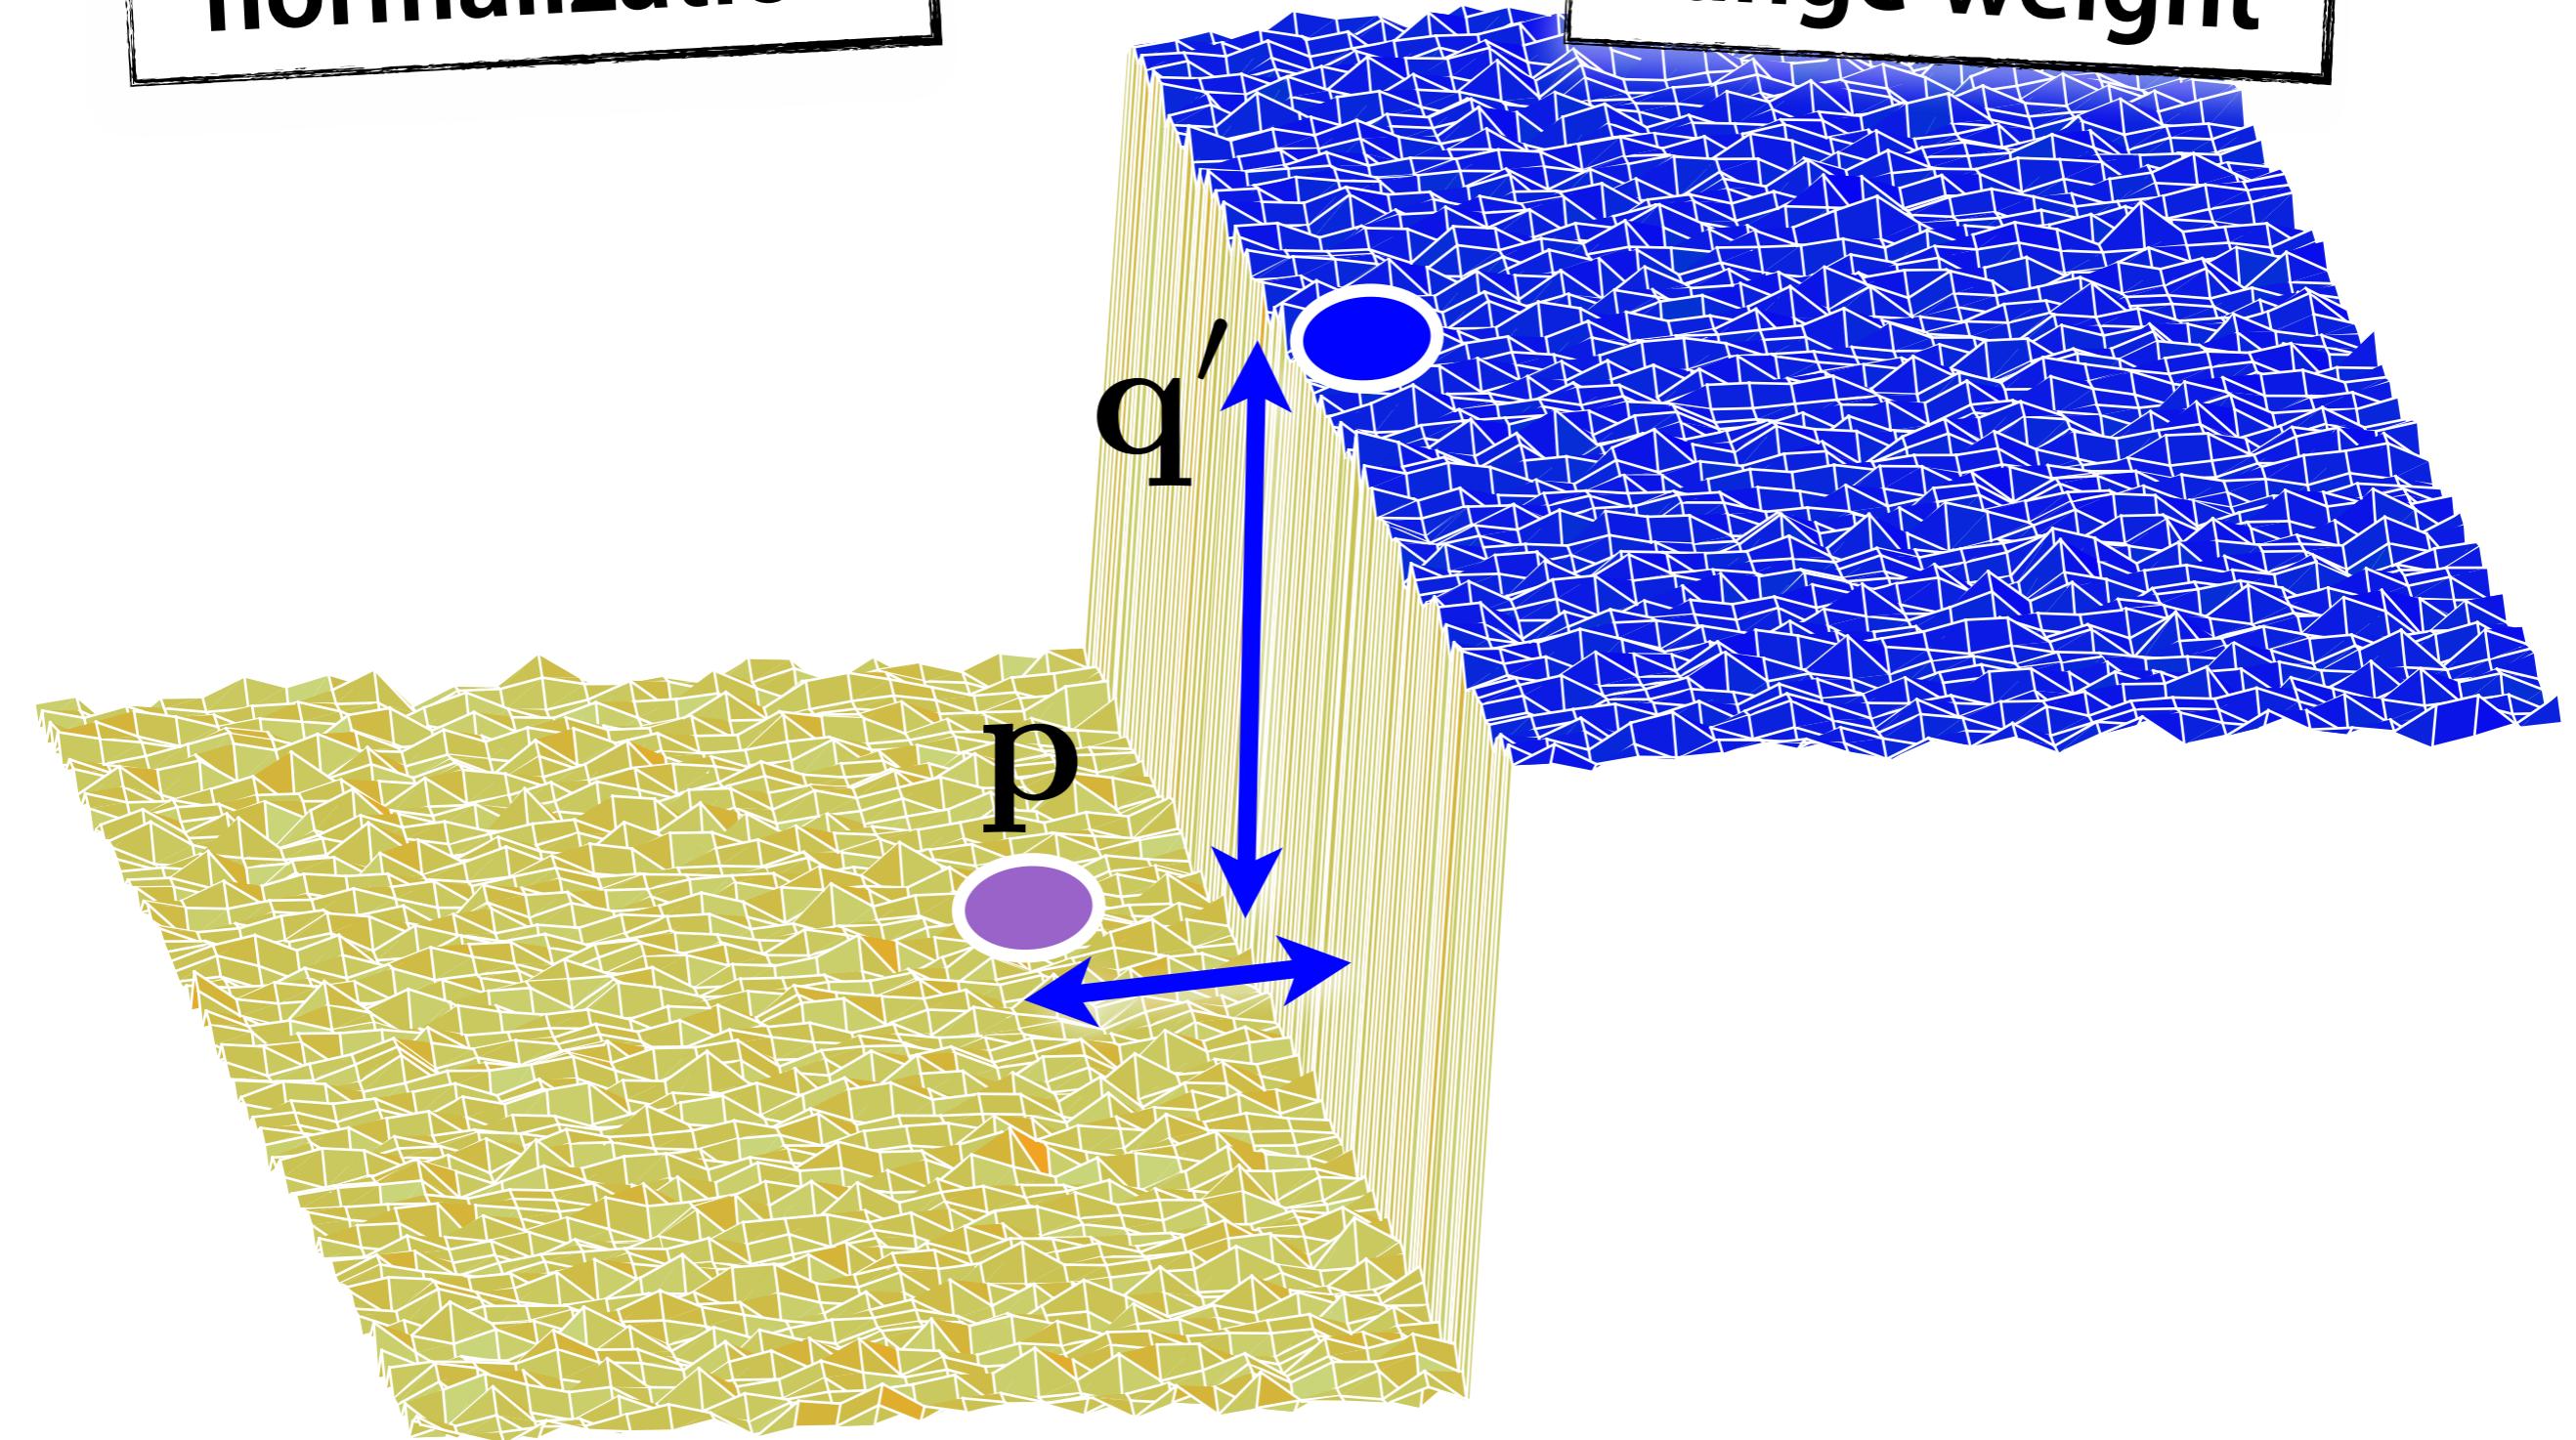


**spatial weight**

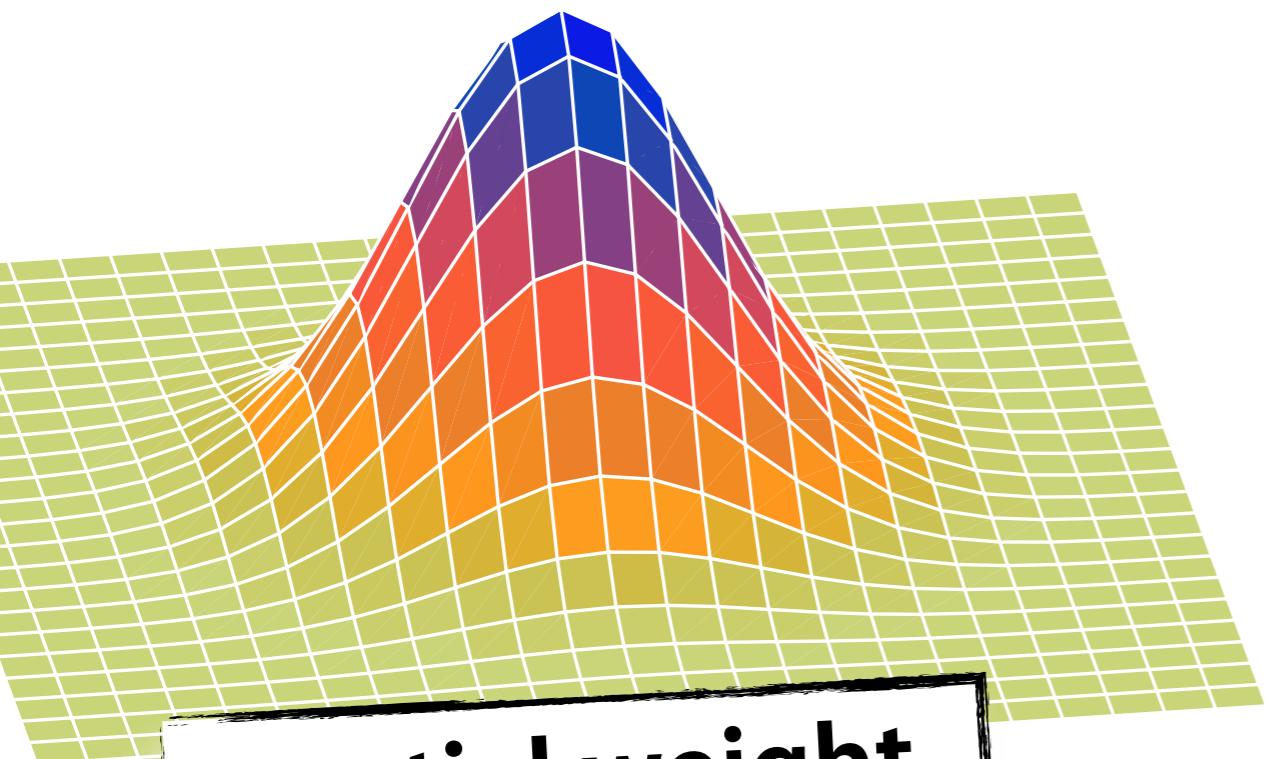
$$I'(\mathbf{p}) = w(\mathbf{p}) \sum G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|) I(\mathbf{q})$$

**normalization**

**range weight**

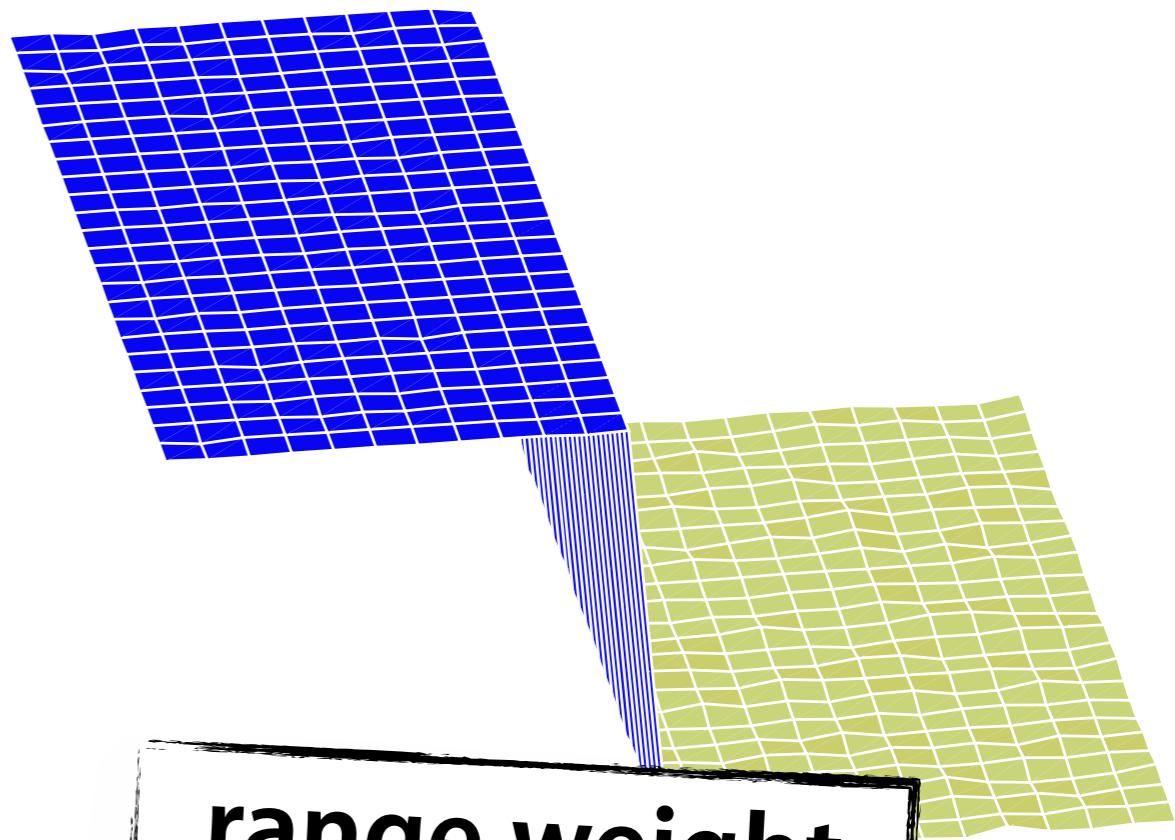


$$G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|)$$



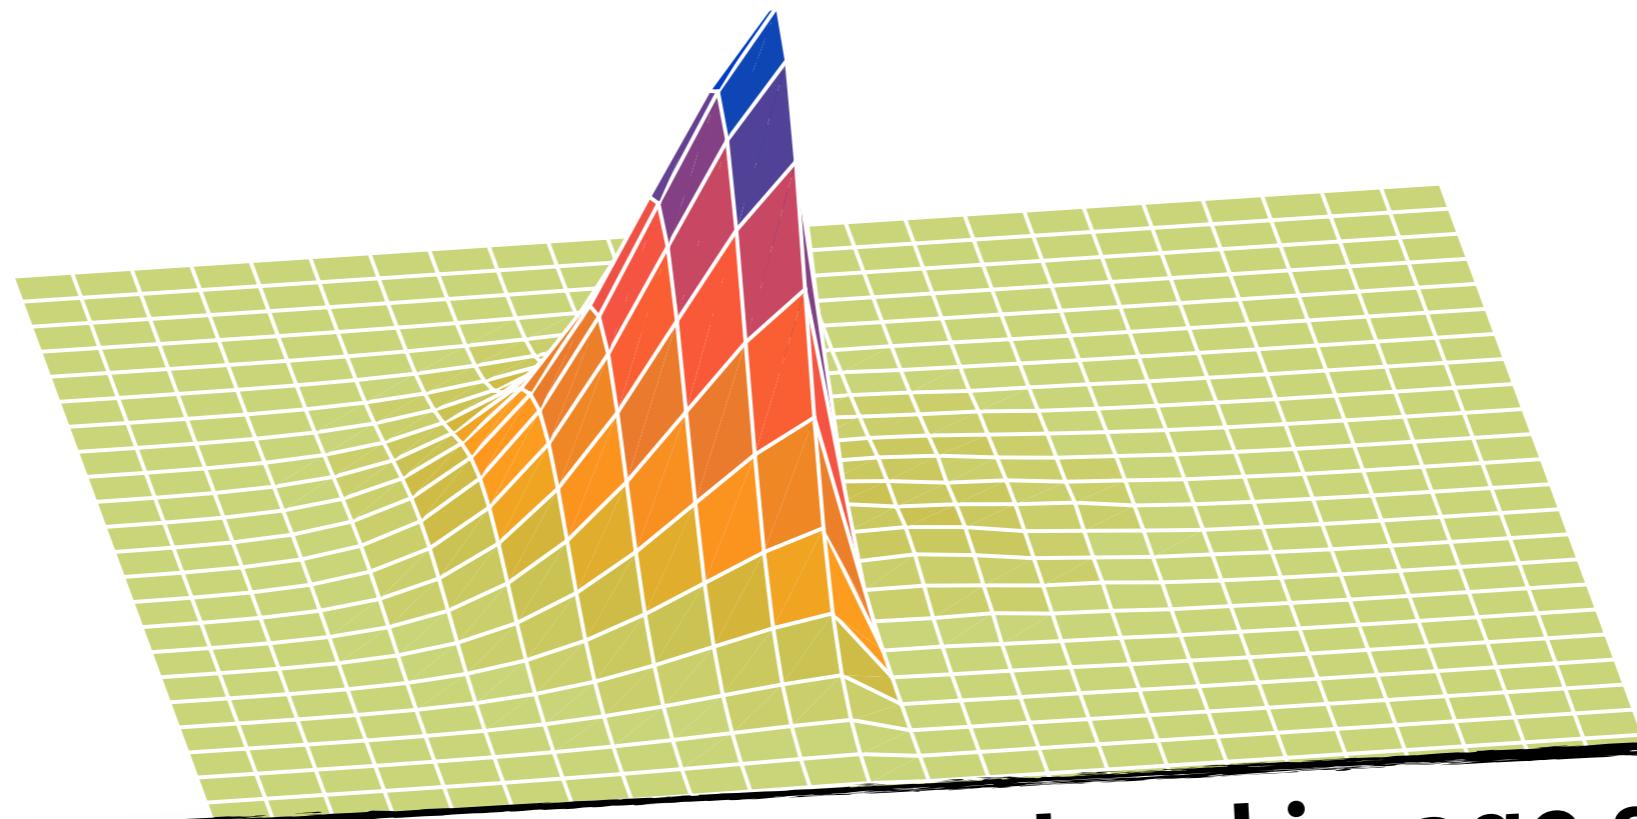
**spatial weight**

$$G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|)$$



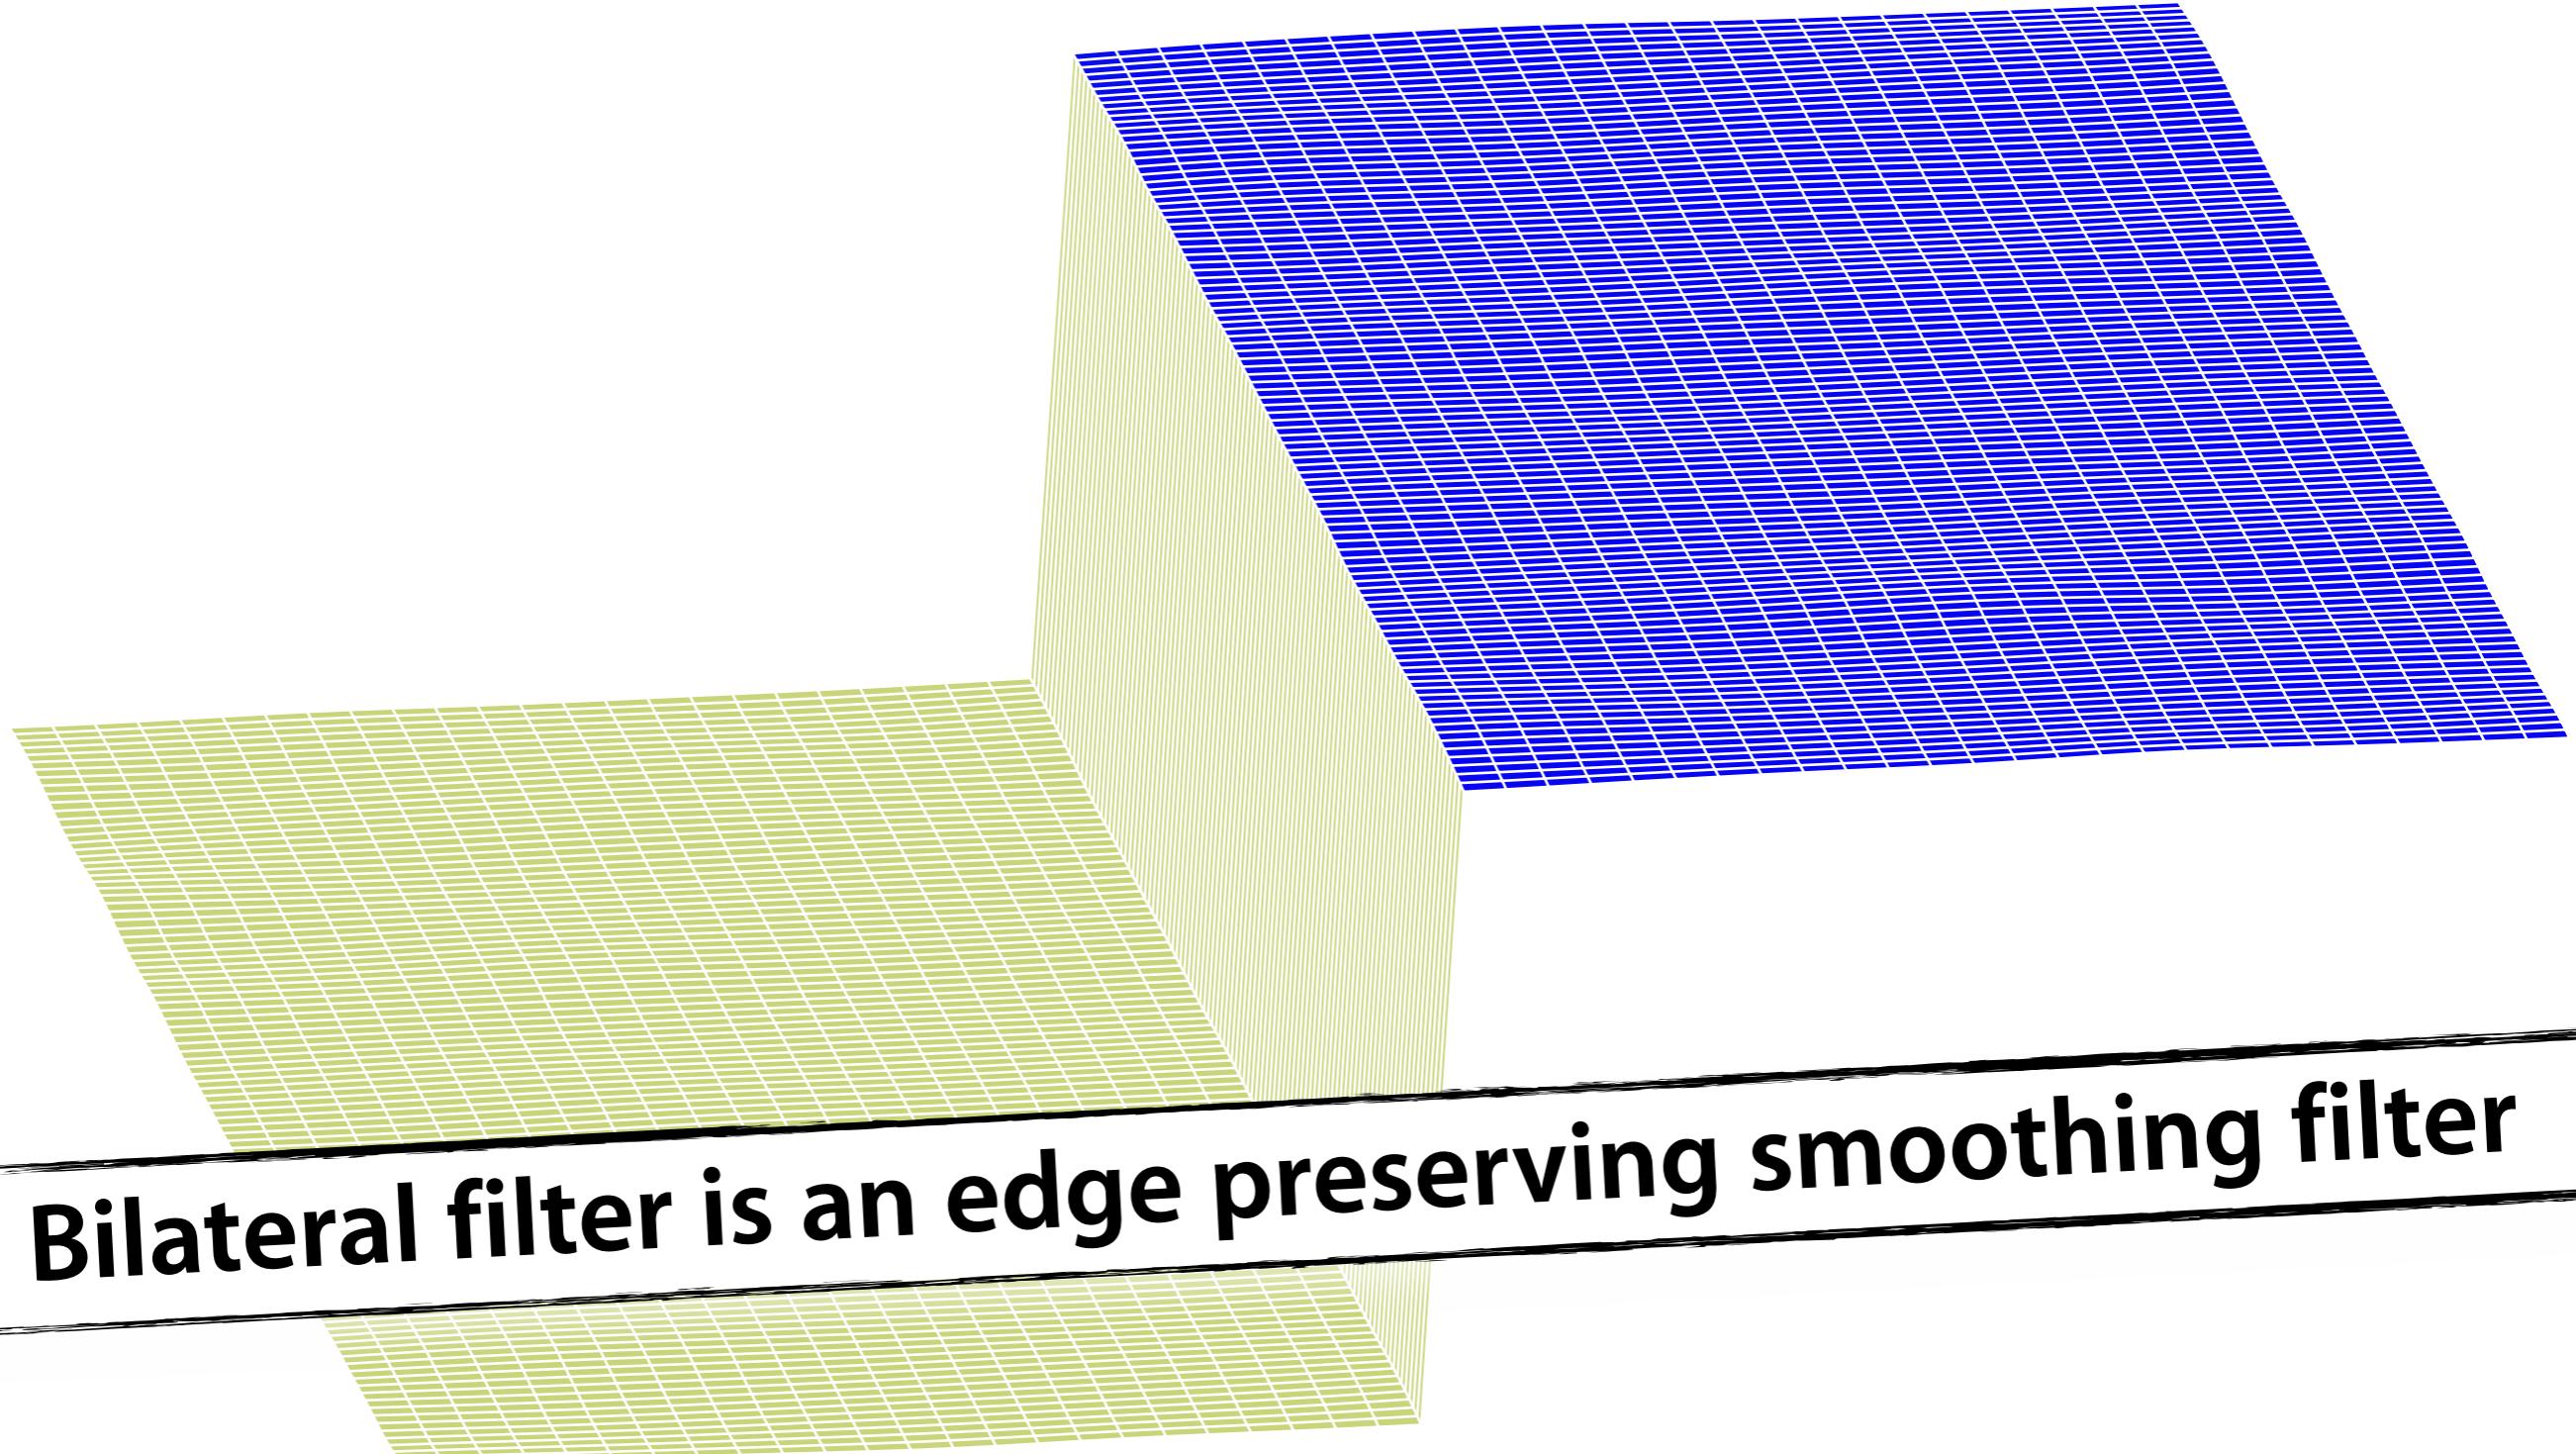
**range weight**

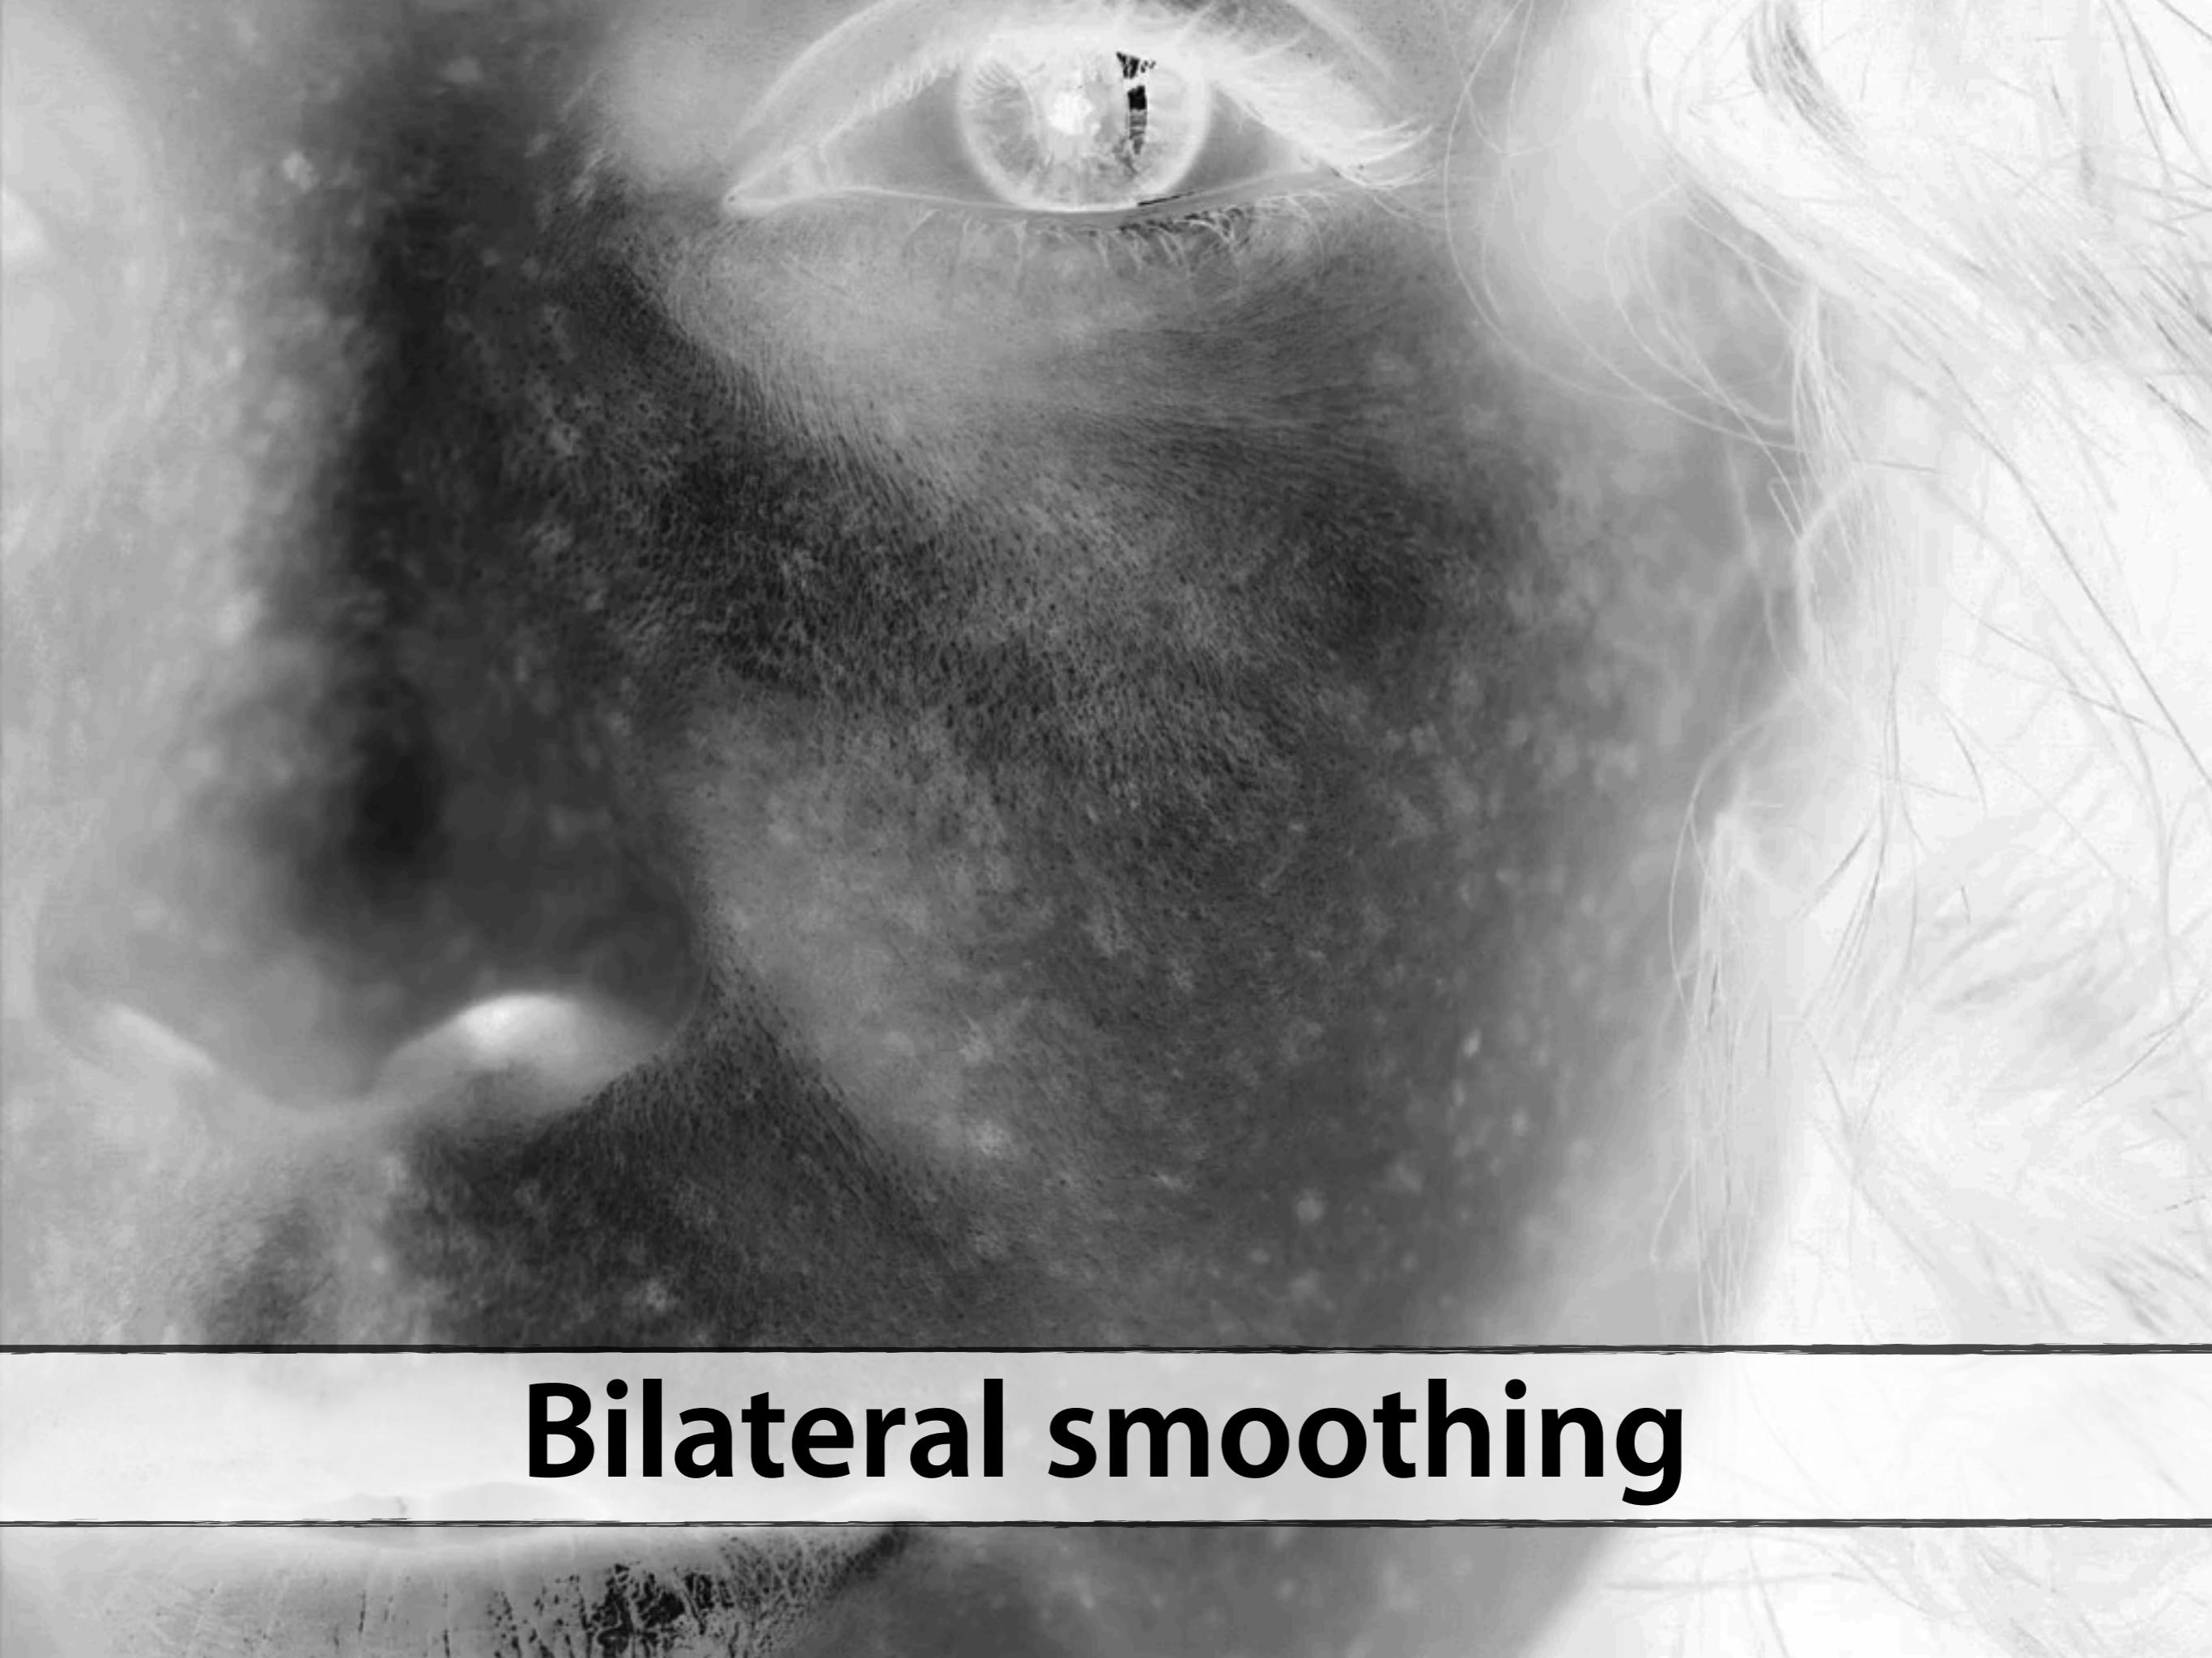
$$w(\mathbf{p}) G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|)$$



**weight filter shape conforms to local image structure**

$$I'(\mathbf{p}) = w(\mathbf{p}) \sum_{\mathbf{q} \in S} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I(\mathbf{p}) - I(\mathbf{q})\|) I(\mathbf{q})$$





# Bilateral smoothing