

**Intro to**

# **Computer Vision**

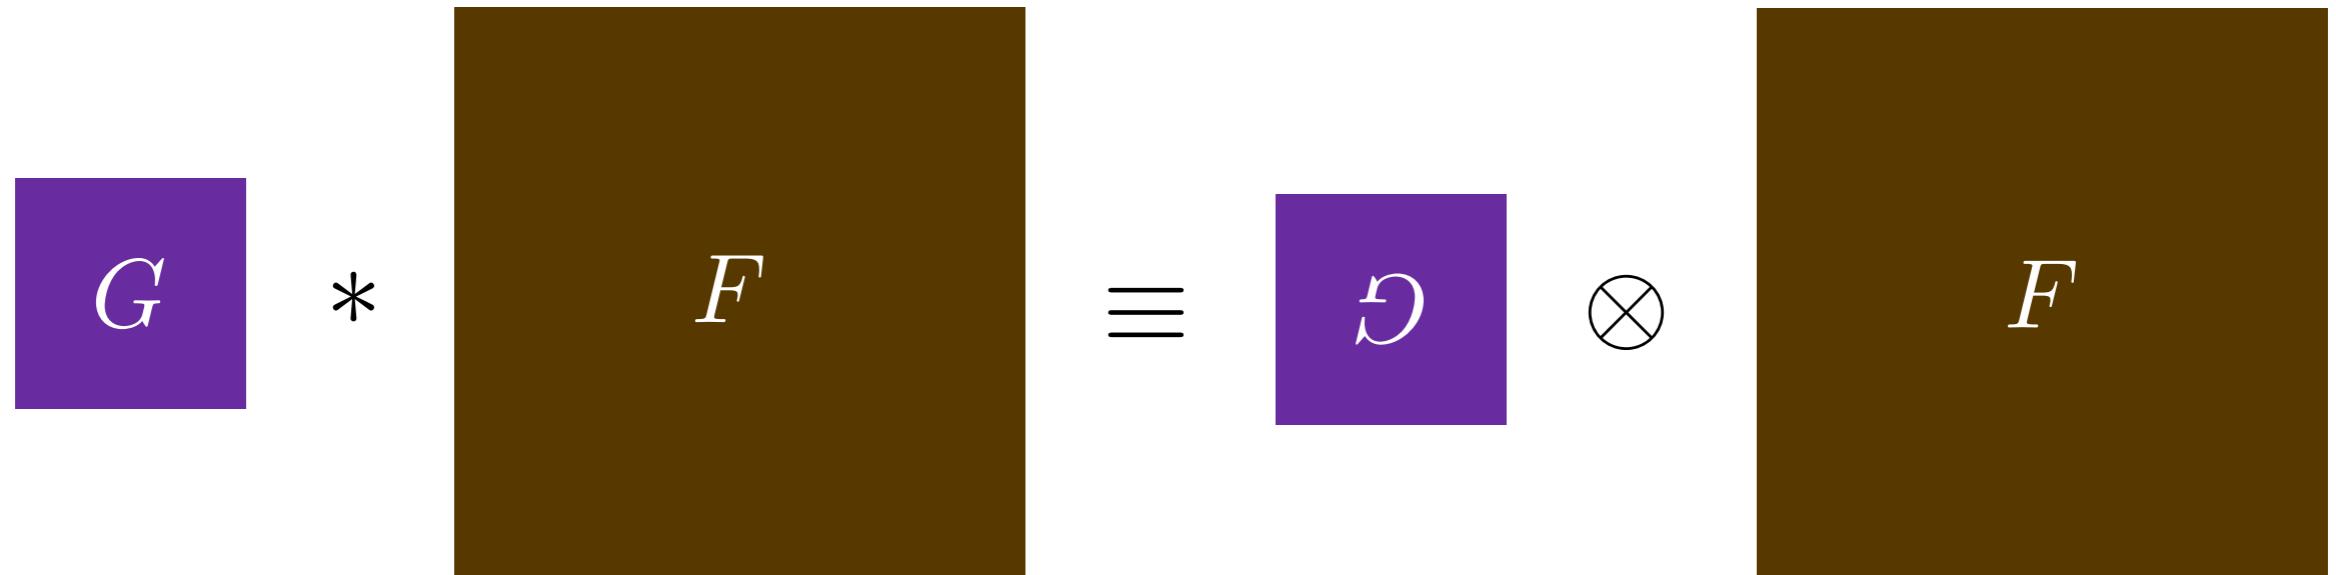
**with Prof. Kosta Derpanis**

## **Edge detection**

# Review

**Correlation**  $H = G \otimes F$

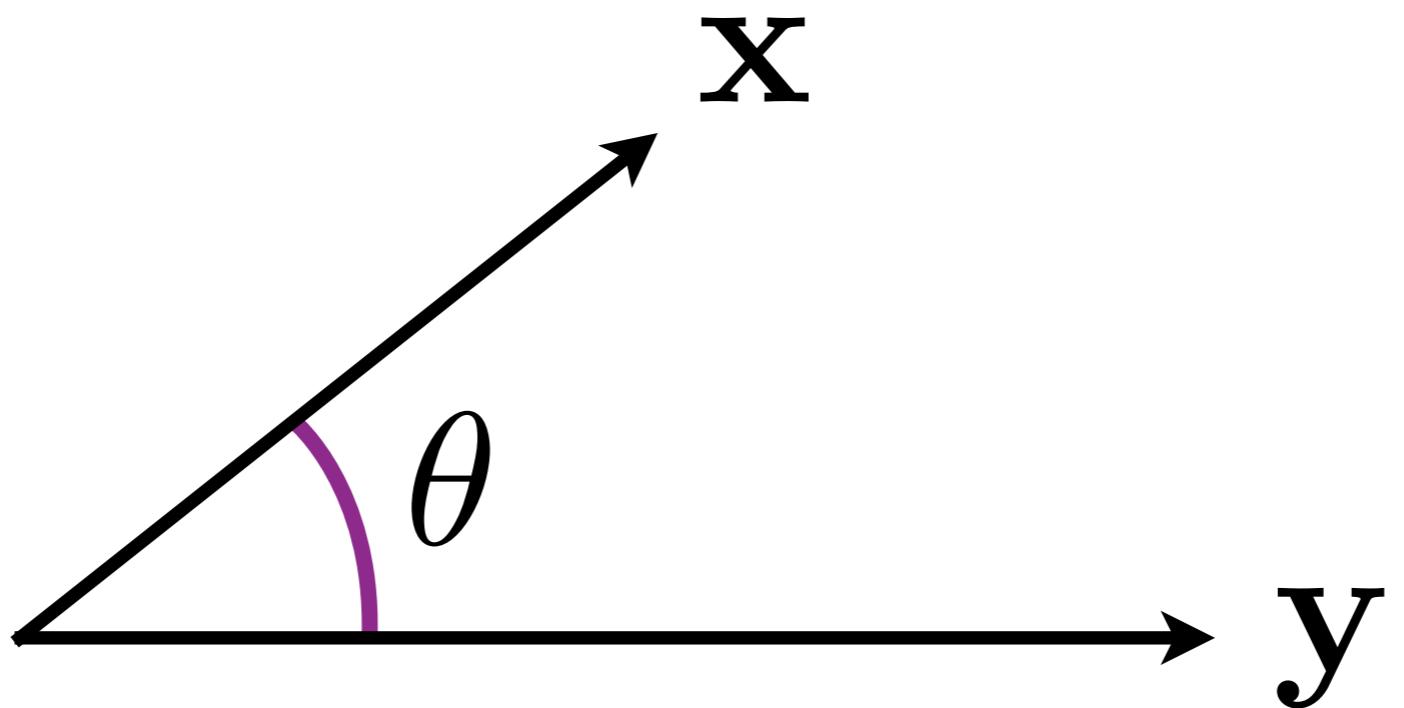
$$H[x, y] = \sum_{u=-K}^K \sum_{v=-K}^K G[u, v] F[x + u, y + v]$$



**Convolution**  $H = G * F$

$$H[x, y] = \sum_{u=-K}^K \sum_{v=-K}^K G[u, v] F[x - u, y - v]$$

$$\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^N x_i y_i = \|\mathbf{x}\| \|\mathbf{y}\| \cos \theta$$



C

template



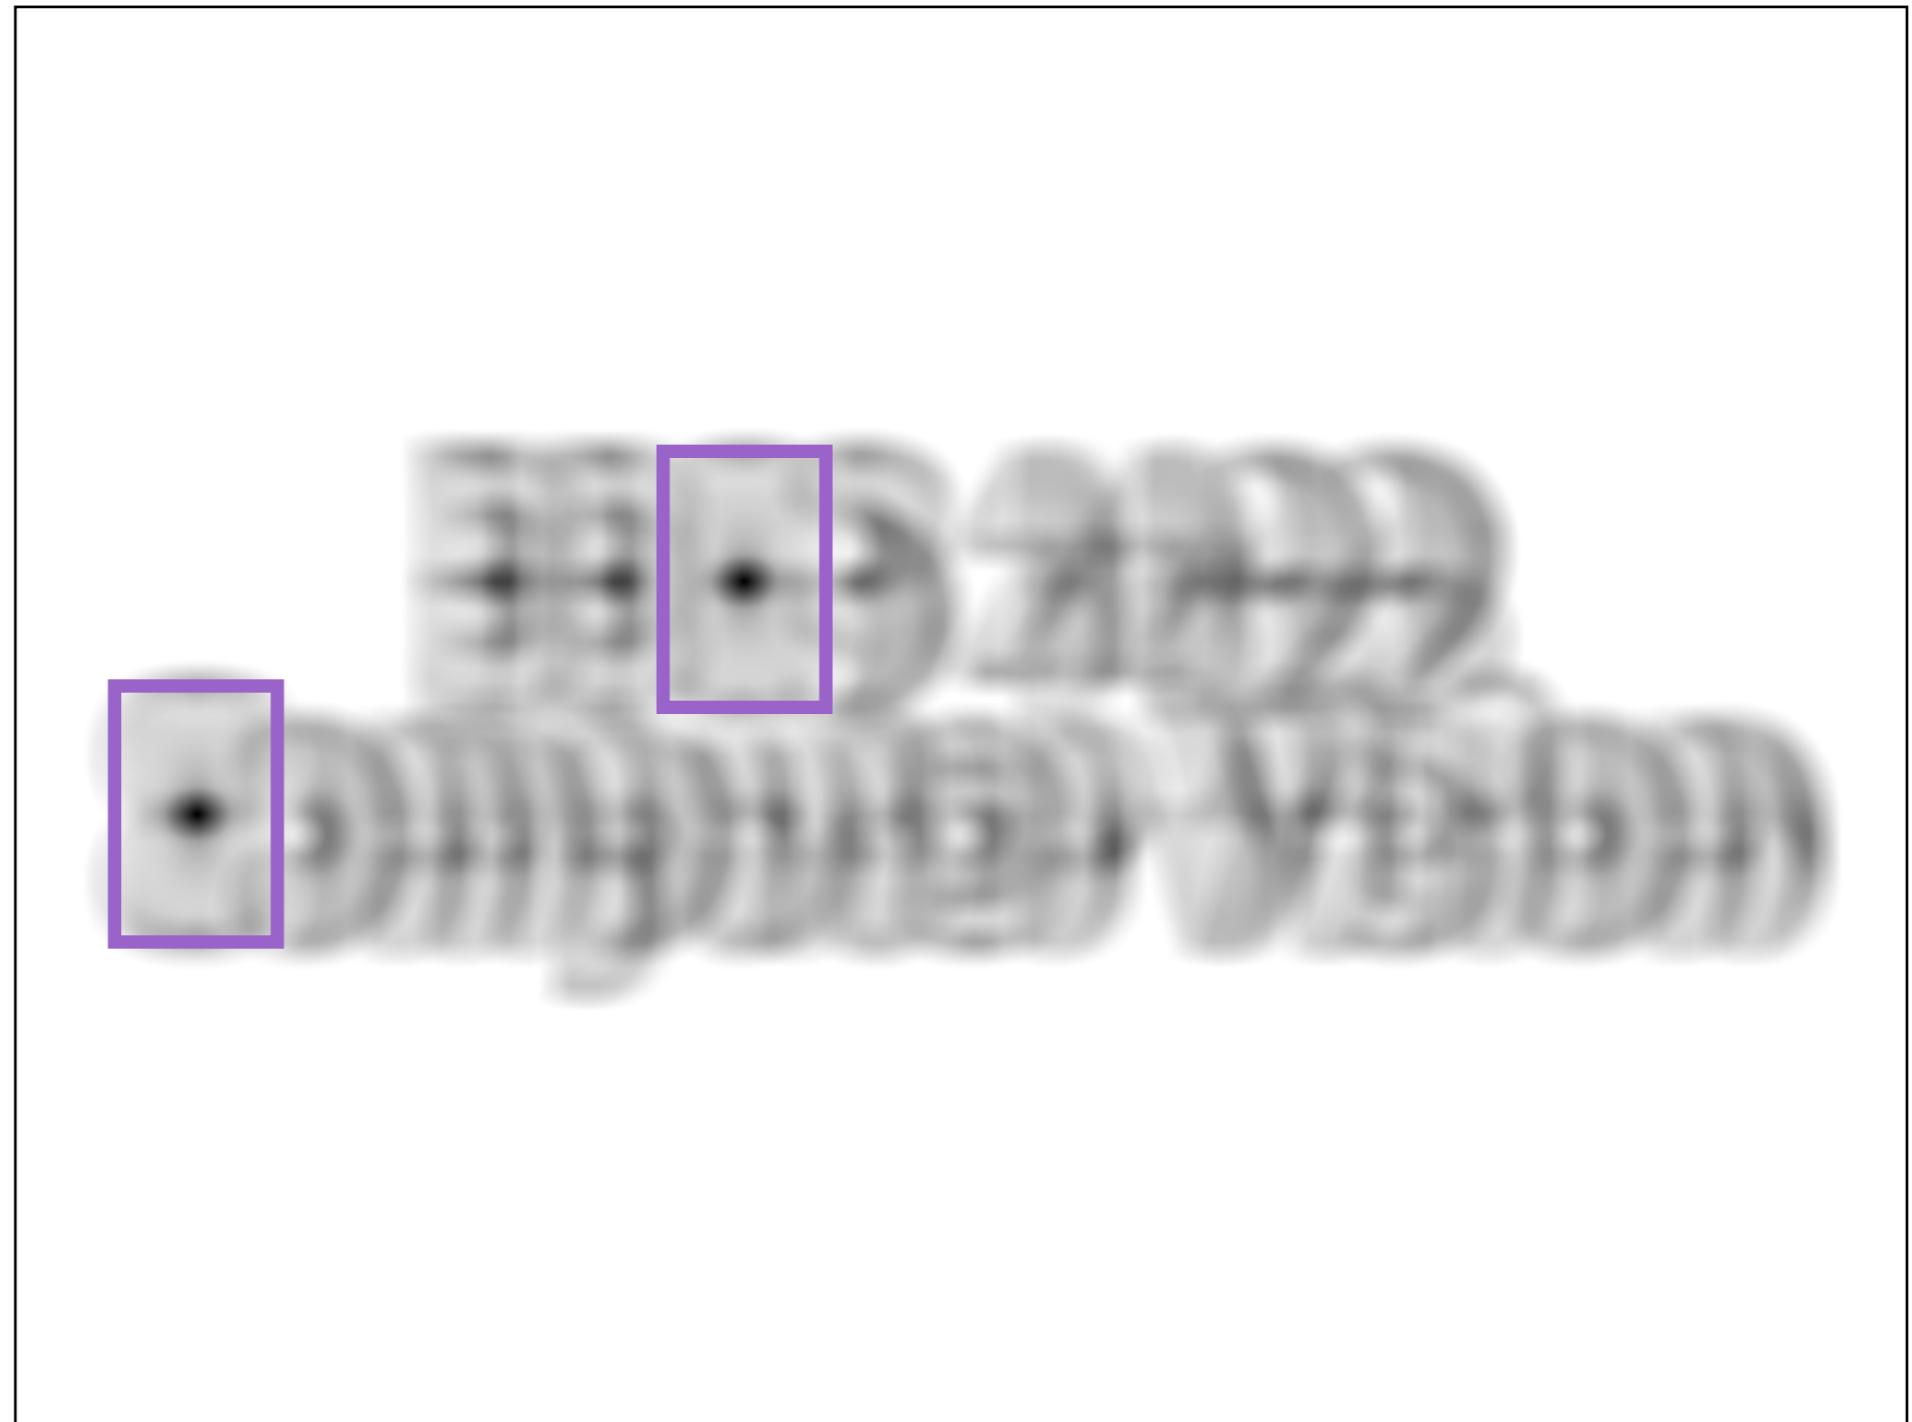
# EECS 4422

# Computer Vision

search image

**n**

=



**correlation map**

n

=

EECS 4422  
Computer Vision

detected templates

**C**  $\otimes$   
template

subimage

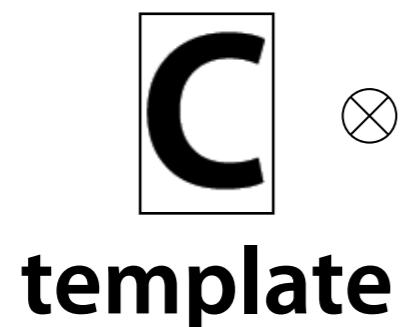
EE**CS** 4422

# Computer Vision

search image

**correlation score** =  $t \cdot s$

**scalar product between vectors**



subimage

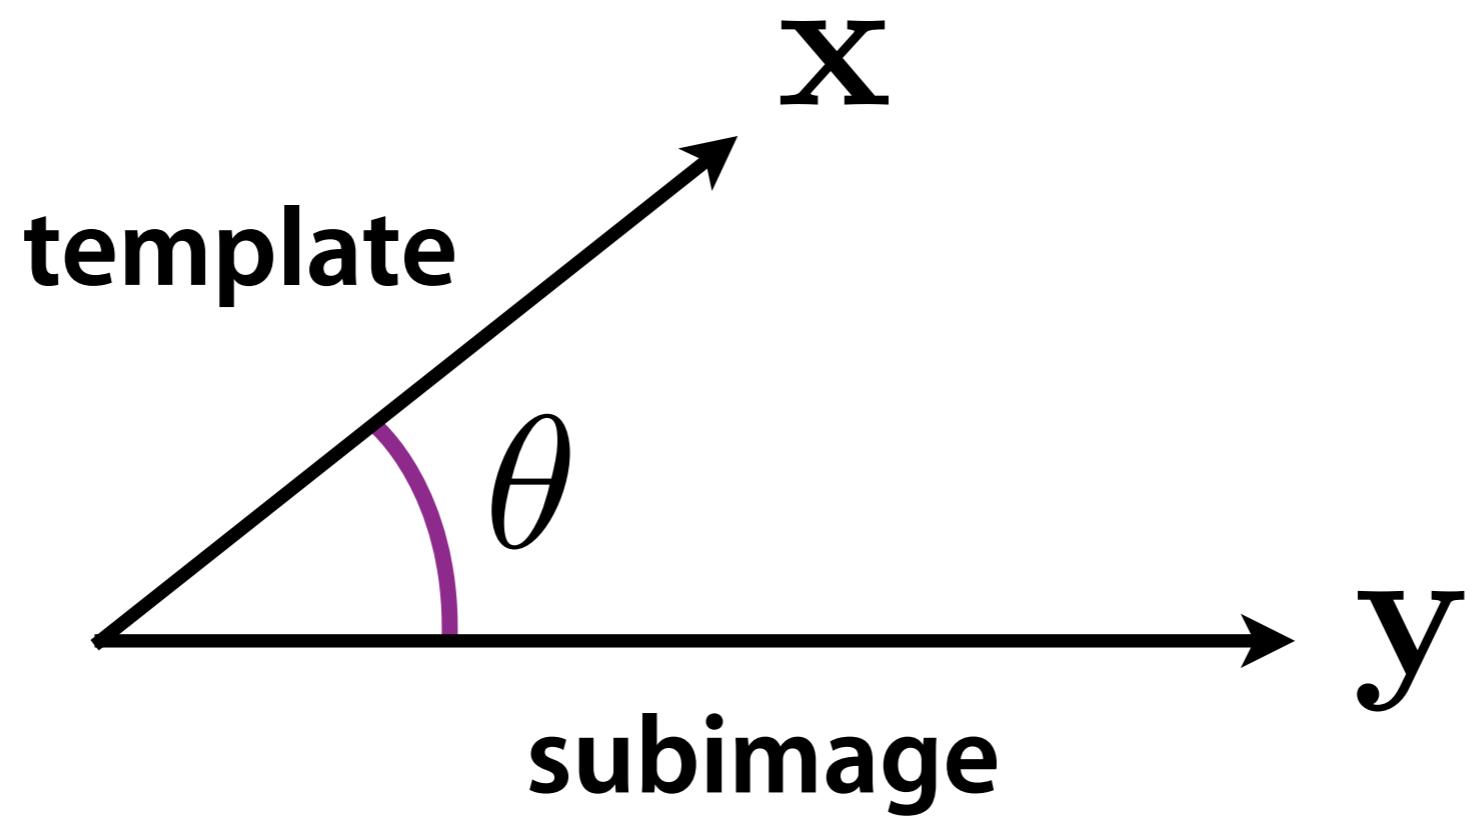
EECS 4422

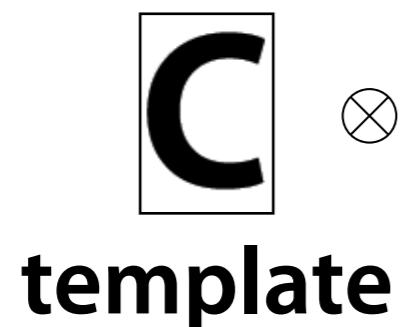
Computer Vision

search image

$$\text{correlation score} = \mathbf{t} \cdot \mathbf{s} = \|\mathbf{t}\| \|\mathbf{s}\| \cos \theta$$

$$\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^N x_i y_i = \|\mathbf{x}\| \|\mathbf{y}\| \cos \theta$$





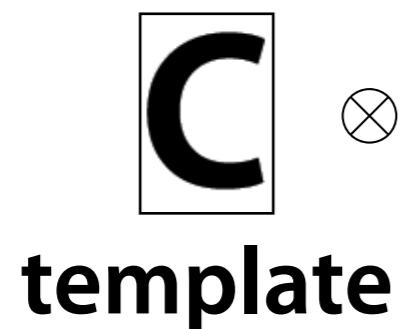
subimage

EECS 4422

Computer Vision

search image

$$\text{correlation score} = \mathbf{t} \cdot \mathbf{s} = \|\mathbf{t}\| \|\mathbf{s}\| \cos \theta$$



subimage

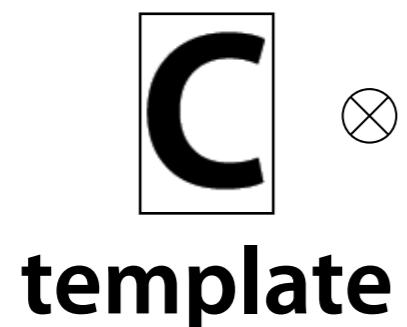
EECS 4422

# Computer Vision

search image

$$\text{correlation score} = \mathbf{t} \cdot \mathbf{s} = \|\mathbf{t}\| \|\mathbf{s}\| \cos \theta$$

score is sensitive to multiplicative brightness scaling



subimage

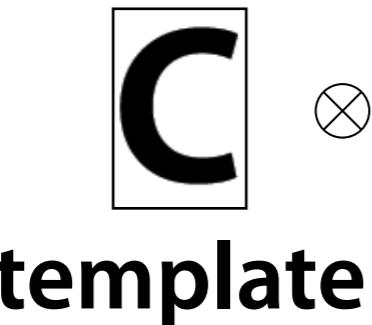
EECS 4422

Computer Vision

search image

**correlation score** =  $t \cdot s = \|t\| \|s\| \cos \theta$

**normalized score** =  $\frac{t}{\|t\|} \cdot \frac{s}{\|s\|} = \cos \theta$



template

# EECS 4422 Computer Vision

search image

$$\text{correlation score} = \mathbf{t} \cdot \mathbf{s} = \|\mathbf{t}\| \|\mathbf{s}\| \cos \theta$$

$$\text{normalized score} = \frac{\mathbf{t}}{\|\mathbf{t}\|} \cdot \frac{\mathbf{s}}{\|\mathbf{s}\|} = \cos \theta$$

How can we make it **invariant** to  
multiplicative and additive brightness change?



⊗

template

EECS 4422

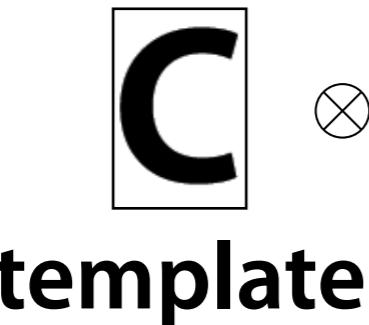
# Computer Vision

search image

**correlation score** =  $t \cdot s = \|t\| \|s\| \cos \theta$

**normalized score** =  $\frac{t}{\|t\|} \cdot \frac{s}{\|s\|} = \cos \theta$

**normalized cross-correlation** =  $\frac{(t - \bar{t})}{\|(t - \bar{t})\|} \cdot \frac{(s - \bar{s})}{\|(s - \bar{s})\|}$



# EECS 4422

## Computer Vision

search image

$$\text{correlation score} = \mathbf{t} \cdot \mathbf{s} = \|\mathbf{t}\| \|\mathbf{s}\| \cos \theta$$

$$\text{normalized score} = \frac{\mathbf{t}}{\|\mathbf{t}\|} \cdot \frac{\mathbf{s}}{\|\mathbf{s}\|} = \cos \theta$$

$$\text{normalized cross-correlation} = \frac{(\mathbf{t} - \bar{\mathbf{t}}) \cdot (\mathbf{s} - \bar{\mathbf{s}})}{\|(\mathbf{t} - \bar{\mathbf{t}})\| \|(\mathbf{s} - \bar{\mathbf{s}})\|}$$

average brightness

*y*



**steep cliffs**

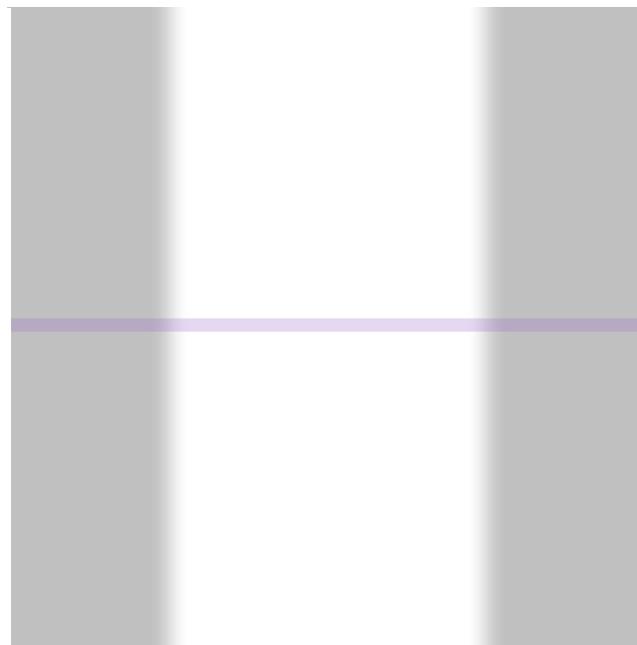
*x*

image( $x, y$ )

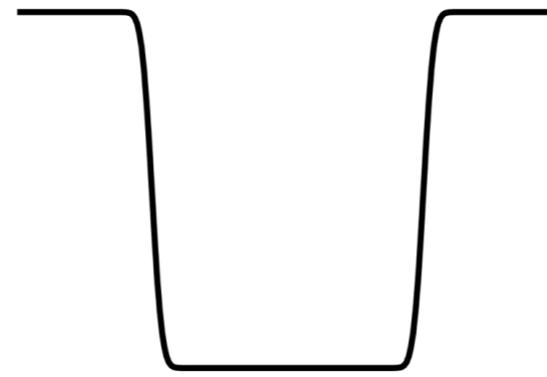




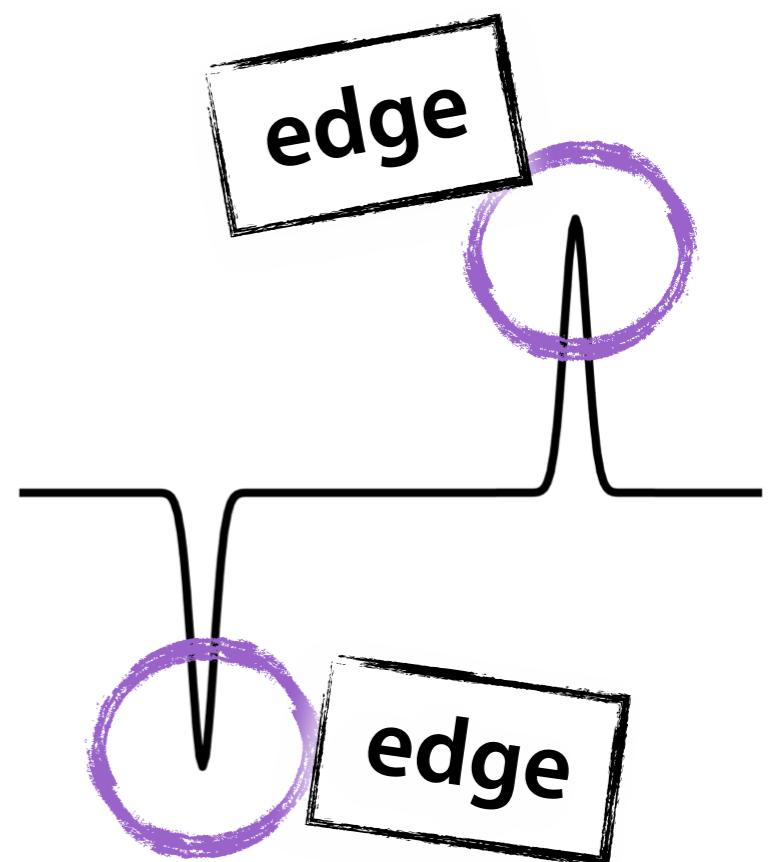
# Edge Detection



image

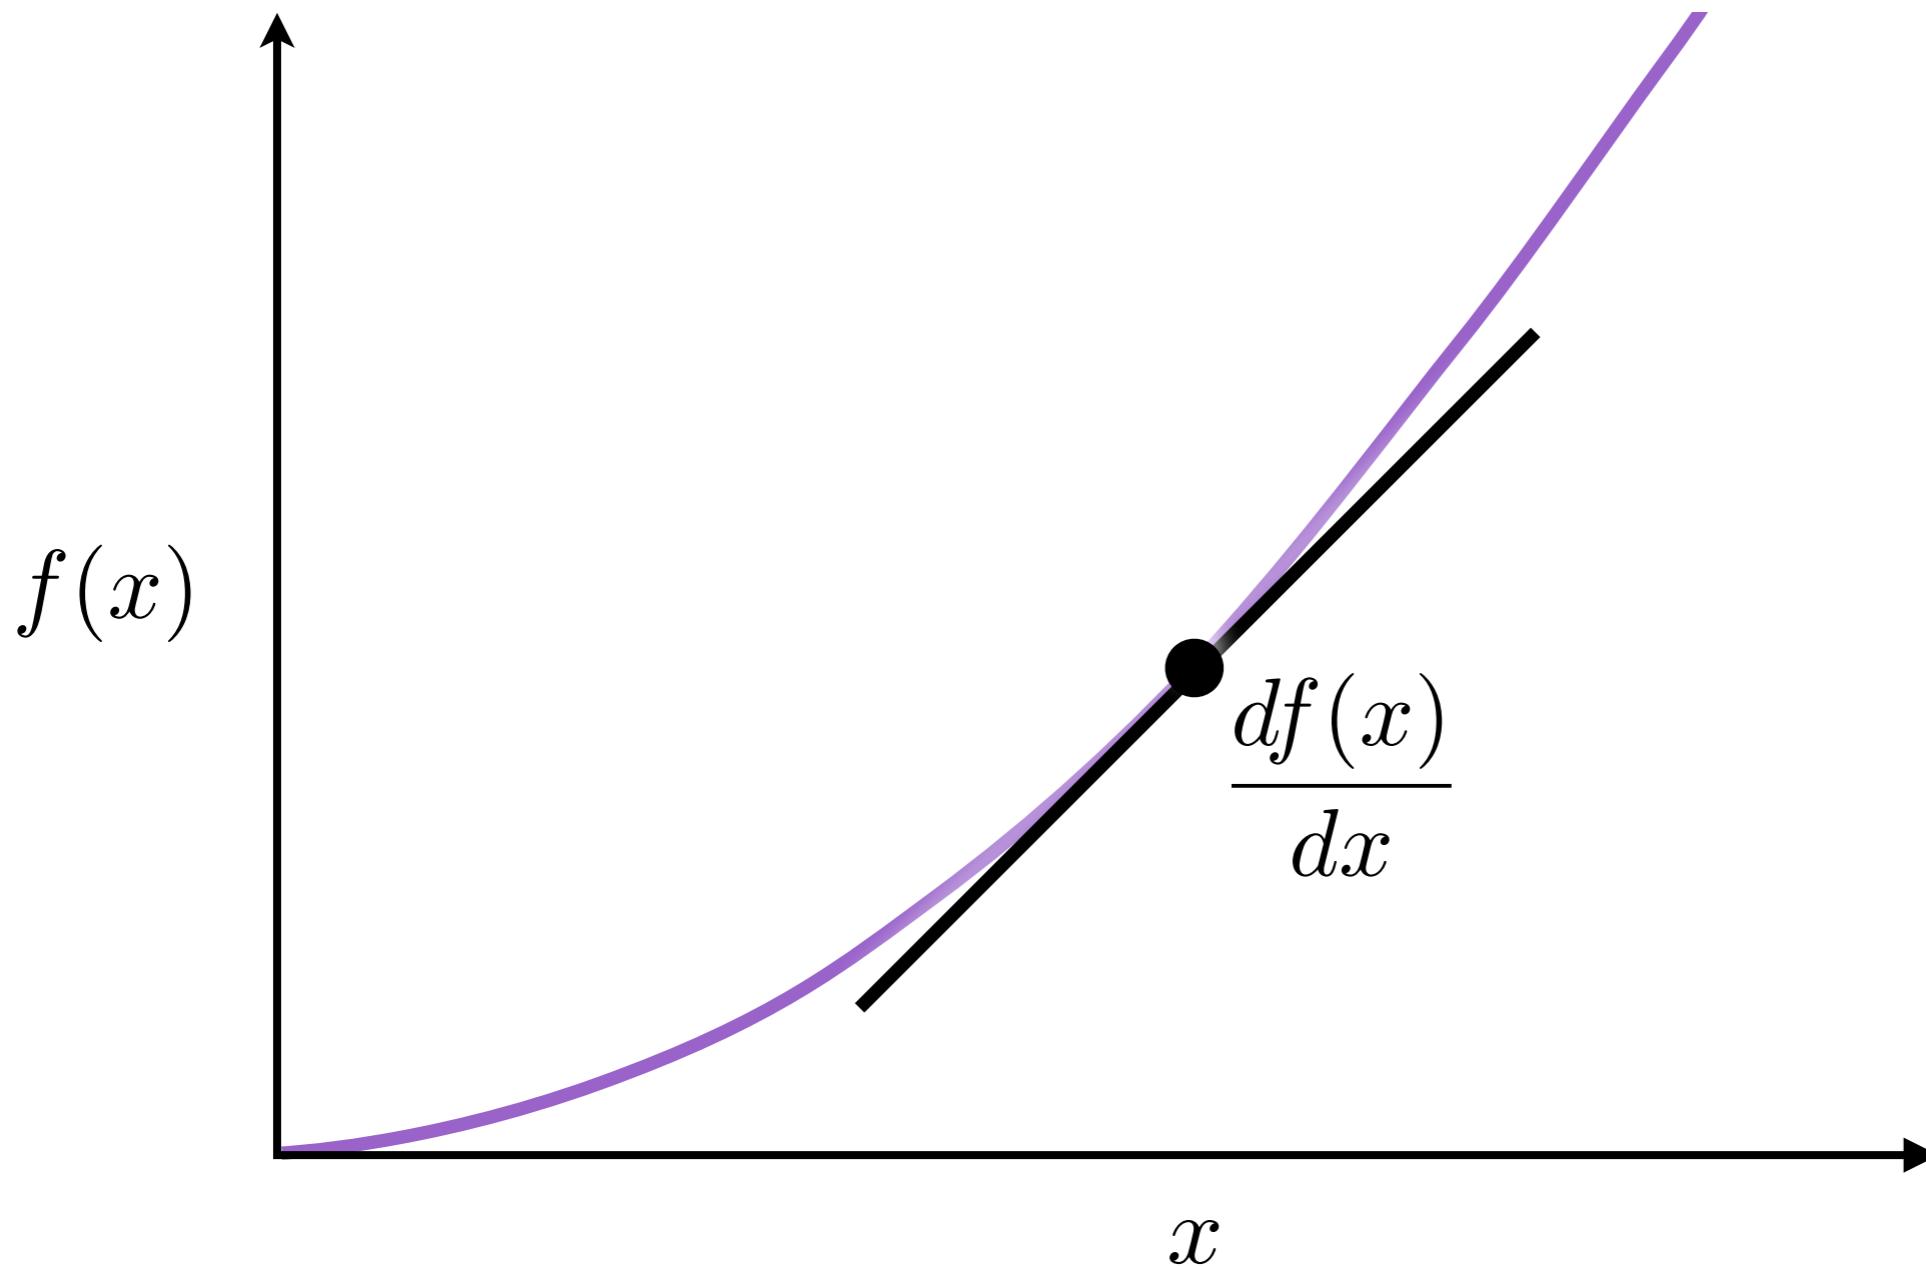


intensity  
function  
(slice)



first derivative

$$\frac{df(x)}{dx} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$



# Partial Derivatives

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon, y) - f(x, y)}{\epsilon}$$

$$\frac{\partial f(x, y)}{\partial y} = \lim_{\epsilon \rightarrow 0} \frac{f(x, y + \epsilon) - f(x, y)}{\epsilon}$$

# Discrete Derivatives

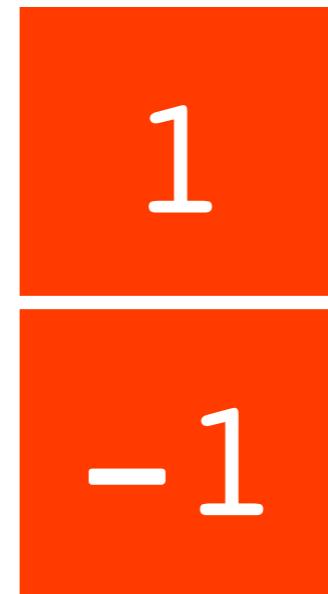
$$\begin{aligned}\frac{\partial f(x, y)}{\partial x} &\approx \frac{f(x + 1, y) - f(x, y)}{1} \\&= f(x + 1, y) - f(x, y)\end{aligned}$$

$$\frac{\partial f(x, y)}{\partial x} \approx f(x + 1, y) - f(x, y)$$

$$\frac{\partial f(x, y)}{\partial y} \approx f(x, y + 1) - f(x, y)$$

What are the associated convolution filters?

1	-1
---	----



$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial y}$$

# finite differences

$$\frac{\partial f(x, y)}{\partial x}$$

# finite differences

$$\frac{\partial f(x, y)}{\partial y}$$

*separable  
filters*

$$F[x, y] = U[x]V[y]$$

$$\begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix} = \begin{matrix} 1 \\ 0 \\ -1 \end{matrix} \quad \begin{matrix} 1 & 2 & 1 \\ 1 & 2 & 1 \end{matrix}$$

**Sobel Filter**

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

# Sobel Filter

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

Other  
Derivative Filters

## Sobel Filter

1	0	-1
1	0	-1
1	0	-1

1	1	1
0	0	0
-1	-1	-1

## Prewitt

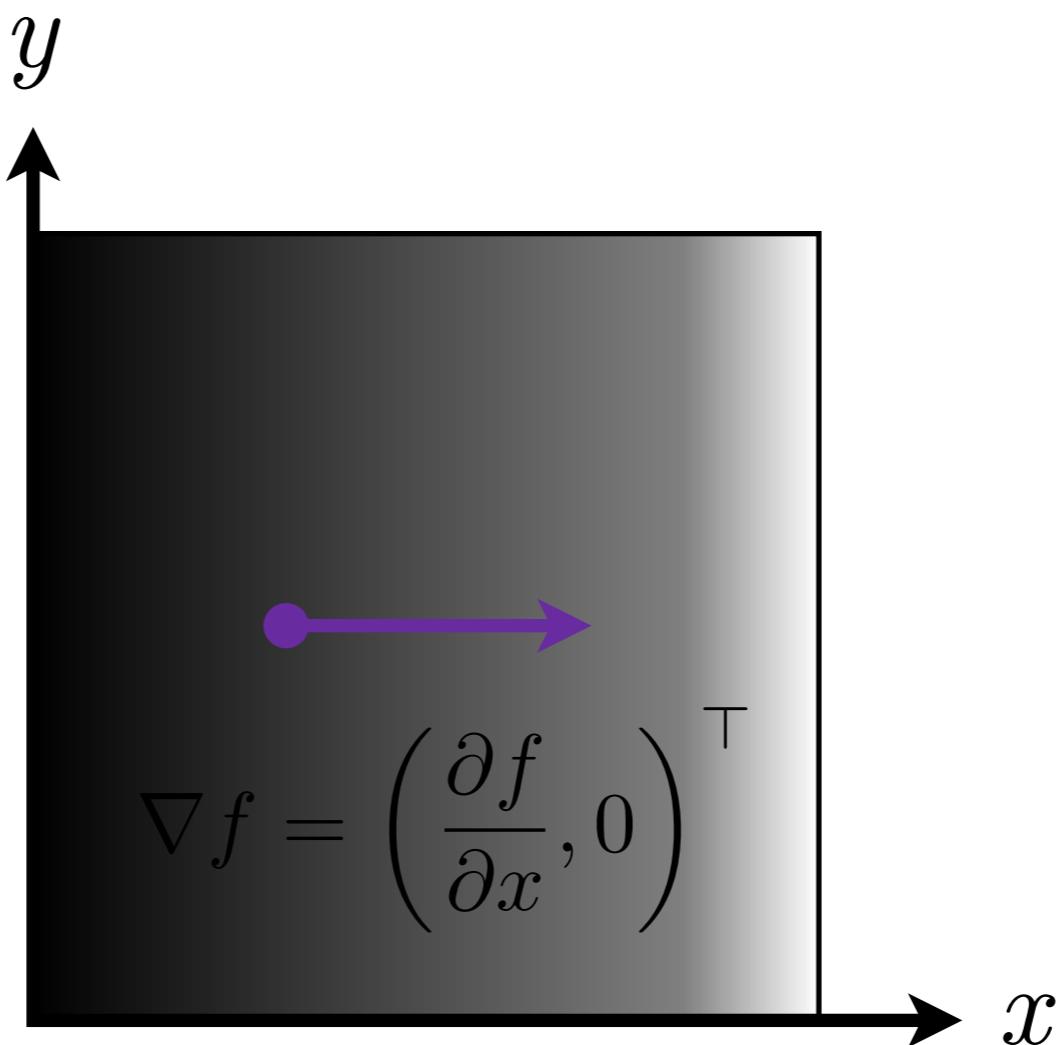
0	-1
1	0

-1	0
0	1

## Roberts Cross

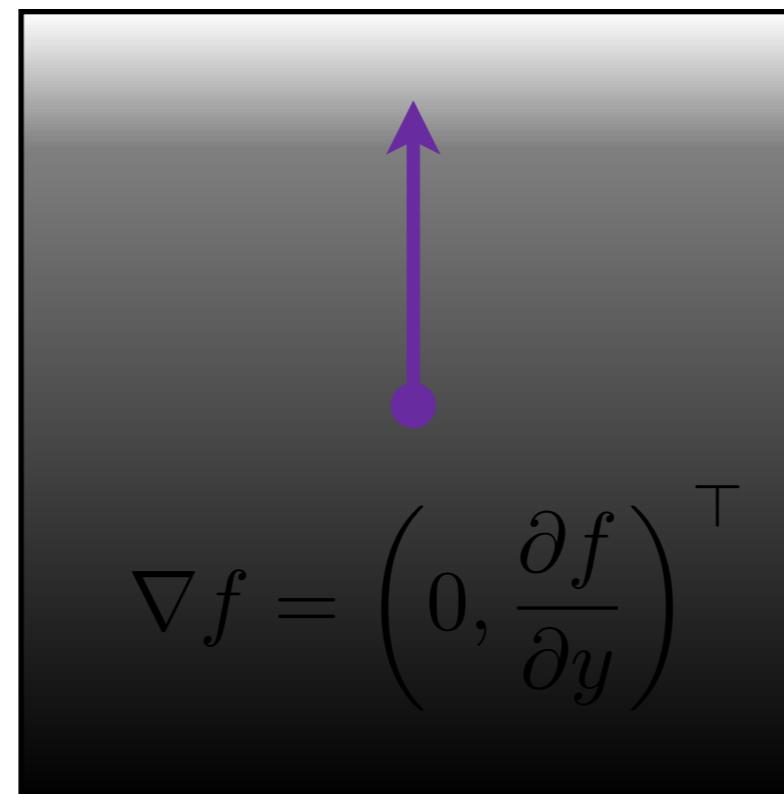
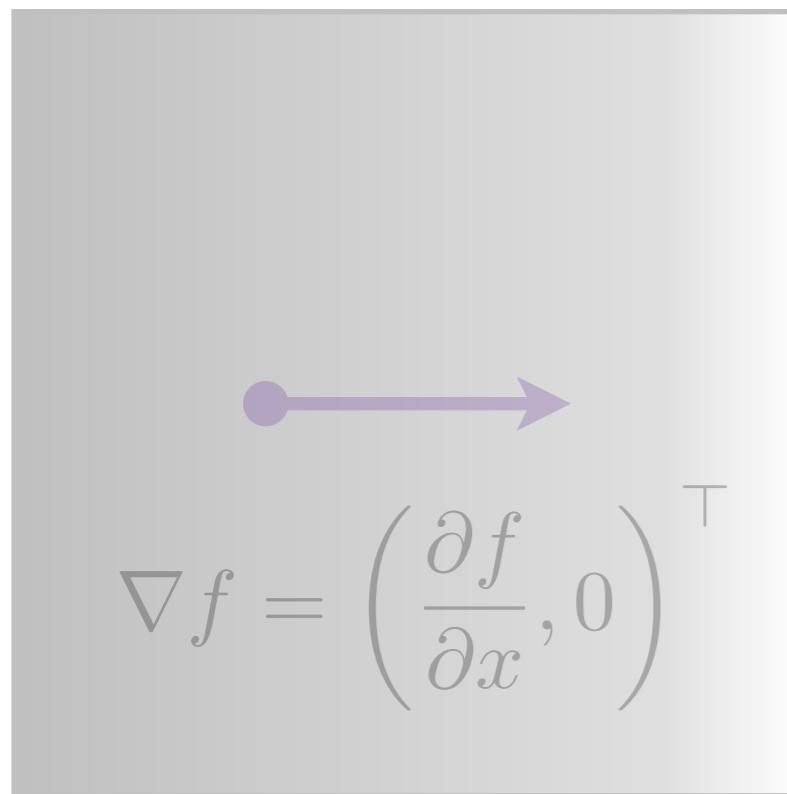
# gradient

vector that points in the direction of the greatest rate of increase of a function and whose magnitude is that rate of increase



# gradient

vector that points in the direction of the greatest rate of increase of a function and whose magnitude is that rate of increase



# gradient

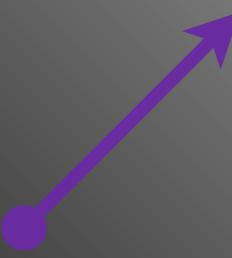
vector that points in the direction of the greatest rate of increase of a function and whose magnitude is that rate of increase



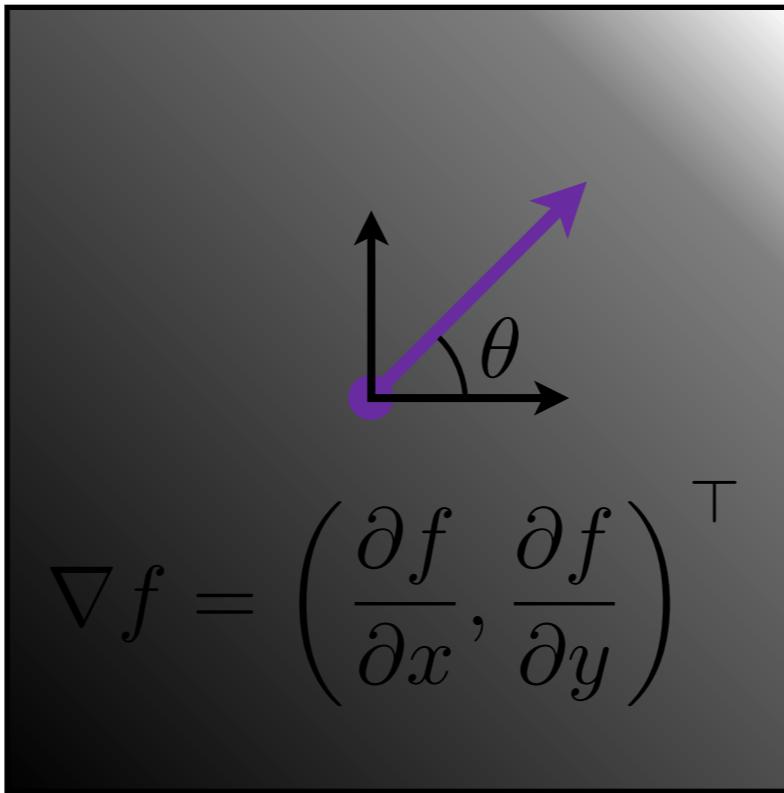
$$\nabla f = \left( \frac{\partial f}{\partial x}, 0 \right)^\top$$



$$\nabla f = \left( 0, \frac{\partial f}{\partial y} \right)^\top$$



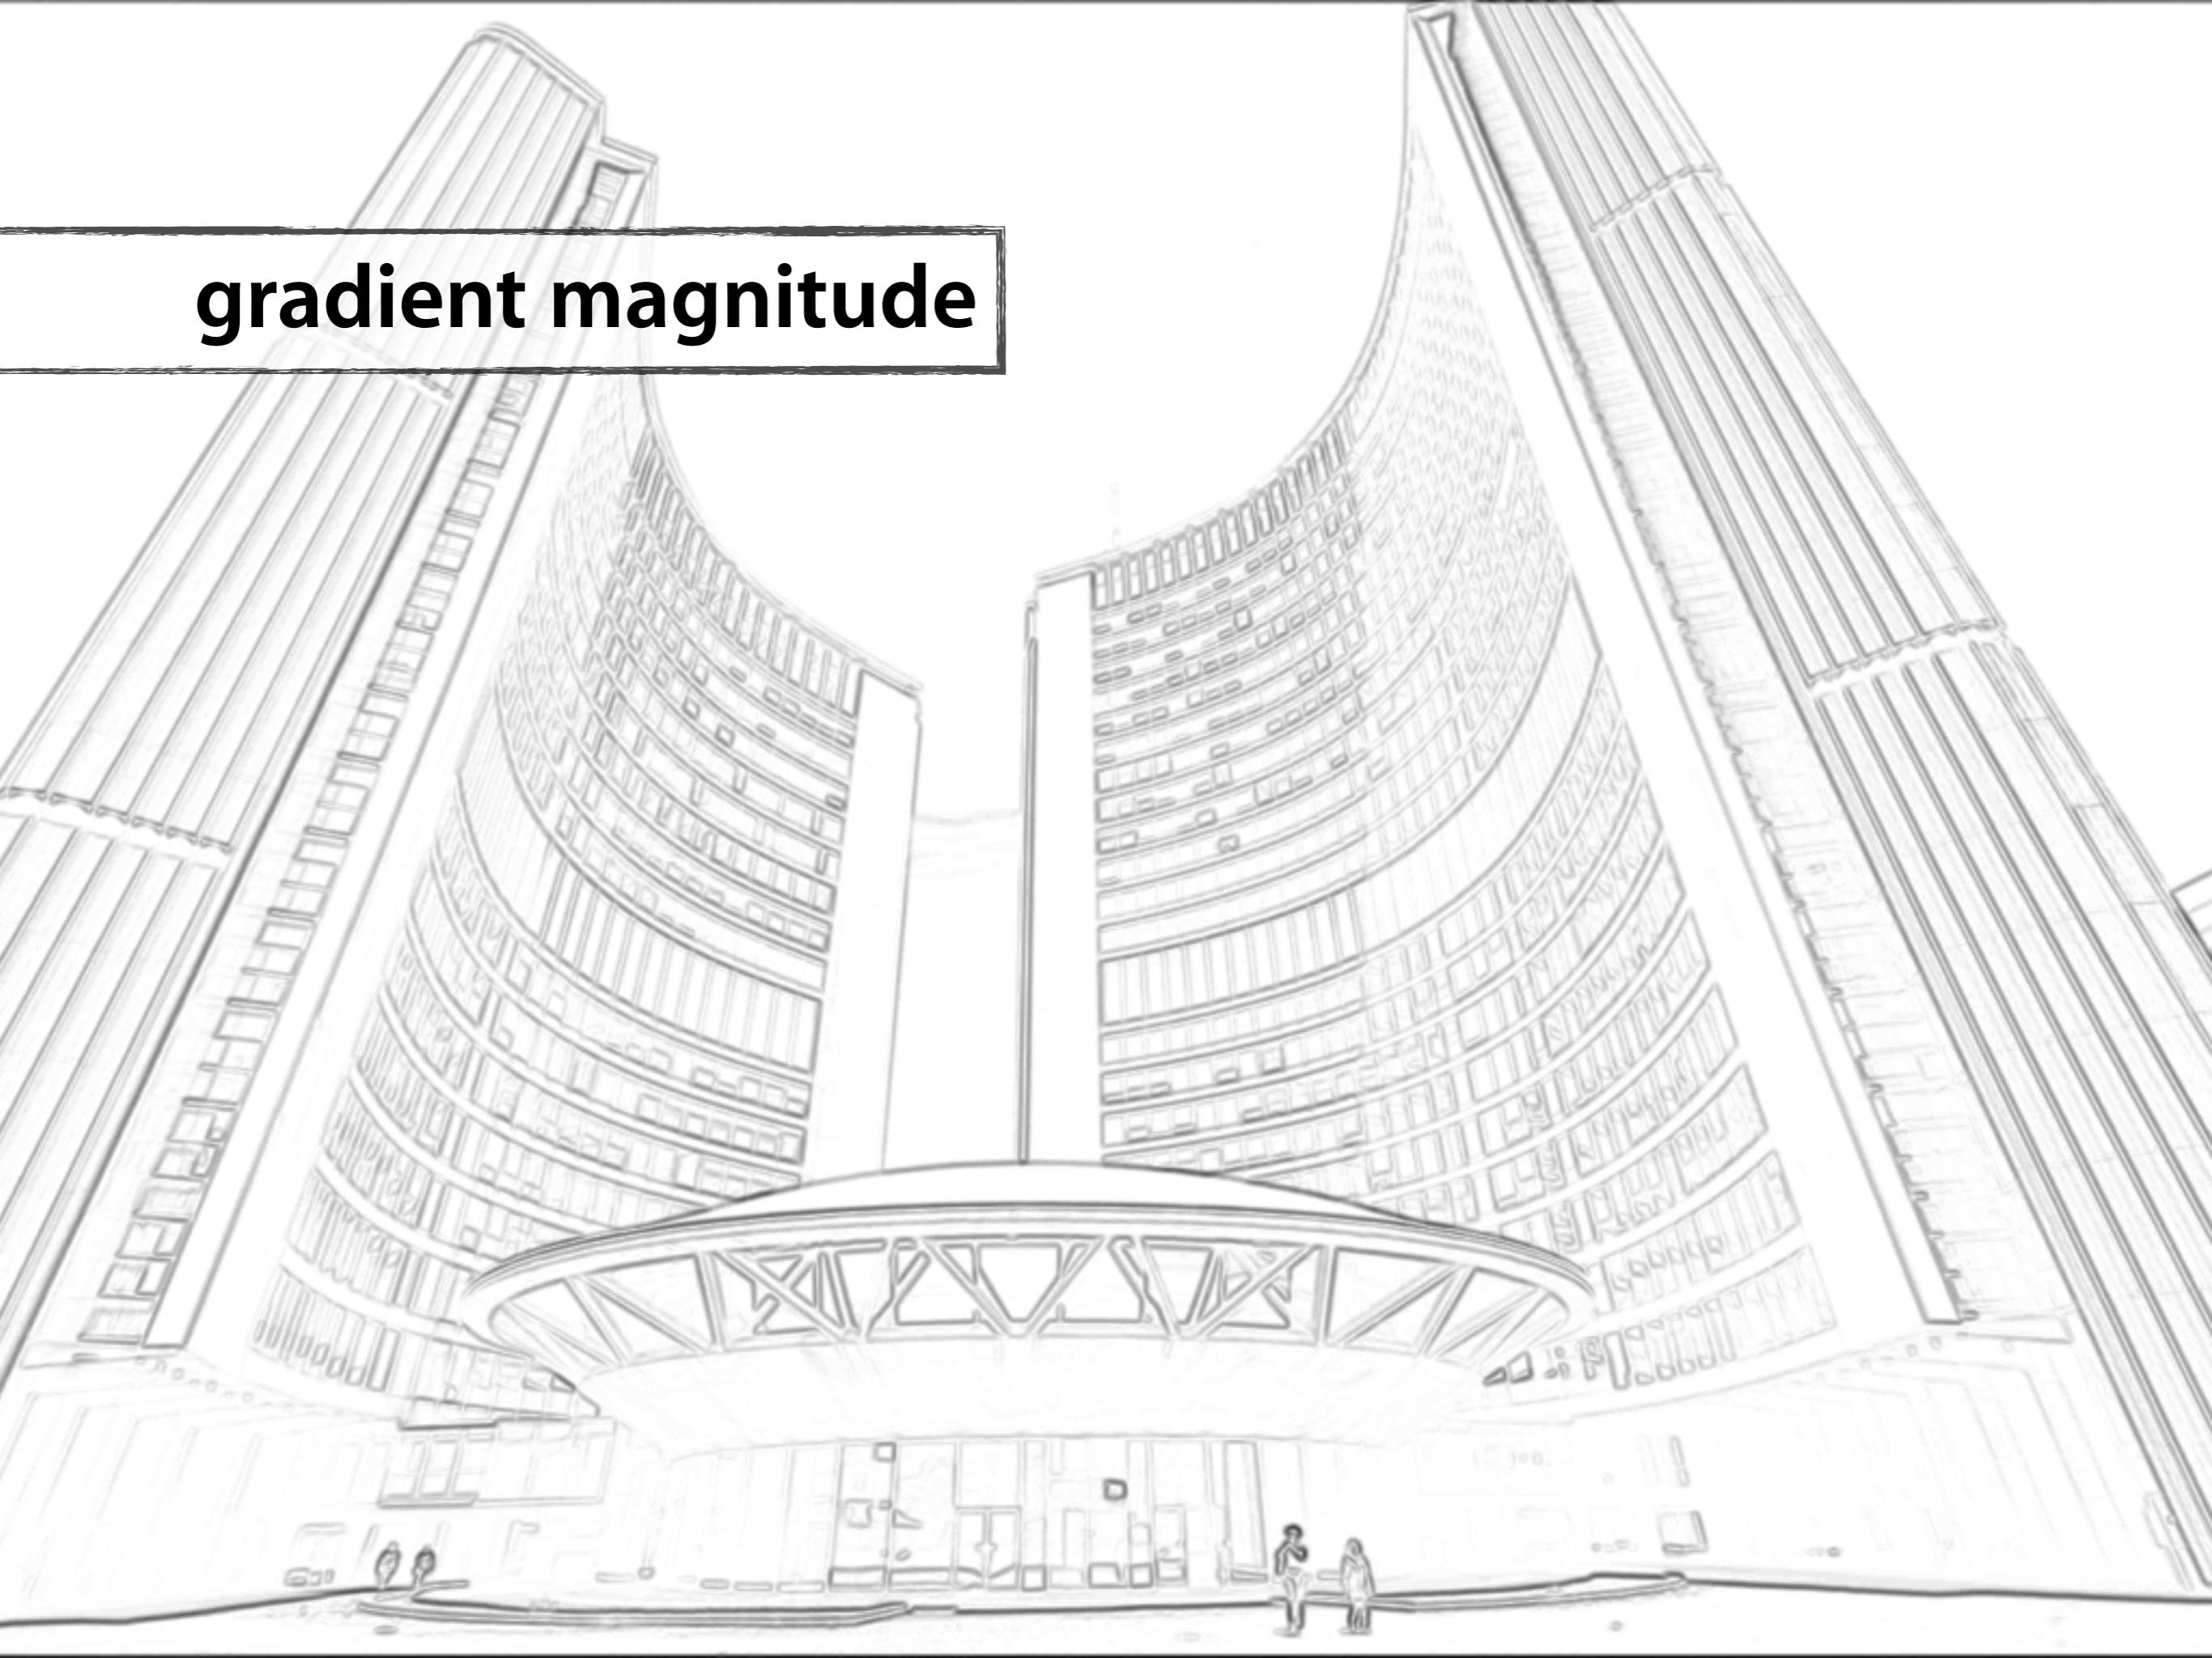
$$\nabla f = \left( \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)^\top$$



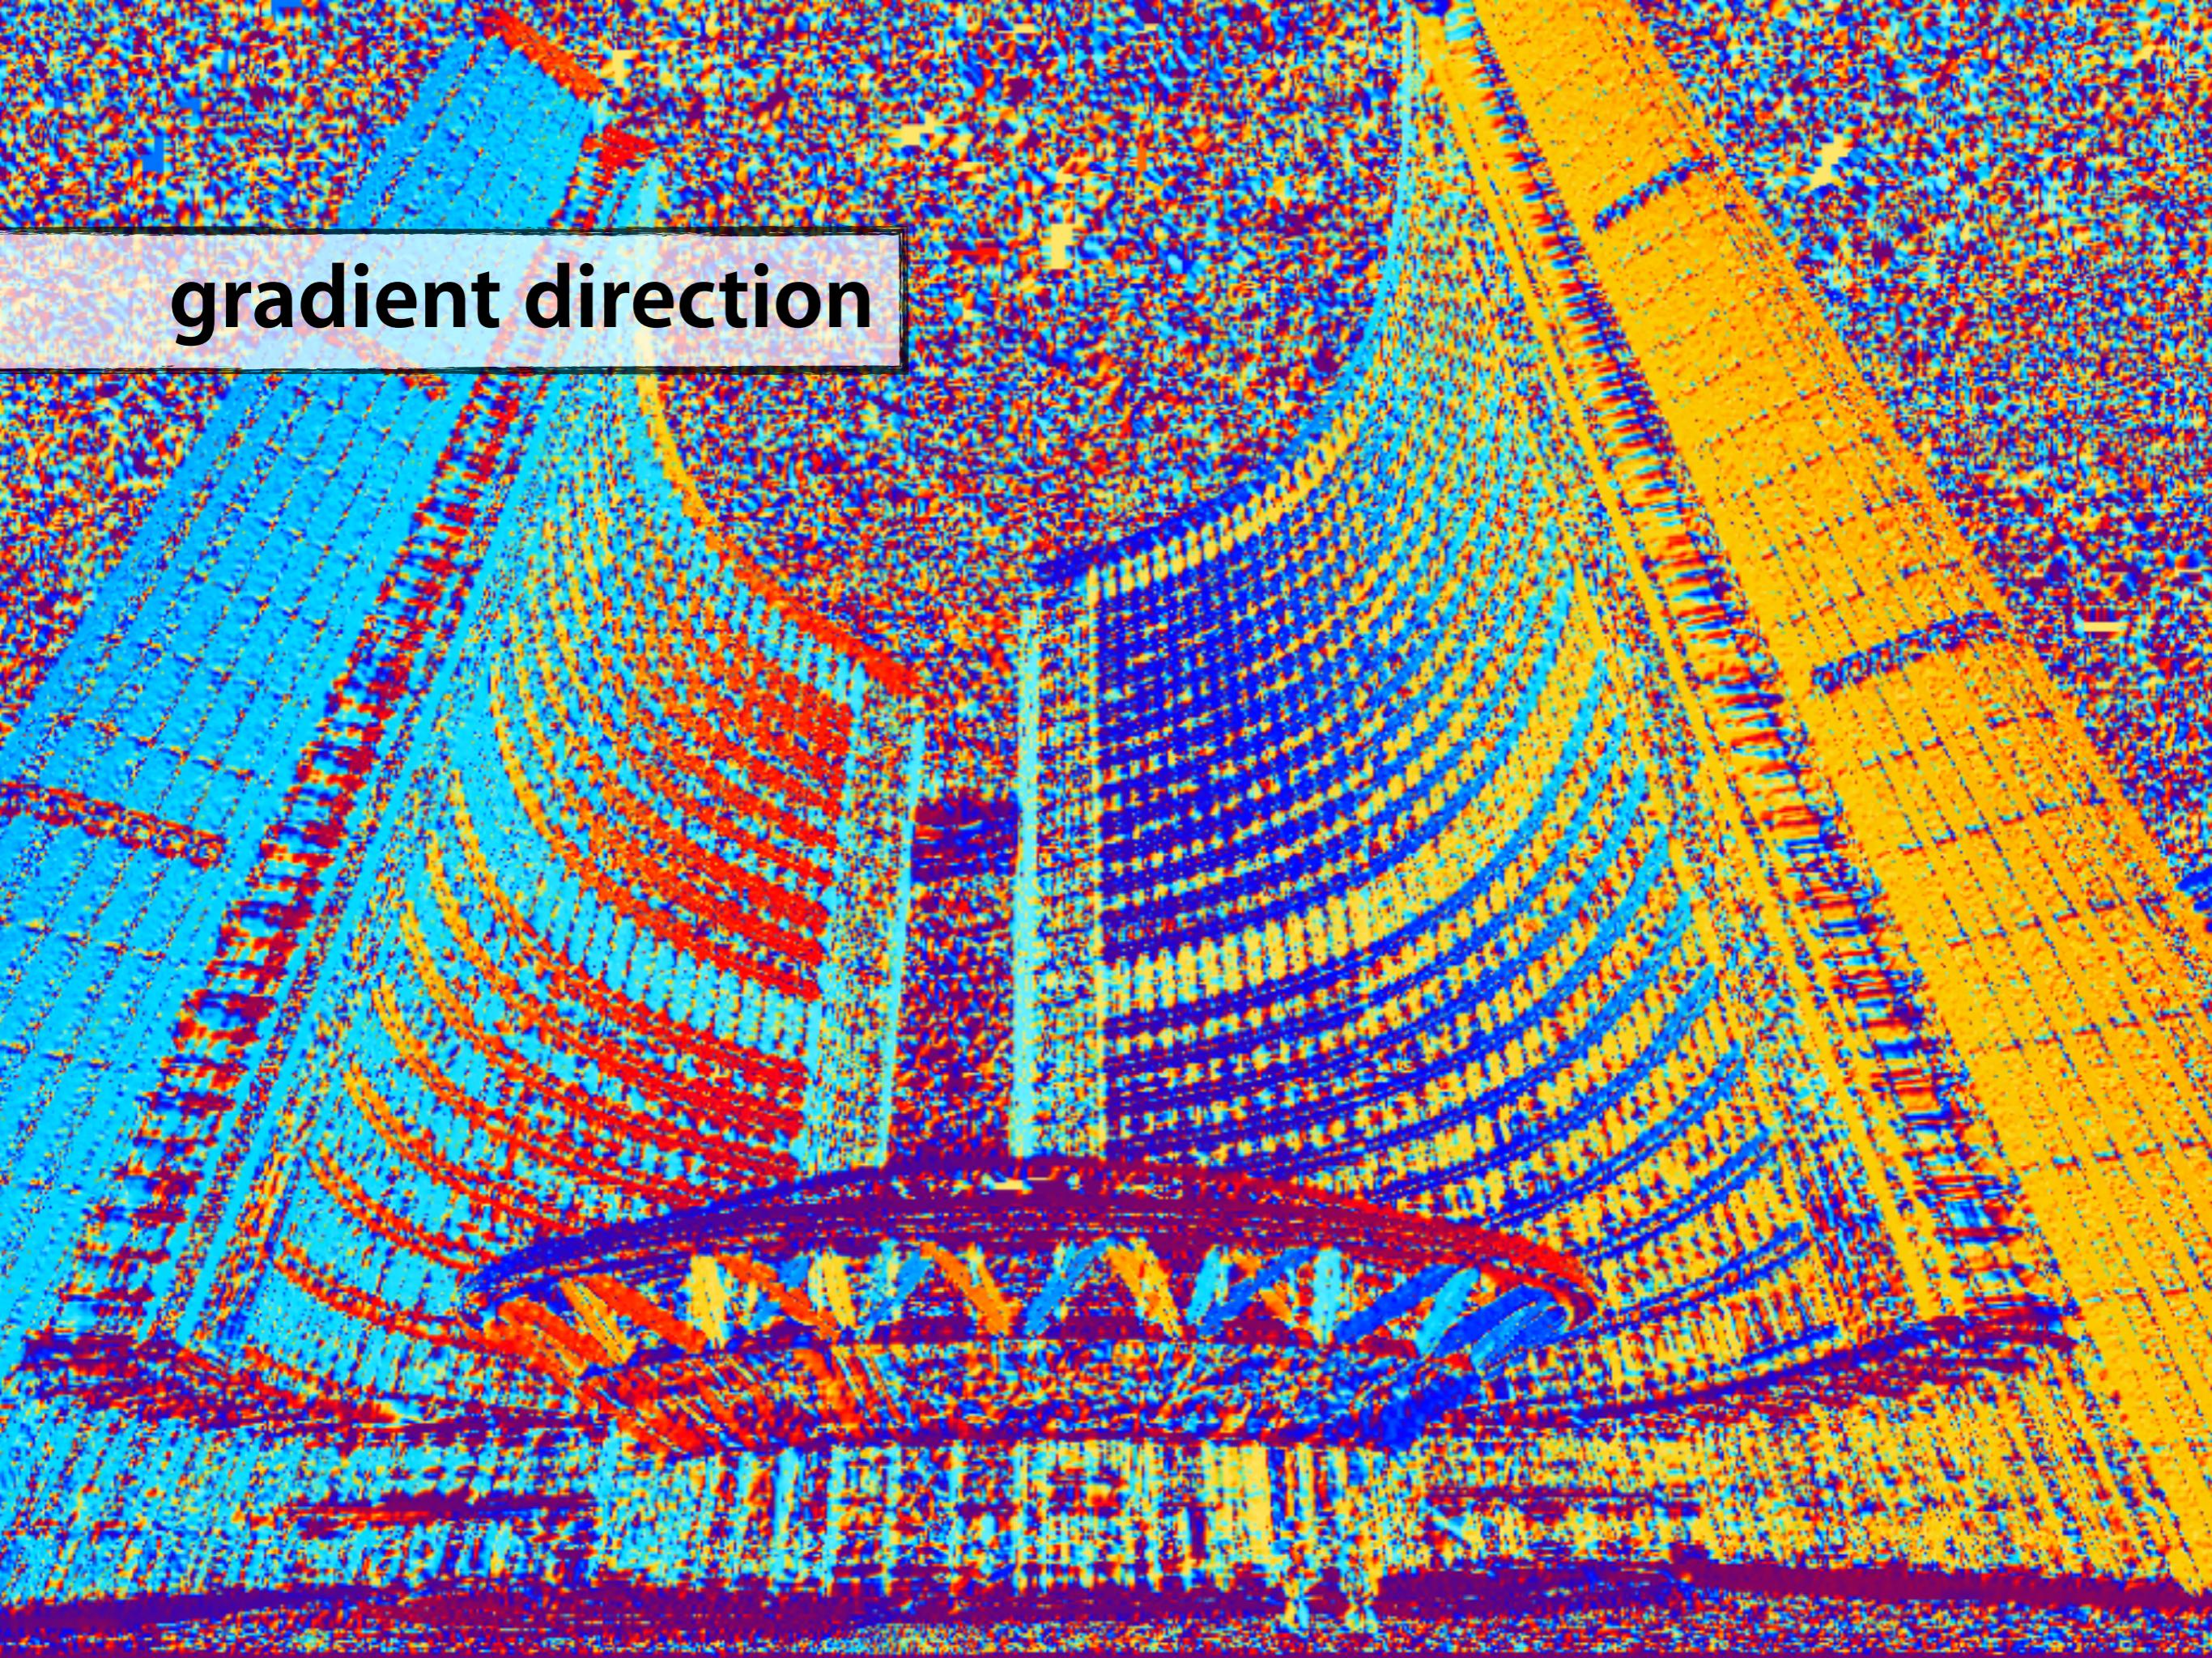
$$\nabla f = \left( \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)^\top$$

**Gradient magnitude**     $\|\nabla f\| = \sqrt{\frac{\partial f^2}{\partial x} + \frac{\partial f^2}{\partial y}}$

**Gradient direction**     $\theta = \tan^{-1} \left( \frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}} \right)$

A detailed black and white architectural rendering of a modern urban landscape. In the foreground, a long, low building features a series of arched windows along its side. A curved walkway or bridge connects this building to a larger, more complex structure in the background. This background structure consists of several tall, curved towers with intricate, grid-like patterns on their facades. The perspective is from a low angle, looking up at the towering buildings.

**gradient magnitude**



**gradient direction**



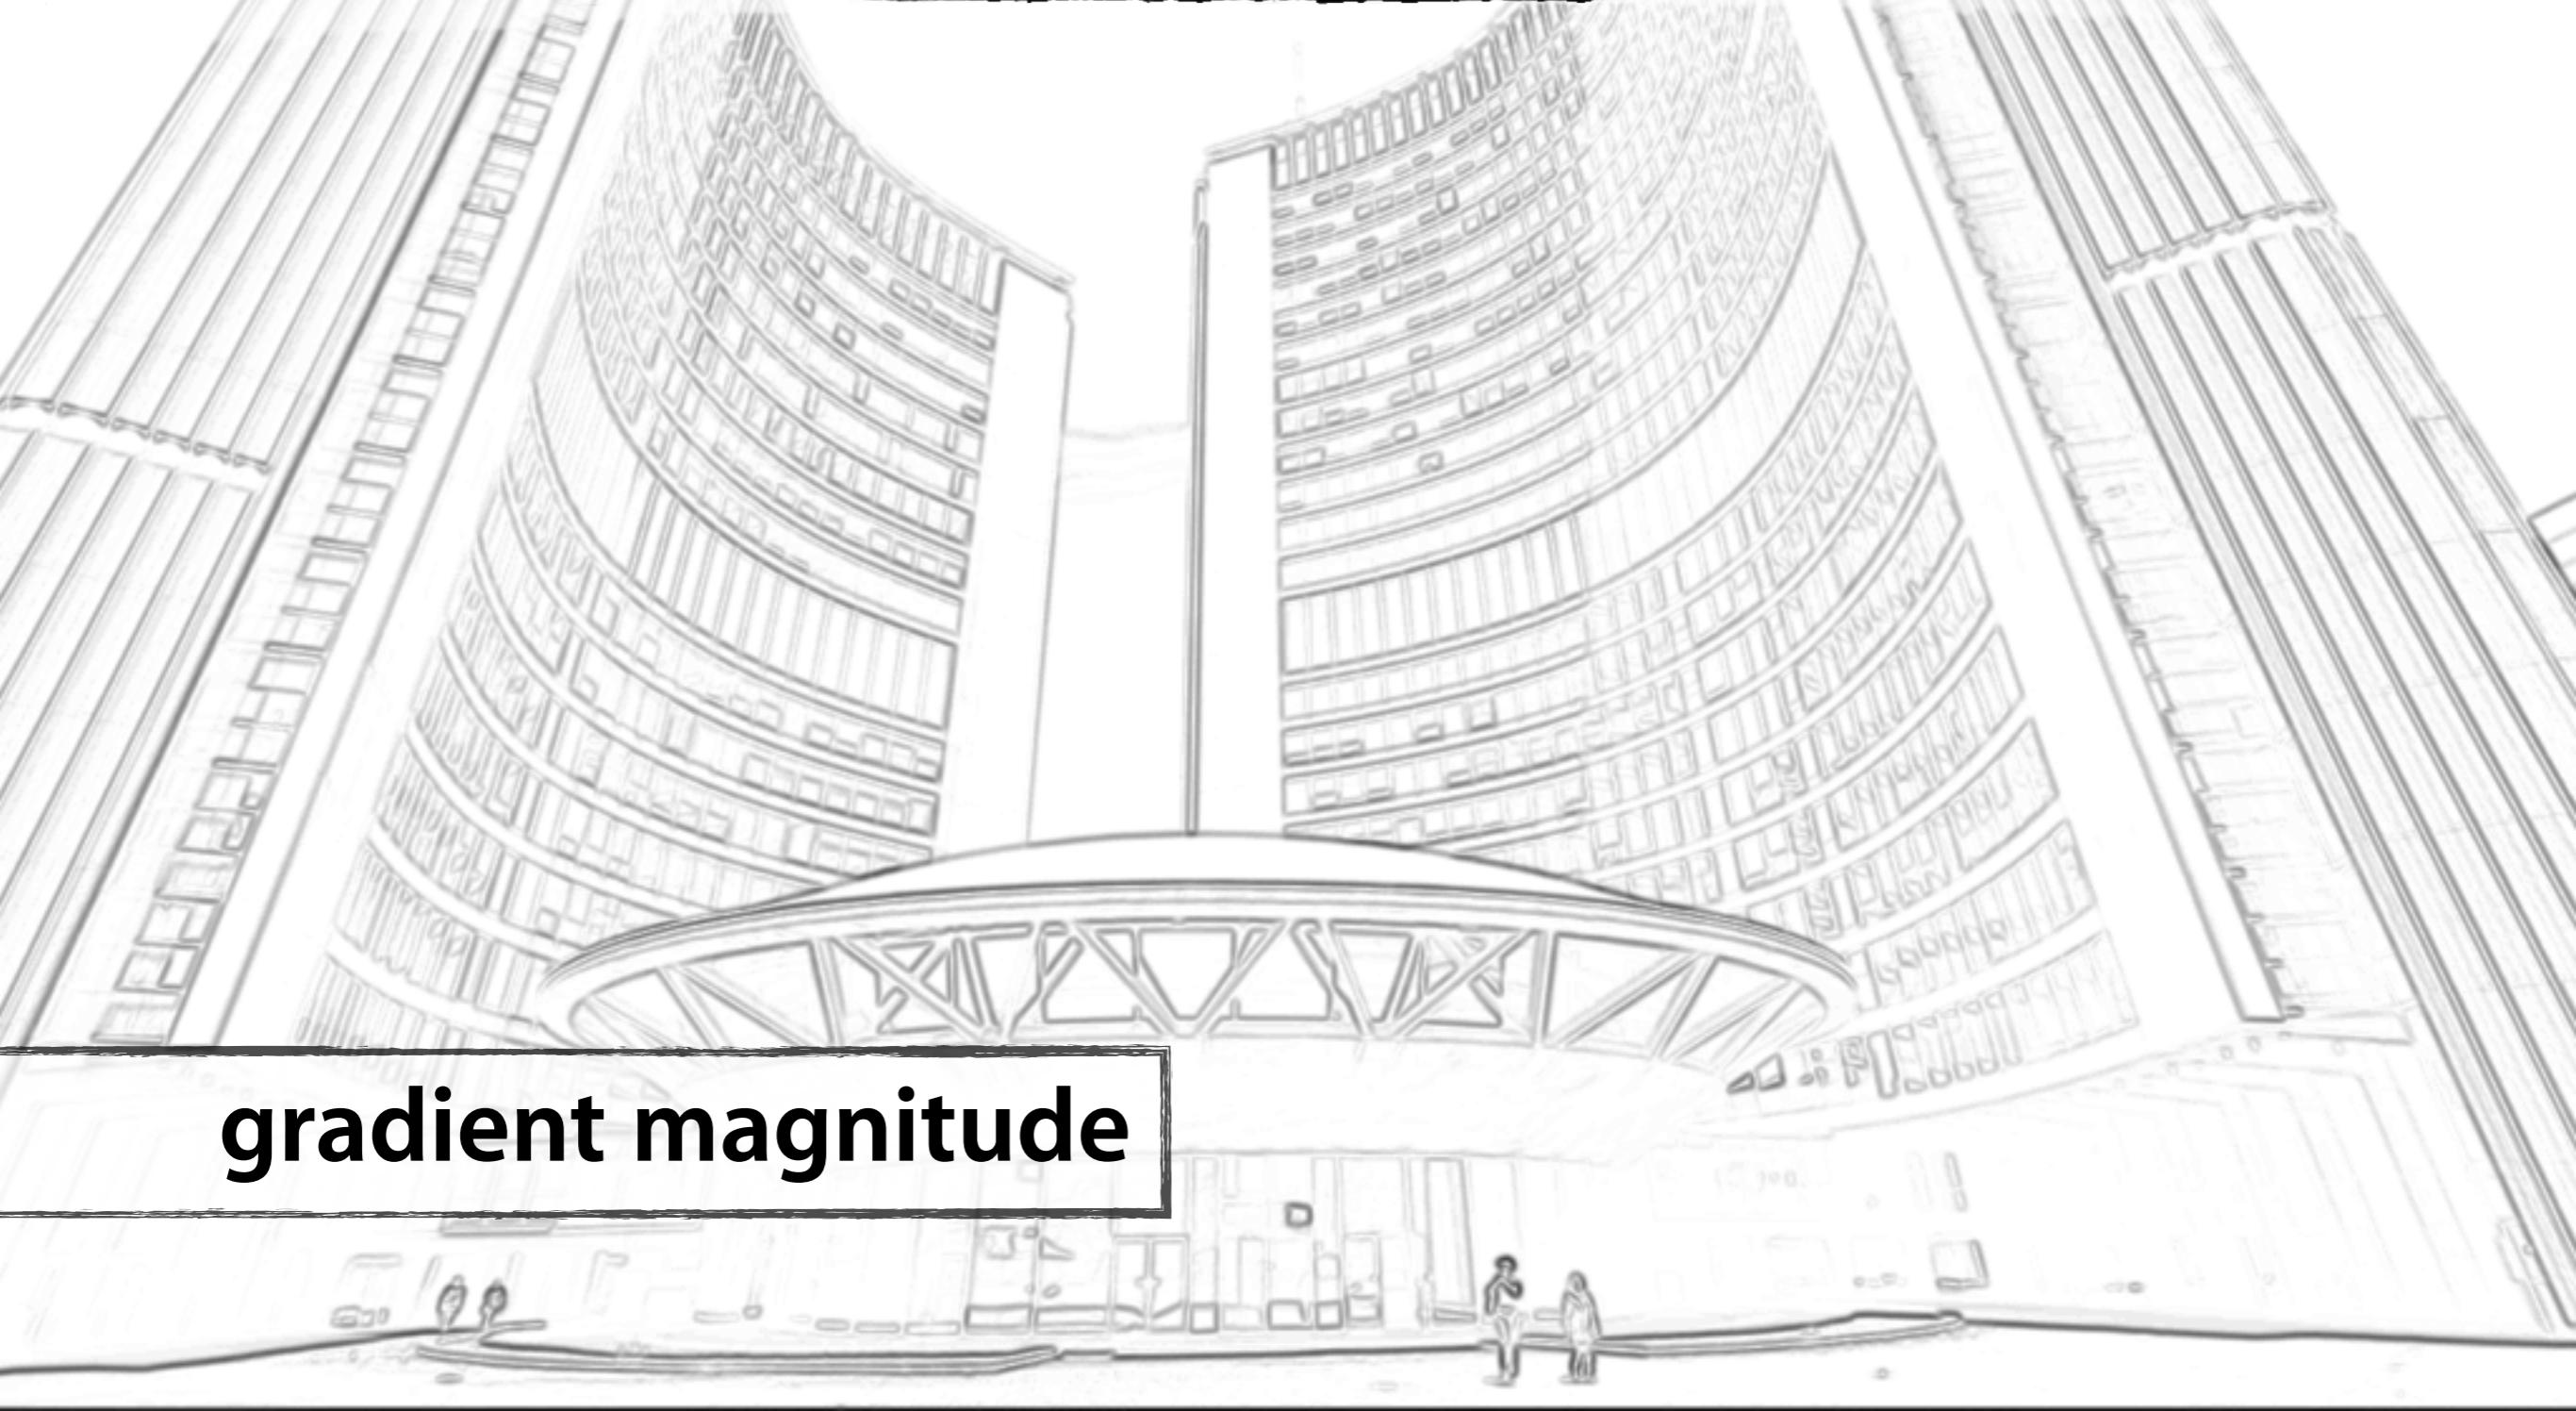
```
>> h = fspecial('sobel');
>> im_dy = imfilter(double(im), h, 'conv');
>> imshow(im_dy, [ ])
```



```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> imshow(im_dy, [ ])
```

correlation is default

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im), h', 'conv');  
>> imshow(sqrt(img_dx.^2 + img_dy.^2), [])
```



**gradient magnitude**

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im) h', 'conv');  
>> imshow(sqrt(img_dx.^2 + img_
```

transpose

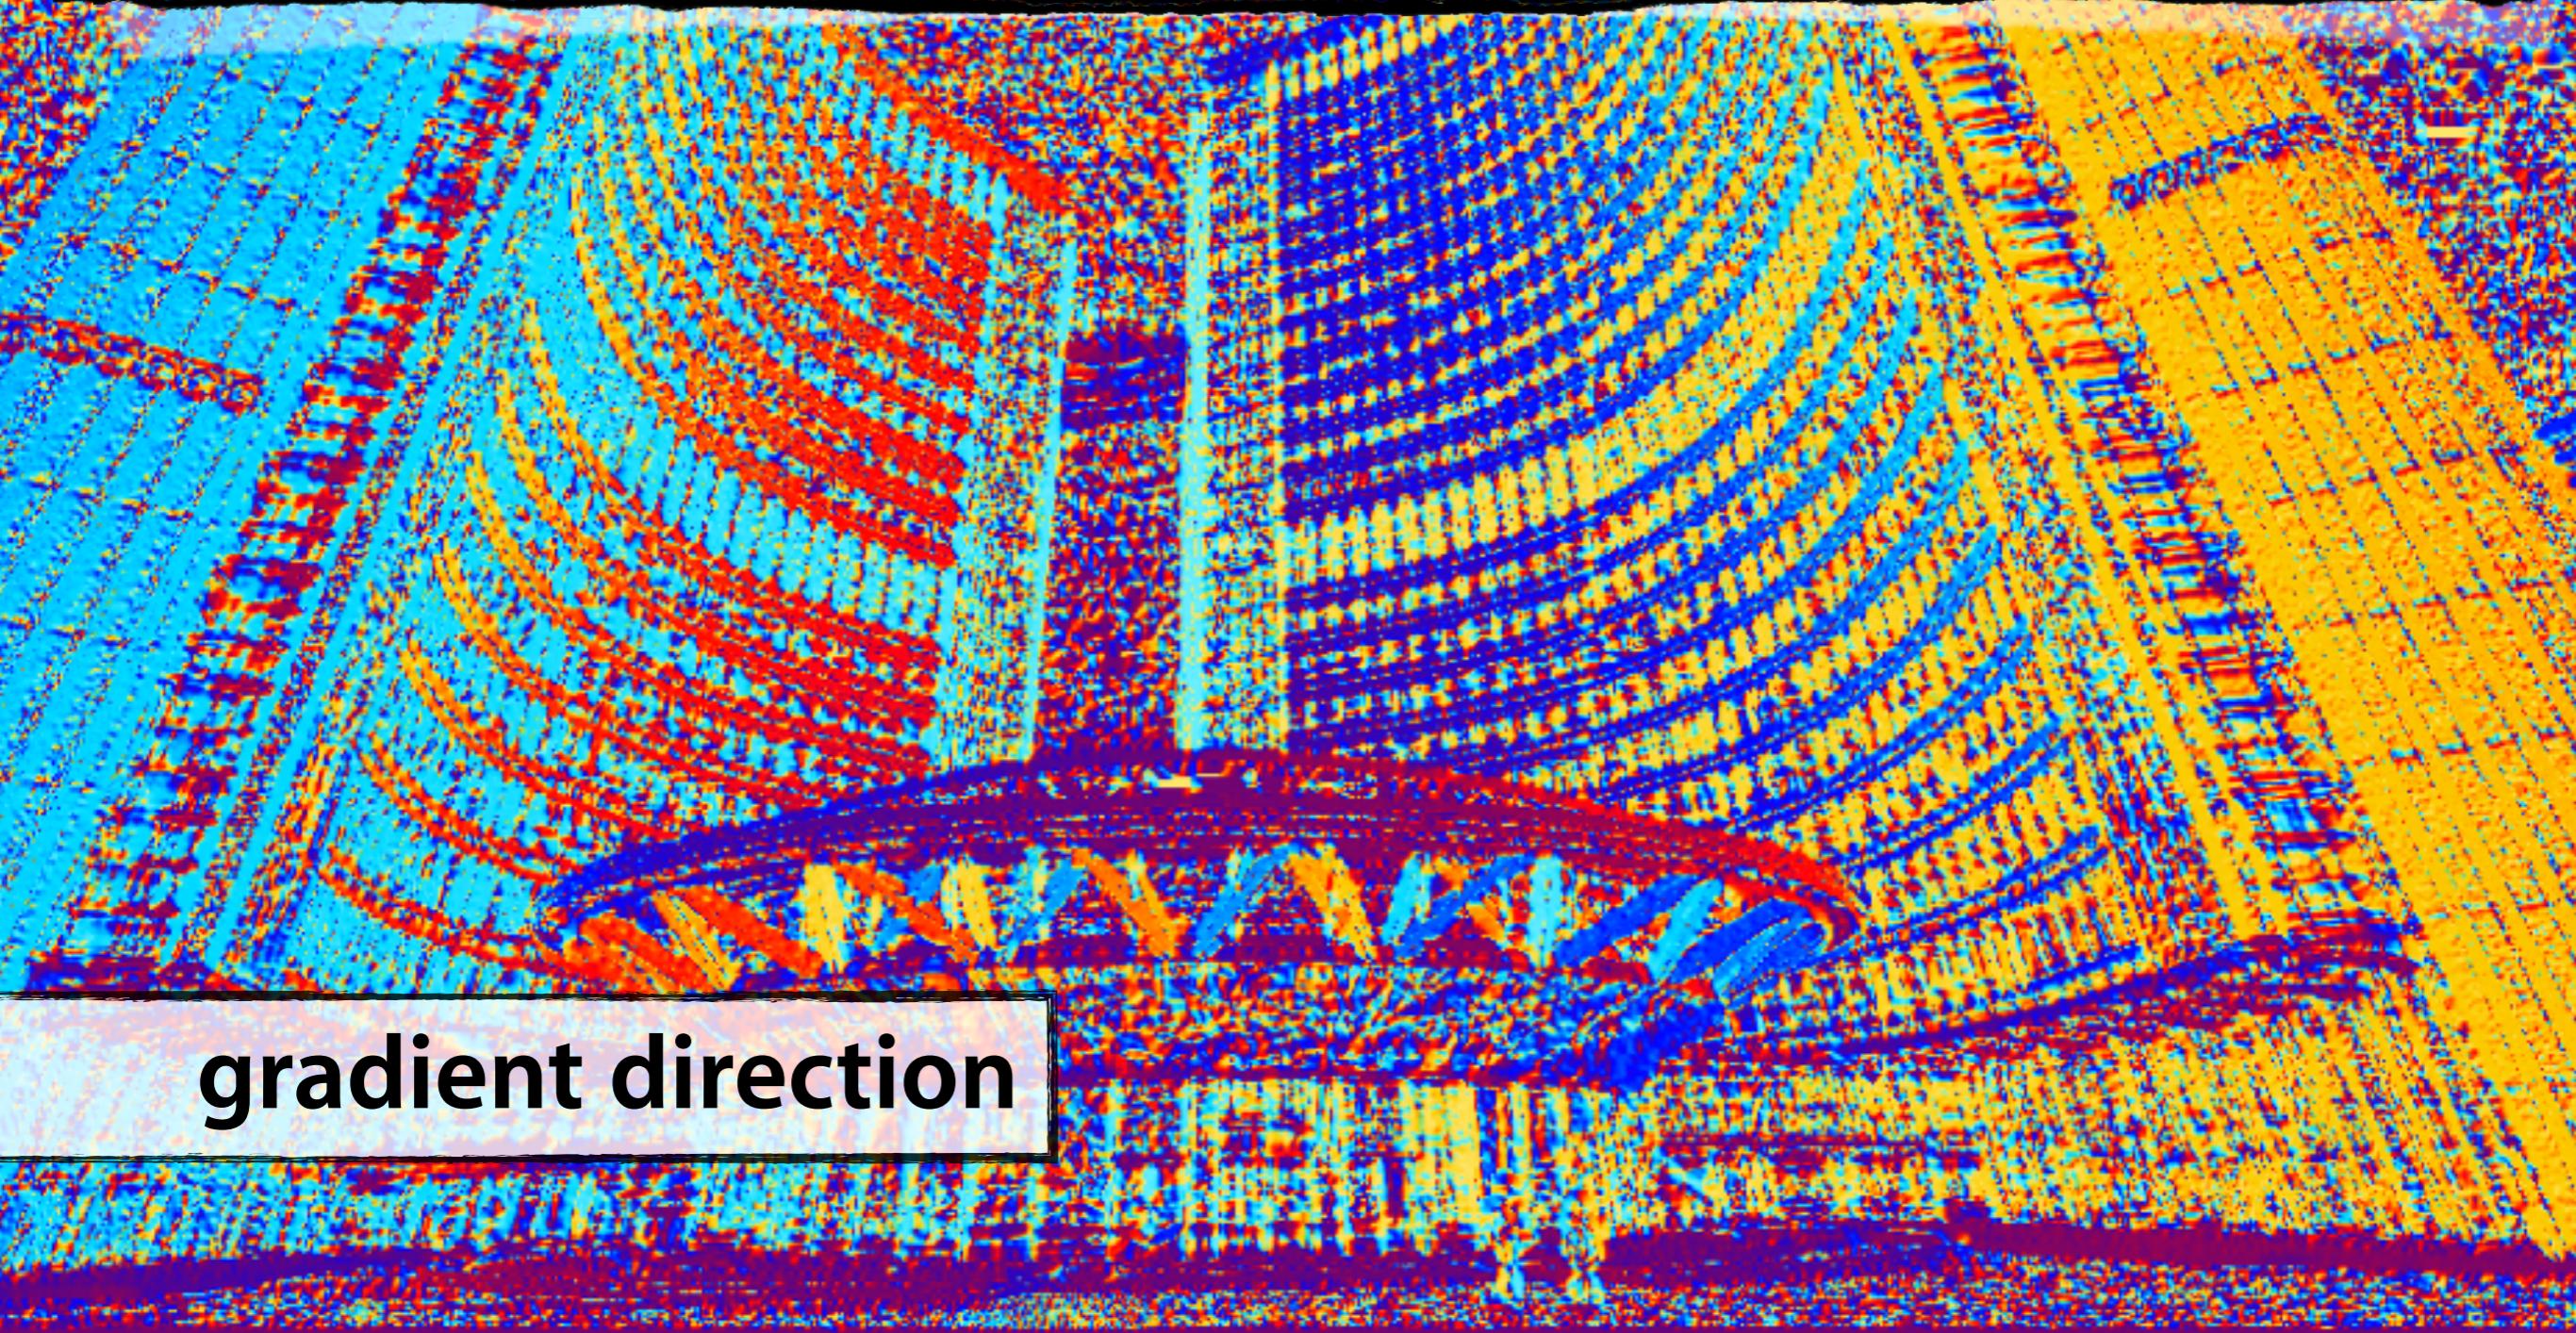
gradient magnitude

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im), h', 'conv');  
>> imshow(sqrt(img_dx.^2 + img_dy.^2), [])
```

**pointwise product**

**gradient magnitude**

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im), h', 'conv');  
>> imshow(atan2(im_dy, im_dx), [ ])  
>> colormap jet
```



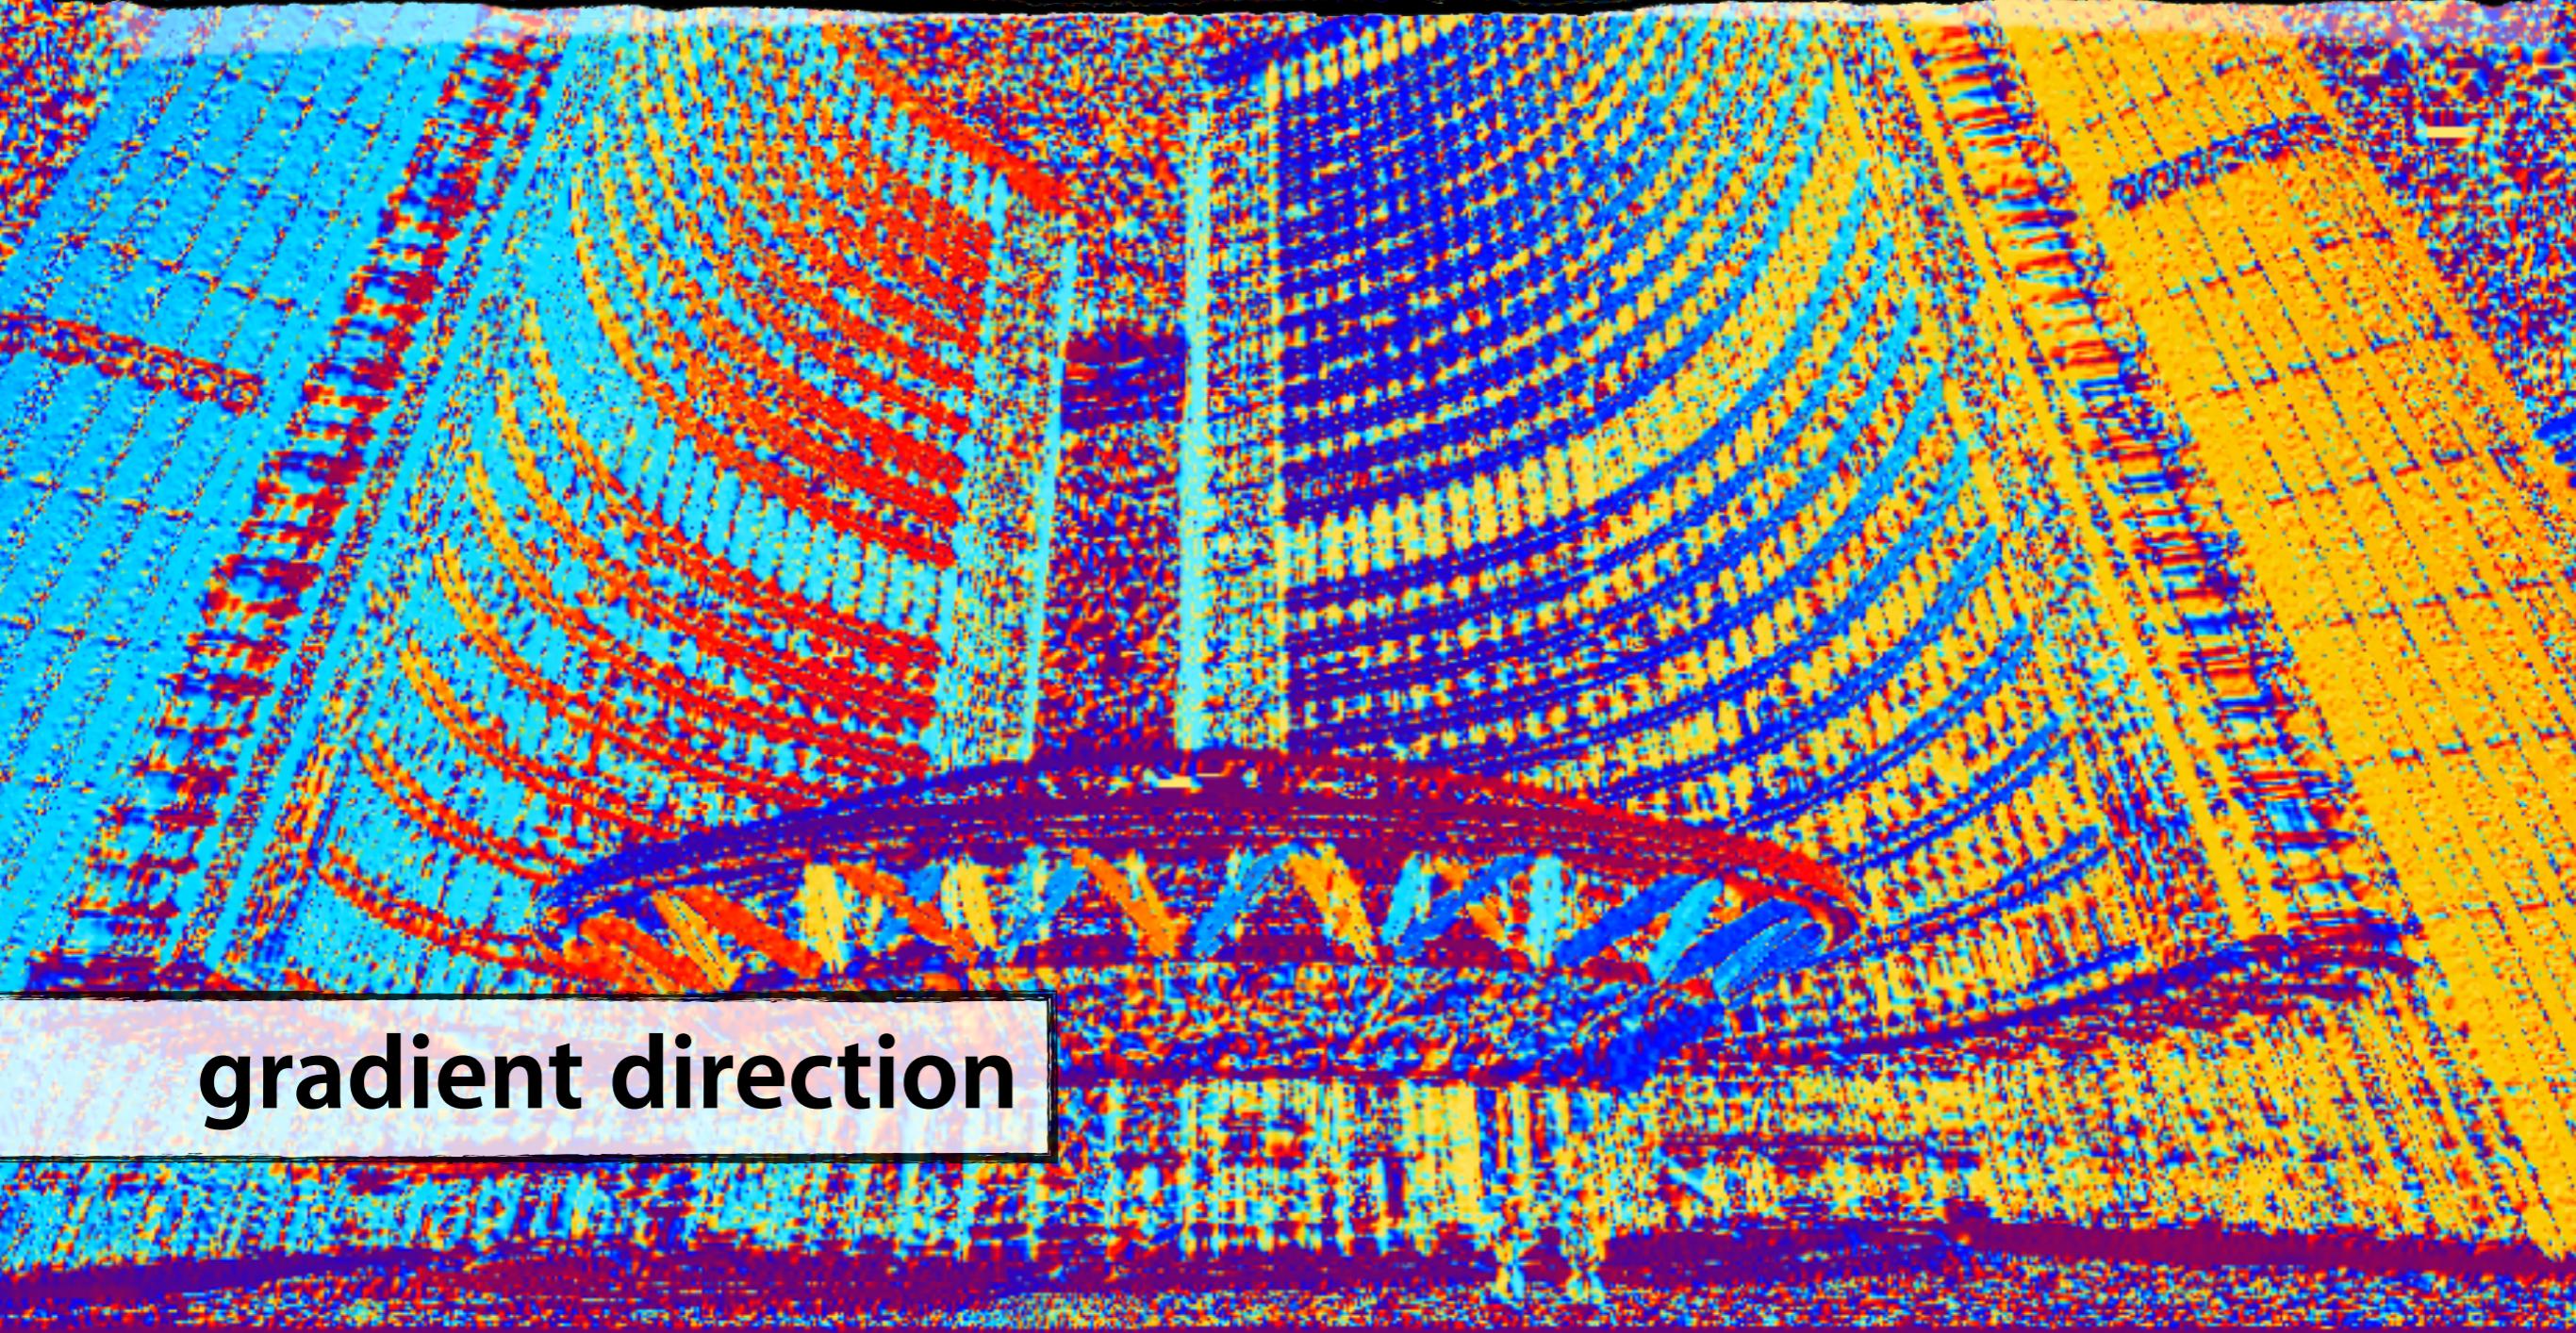
gradient direction

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im), h', 'conv');  
>> imshow(atan2(im_dy, im_dx), [ ])  
>> colormap jet
```

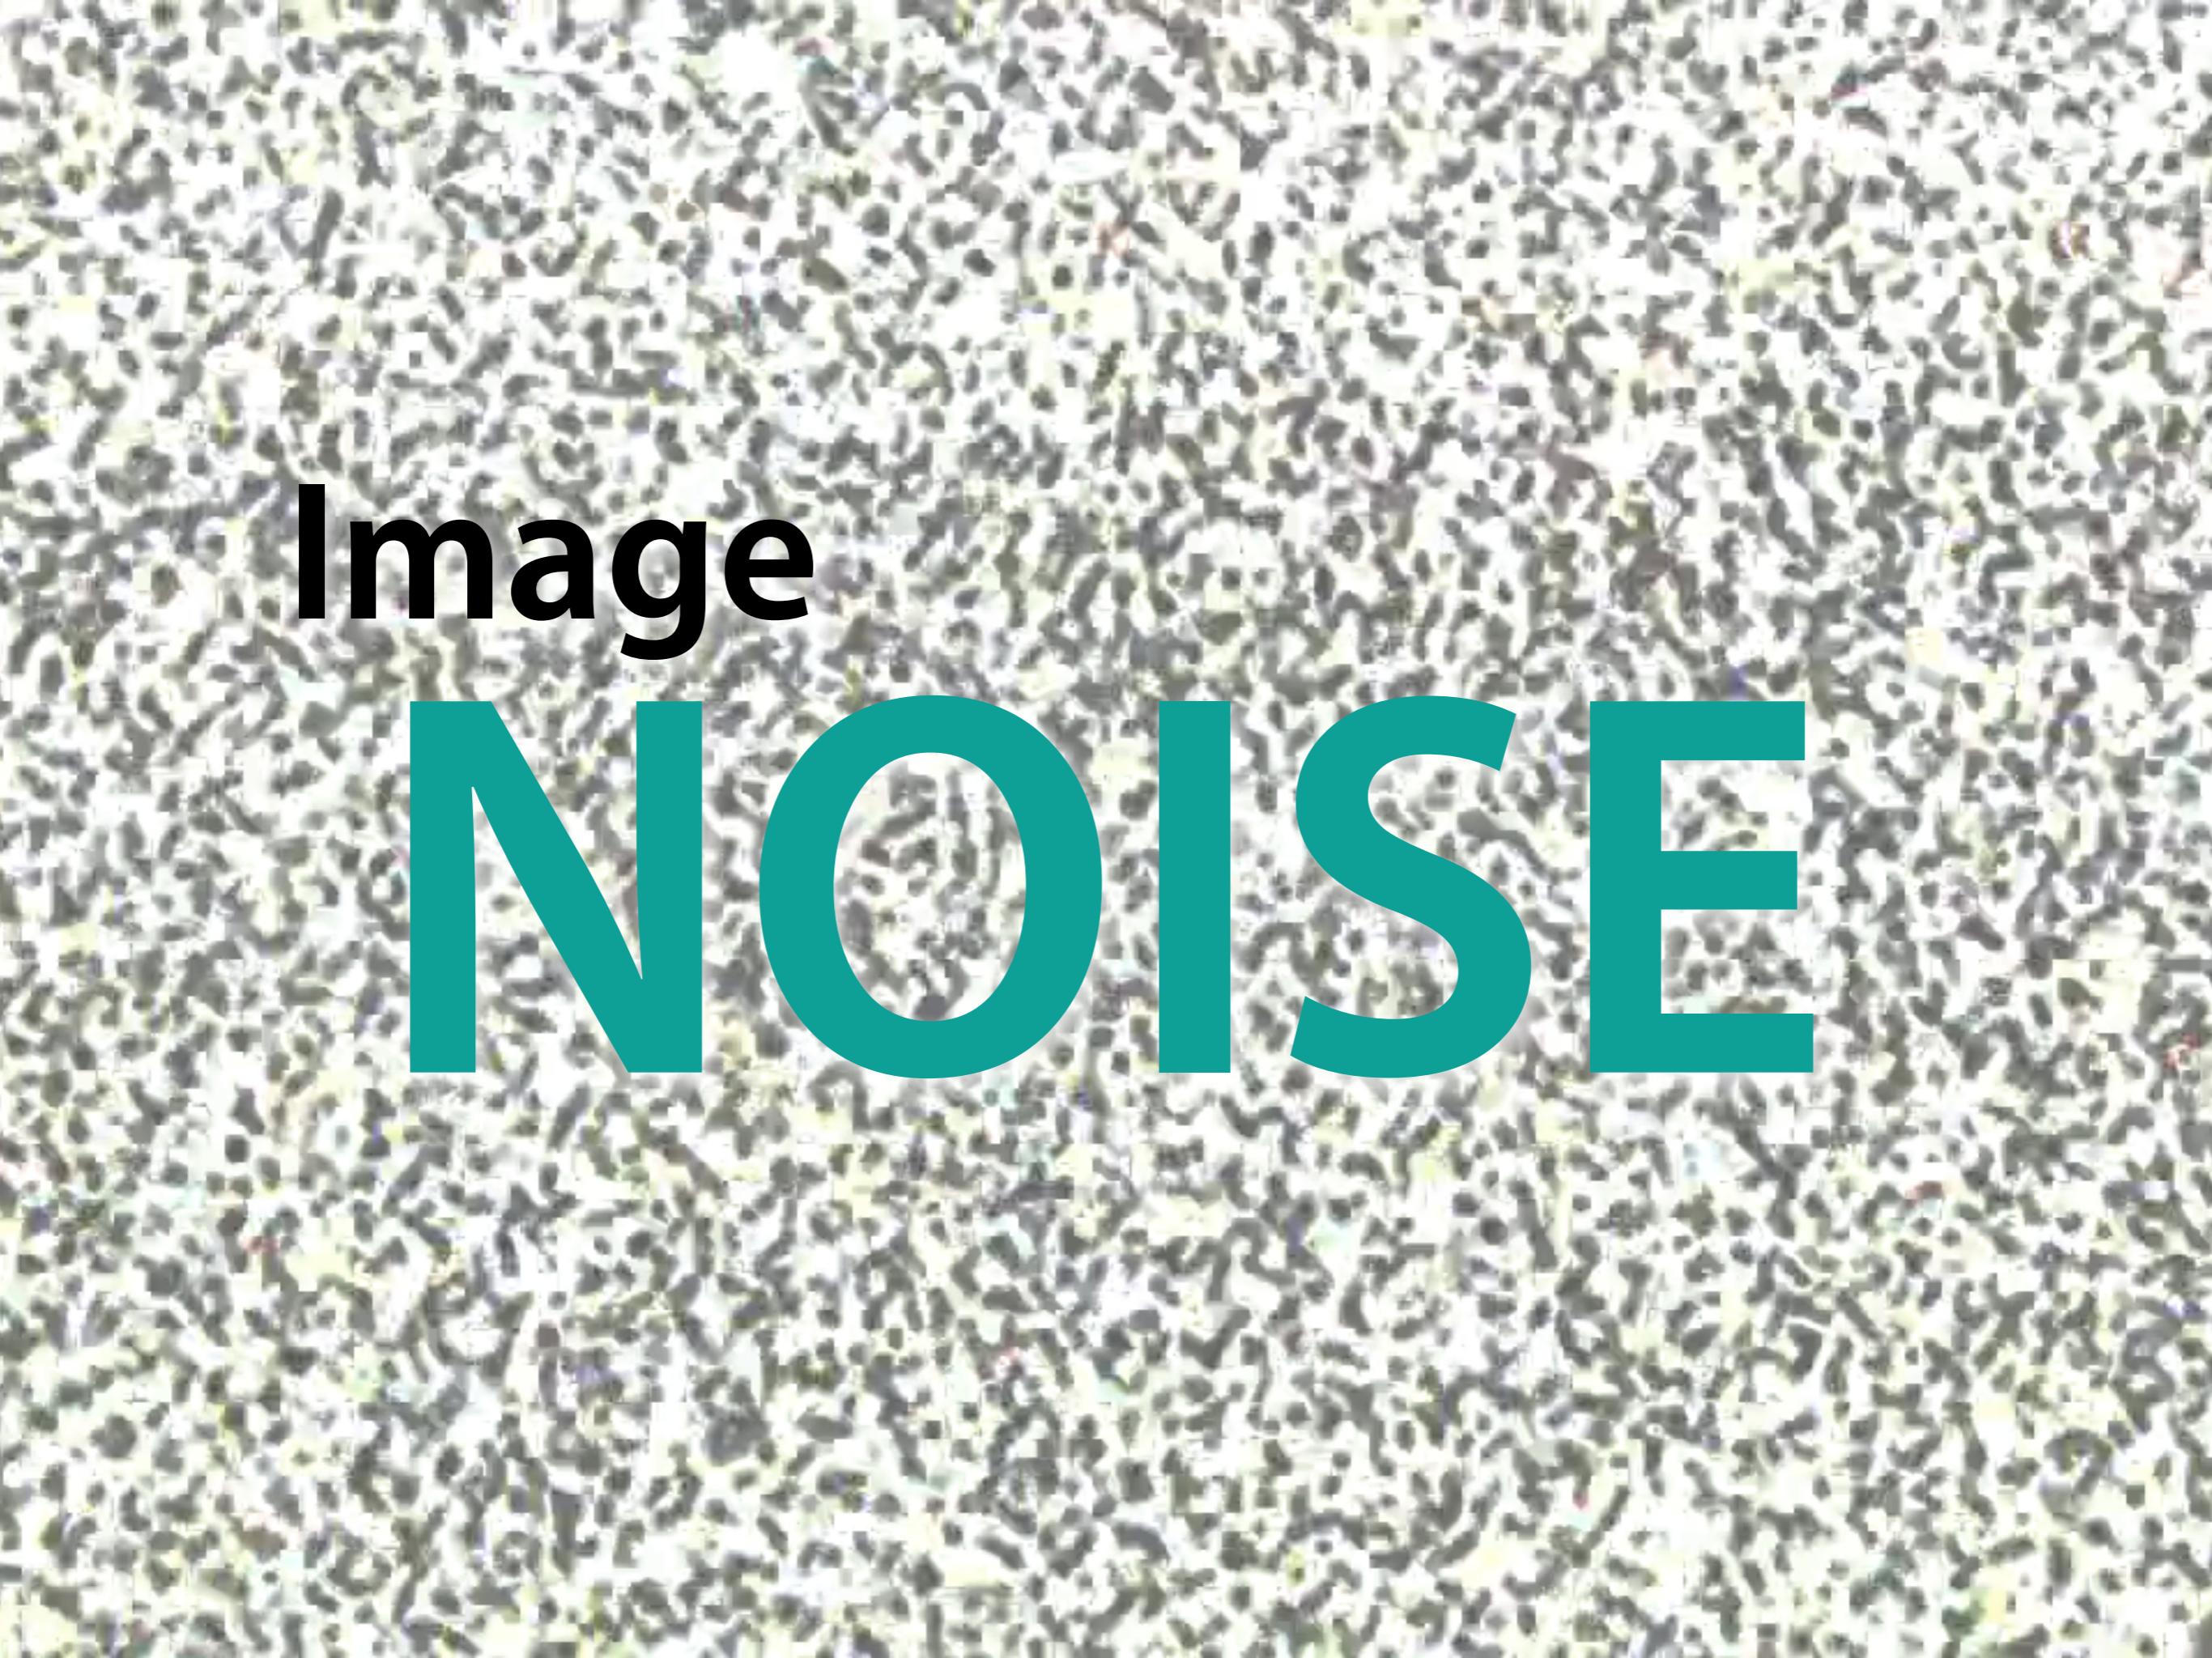
**change output colour space**

**gradient direction**

```
>> h = fspecial('sobel');  
>> im_dy = imfilter(double(im), h, 'conv');  
>> im_dx = imfilter(double(im), h', 'conv');  
>> imshow(atan2(im_dy, im_dx), [ ])  
>> colormap jet
```



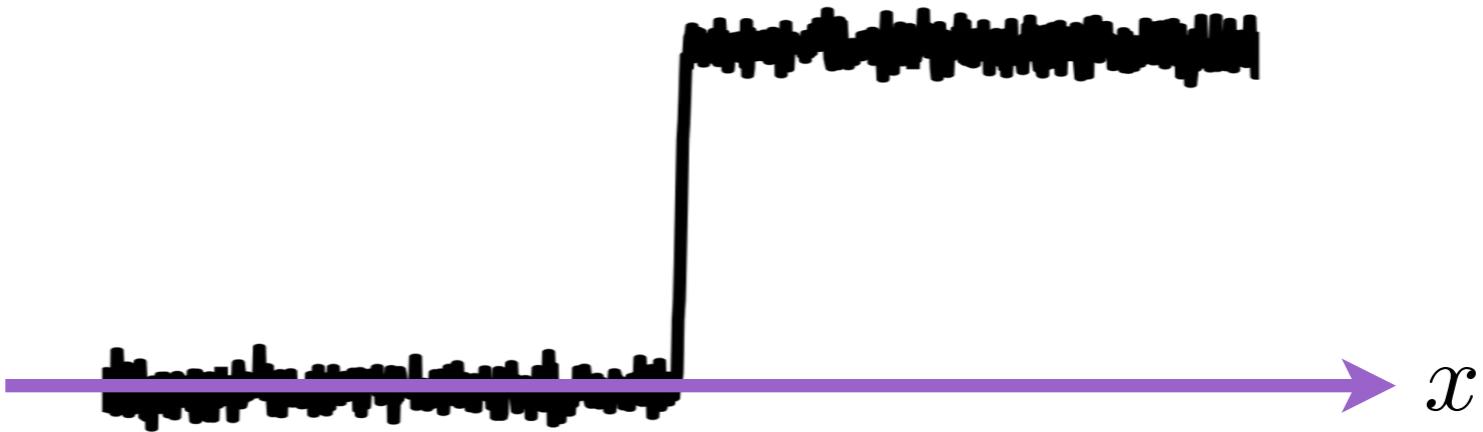
gradient direction



A high-contrast, black-and-white image showing a dense, noisy texture. The noise is grainy and irregular, resembling a heavily overexposed photograph or a signal with high levels of electronic noise. There are faint, darker shapes that suggest the presence of trees and possibly a building in the background, but they are completely obscured by the noise.

Image  
**NOISE**

$$f(x)$$

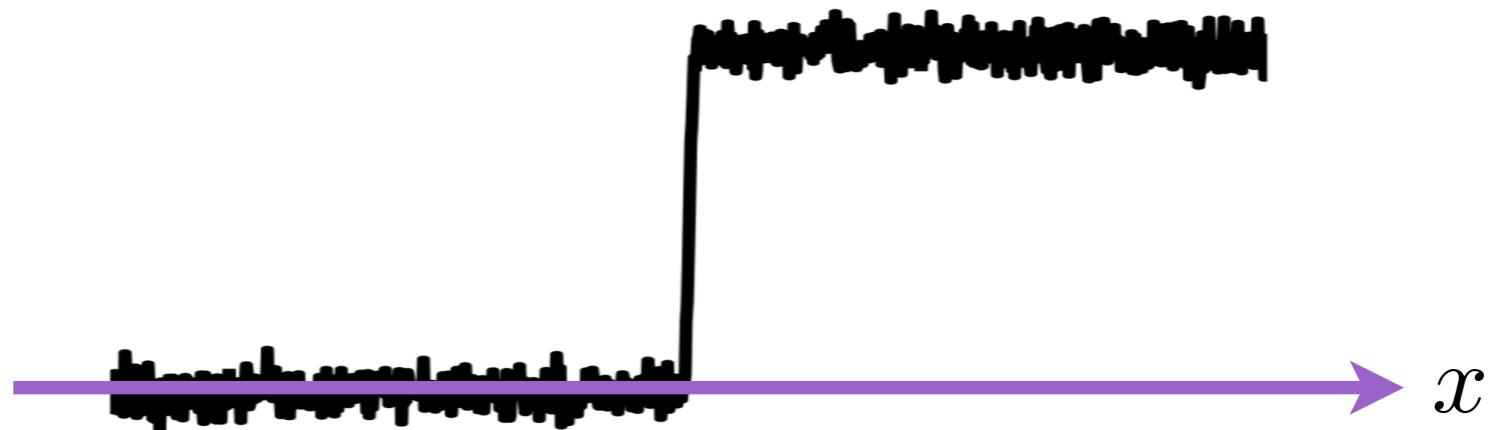


$$\frac{d}{dx} f(x)$$



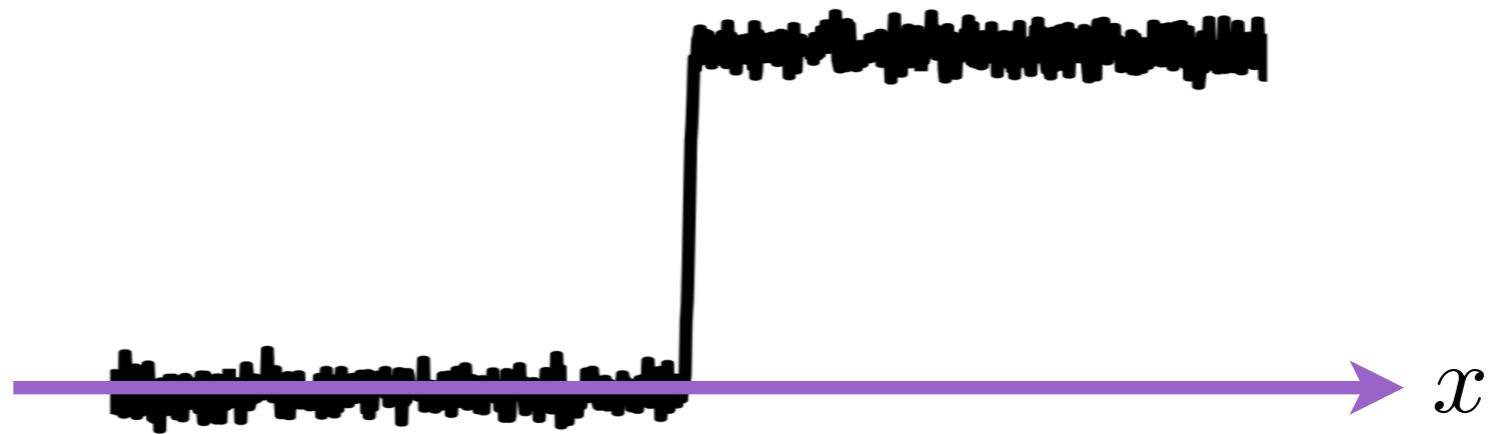
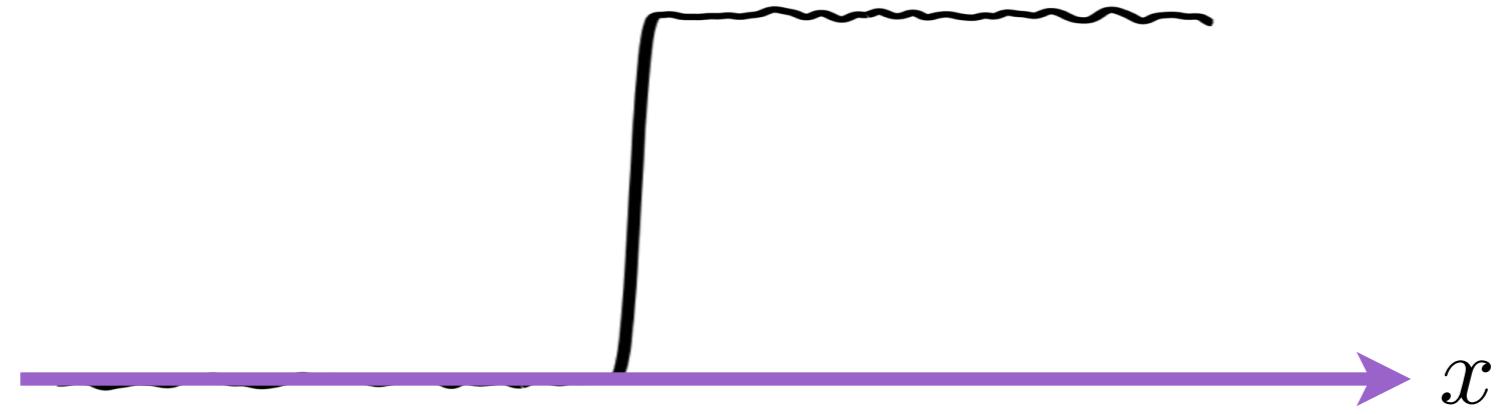
Noise is **AMPLIFIED!**

$f(x)$

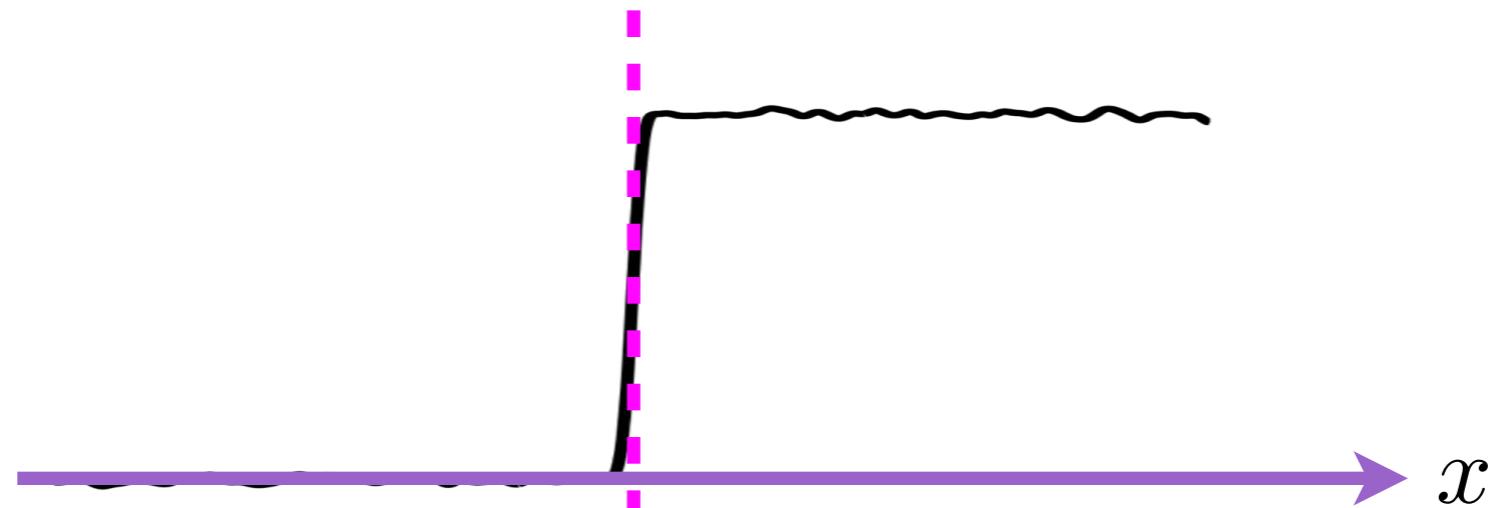


**How do we avoid the issues caused by noise?**

**smooth the image**

$f(x)$  $g(x)$  $f * g$ 

$$f * g$$



$$\frac{d}{dx}(f * g)$$



**What about 2D?**

*discrete  
convolution*

$$H[x, y] = \sum_{u=-K}^{K} \sum_{v=-K}^{K} G[u, v] F[x - u, y - v]$$

continuous  
convolution

$$H(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} G(u, v) F(x - u, y - v) du dv$$

$$\frac{\partial}{\partial x} (g * I)$$

Gaussian filter

$$\frac{\partial}{\partial x} (g * I)$$

image

derivative of Gaussian

$$\frac{\partial}{\partial x} (g * I) \equiv \frac{\partial g}{\partial x} * I$$

$$\frac{\partial}{\partial x} (g * I) \equiv \frac{\partial g}{\partial x} * I$$

**derivative theorem of convolution**

$$\frac{\partial}{\partial x}(g * I) \equiv \frac{\partial g}{\partial x} * I$$

**How do we perform discrete filtering?**

$$\frac{\partial}{\partial x}(g * I) \equiv \frac{\partial g}{\partial x} * I$$

Solution  
1

1	-1
---	----

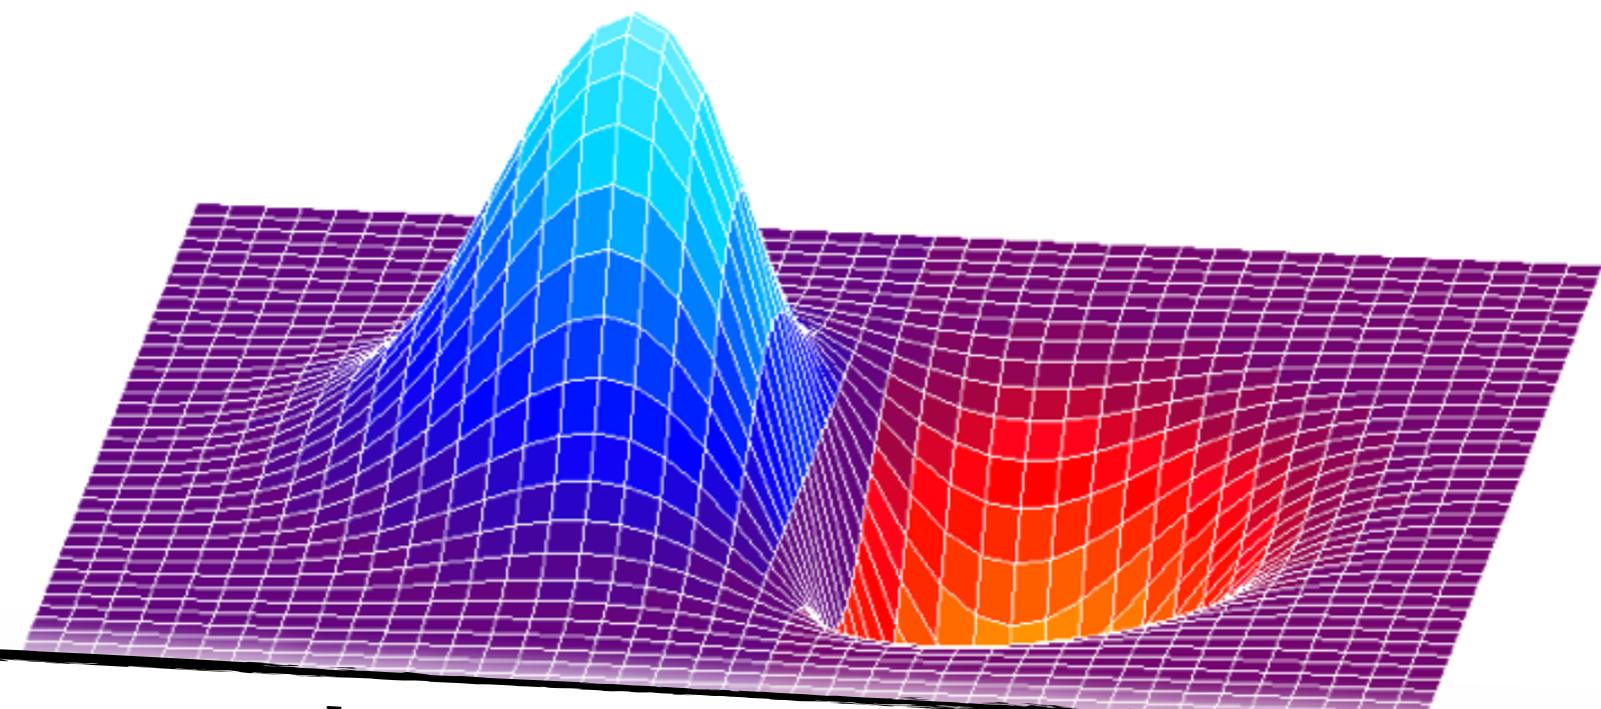
\*

0.0113	0.0838	0.0113
0.0838	0.6193	0.0838
0.0113	0.0838	0.0113

## Solution 2

$$\frac{\partial}{\partial x} (g * I) \equiv \frac{\partial g}{\partial x} * I$$

$$\frac{\partial}{\partial x} g(x, y; \sigma) = -\frac{x}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$



**sampled derivative of Gaussian**

# **Steerable Filters**

# steerable filter

filter class in which a filter of arbitrary orientation  
can be expressed via a linear combination of a  
“basis” set of filters

# basis

A set of vectors in a vector space  $V$  is called a **basis** if every element in  $V$  can be written in a unique way as a finite linear combination of elements of  $B$ .

## Definition:

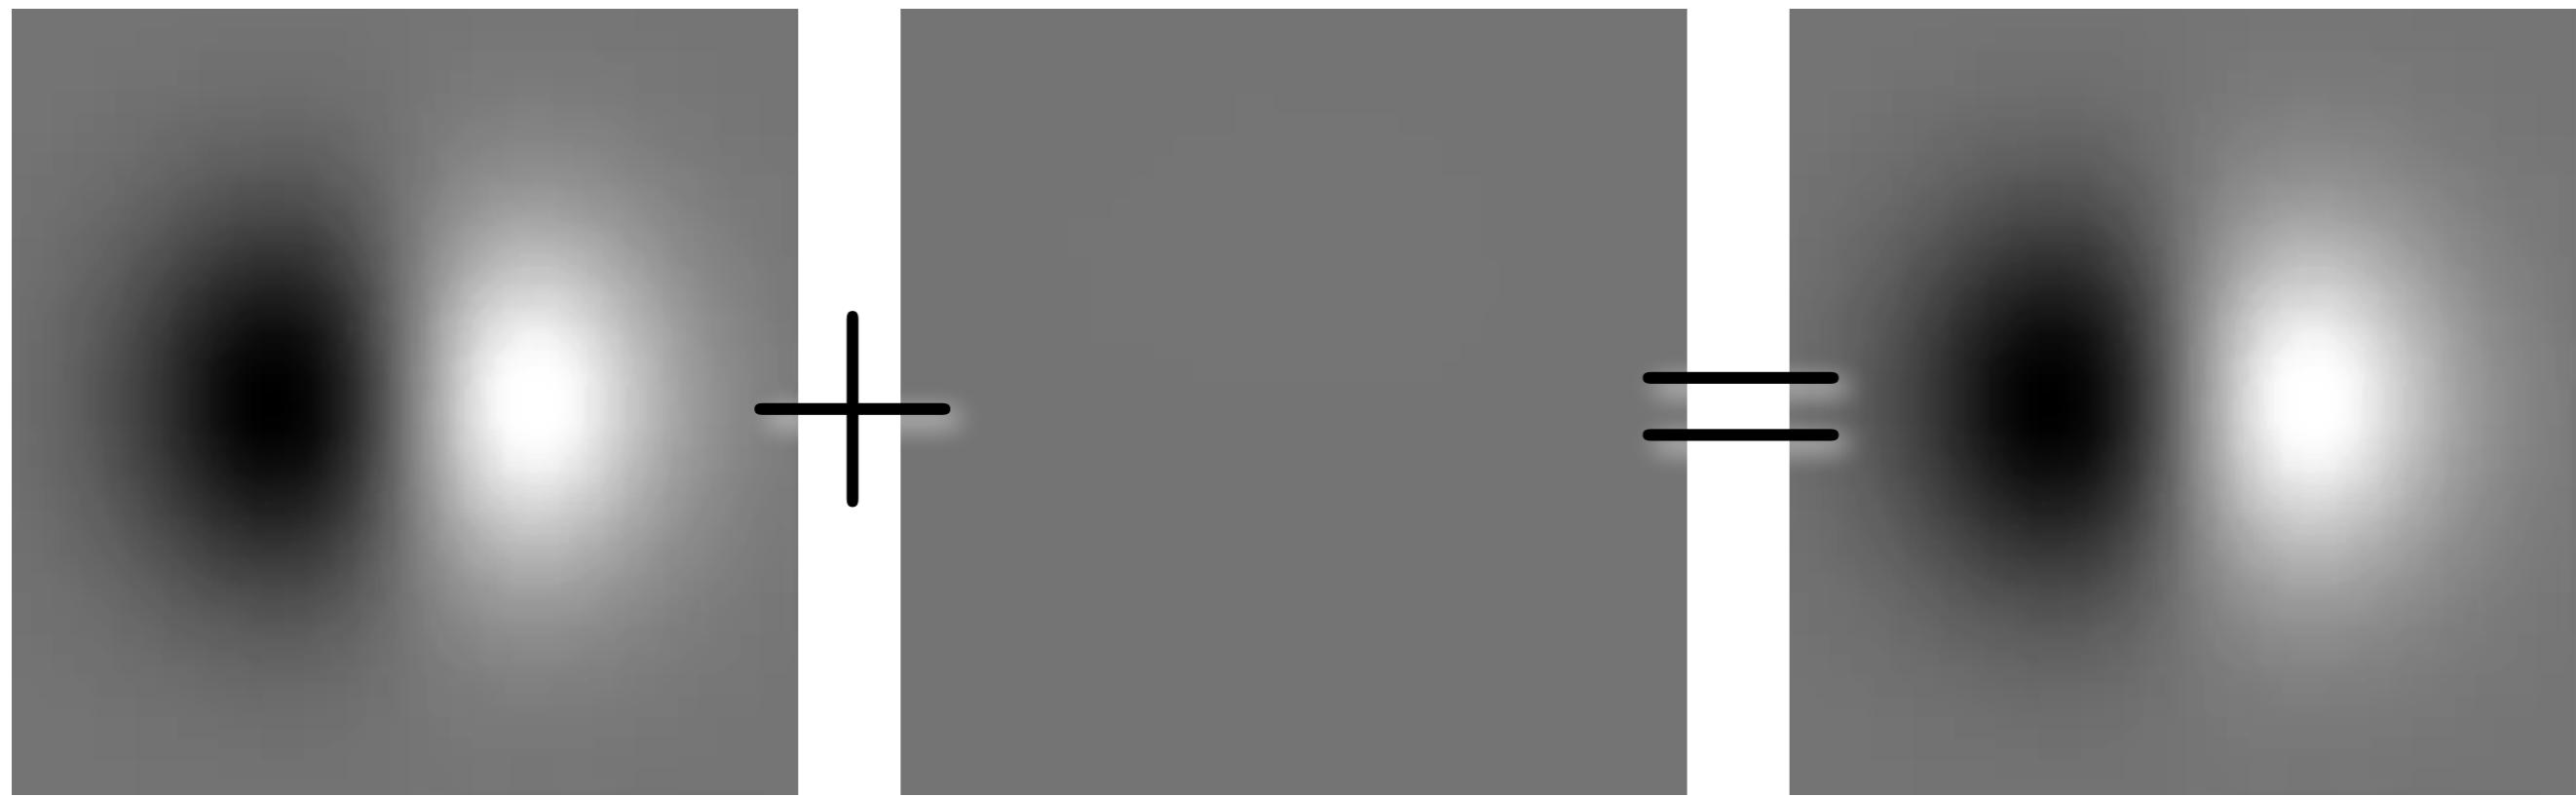
The **directional derivative**,  $\nabla_{\mathbf{u}} f(\mathbf{x}_0)$ , is the rate that a function  $f(\mathbf{x}_0)$  changes at a point  $\mathbf{x}_0$  in the direction  $\hat{\mathbf{u}}$ . It can be defined as:

## Definition:

The **directional derivative**,  $\nabla_{\mathbf{u}} f(\mathbf{x}_0)$ , is the rate that a function  $f(\mathbf{x}_0)$  changes at a point  $\mathbf{x}_0$  in the direction  $\hat{\mathbf{u}}$ . It can be defined as:

$$\begin{aligned}\nabla_{\hat{\mathbf{u}}} f(\mathbf{x}_0) &= \lim_{h \rightarrow 0} \frac{f(\mathbf{x}_0 + h\hat{\mathbf{u}}) - f(\mathbf{x}_0)}{h} \\ &= \nabla f(\mathbf{x}_0) \cdot \hat{\mathbf{u}} \\ &= \left( \frac{\partial f(\mathbf{x}_0)}{\partial x}, \frac{\partial f(\mathbf{x}_0)}{\partial y} \right)^{\top} \cdot (u_0, u_1)^{\top} \\ &= u_0 \frac{\partial f(\mathbf{x}_0)}{\partial x} + u_1 \frac{\partial f(\mathbf{x}_0)}{\partial y}\end{aligned}$$

$$\cos(\theta)\frac{\partial G(x,y)}{\partial x} \; + \; \sin(\theta)\frac{\partial G(x,y)}{\partial y} \; = \; \frac{\partial G(x,y)}{\partial\theta}$$



$$\cos(\theta) \frac{\partial G(x, y)}{\partial x}$$

$$\sin(\theta) \frac{\partial G(x, y)}{\partial y}$$

$$\frac{\partial G(x, y)}{\partial \theta}$$

$$I_\theta(x, y) = [\cos(\theta) \frac{\partial G(x, y)}{\partial x} + \sin(\theta) \frac{\partial G(x, y)}{\partial y}] * I(x, y)$$

**Can we compute this more efficiently?**

*distributive*

$$E[x, y] * (F[x, y] + G[x, y])$$

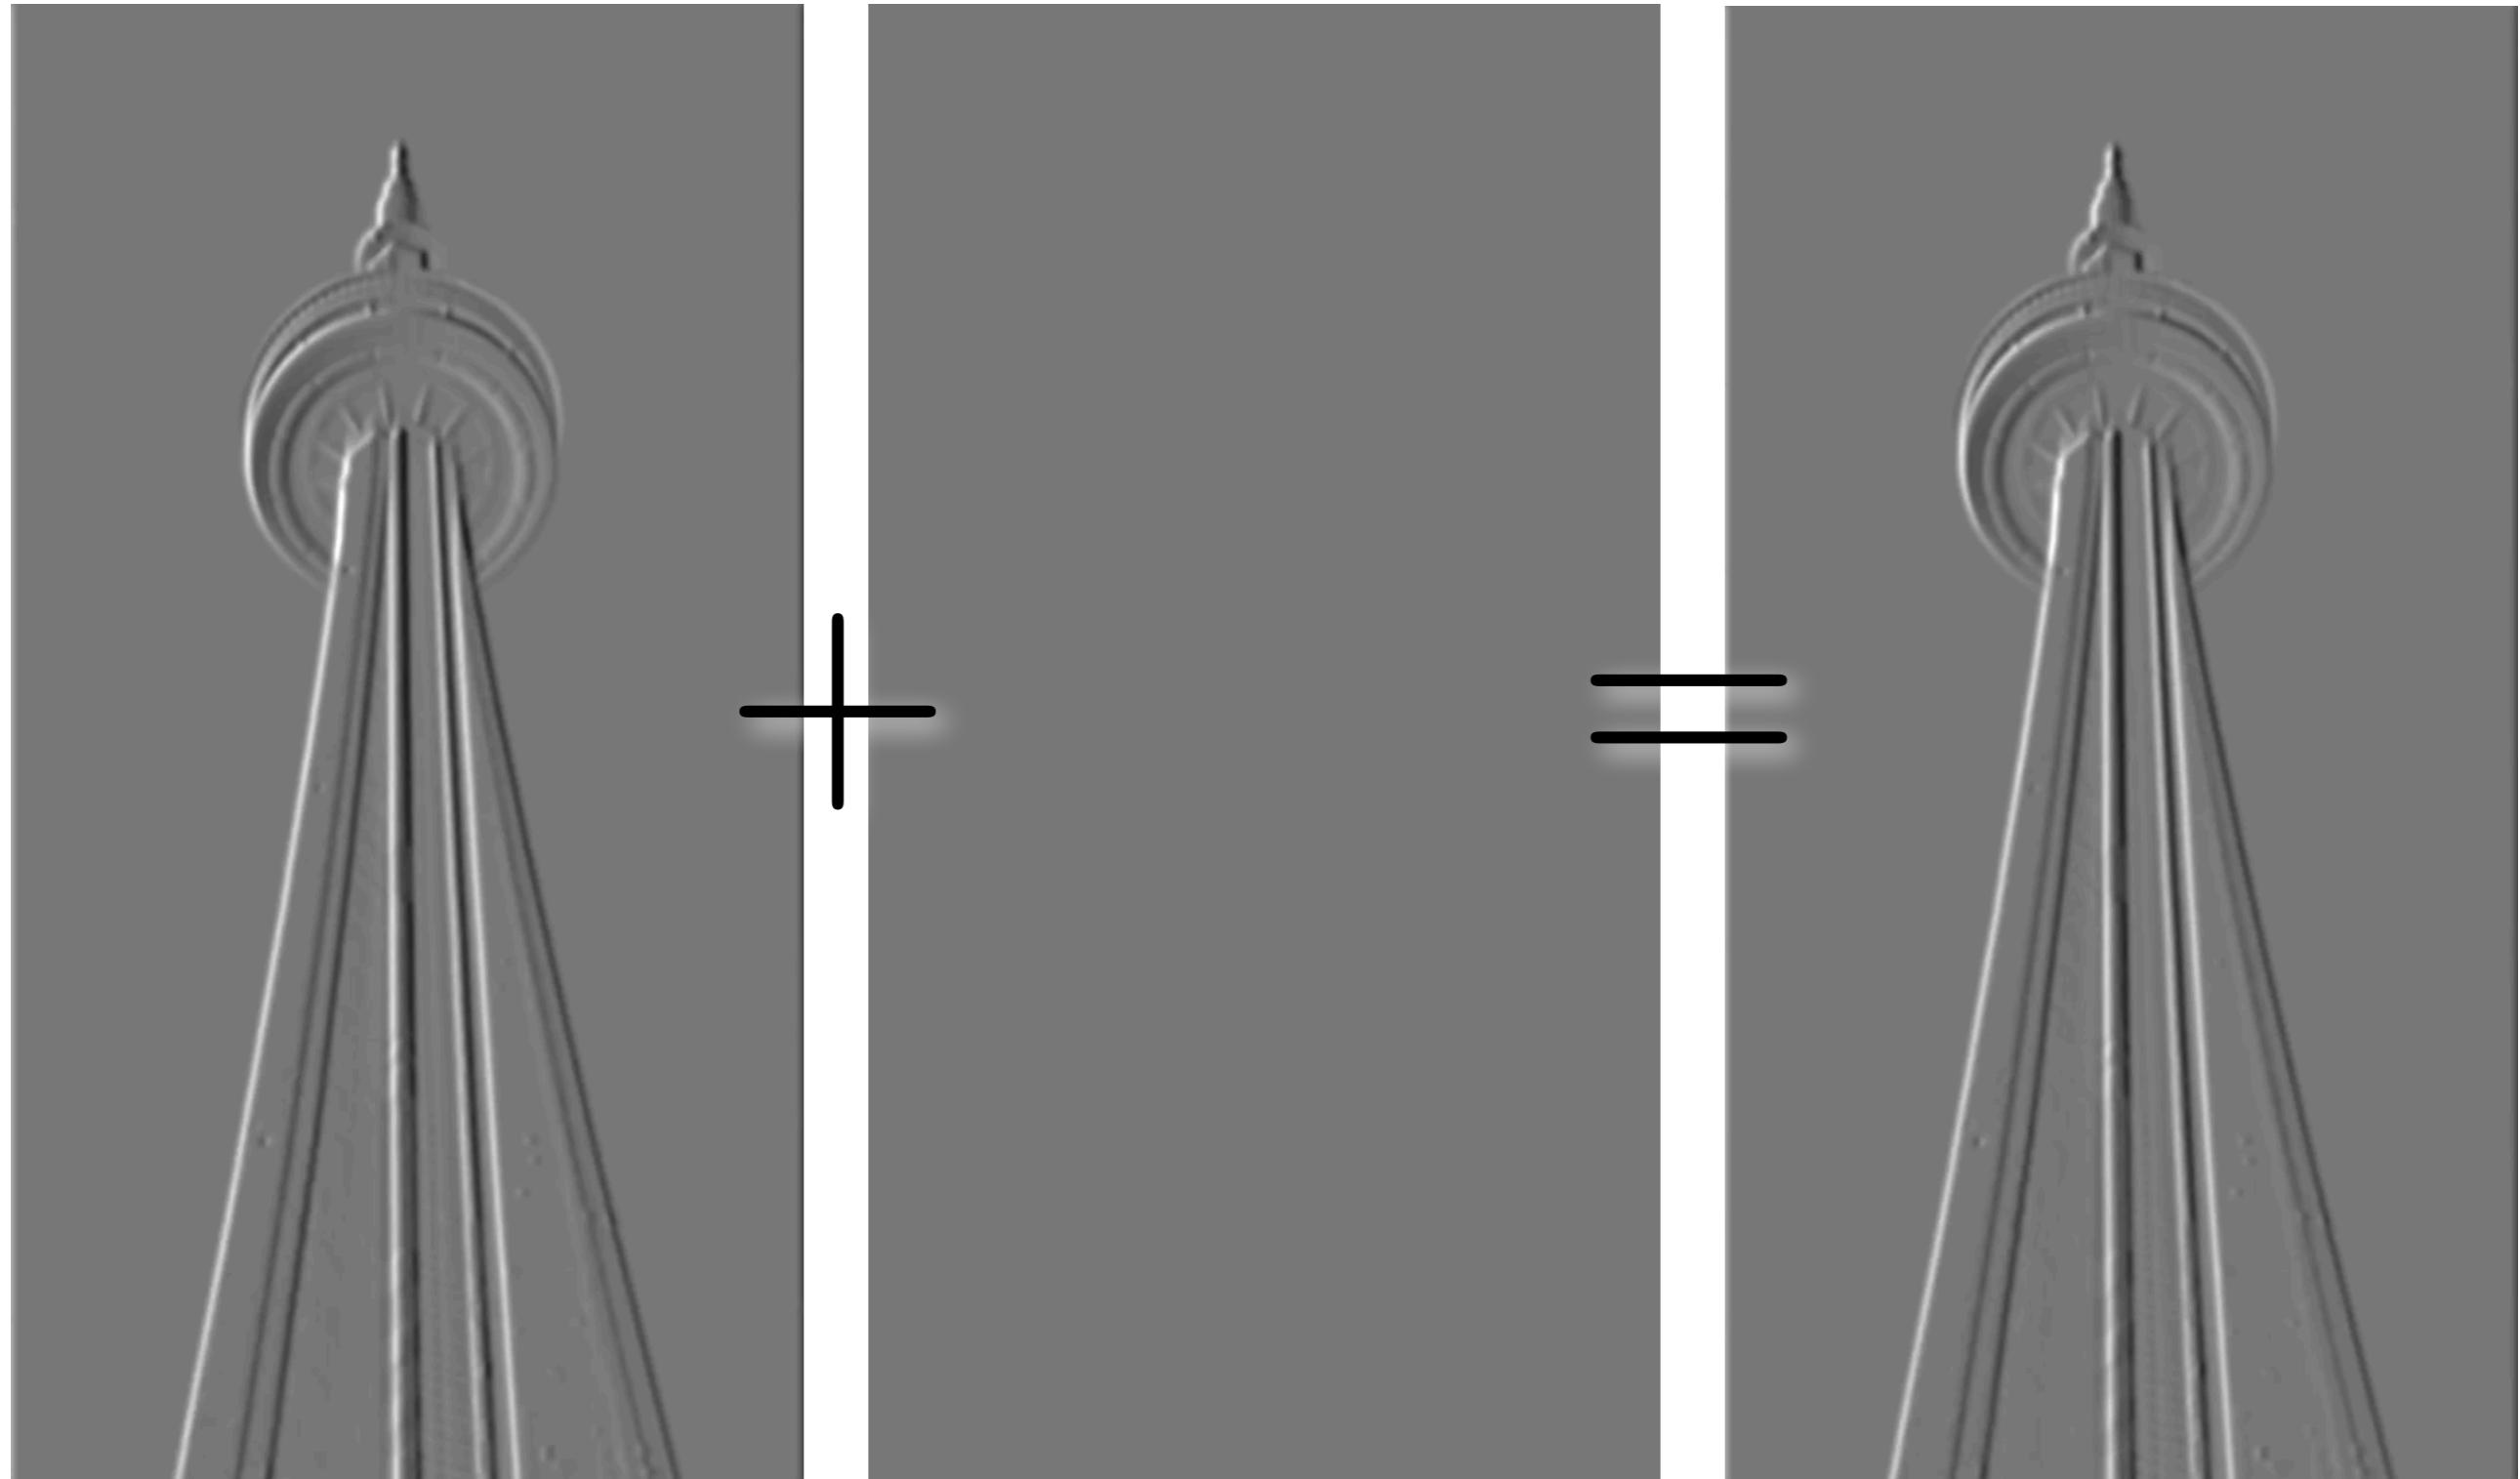
$$= (E[x, y] * F[x, y]) + (E[x, y] * G[x, y])$$

$$I_\theta(x, y) = [\cos(\theta) \frac{\partial G(x, y)}{\partial x} + \sin(\theta) \frac{\partial G(x, y)}{\partial y}] * I(x, y)$$

**Can we compute this more efficiently?**

$$= \cos(\theta) \left[ \frac{\partial G(x, y)}{\partial x} * I(x, y) \right] + \sin(\theta) \left[ \frac{\partial G(x, y)}{\partial y} * I(x, y) \right]$$

**distributive property leads to efficient implementation**



$$\cos(\theta) \left[ \frac{\partial G}{\partial x} * I \right]$$

$$\sin(\theta) \left[ \frac{\partial G}{\partial y} * I \right]$$

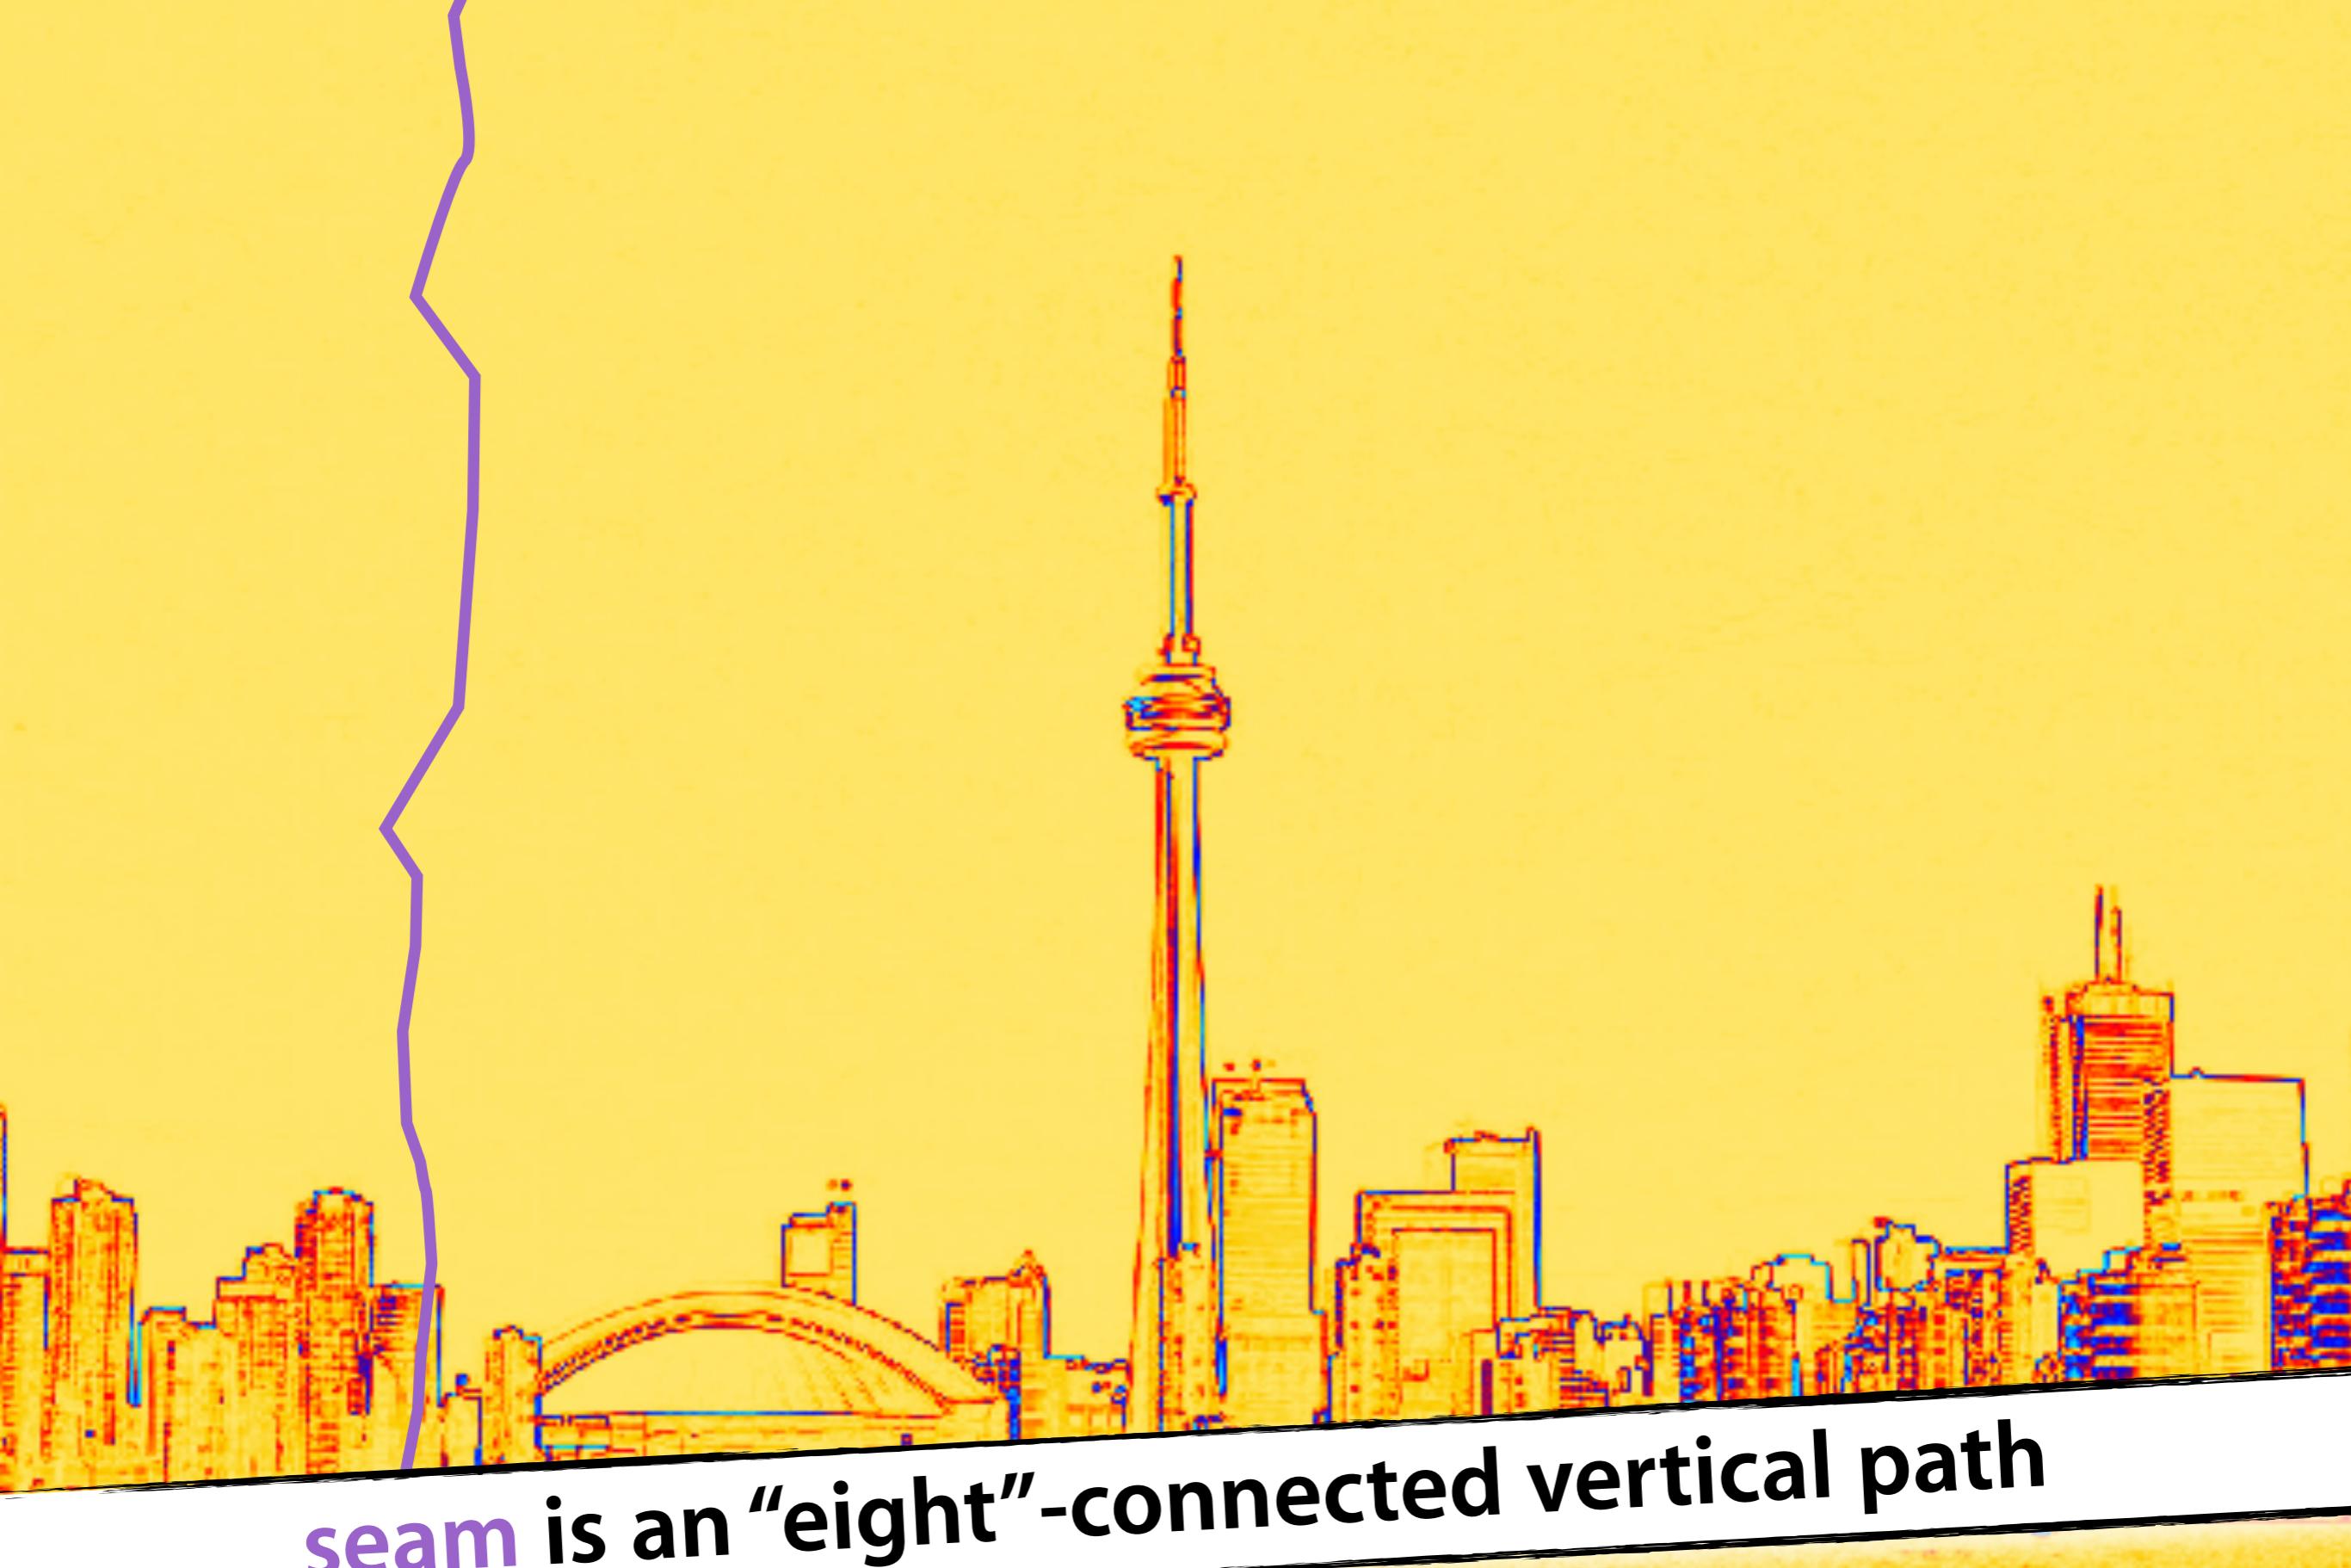
$$\frac{\partial [G * I]}{\partial \theta}$$



**reduce image size by removing non-salient seams**



$$\text{Energy}[x, y] = \sqrt{I_x^2 + I_y^2}$$



**seam** is an “eight”-connected vertical path



**Goal:** Find the seam with the least cumulative energy

$$\mathbf{s}^* = \arg \min_{\mathbf{s}} \text{Cost}[\mathbf{s}]$$

where

$$\text{Cost}[\mathbf{s}] = \sum_{i=0}^{\text{height}} \text{Energy}[\mathbf{s}_i]$$

# Dynamic Programming

$$M[i, j] = \text{Energy}[i, j] +$$

$$\min(M[i - 1, j - 1], M[i, j - 1], M[i + 1, j - 1])$$

**minimum cumulative energy ending at pixel  $[i, j]$**

$$\begin{aligned}\mathbf{M}[i, j] = & \text{Energy}[i, j] + \\& \min(\mathbf{M}[i - 1, j - 1], \mathbf{M}[i, j - 1], \mathbf{M}[i + 1, j - 1])\end{aligned}$$

---

Energy[i, j]

0	5	1
6	8	4
9	1	7

optimal solution

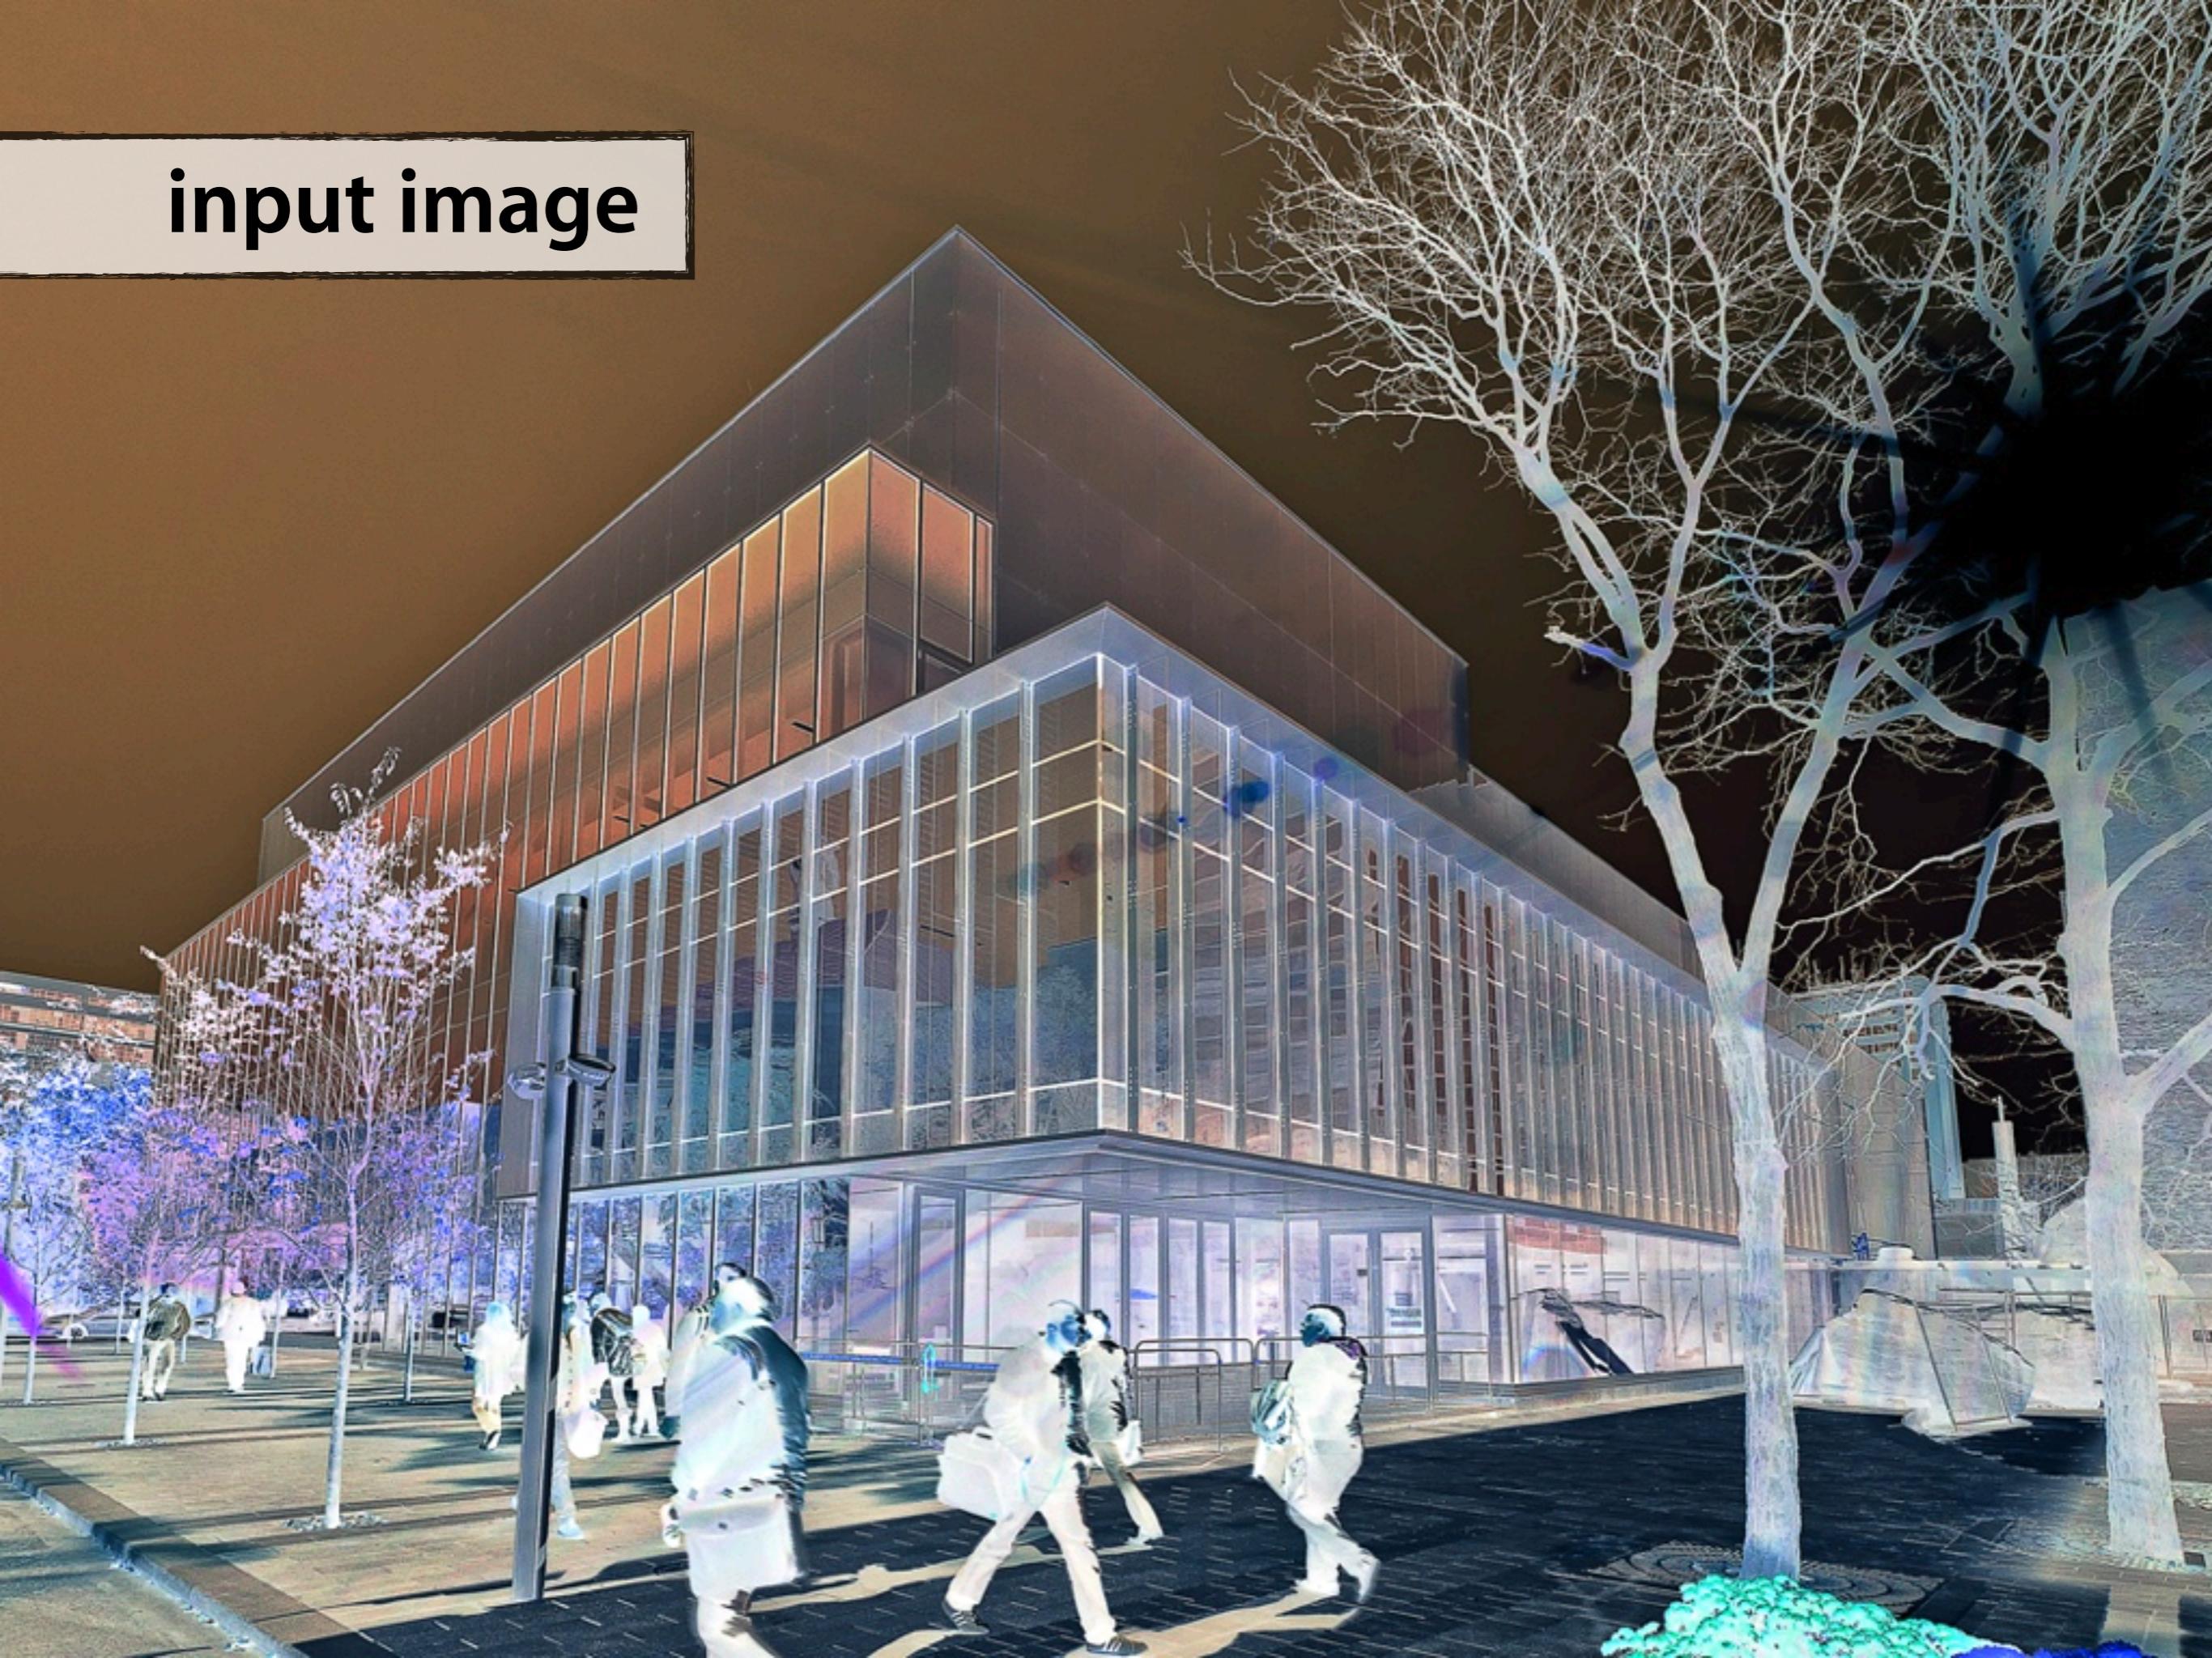
$\mathbf{M}[i, j]$

0	5	1
6	8	5
15	6	12

partial optimal solution

To find the optimal seam, backtrack from minimum

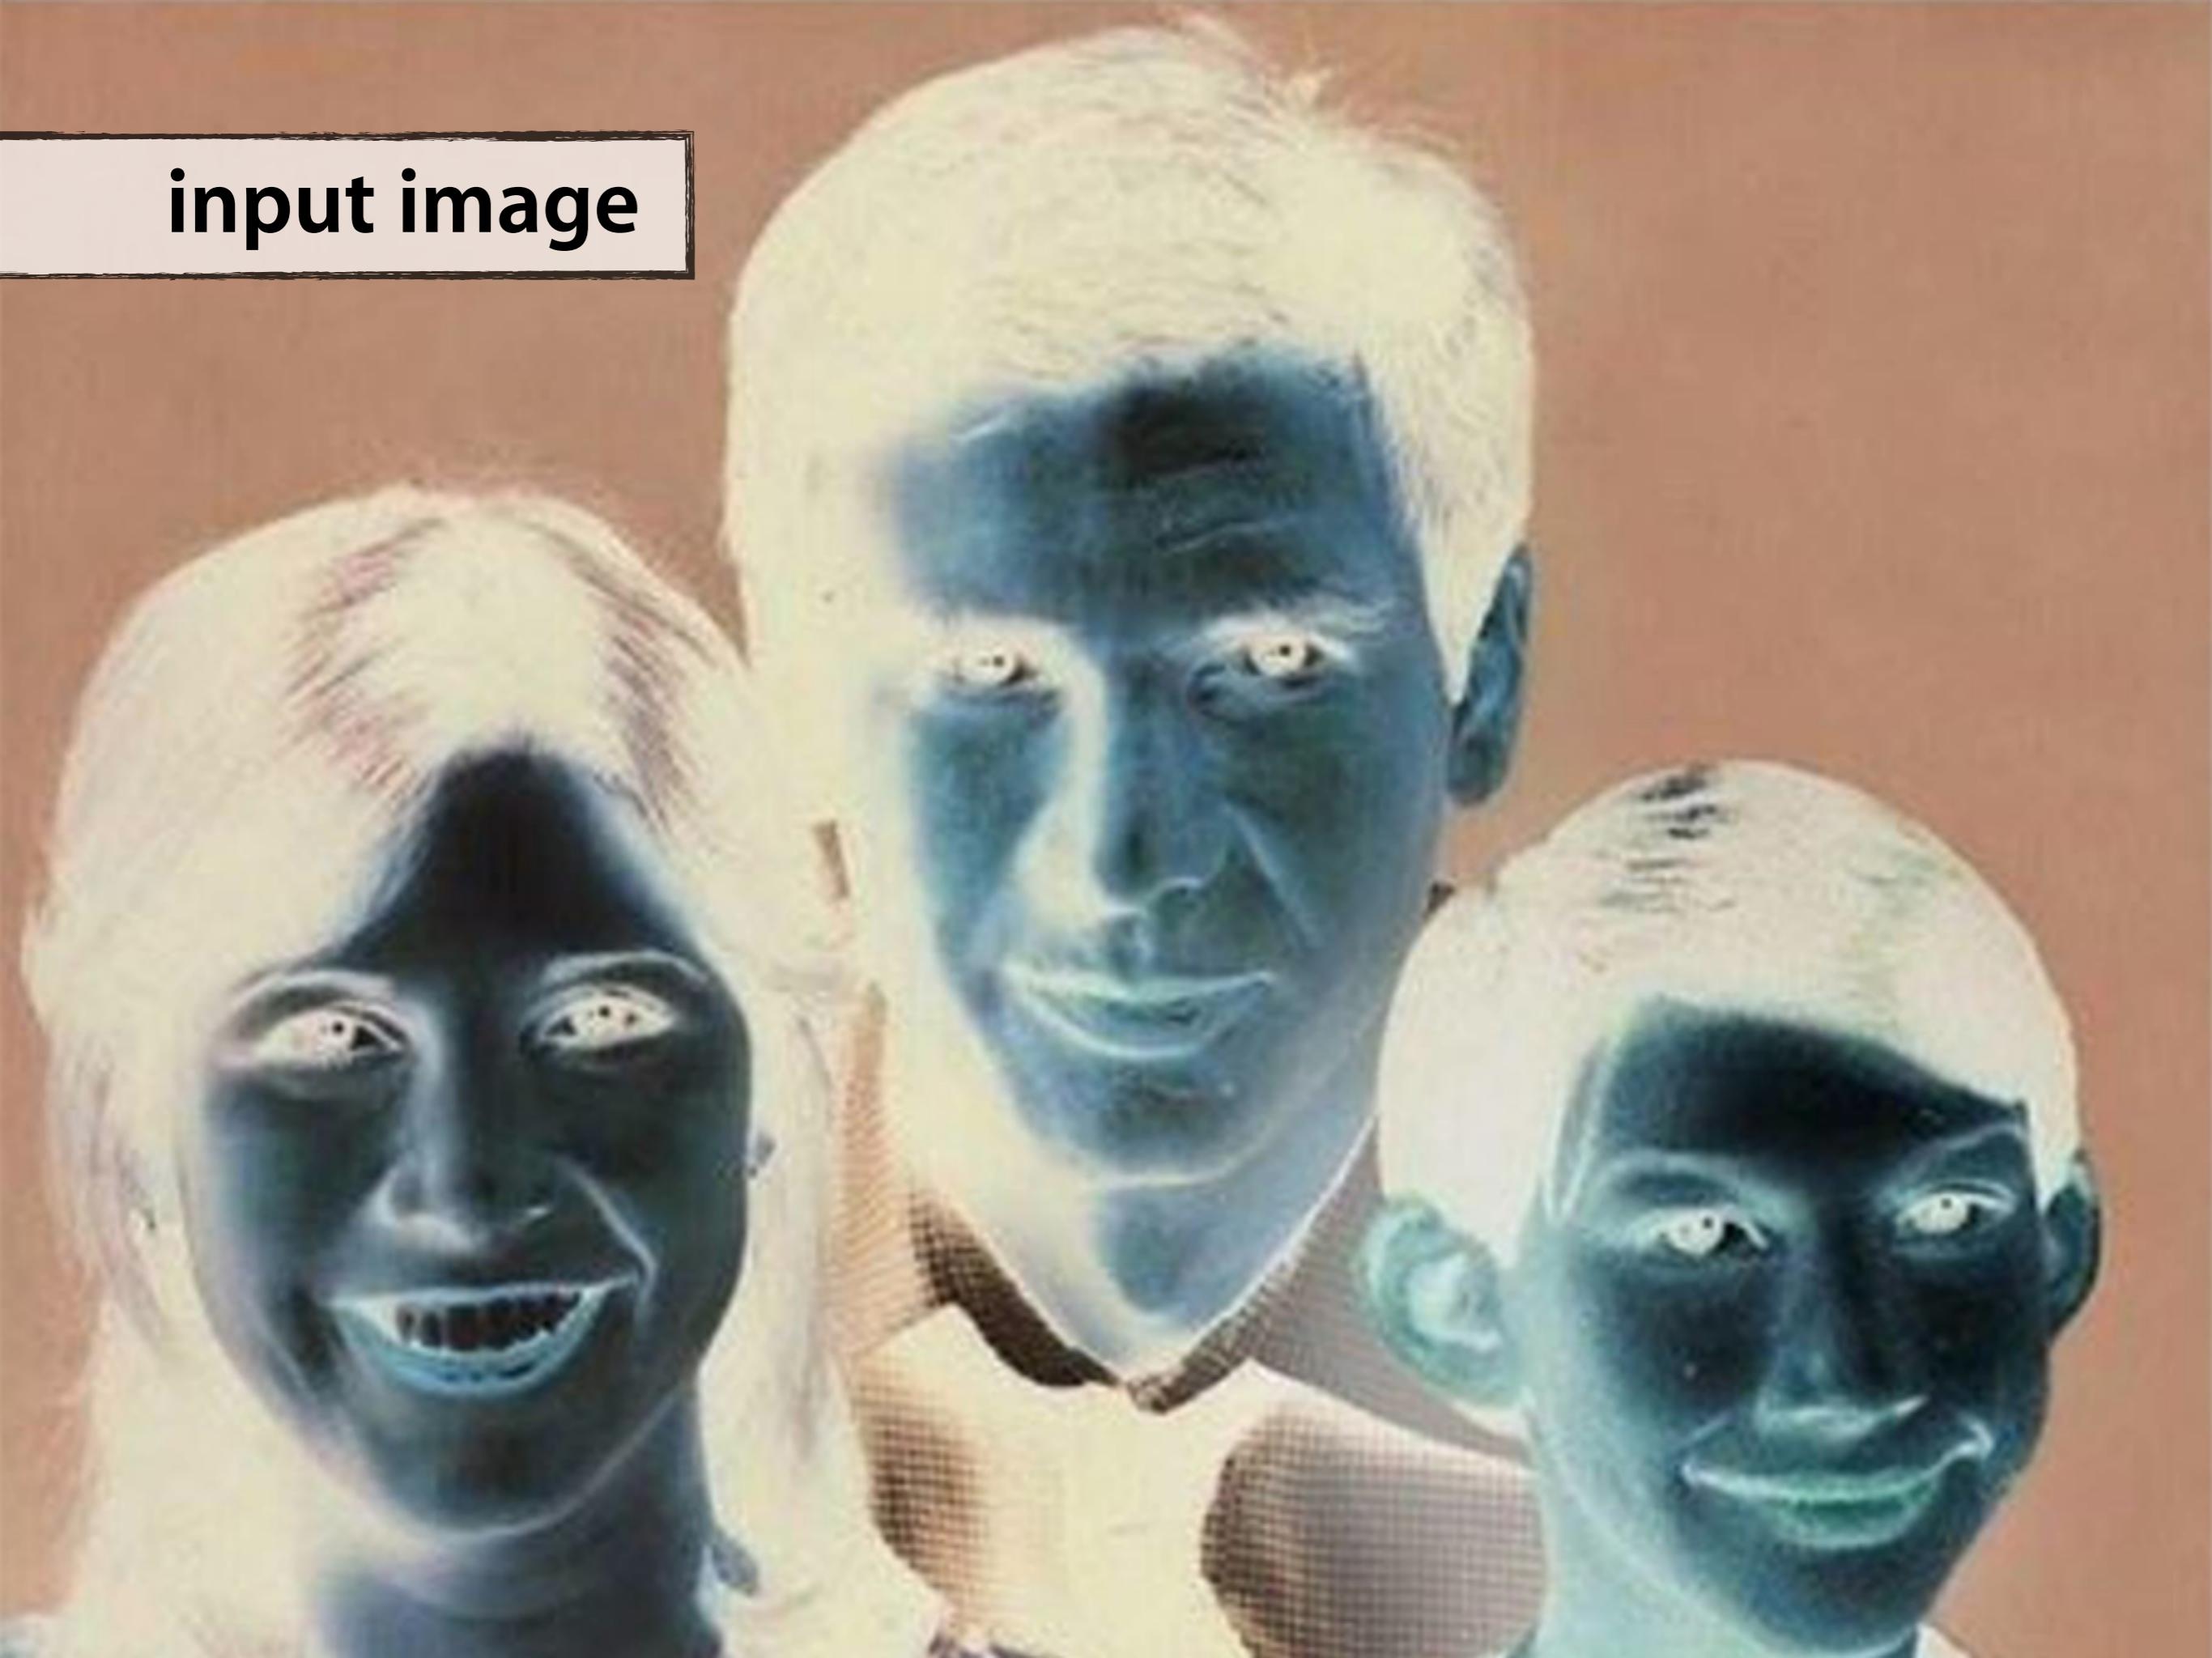
**input image**





seam carving result

**input image**





**seam carving result**

# A Computational Approach to Edge Detection

JOHN CANNY, MEMBER, IEEE

**Abstract**—This paper describes a computational approach to edge detection. The success of the approach depends on the definition of a comprehensive set of goals for the computation of edge points. These goals must be precise enough to delimit the desired behavior of the detector while making minimal assumptions about the form of the solution. We define detection and localization criteria for a class of edges, and present mathematical forms for these criteria as functionals on the

detector as input to a program which could isolate simple geometric solids. More recently the model-based vision system ACRONYM [3] used an edge detector as the front end to a sophisticated recognition program. Shape from motion [29], [13] can be used to infer the structure of three-dimensional objects from the motion of edge con-

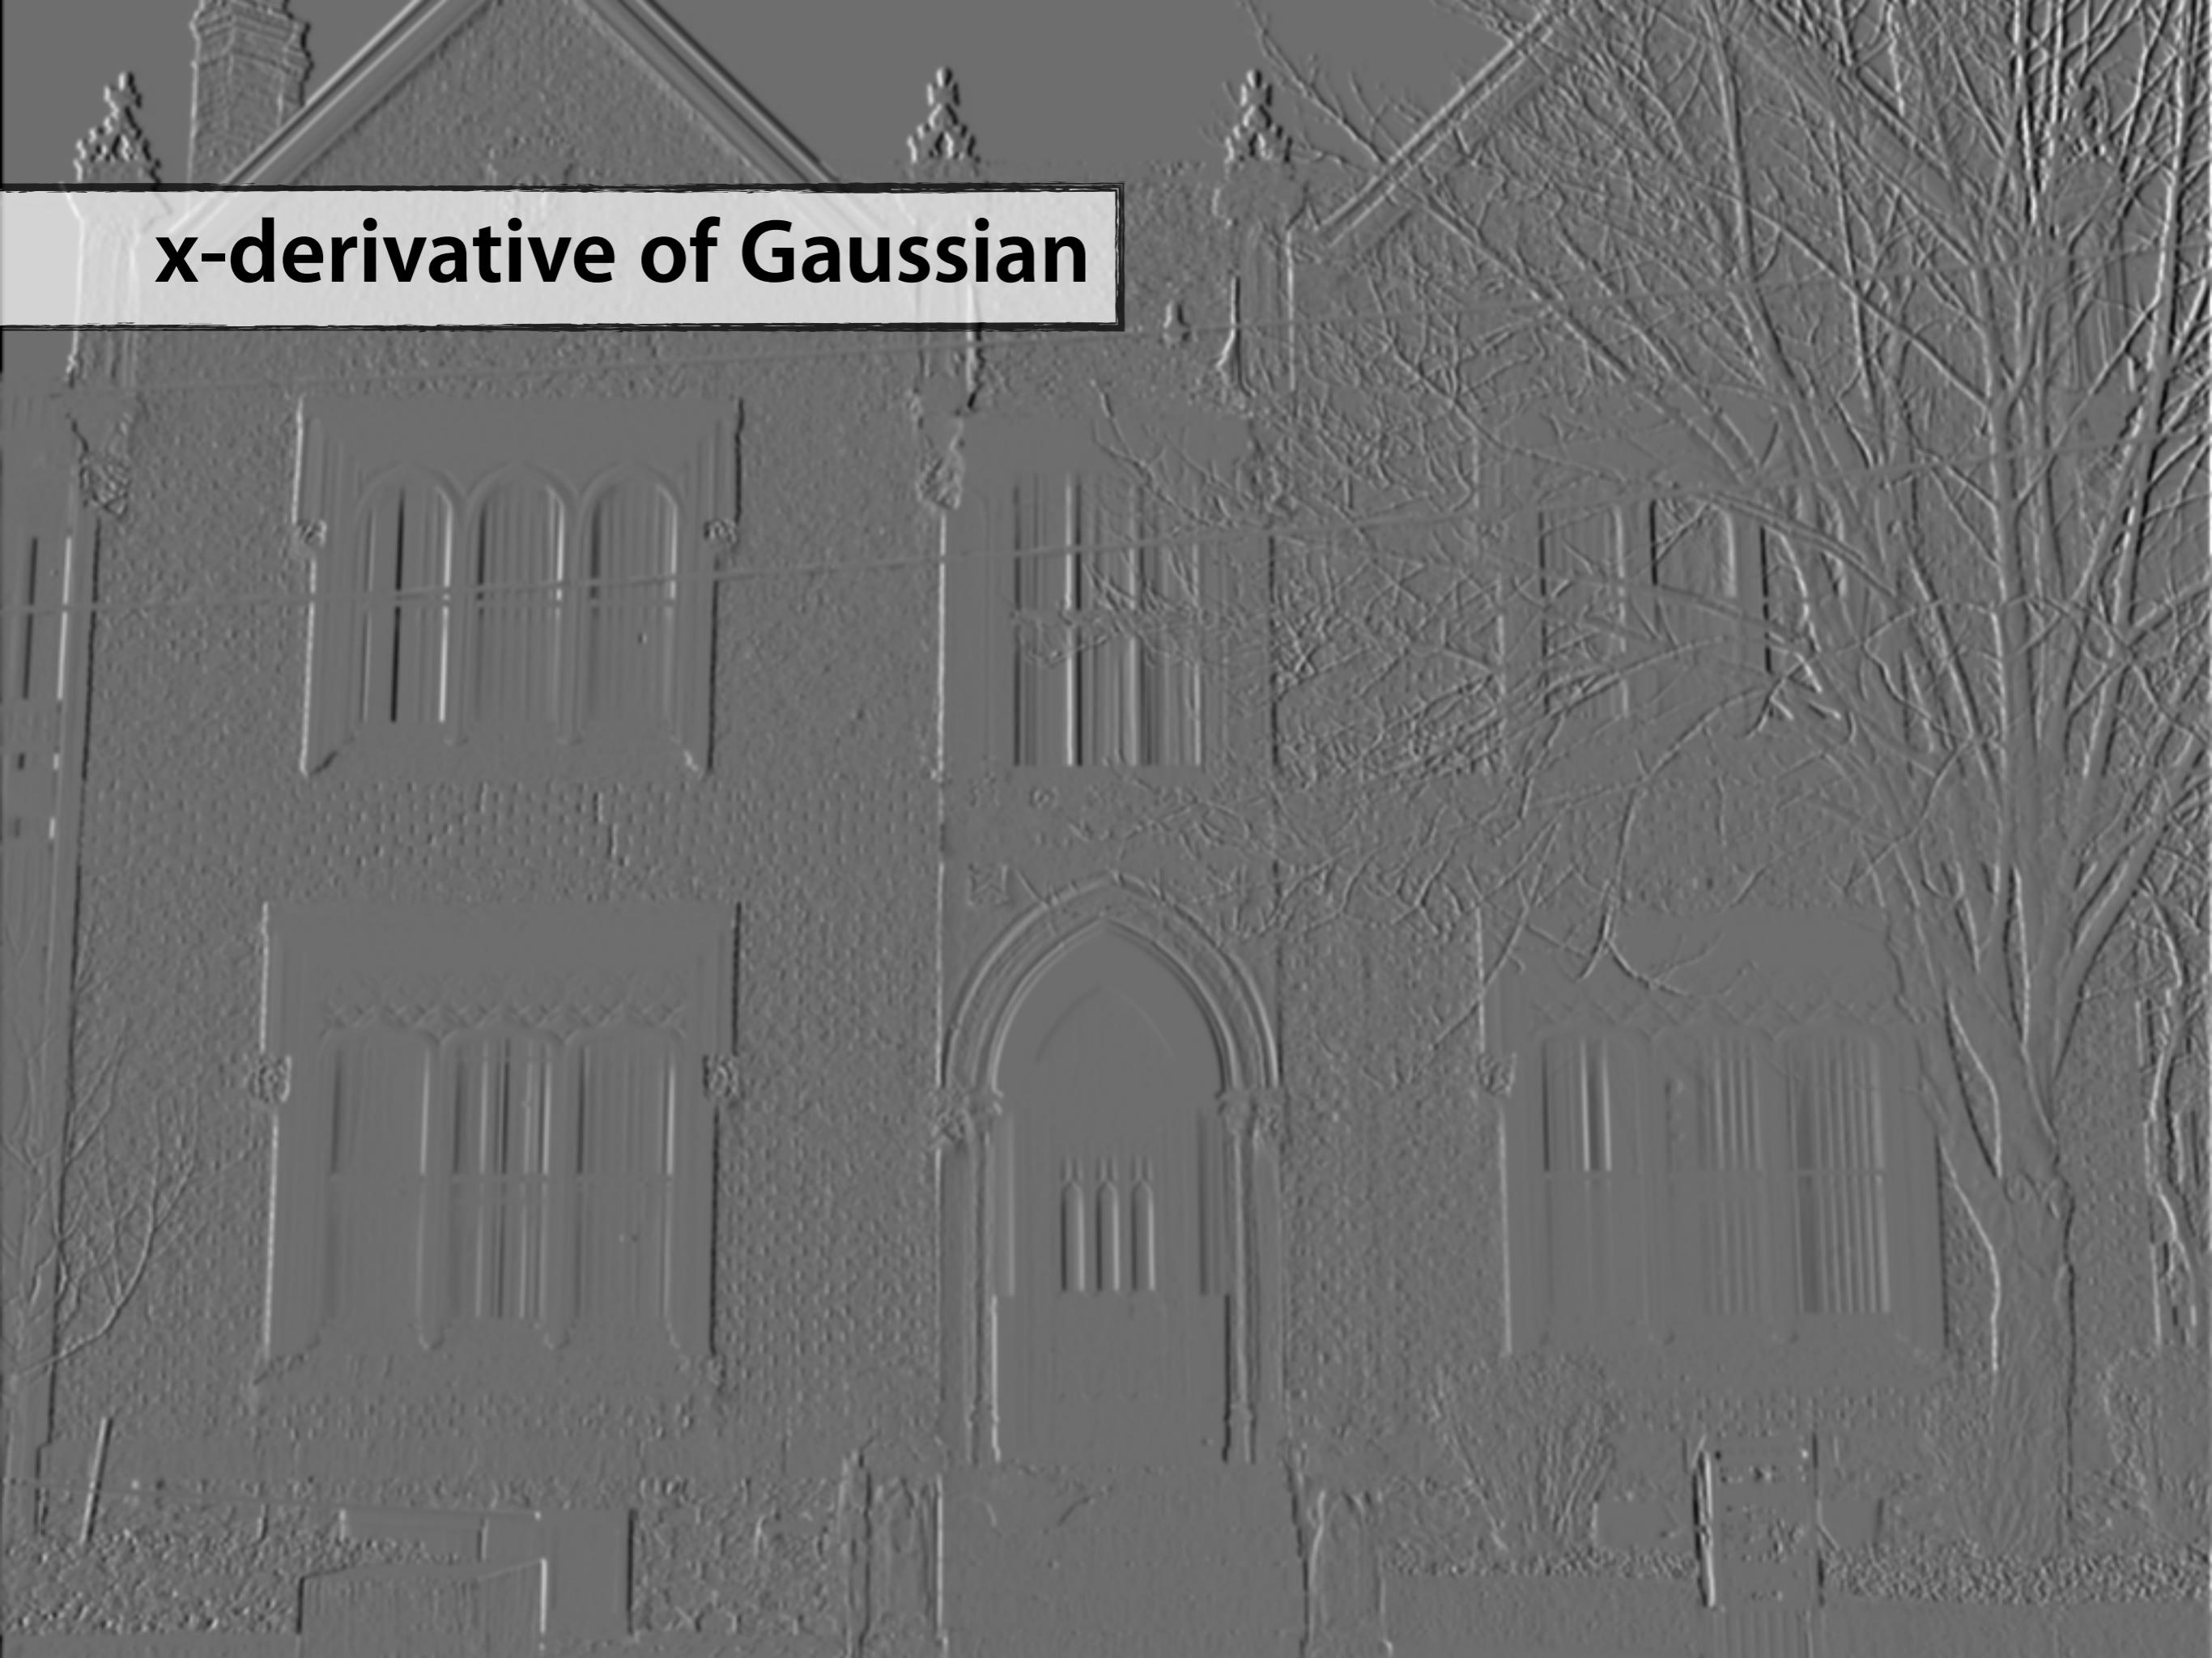
# 4

**major steps**

# Step 1

**Filter image with x and y derivatives of Gaussian**

# **x-derivative of Gaussian**



# y-derivative of Gaussian



# Step 2

**Find magnitude and direction of gradient**

# gradient magnitude

```
img_grad_mag = sqrt(img_dx.^2 + img_dy.^2);
```

# gradient orientation

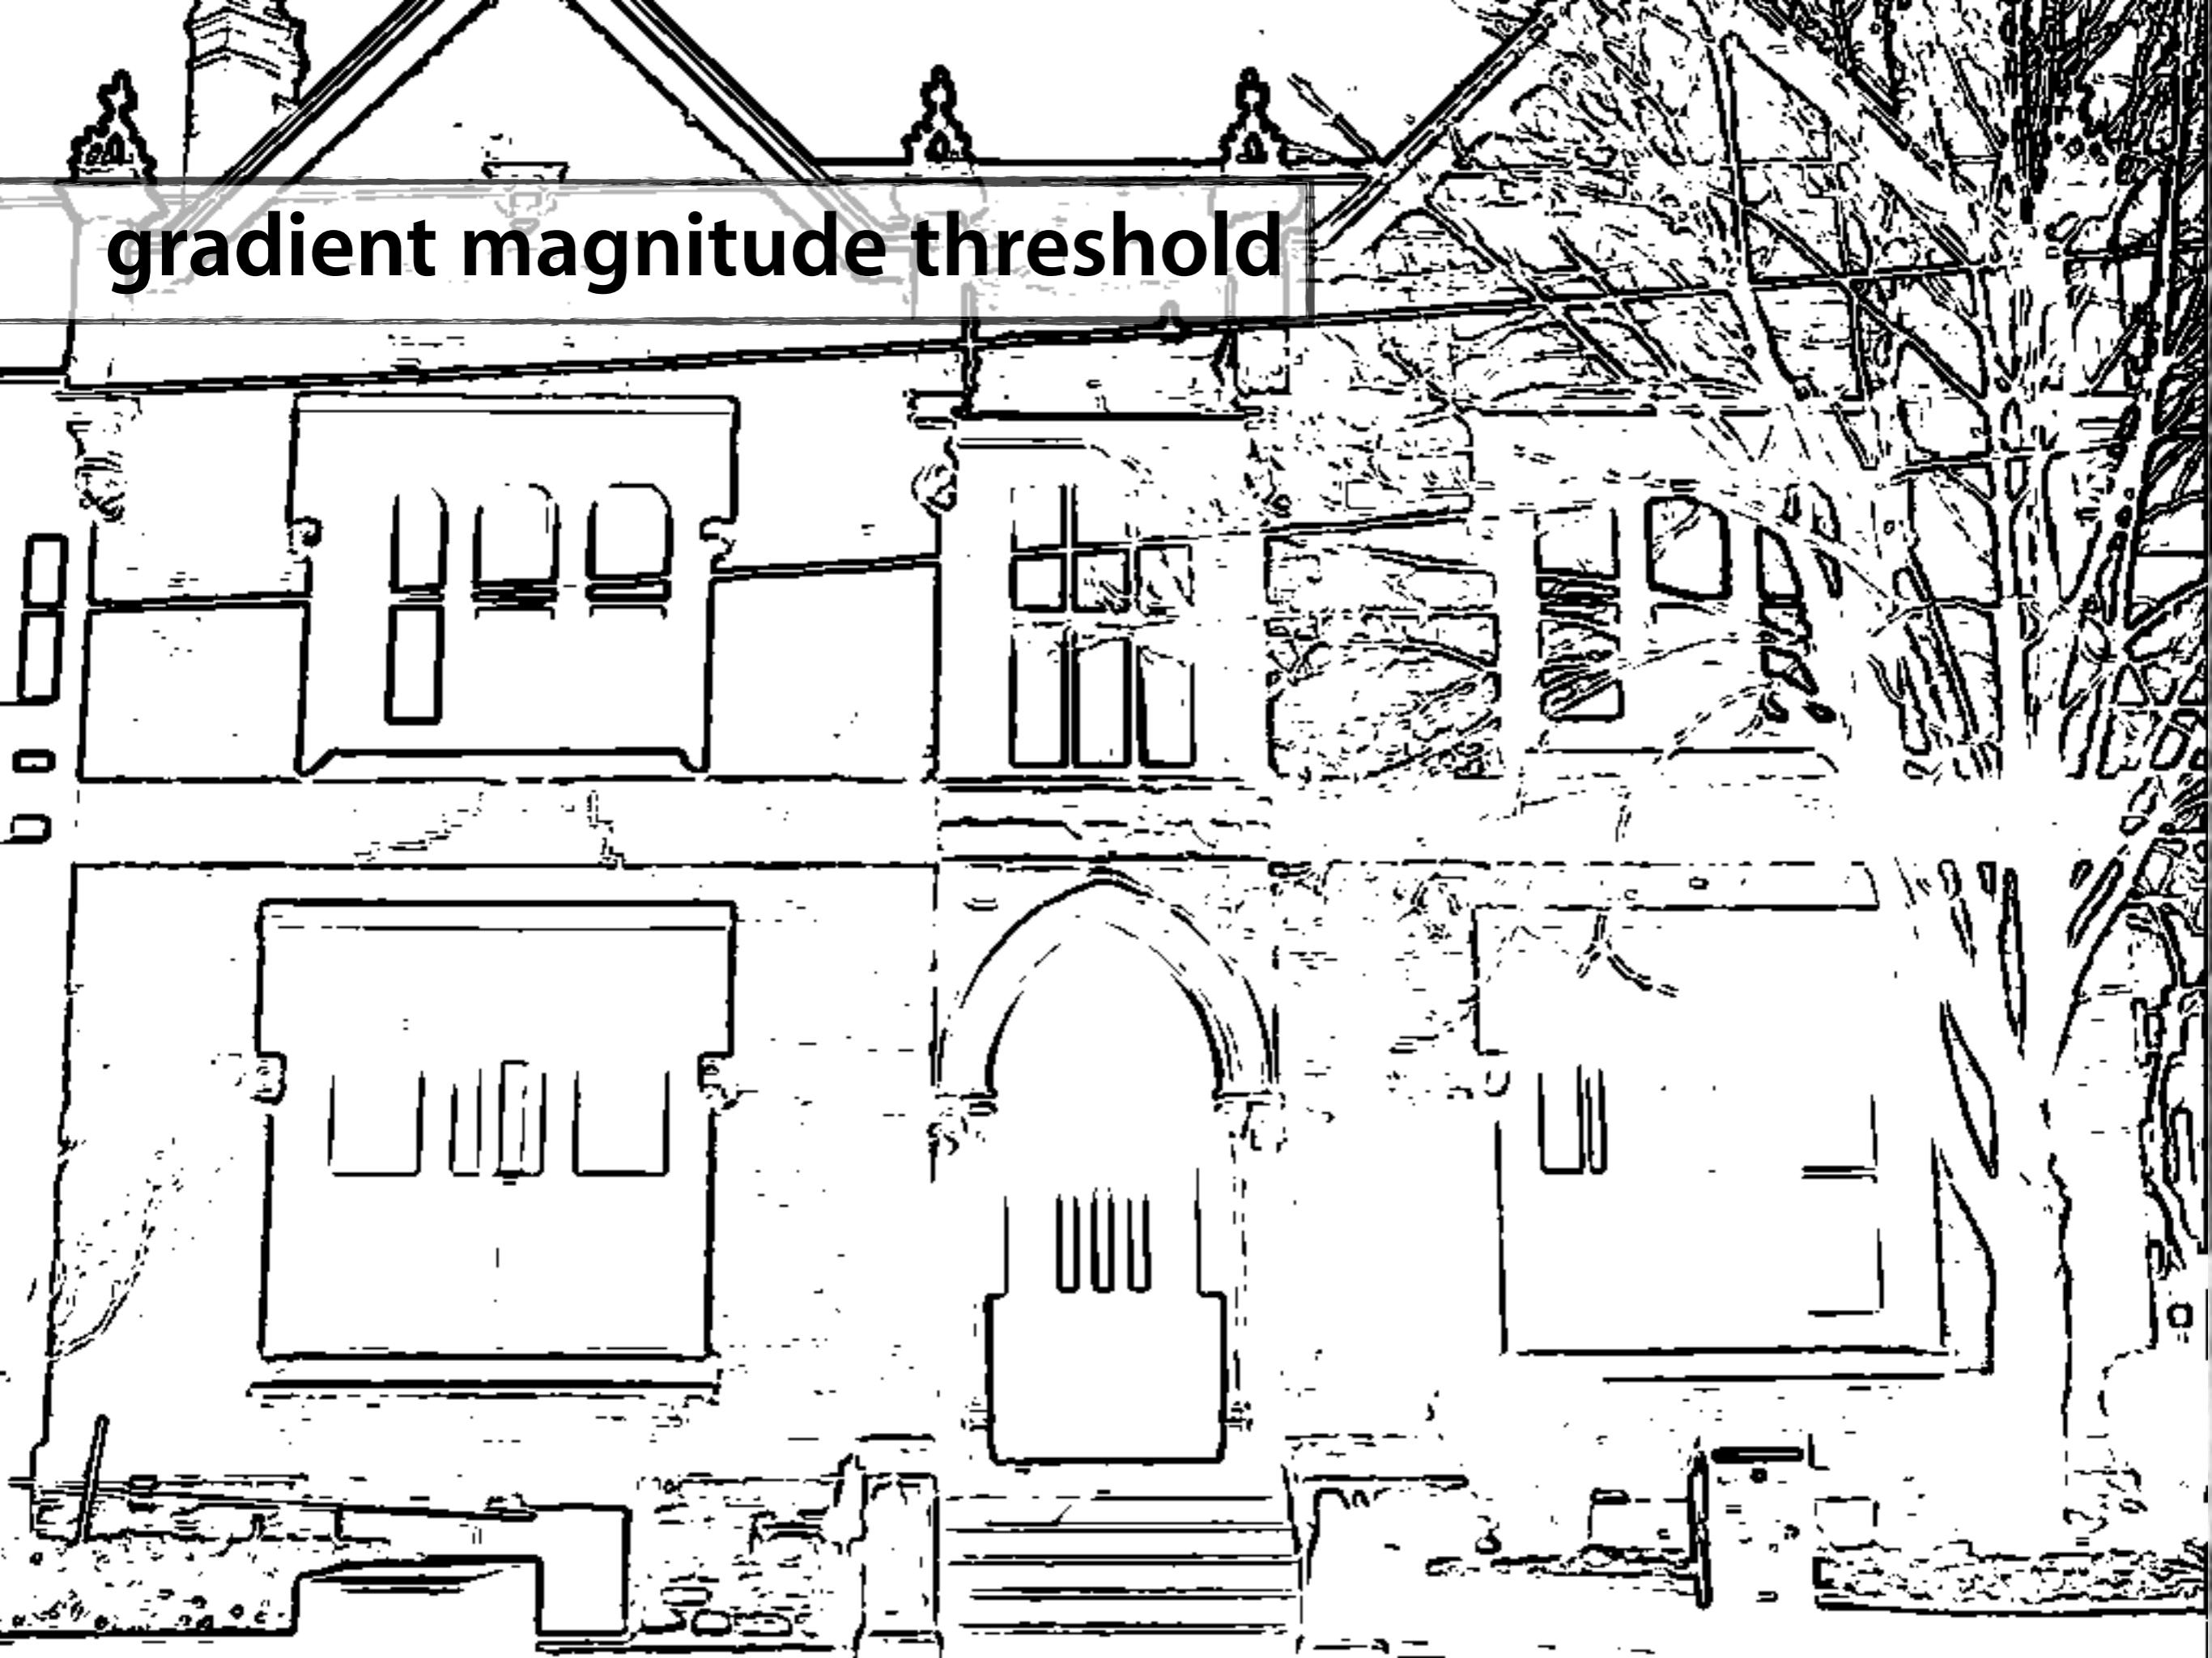
```
img_grad_or = atan2(img_dy, img_dx);
```

# gradient magnitude

Threshold gradient magnitude and done?

```
>> img_thresh = im_grad_mag > 15;
```

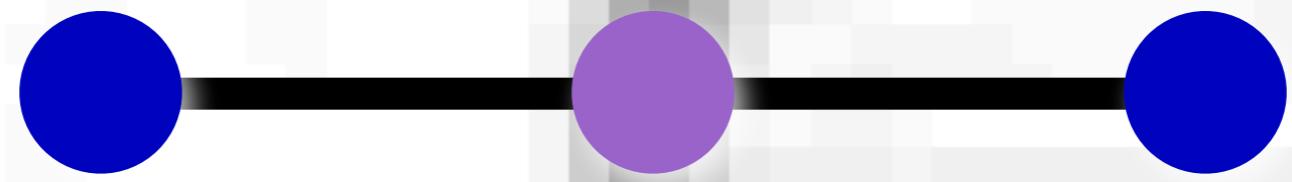
# gradient magnitude threshold



**End up with thick edges**

# Step 3

**Perform non-maximum suppression**



Is the gradient magnitude greater than its neighbours  
along the gradient orientation?

If yes, keep, otherwise suppress.

# gradient magnitude



# non-max suppression



# **Step 4**

**Threshold and link  
also known as Hysteresis Thresholding**

**low threshold**



```
>> img_thresh = img_grad_mag > 5;
```



**high threshold**

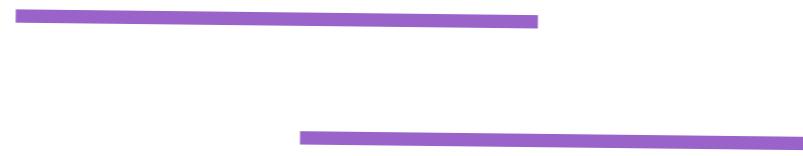
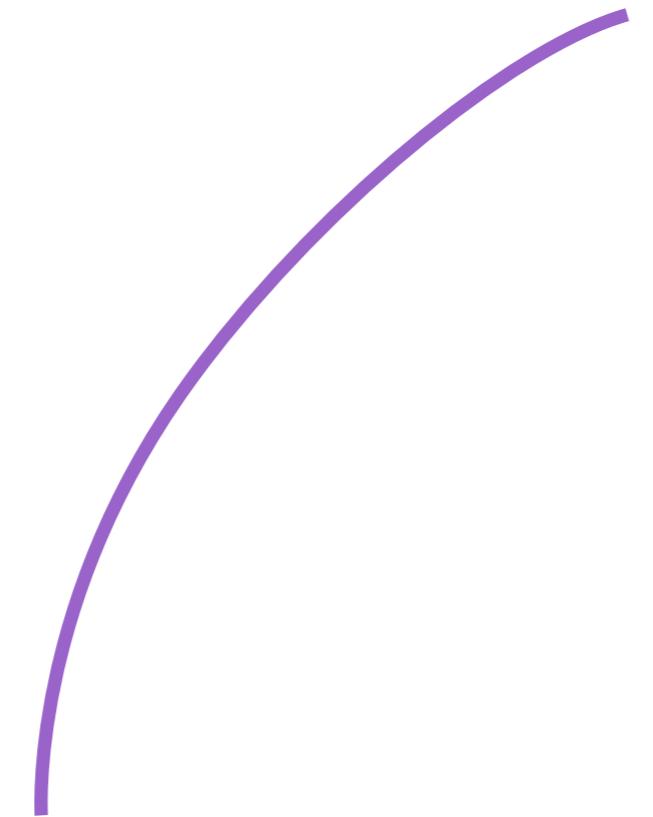
```
>> img_thresh = img_grad_mag > 30;
```

**non-max suppression**



# non-max suppression





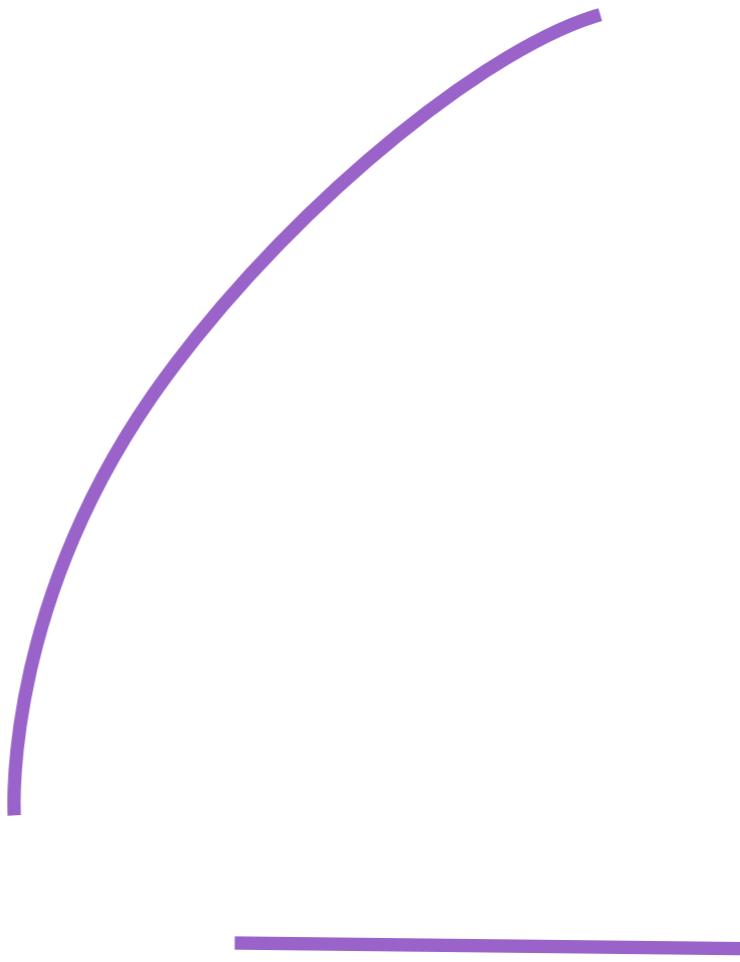
**Define two thresholds, low and high**



**Define two thresholds, low and high**



**Starting from a pixel passing a high threshold, keep only pixels on the path that pass the low threshold**



**Starting from a pixel passing a high threshold, keep only pixels on the path that pass the low threshold**

**non-max suppression**



# **hysteresis thresholding**



Canny Detector  
tl;dl

1. Filter image with x and y derivatives of Gaussian
2. Find magnitude and direction of gradient
3. Perform non-maximum suppression
4. Threshold and link



```
>> % type 'help edge' to see different  
    edge-finding methods  
>> img_edge = edge(img, 'canny');  
>> imshow(img_edge, [ ])
```

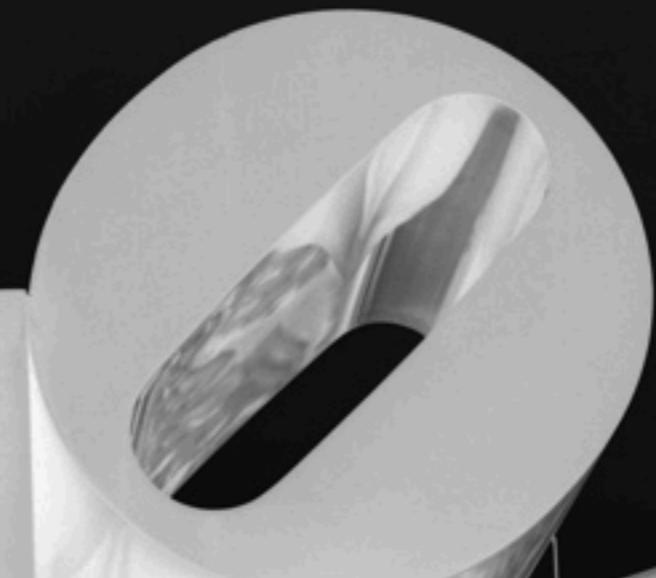
**What is the right  
SCALE?**

Scale?



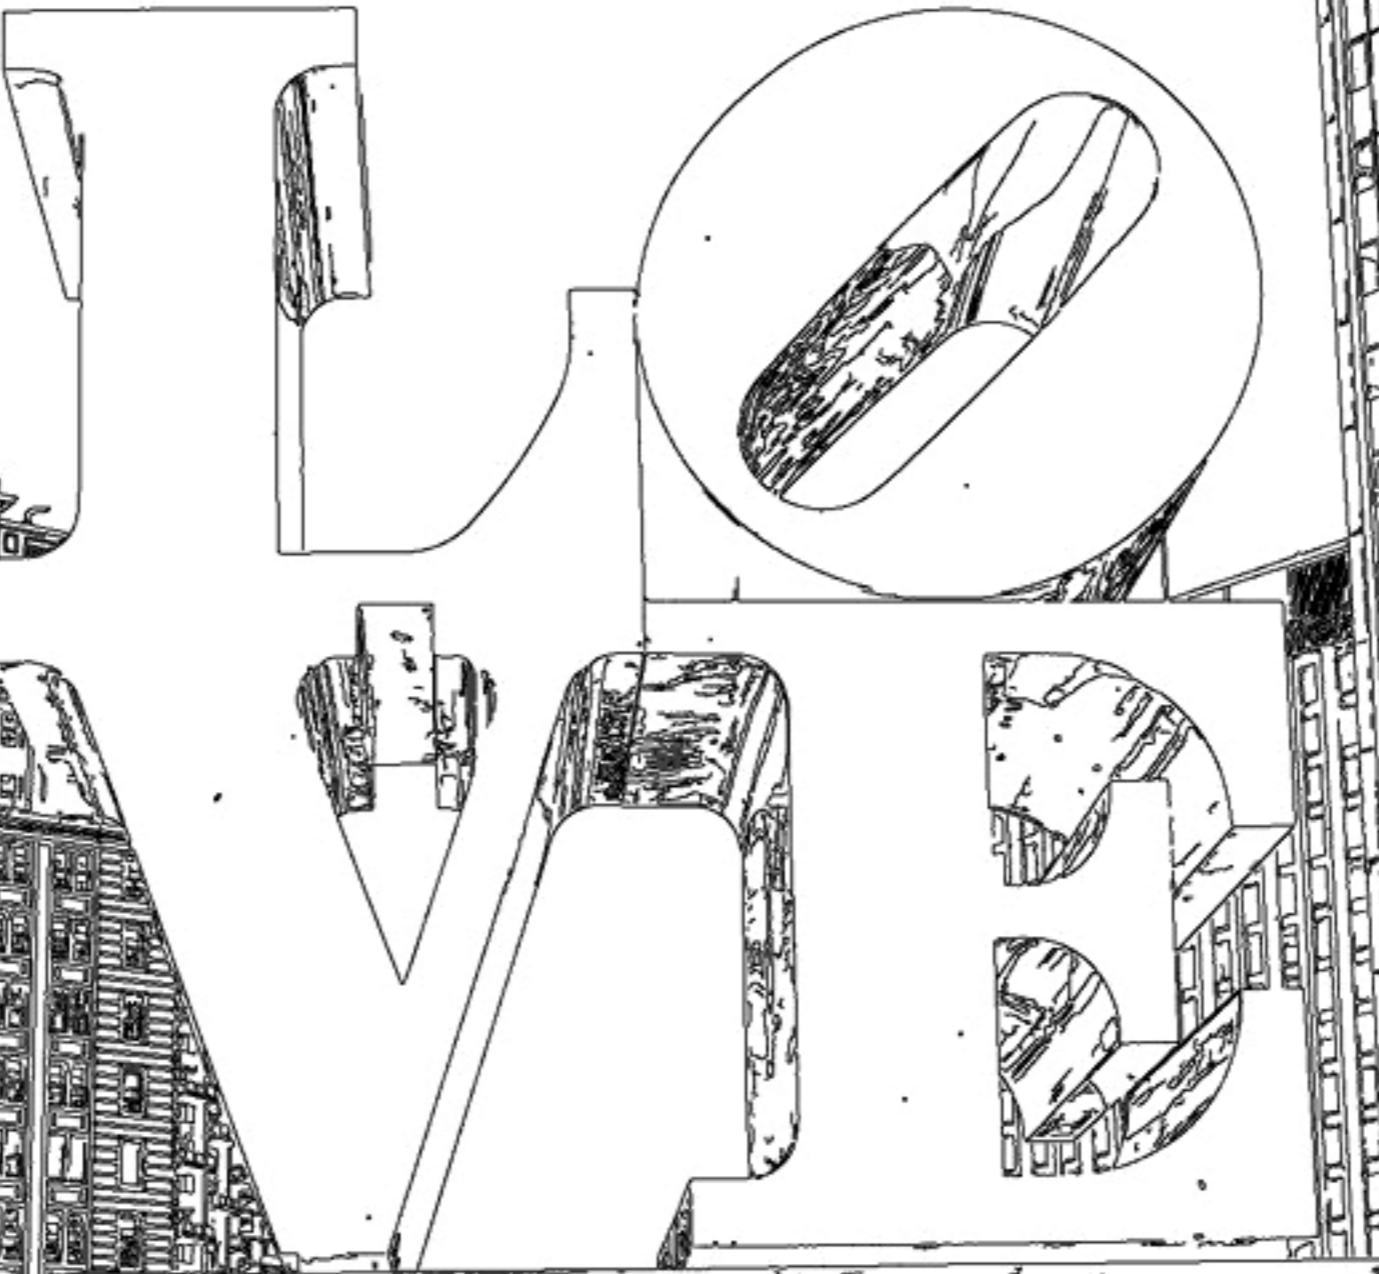


input



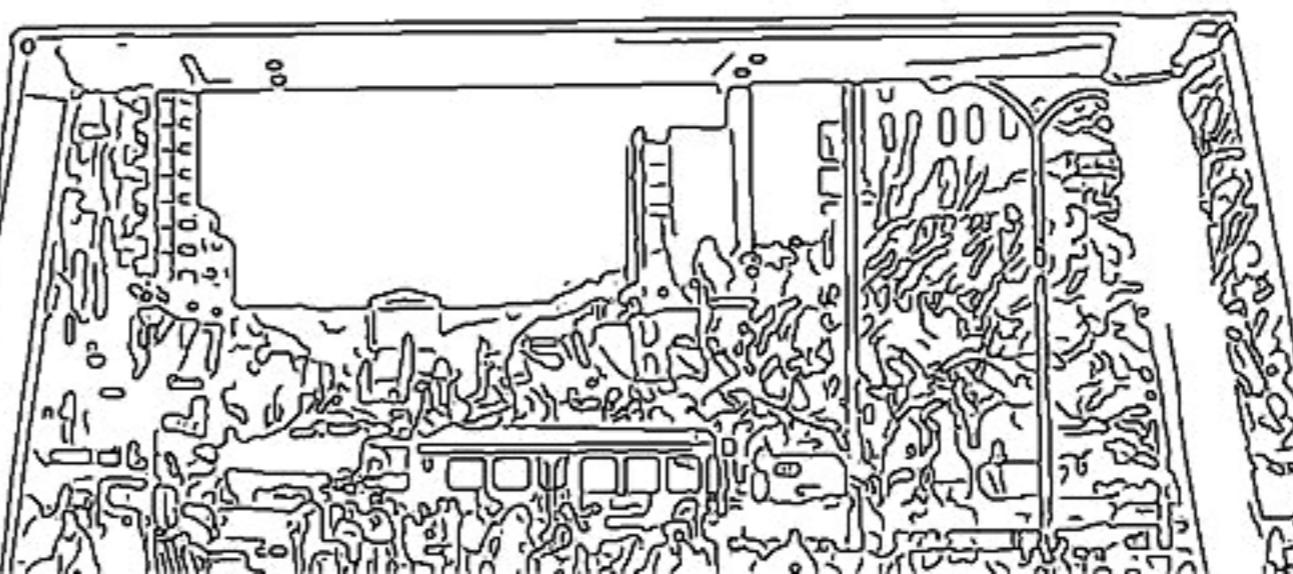
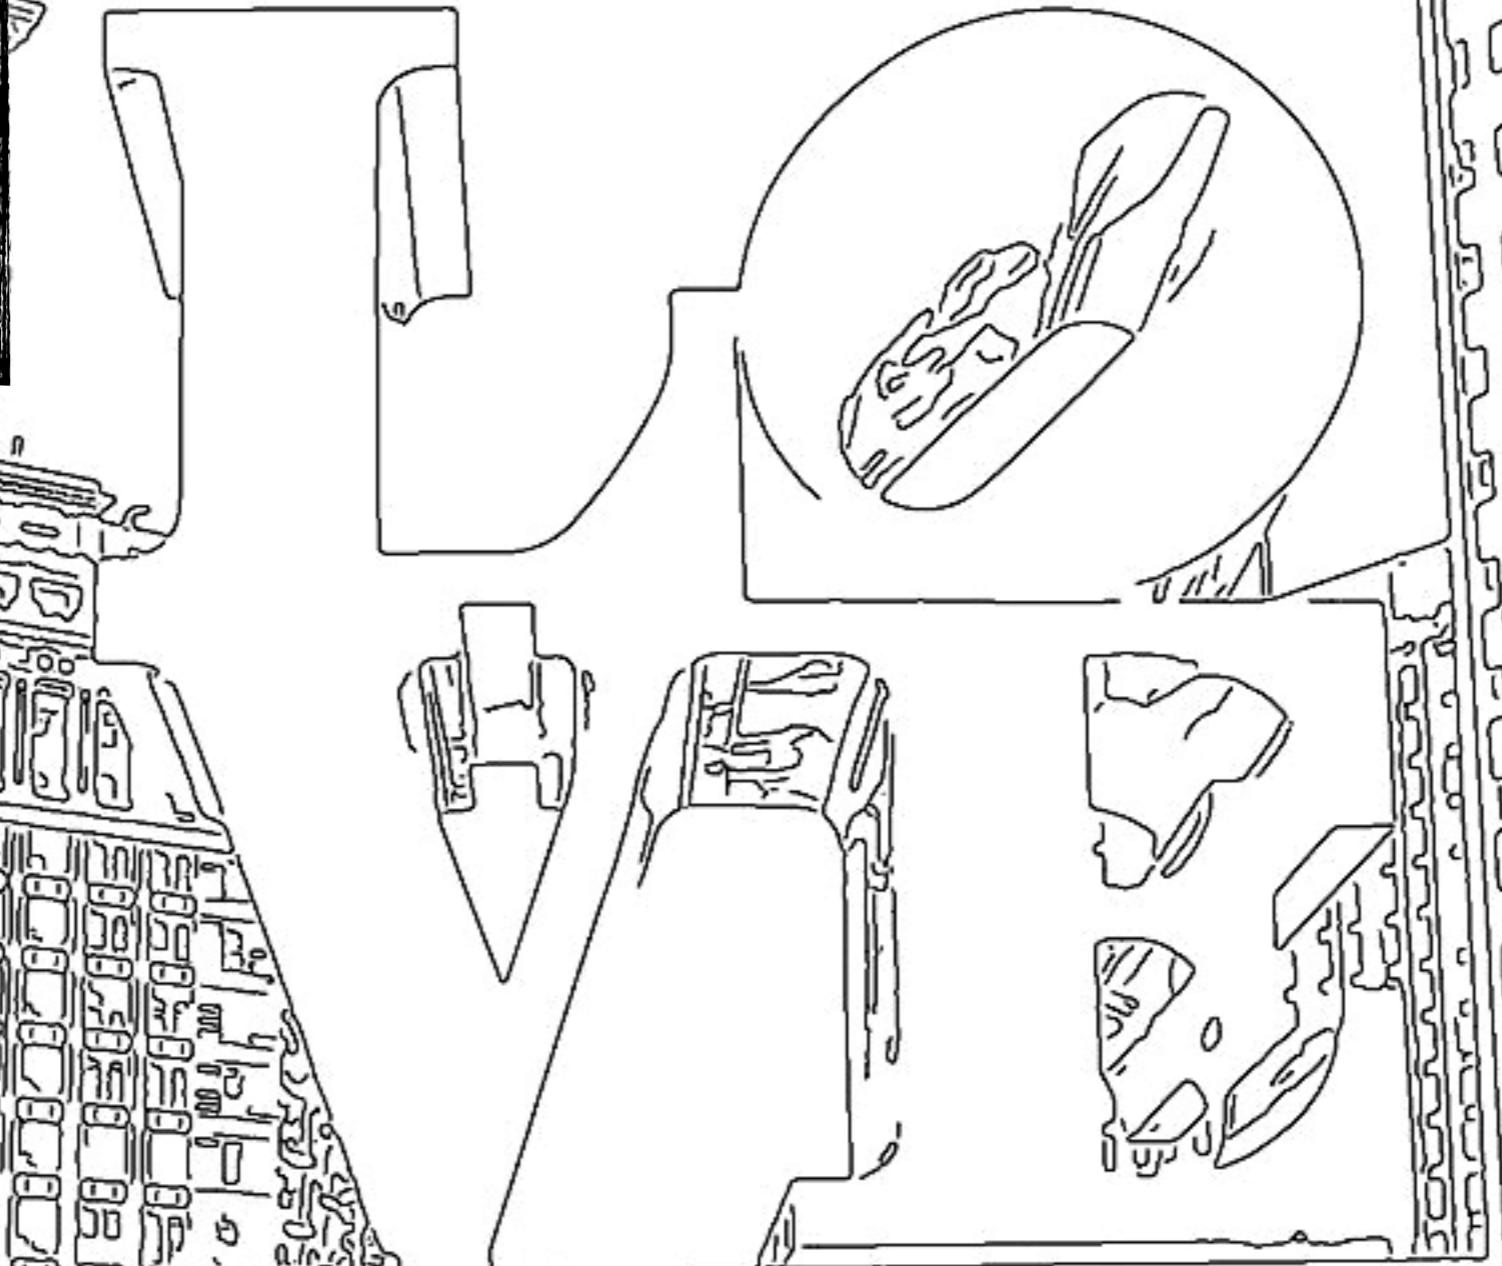
# Canny

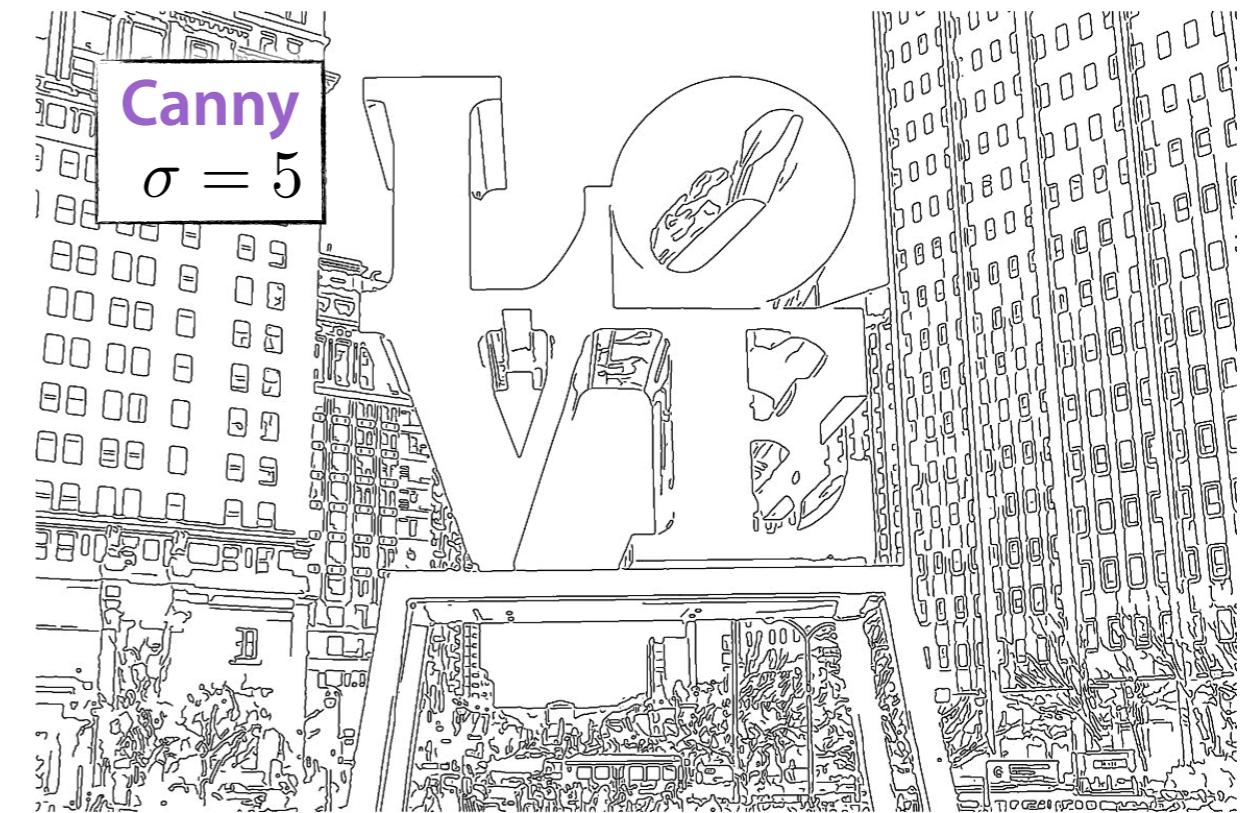
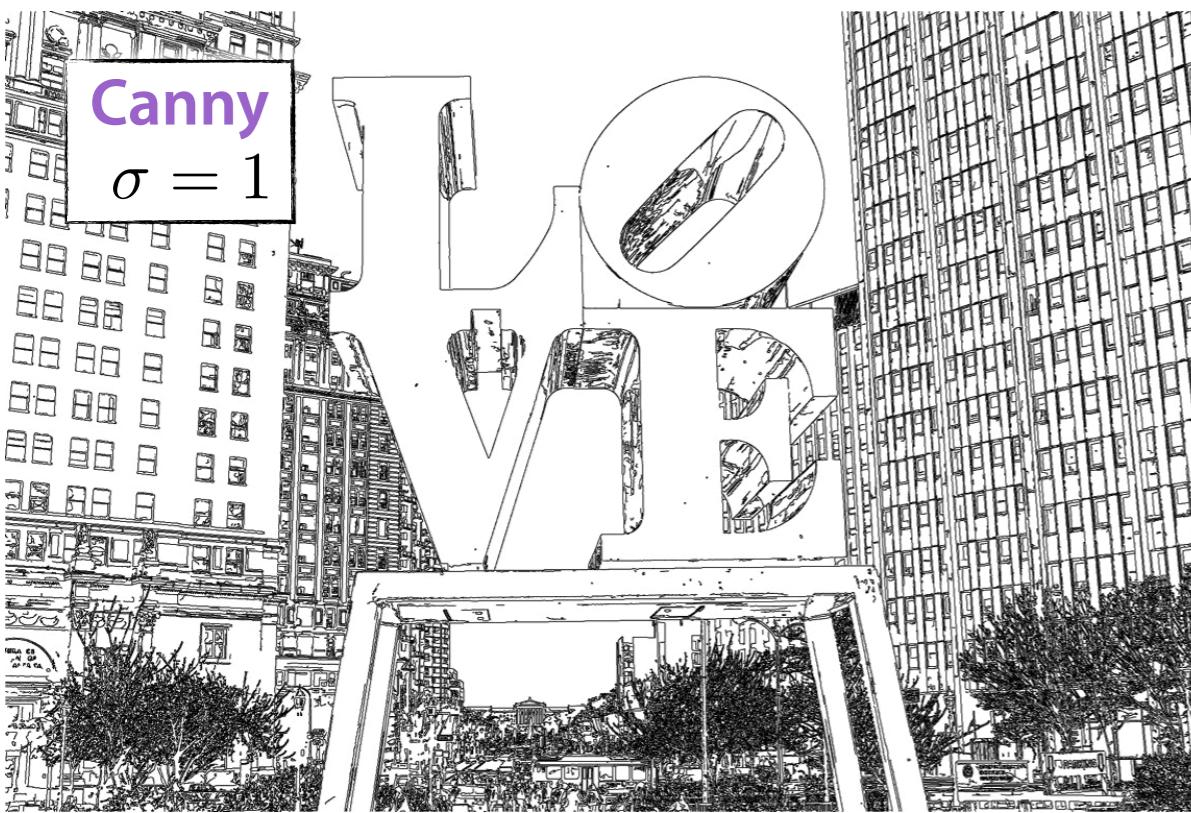
$\sigma = 1$



# Canny

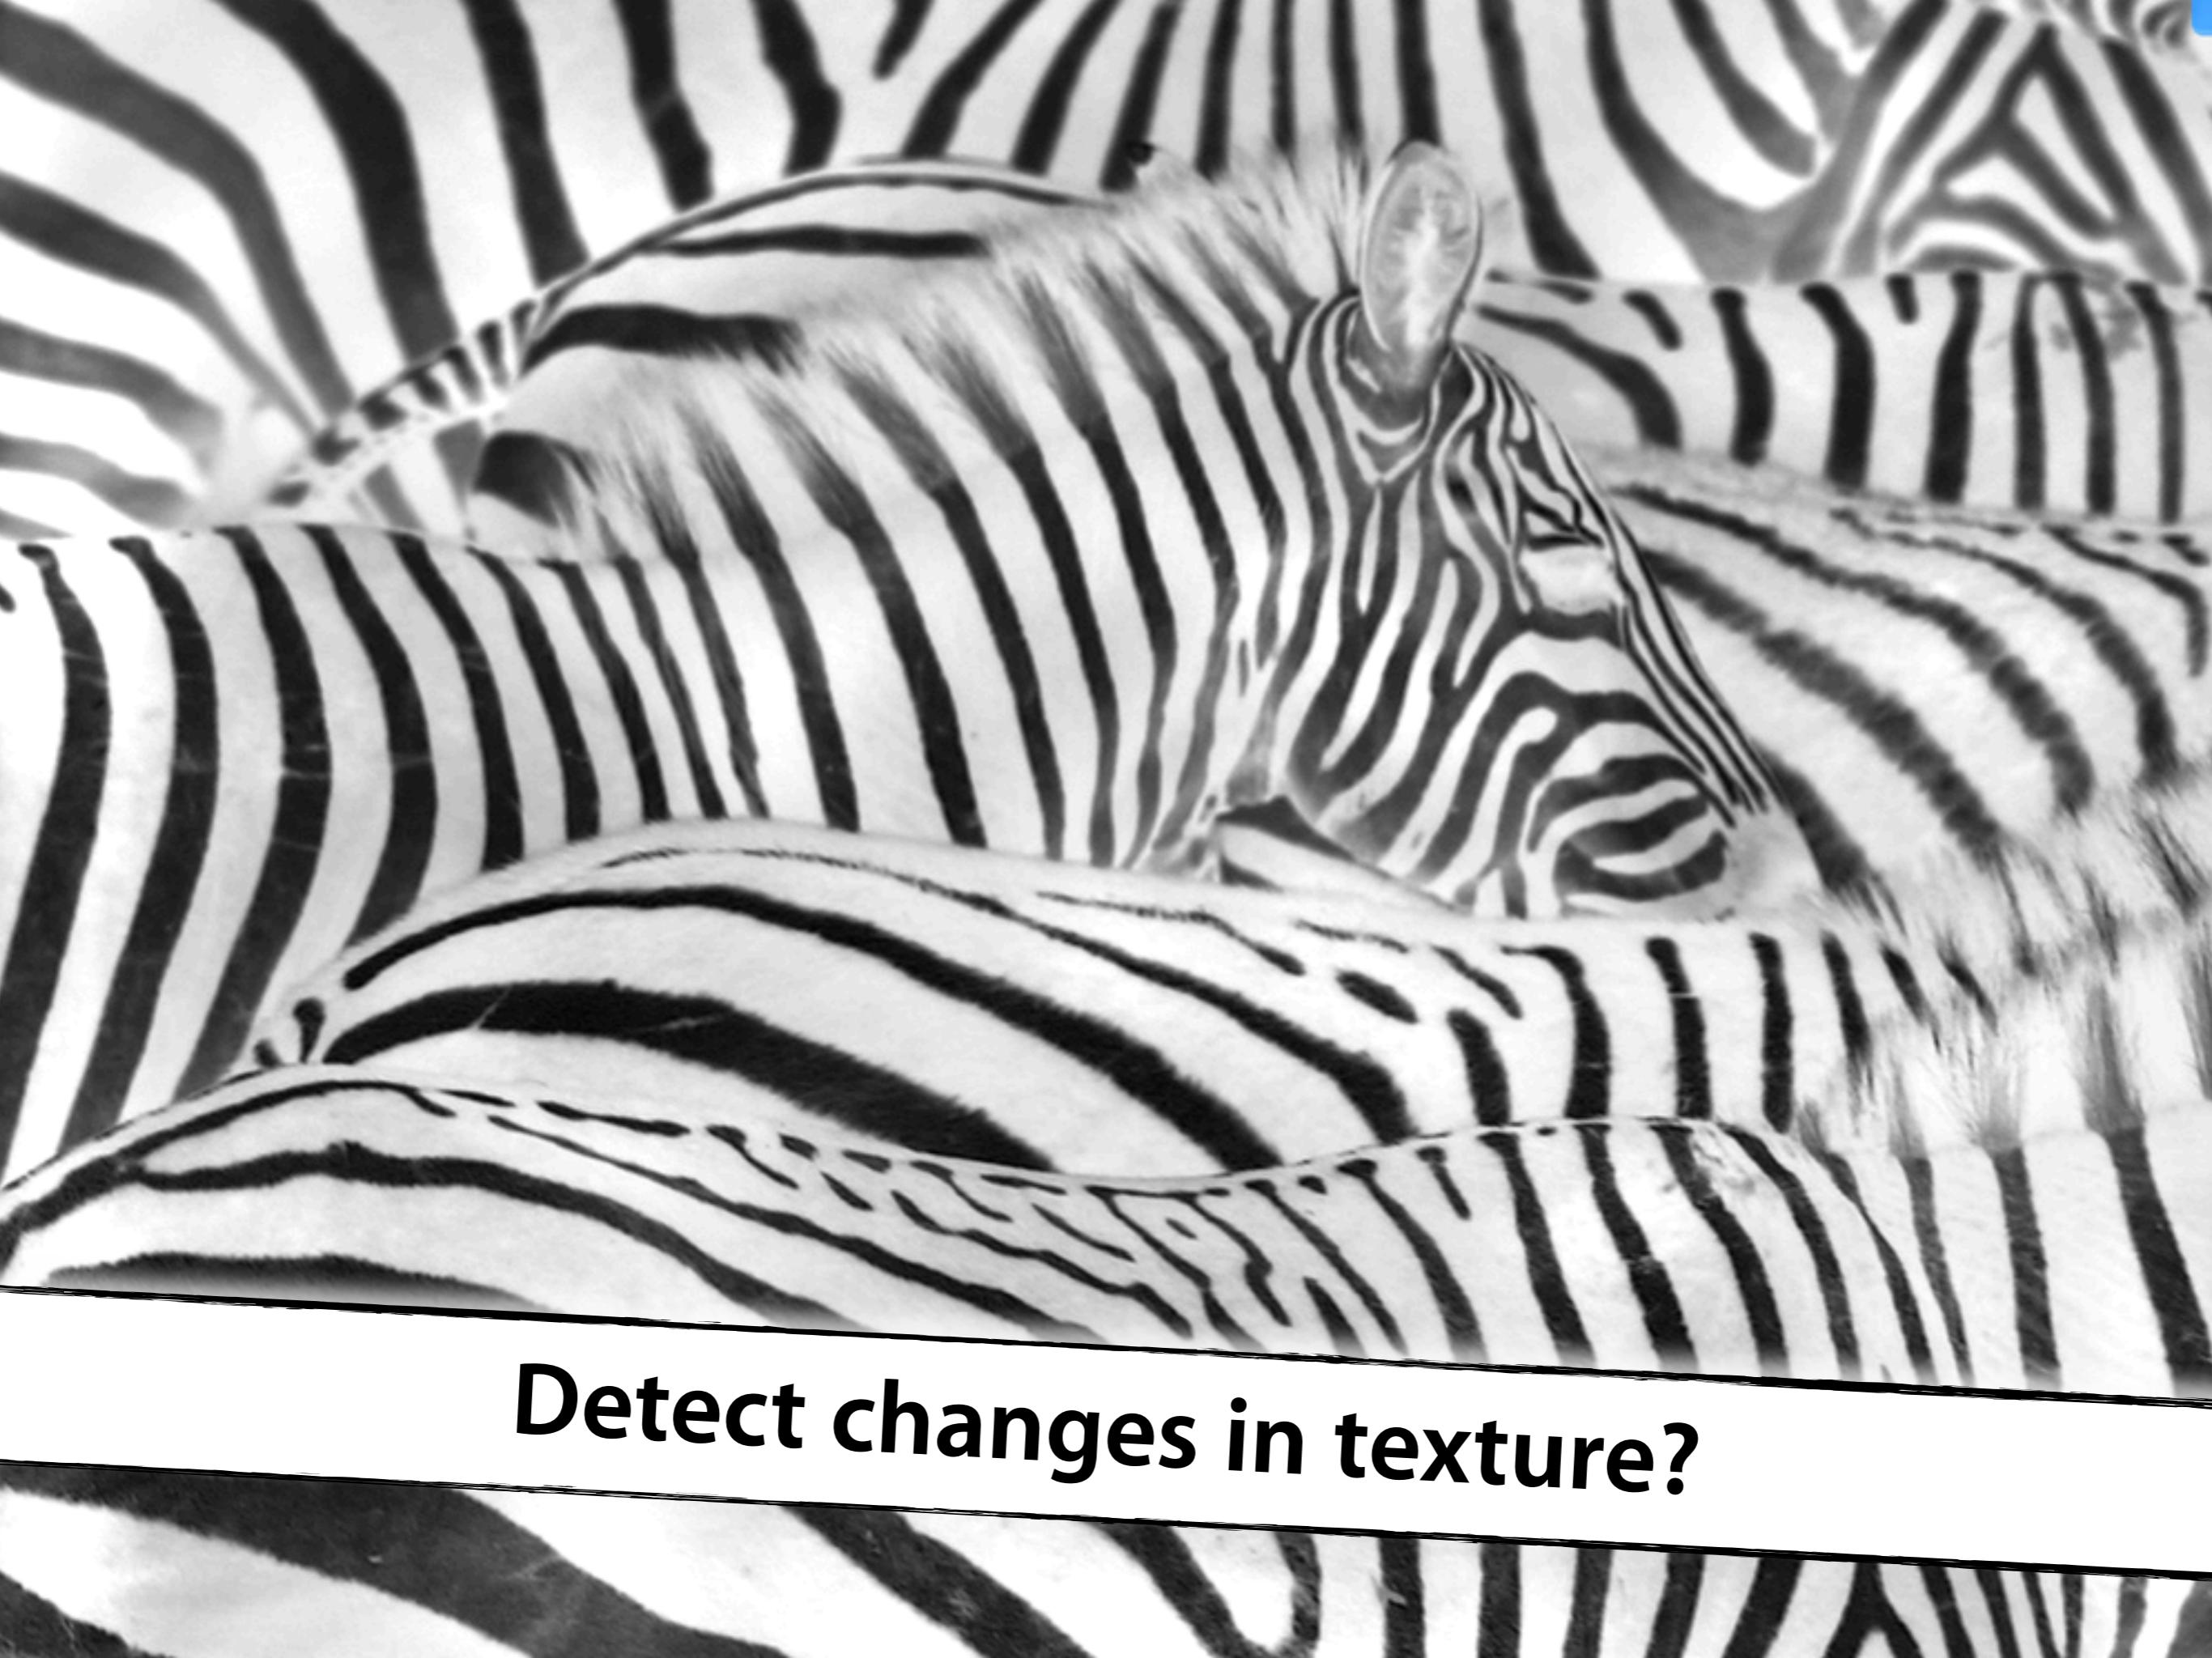
$\sigma = 5$







**Detect changes in colour?**



**Detect changes in texture?**

**How can we combine cues to detect object contours?**

A white humanoid robot with a metallic finish and glowing blue eyes is shown from the waist up, holding a yellow book open with both hands and looking down at it intently. It is positioned in front of several stacks of books of various colors, including blue, green, pink, and brown. The background is black.

# Machine learning

$x$

feature

$\hat{f}$

$y$

class label

Given training data, estimate the function  $f$

$$\{(x_i, y_i)\}$$



